

Poll: Why is multiplying numbers by 100 easier than multiplying by 128?

EECS 370 - Lecture 2

Binary and
Instruction Set Architecture (ISA)



Live Poll + Q&A: [slido.com #eecs370](https://slido.com/#eecs370)

Poll and Q&A Link

Announcements

- P1a
 - Posted today
- Labs
 - No lab Monday (holiday)
 - Lab 1 due next Wednesday
 - Attendance starts next Friday
- OH
 - Started

My Office Hours

- 2 types:
- Group:
 - In-person
 - 30 minutes right after class (2901 BBB)
 - Prioritize group questions over individual debugging
 - Starting today
- Individual:
 - Some in-person, some virtual
 - See Google calendar for details
 - One-on-one: any questions welcome

Extra Resources

- Want more examples on binary? Two's complement?
 - See "resources tab" on website
 - Extra videos, review sheets

The screenshot displays the 'Course Resources' page for EECS 370. On the left is a dark sidebar with a list of navigation links: 'EECS 370', 'Calendar', 'Lecture', 'Discussion', 'Assignments', 'Exams', 'Admin Requests', 'Schedule', and 'Course Resources' (which is highlighted in teal). The main content area is titled 'Course Resources' and is organized into four columns, each with a header and a list of resource links in purple boxes:

- Review Content**
 - EECS 370 Youtube Channel
 - Binary, Hex, and 2's complement Review Sheet
- Simulators**
 - Cache Simulator
 - Pipeline Simulator
- Reference Material**
 - Green LEGv8 Cheat Sheet
 - C for C++ users by Ian Cooke
 - Symbol Table and Relocation Table for EECS 370
- GDB Content**
 - GDB Tutorial
 - GDB Reference Card

Instruction Set Architecture (ISA) Design Lectures

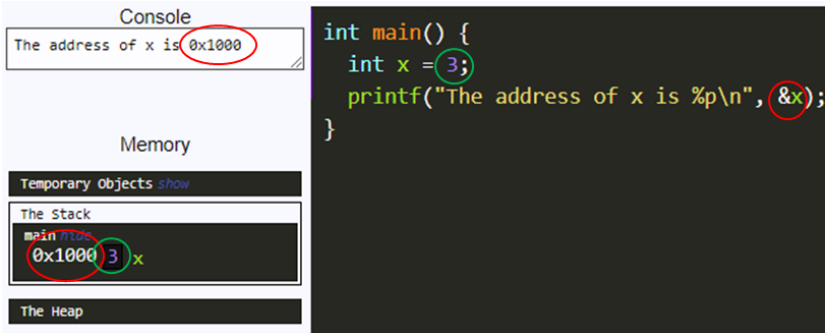
- **Lecture 2: ISA - storage types, binary and addressing modes**
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout

Agenda

- **Computer Model and Binary**
- ISAs
 - Registers
 - Control Flow
 - Representing Different Values

Basic Computer Model

- You know from 280 that computers have "memory"
 - Abstractly, a long array that holds values
- Every piece of data in a running program lives at a numerical **address** in memory
 - You can see the address in C by using the "&" operator



The image shows a debugger interface with three main components:

- Console:** Displays the output "The address of x is 0x1000". The address "0x1000" is circled in red.
- Memory:** Shows the stack frame for "main". It lists the address "0x1000" (circled in red), the value "3" (circled in green), and the variable "x".
- Code:** A C program snippet:

```
int main() {  
    int x = 3;  
    printf("The address of x is %p\n", &x);  
}
```

The value "3" is circled in green, and the "&x" is circled in red.

- Most programs work by loading values from memory to the processor, operating on those values, and writing values back into memory

Basic Memory Model

- 1st question in understanding how programs run on computers:
 - How are values actually represented in memory?
- Answer: binary

Aside: Decimal and Binary



- Humans represent numbers in base-10 (decimal) because we have 10 fingers (or "digits")
- The n^{th} digit corresponds to 10^n

$$\begin{aligned} & \text{1407} \\ &= 1 \cdot 10^3 + 4 \cdot 10^2 + 0 \cdot 10^1 + 7 \cdot 10^0 \\ &= 1000 + 400 + 00 + 7 \end{aligned}$$

Collection of 8
bits is called a
byte

- Computers are made of wires with either high or low voltages
- Internally represents values in base-2 (binary) since it has "binary digits"
 - (or bits for short)
- In binary, the n^{th} bit corresponds to 2^n

$$\begin{aligned} & \text{1101} \\ &= 1 \cdot 2^3 + 1 \cdot 2^2 + 0 \cdot 2^1 + 1 \cdot 2^0 \\ &= 8 + 4 + 0 + 1 \\ &= 13 \end{aligned}$$



*Does Bart Simpson
count in octal?*

Aside: Hexadecimal

- A bunch of 0s and 1s is hard to read for humans
 - But translating to decimal and back is tricky
- Solution: Bases that are a power of 2 are easy to translate between, since a fixed group of bits corresponds to one digit
- In practice, base-16 or **hexadecimal** is used
 - Digits 0-9, plus letters A-F to represent 10-15

Aside: Hexadecimal

Represent binary using 0b. Hex
using 0x. If not specified, it's
decimal

- Every 4 bits corresponds to 1 hex digit (since $2^4=16$)

(binary)	0b	0010	0101	1010	1011
(hexadecimal)	0x	2	5	A	B

eg: 52 in binary? Hex?

$\frac{1}{32}$	$\frac{1}{16}$	$\frac{0}{8}$	$\frac{1}{4}$	$\frac{0}{2}$	$\frac{0}{1}$	0b 110100	0x 25AB
0	0	1	1	0	1	0	0
0x 34							

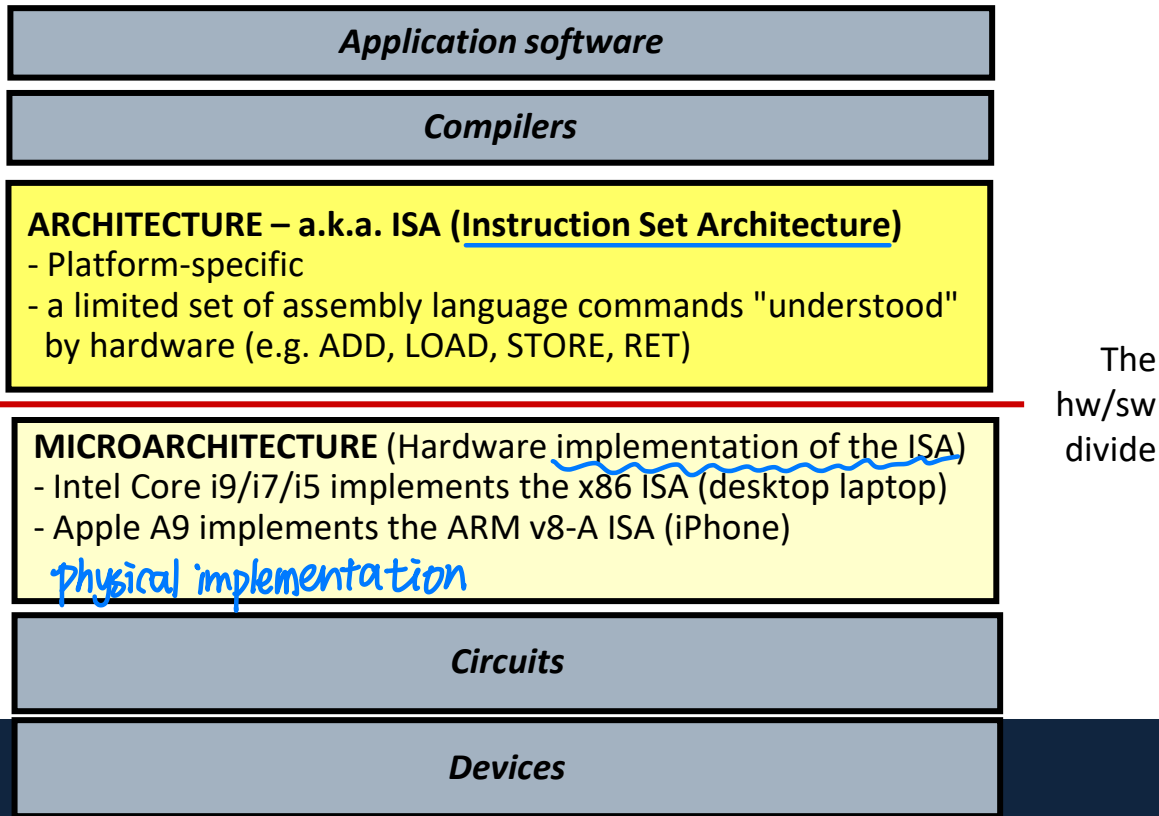
Other Units in this Class

Unit	Number of Bytes <i>8 bits.</i>
word	4 (in this class) <i>32 bits.</i>
Kilobyte (KB)	$2^{10} = 1,024$
Megabyte (MB)	$2^{20} = 1,048,576$
Gigabyte (GB)	$2^{30} = \text{About a billion}$

Agenda

- Computer Model and Binary
- **ISAs**
 - **Registers**
 - Control Flow
 - Representing Different Values

Where do ISAs come into the game ?



How is Assembly Different from C/C++?

- C/C++ instructions operate on **variables**

- e.g.

$x = i + j;$

- Practically unlimited

- We might guess that assembly instructions act on addresses, e.g.

$0x10000100 = 0x10000200 + 0x10000300$

- Problems:

1. This makes the instructions really long
2. As we'll see later in the course, memory is slow
 - We don't want to go multiple times for every instruction

How is Assembly Different from C/C++?

- Modern ISAs define **registers**
 - Basically a small number (~8-32) of fixed-length, hardware variables that have simple names like "r5"
- In a **load-store architecture** (what we'll assume in this class):
 - **load** instructions bring values from memory into a register
 - Other instructions specify register indices (compact and fast)
 - **store** instructions send them back to memory

Example Assembly Code

*Example ISA
(simplified)*

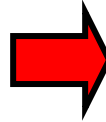
We'll talk more about how
loads / stores work later

0x1000
0x1004
0x1008

```
int a, b, c;  
main()  
{  
    a = a + b + c;  
}
```

C program

Compile



a r1 ← load(0x1000);
b r2 ← load(0x1004);
c r3 ← load(0x1008);
r1 ← add(r1, r2); } *r1 = r1 + r2 + r3*
r1 ← add(r1, r3)
r1 → store(0x1000)
*take the new value in r1 and write
back to 0x1000 address*

Assembly code

Example Architectures

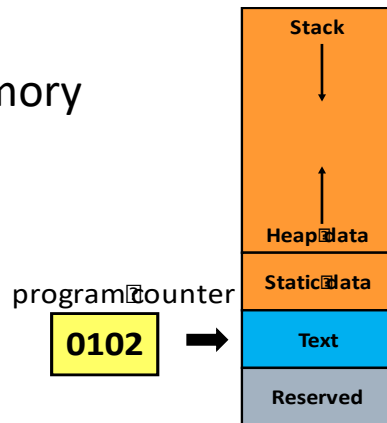
- ARMv8—LEGV8 subset from P+H text book
 - 32 registers (X0 – X31)
 - 64 bits in each register
 - Some have special uses e.g. X31 is always 0—XZR
- Intel x86 (not discussed much in this class)
 - 4 general purpose registers (eax, ebx, ecx, edx) 32 bits
 - Special registers: 3 pointer registers (si, di, ip), 4 segment (cs, ds, ss, es), 2 stack (sp, bp), status register (flags)
- LC2K (simple architecture made up for this class)
 - 8 registers, 32 bits each *≈ 4 billion*

Agenda

- Computer Model and Binary
- ISAs
 - Registers
 - **Control Flow**
 - Representing Different Values


How is Assembly Different from C/C++?

- C/C++: next line of code is executed until you get to:
 - function call
 - return statement
 - if statement or for/while loop
 - etc
- Assembly: a program counter (PC) keeps track of which memory address has the next instruction, gets incremented until
 - a "branch" or "jump" instruction
 - Used to change control flow (more later)
 - This model is called a **von Neumann Architecture**



Traditional (von Neumann) Architecture

Here's the (endless) loop that hardware repeats forever:

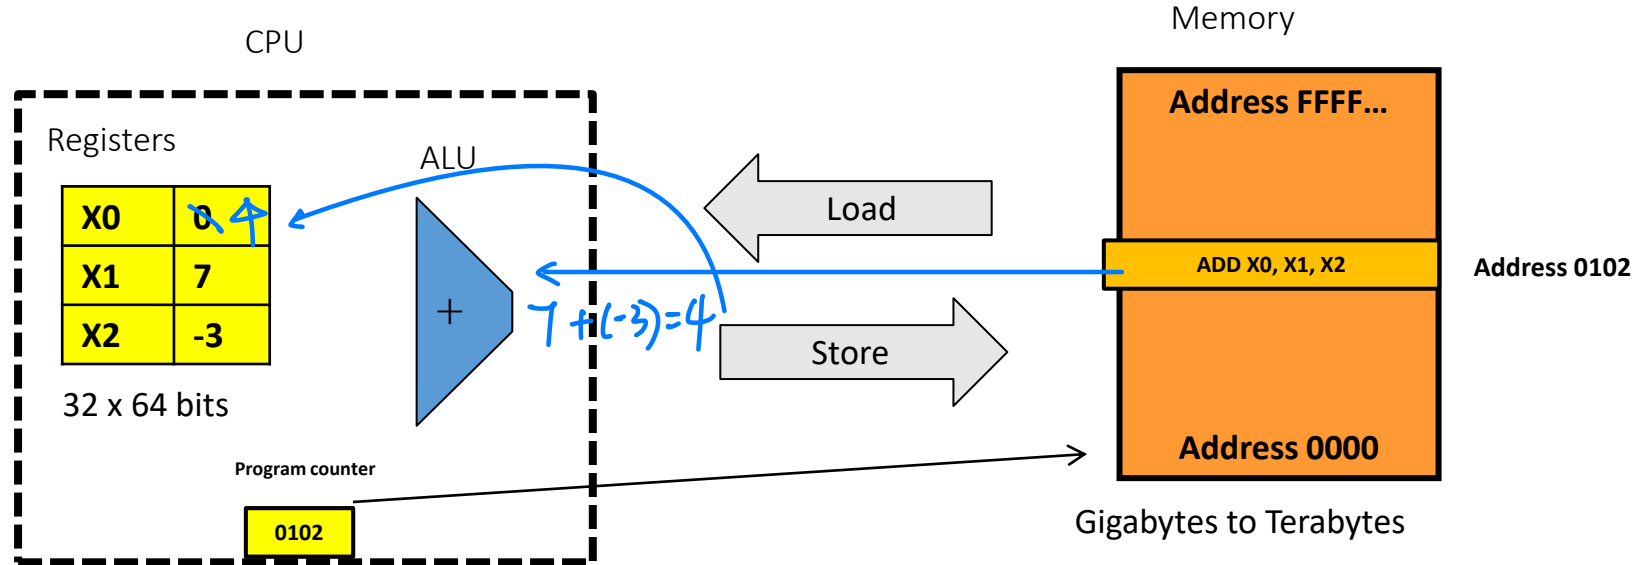
- 
1. Fetch—get next instruction—use PC to find where it is in memory and place it in instruction register (IR)
 - PC is changed to “point” to the next instruction in the program
 2. Decode—control logic examines the contents of the IR to decide what instruction it should perform
 3. Execute—the outcome of the decoding process dictates
 - an arithmetic or logical operation on data
 - an access to data in the same memory as the instructions
 - OR a change to the contents of the PC

Let's execute this short program
(destination register listed first):

```
ADD X0, X1, X2  
SUB X1, X2, X0
```

Destination

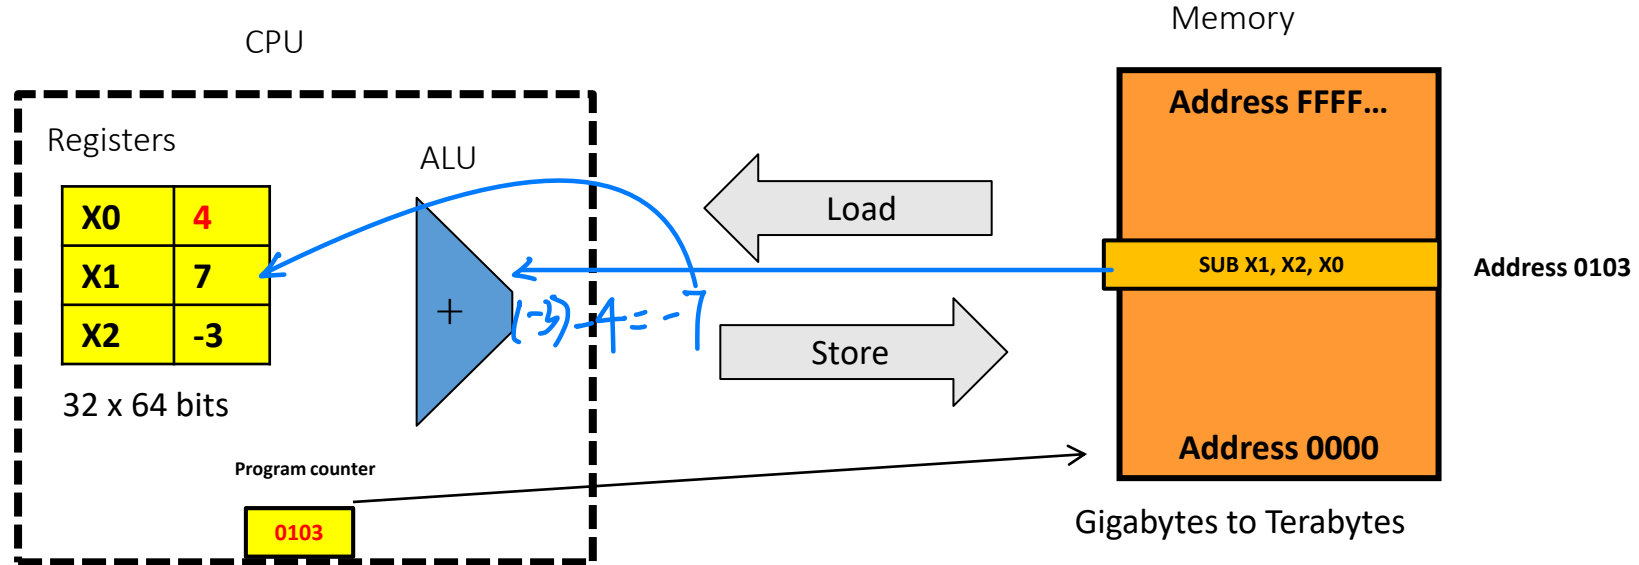
(Simplified) System Organization



Let's execute this short program
(destination register listed first):

ADD X0, X1, X2
SUB X1, X2, X0

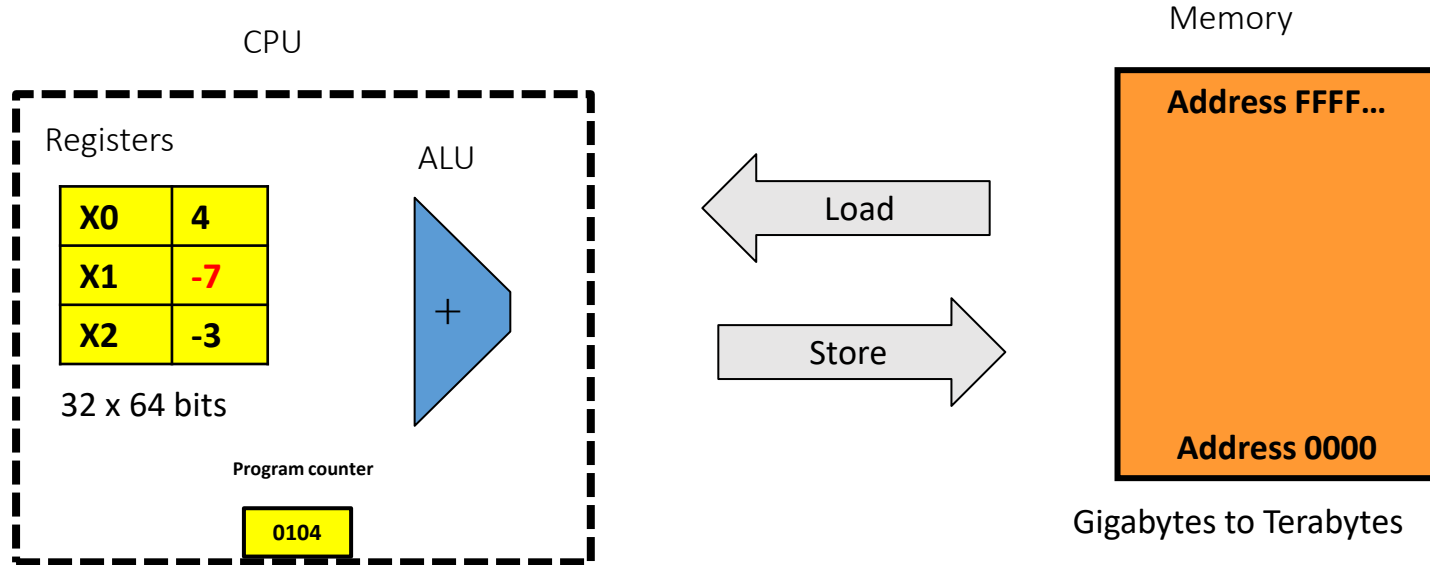
(Simplified) System Organization



Let's execute this short program
(destination register listed first):

ADD X0, X1, X2
SUB X1, X2, X0

(Simplified) System Organization



Assembly Code – ARM Example

Poll: What are the final contents of X1, X2, and X3?

- What are the contents of the registers after executing the given assembly code (destination register is listed first in ARM)?

Program:

opcode	d	s1	s2imm
ADD	X3	X1	X2
ADDI	X3	X3	#3
SUB	X2	X3	X1

ADDI means "add immediate", the last field is a literal value, not a register index

Initial register file:

X1	25
X2	-4
X3	57

		(1) ADD X3, X1, X2	(2) ADDI X3, X3, #3	(3) SUB X2, X3, X1
X1	25	25	25	25
X2	-4	-4	-4	-1
X3	57	21	24	24

Assembly Code – ARM Example

Poll: What are the final contents of X1, X2, and X3?

- What are the contents of the registers after executing the given assembly code (destination register is listed first in ARM)?

Program:

opcode	d	s1	s2imm
ADD	X3	X1	X2
ADDI	X3	X3	#3
SUB	X2	X3	X1

ADDI means "add immediate", the last field is a literal value, not a register index

Initial register file:

X1	25
X2	-4
X3	57

		(1) ADD X3, X1, X2	(2) ADDI X3, X3, #3	(3) SUB X2, X3, X1
X1	25	X1 25	X1 25	X1 25
X2	-4	X2 -4	X2 -4	X2 -1
X3	57	X3 21	X3 24	X3 24

Agenda

- Computer Model and Binary
- ISAs
 - Registers
 - Control Flow
 - **Representing Different Values**

Different Data Types

- How does memory distinguish between different data types?
 - E.g. int, int *, char, float, double *Below C. there are no*
- It doesn't! It's all just 0s and 1s! *data type*
- We'll see how to encode each of these later
- Exact length depends on architectures

How is Assembly Different from C/C++?

- No data types in assembly
- Everything is 0s and 1s: up to the programmer to interpret whether these bits should be interpreted as ints, bools, chars... or even instructions themselves!

```
char c = 'a';  
c++; // c is now 'b'
```

// results in the same assembly as

```
int x = 97;  
x++; // c is now 98
```

```
x = (int) c; // this instruction has no effect... why?
```

Minimum Datatype Sizes

Type	Minimum size (bits)
char	8
int (most) 32	16
long int	32
float	32
double	64

Representing Values in Hardware

- Unsigned integers represented as we've seen
- Chars are represented as ASCII values
 - e.g. 'a' -> 97, 'b' -> 98, '#' -> 35
- What about negative numbers?
- Fractional numbers?

Negative Numbers

- There are many ways we could represent negative numbers
- Because it will eventually make our hardware simpler, the most common representation is 2's complement



No, not 2's *compliment*!

Two's Complement Representation

- Recall that 1101 in binary is 13 in decimal.

$$\begin{array}{cccc} 1 & 1 & 0 & 1 \\ 2^3 & 2^2 & 2^1 & 2^0 \end{array} = 8 + 4 + 1 = 13$$

- 2's complement numbers are very similar to unsigned binary numbers.
 - The only difference is that the first number is now negative.

$$\begin{array}{cccc} 1 & 1 & 0 & 1 \\ -2^3 & 2^2 & 2^1 & 2^0 \end{array} = -8 + 4 + 1 = -3$$

Fun with 2's Complement Numbers

- What is the range of representation of a 4-bit 2's complement number?
 - [-8, 7] (corresponding to 1000 and 0111)
- What is the range of representation of an n-bit 2's complement number?
 - $[-2^{(n-1)}, 2^{(n-1)} - 1]$
- Useful trick: You can negate a 2's complement number by inverting all the bits and adding 1.
 - 5 is represented as **0101**
 - Negate each bit: **1010**
 - Add 1: **1011** = $-8 + 2 + 1 = -5$

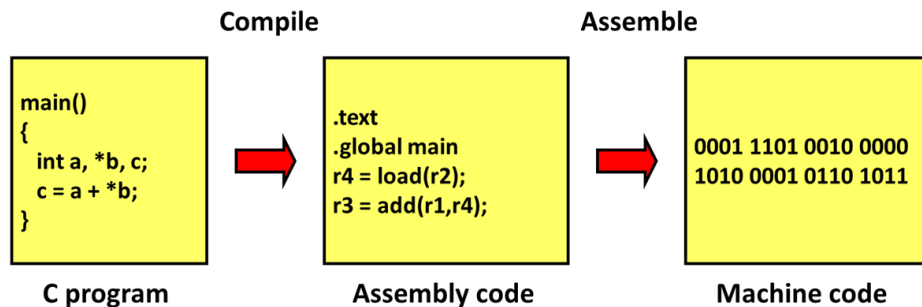
What about fractional numbers?

- One idea: fixed point notation
 - Have some bits represent numbers before decimal point, some bits represent numbers after decimal point
- Better idea: floating point notation
 - Inspired by scientific notation (e.g. 1.3×10^{-3})
 - Allows for larger range of numbers
 - We'll come back to this in a few lectures

Representing Instructions?

- Instructions, not just data, are stored in memory
- So, they must be expressible as numbers
- We'll look at how to encode instructions next time

*Example ISA
(simplified)*



Next Time

- Finish Up ISAs
- LC2K details
- Lingering questions / feedback? I'll include an anonymous form at the end of every lecture: <https://bit.ly/3oXr4Ah>



Extra Slides

Addressing Modes

- Direct addressing
- Register indirect
- Base + displacement
- PC-relative

Direct Addressing

- Consider this code:

```
const double PI = 3.14;  
  
double two_pi() {  
    return 2*PI;  
}
```

Not practical in modern ISAs...
if we have 32 bit instructions
and 32 bit addresses, the entire
instruction is the address!

- When we load PI, it's ALWAYS the same address
 - If the ISA supports it, we can just hardcode that address in the instruction
 - Like register addressing
 - Specify address as immediate constant
- ```
load r1, mem[1500] ; r1 ← contents of location 1500
jump mem[3000] ; jump to address 3000
```
- Useful for addressing locations that don't change during execution
    - Branch target addresses
    - Global/static variable locations



# Register indirect

- Consider this code:

```
int my_arr[2] = {6666, 7777};
int* ptr = &my_arr[0];
for(int i=0; i<2; i++) {
 int x = *ptr;
 ptr++;
}
```

- Everytime we load into x, it's a different address
- But the address is always stored in another variable
- If ISA supports it, we could use a load like this
  - **load r1, mem[r2]**

# Register indirect

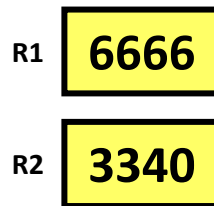
- Consider this code:

```
int my_arr[2] = {6666, 7777};
int* ptr = &my_arr[0];
for(int i=0; i<2; i++) {
 int x = *ptr;
 ptr++;
}
```

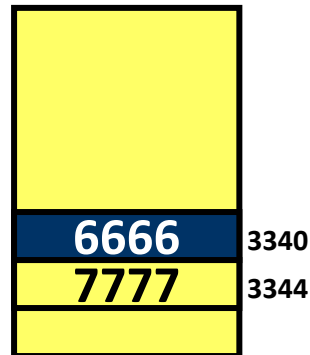


→ `load r1, mem[ r2 ]`  
`add r2, r2, #4`  
`load r1, mem[ r2 ]`

register file



memory



# Register indirect

- Consider this code:

```
int my_arr[2] = {6666, 7777};
int* ptr = &my_arr[0];
for(int i=0; i<2; i++) {
 int x = *ptr;
 ptr++;
}
```

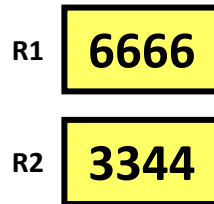


load r1, mem[ r2 ]

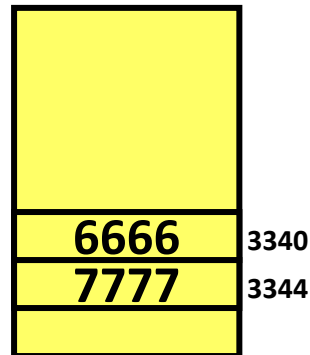
→ add r2, r2, #4

load r1, mem[ r2 ]

register file



memory



3340

3344

# Register indirect

- Consider this code:

```
int my_arr[2] = {6666, 7777};
int* ptr = &my_arr[0];
for(int i=0; i<2; i++) {
 int x = *ptr;
 ptr++;
}
```



load r1, mem[ r2 ]

add r2, r2, #4

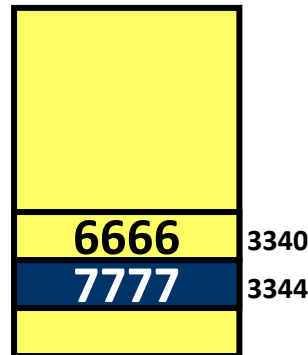
→ load r1, mem[ r2 ]

register file

R1 6666

R2 3344

memory



This is better, but we  
can be more general

# Base + Displacement

- Consider this code:

```
struct My_Struct {
 int tot;
 //...
 int val;
};
```

```
My_Struct a;
//...
a.tot += a.val;
```



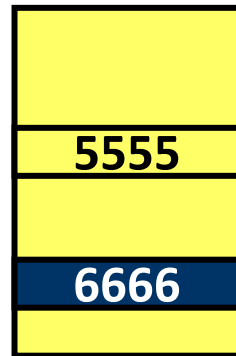
load r1, mem[r2 + 32]

register file

R2

2340

memory



2340

2372

- If a register holds the starting address of "a"...
- Then the specific values needed are a slight **offset**
- Base + Displacement**
  - reg value + immed

Very general, most  
common addressing  
mode today

# Class Problem

- a. What are the contents of register/memory after executing the following instructions

```
r2 = load mem[r3]
r3 = load mem[r2+4]
store mem[r2+8], r3
```

Poll: What are the contents of register / memory?

| register file |     | memory |     |
|---------------|-----|--------|-----|
| R1            | 0   | 108    | 100 |
| R2            | 10  | -1     | 104 |
| R3            | 108 | 100    | 108 |

# PC-relative addressing

- **Relevant for P1.a!**
- Variant on base + displacement
- Remember PC is "Program Counter", keeps track of which line (memory address) of the program we're at
- PC register is base, longer displacement possible since PC is assumed implicitly (more bits available)
  - Used for branch instructions
    - `jump [ - 8 ] ;` jump back 2 instructions (32-bit instructions)

# ISA Types

## Reduced Instruction Set Computing (RISC)

- Fewer, simpler instructions
- Encoding of instructions are usually the same size
- Simpler hardware
- Program is larger, more tedious to write by hand
- E.g. LC2K, RISC-V, ARM (kinda)
- More popular now

## Complex Instruction Set Computing (CISC)

- More, complex instructions
- Encoding of instructions are different sizes
- More complex hardware
- Short, expressive programs, easier to write by hand
- E.g. x86
- Less popular now

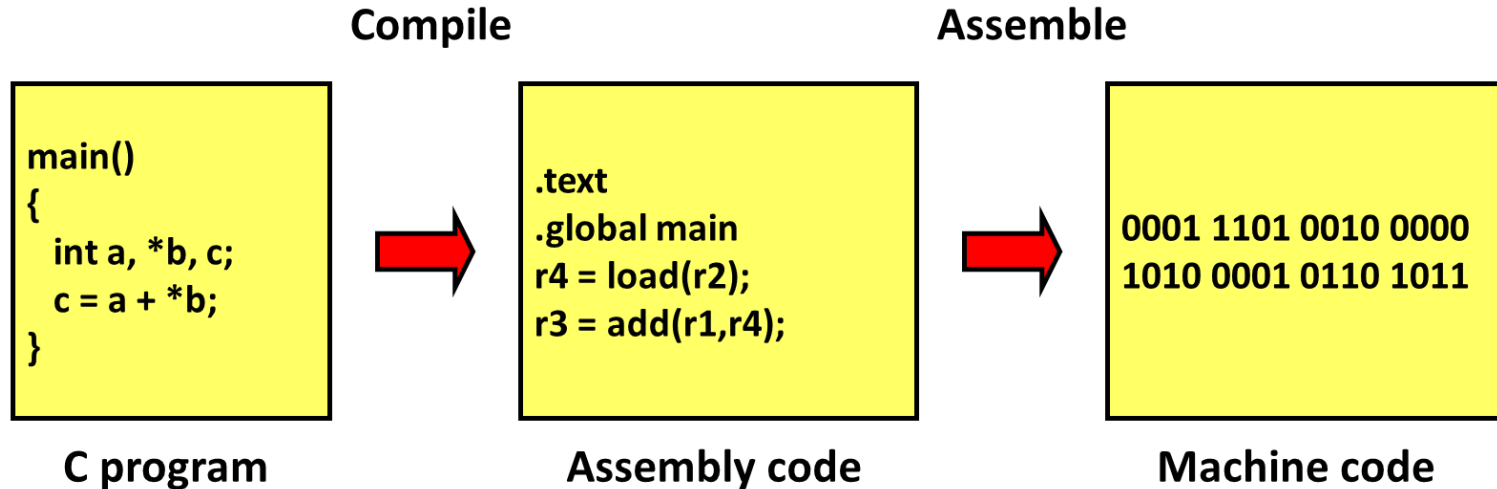


# Encoding Instructions

- So binary numbers can represent signed and unsigned numbers, chars, and fractional numbers
- But they must also represent instructions themselves!
  - After all, memory is just a collection of 1s and 0s
- We need a way of ***encoding*** instructions in order to store them in memory

# Software program to machine code

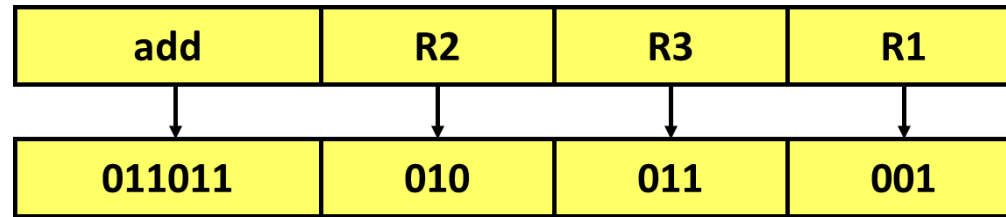
*Example ISA  
(simplified)*



# Assembly Instruction Encoding

- Since the EDSAC (1949) almost all computers stored program instructions the same way they store data.
- Each instruction is encoded as a number

*Example ISA  
(simplified)*



$$\begin{aligned} 011011010011001 &= 2^0 + 2^3 + 2^4 + 2^7 + 2^9 + 2^{10} + 2^{12} + 2^{13} \\ &= 13977 \end{aligned}$$

- This is the number stored in memory (in binary)!

**Poll:** How many different "operation codes" could be supported by this ISA? How many registers?

# Operating on Binary Values

- All values are stored in binary, even when you specify the number in decimal
- It is often convenient to treat values as sequences of bits, rather than values
  - You will need to do this in P1a
- C provides "bitwise operators" to do this
  - Shift (" $<<$ " and " $>>$ ")
  - Bitwise boolean (" $\&$ ", " $\|$ ", " $\wedge$ ", and " $\sim$ ")

# Shift Operators

- Shift a value x bits to the left via "<<"
- Inserts "x" zeros to the right (least significant)
- E.g.

```
int a = 60;
```

```
int s = a << 2;
```

# Shift Operators

- Shift a value  $x$  bits to the left via " $<<$ "
- Inserts " $x$ " zeros to the right (least significant)
- E.g.

```
int a = 60; // 0b0011_1100
```

```
int s = a << 2; // 0b1111_0000
```

- " $a$ " is still 60, " $s$ " is 240
- Same idea for " $>>$ ", but to the right

shifting  $x$  to the left in  
decimal  $\rightarrow$  multiplying  
by  $10^x$

shifting  $x$  to the left in  
binary  $\rightarrow$  multiplying by  
 $2^x$

# Bitwise operations

- Bitwise operations apply a Boolean operation on each bit of a value (or each pair of bits across two values)

```
int a = 60; // 0b0011_1100
```

```
int b = 13; // 0b0000_1101
```

```
int o = a | b; // 0b0011_1101
```

- "a" and "b" are the same, "o" is 61
- **&** – and    **|** – or    **^** – xor    **~** – not
- **Very different** from Boolean **&&**, **||**, etc
  - Why?