


# EECS 370

## Cache Introduction



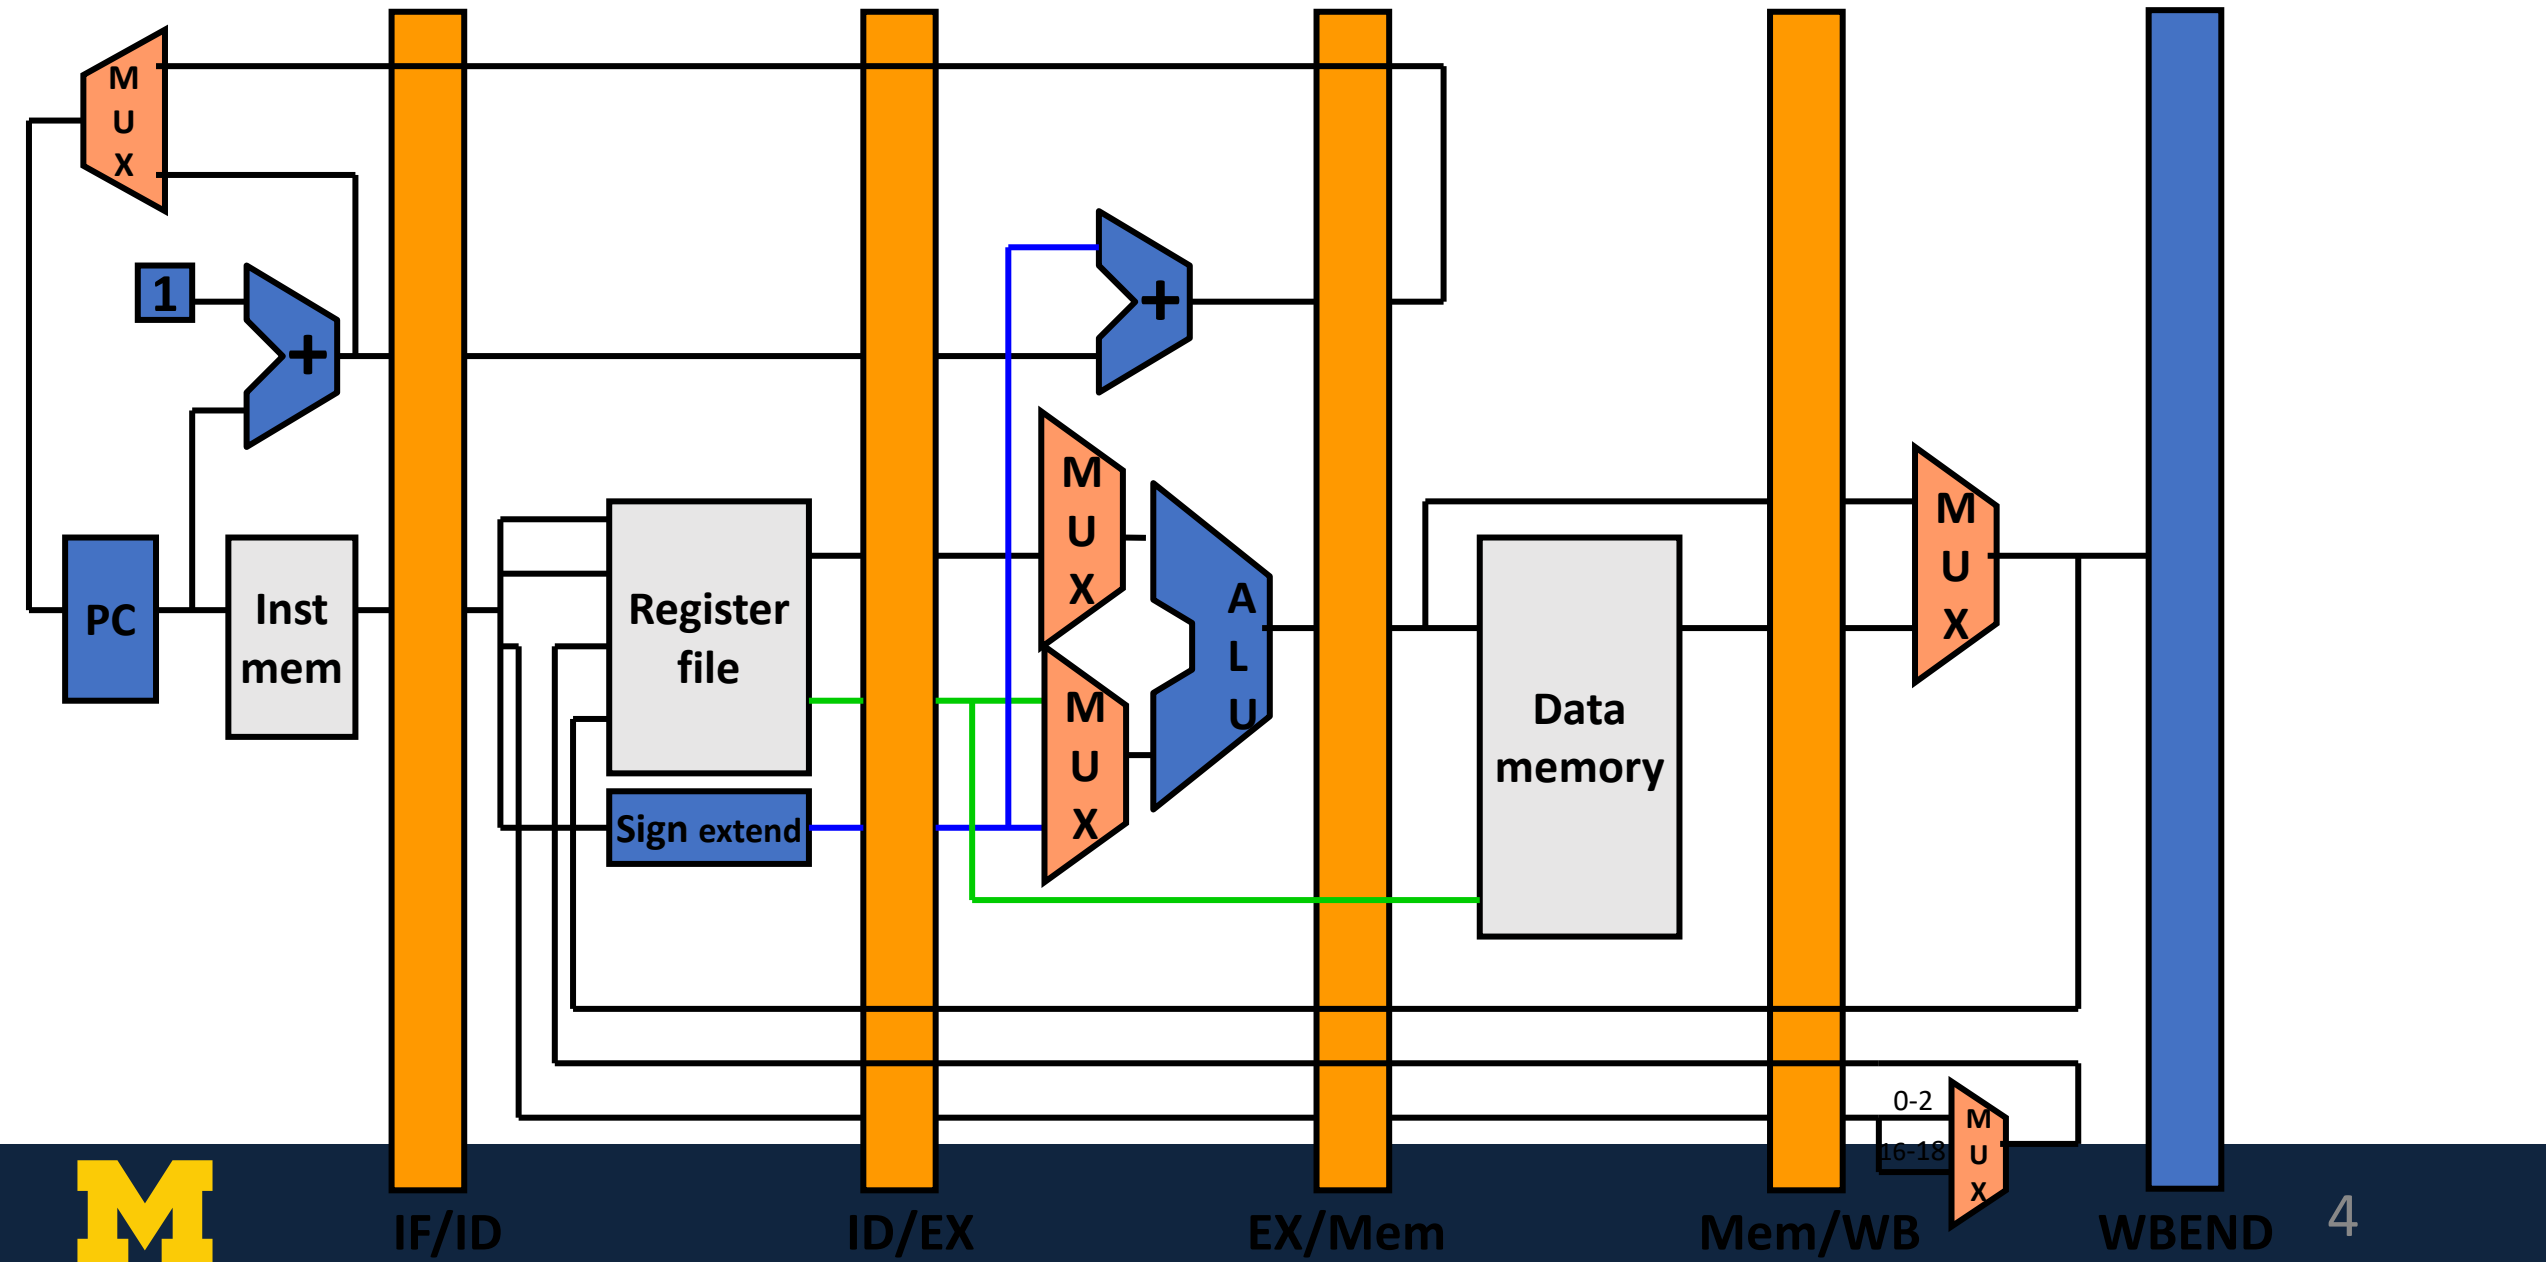
# Announcements

- My office hours on Wed 11/8 are cancelled
- Lab 8 due Wed
  - Lab 9 meets Fr/M
- P3 Checkpoint due Thurs
  - Full P3 due Thu 11/9
- Homework 3 due Mon 11/6

# Project 3 Pipeline

- Design is largely the same, except no "internal forwarding" through the register file
  - I.e. we need to explicitly forward from a new pipeline register WBEND if source instruction is in WB while dependent is in Decode

# Project 3 Pipeline




# Agenda

- **Memory Types**
- Memory Hierarchy and Cache Principles
- Cache example
- How to improve cache

# EECS 370 Overview

- Part I: Software
- Part II: Processor Design
- Part III: Memory Design
  - Starting with: caches
- If we judge a component's value by how much space it takes up on a chip (not a terrible heuristic), caches are **very** valuable



# Cache Aware vs Non-Aware Code

```
#include<stdio.h>
#include<stdlib.h>

#define N 20000
int arrayInt[N][N];

int main(int argc, char **argv)
{
    int i, j;
    int count = 0;

    for(i=0; i< N; i++)
        for(j = 0; j < N; j++ )
        {
            count++;
            arrayInt[i][j] = 10;
        }

    printf("Count :%d\n", count);
}
```

```
#include<stdio.h>
#include<stdlib.h>

#define N 20000
int arrayInt[N][N];

int main(int argc, char **argv)
{
    int i, j;
    int count = 0;

    for(i=0; i< N; i++)
        for(j = 0; j < N; j++ )
        {
            count++;
            arrayInt[j][i] = 10;
        }

    printf("Count :%d\n", count);
}
```

# Memory


- So far, we have discussed two structures that hold data:
  - Register file (little array of words)
  - Memory (bigger array of words)
- How do we build this?
  - We need a lot of memory:  $2^{18}$  for LC2K, a lot more for ARM
  - Bunch of flip-flops? Not practical – too many transistors, would be huge and power hungry
  - Other, clever ways of storing bits *when the bits is small, flip-flop is fast*



# Option 1: SRAM

*Static Random Access memory.*

- Each bit is made of 6 transistors
- Volatile: need constant power to keep data
- Fast: ~1 ns read time
- Still rather large
  - Can only put ~MBs on chip
  - Impractical to scale up to GBs needed for memory



# Option 2: DRAM


↓  
Dynamic

因为 capacitor 里面有的电压  
会不断下降, 所以要不断补充

- Each bit is made of a transistor and capacitor
- Volatile: need constant power to keep data
- Slower: ~50 ns access time
  - Must stall for dozens of cycles on each memory load
- Less expensive for SRAM
  - We can put up to ~16-64 GB in current machines
  - Good for LC2K or 32-bit systems
  - But not for 64-bit systems

同时也是慢的原因。

have voltage different  
(charged) : 1  
Don't have voltage different  
(uncharged) : 0



# Option 3: Disks

- Hard-drives store bits as magnetic charges on spinning disks
  - Non-volatile – holds information even when no power is supplied
  - Obnoxiously slow compared to digital logic: 3,000,000 ns access time
- Recently, solid-state drives have replaced spinning disks with logic gates replacing mechanical systems
  - Also non-volatile
  - Much better speeds (3,000,000 ns), but still too slow to keep up with processor
- Cheap!
  - SSDs cost \$0.0001 per megabyte
  - Scale up to terabytes – practical for modern computing

# Agenda

- Memory Types
- **Memory Hierarchy and Cache Principles**
- Cache example
- How to improve cache

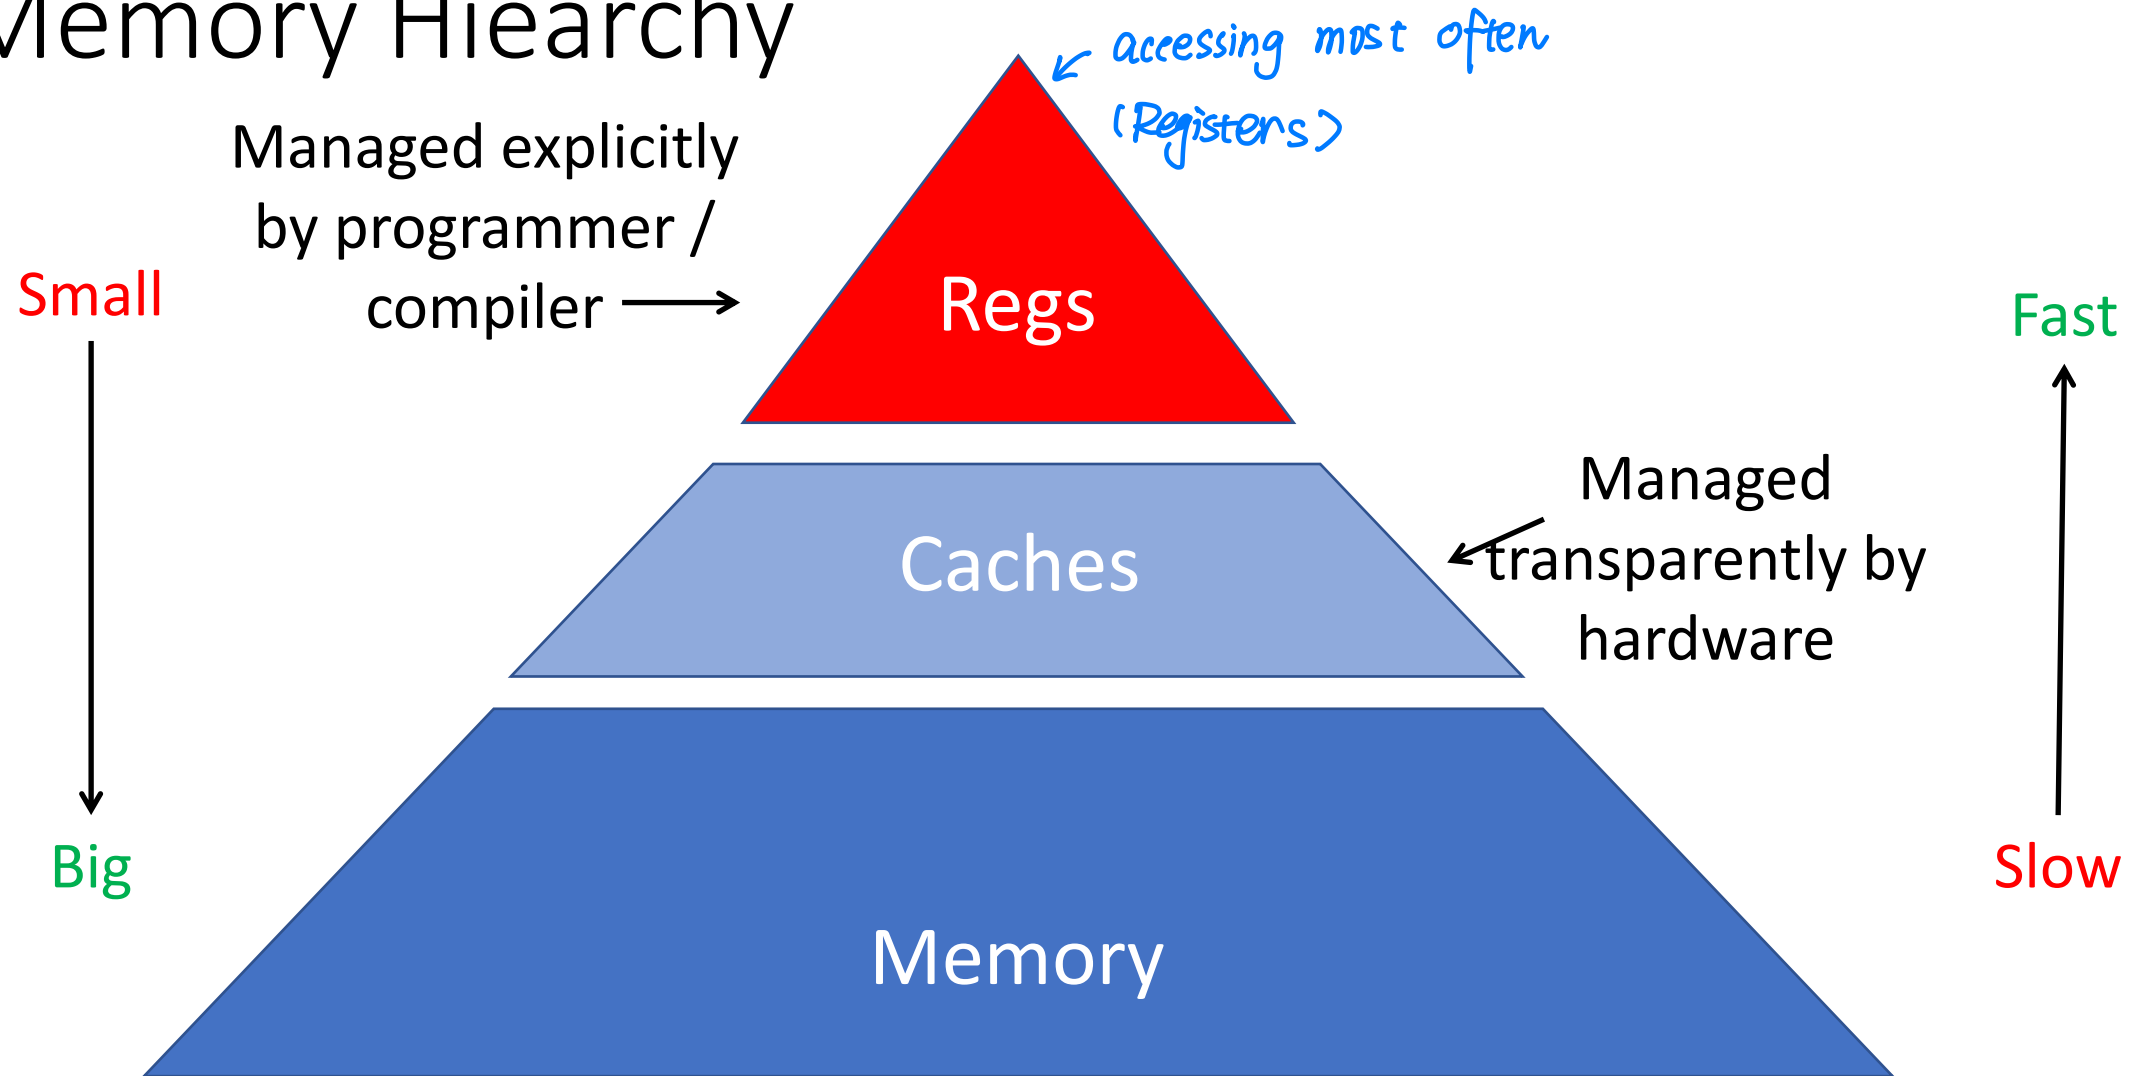
# Memory Goals

- Fast: Ideally run at processor clock speed
  - 1 ns access
- Cheap: Ideally free
  - Not more expensive than rest of system
- DRAM, hard-drives and SSDs are too slow
- SRAM is too expensive
- How to get best properties of multiple memory technologies?

# Memory Hierarchy

- Key observation: we only need to access a small amount of data at a time
- Let's use a small array of SRAM to hold data we need now
  - Call this the **cache**
  - Able to keep up with processor
  - Small (~Kilobytes), so it should be relatively cheap
- Use a large amount of DRAM for **main memory**
  - Can scale up to ~Gigabytes in size
- Everything else, store on disk
  - **Virtual Memory**
- Won't end up building  $2^{64}$  of anything
  - Won't be needed in typical programs
  - Virtual memory (discussed in a couple weeks) will make memory look larger than it is

# Memory Hierarchy




# Cache Analogy

- Studying in the library
  - **Option 1: Every time you grab another book, return current book to shelf and get a new book from shelf**
    - Latency = 5 minutes
  - **Option 2: Keep 10 commonly-used books on shelf above desk**
    - Latency = 1 minute
  - **Option 3: Keep 3 books open on different locations on desk**
    - Latency = 10 seconds



# Caches - Overview




# Function of the Cache

- The cache holds the data we think is **most likely** to be referenced
  - The more often the data we want is in the fast cache, the lower our **average memory access latency** is
  - How do we decide what the most likely accessed memory locations are?


# The simplest cache

- Only worry about load instructions for now
- Word-addressable address space
- Consists of a single, word-size storage location to remember last loaded value



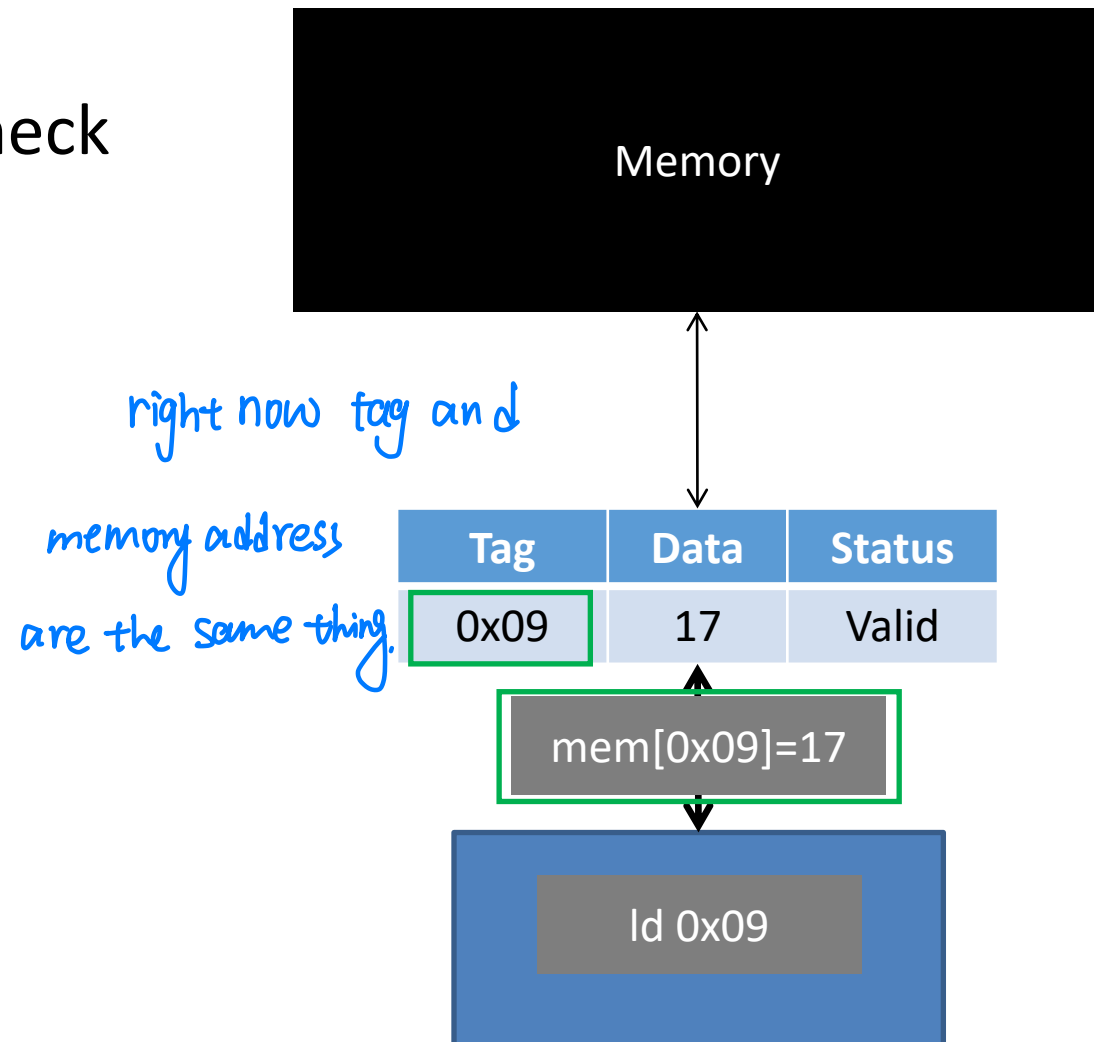
# The simplest cache

- Whenever memory returns data, store it in the cache
- We'll also need to know what address this data corresponds to
  - Store that as “**tag**”
- Also include a “**valid**” status bit



# The simplest cache

- Next memory access, first check if the tag matches address
  - Yes? Return cache data
  - No? Go to memory as before



# Agenda

- Memory Types
- Memory Hierarchy and Cache Principles
- **Cache example**
- How to improve cache

# Scaling Up

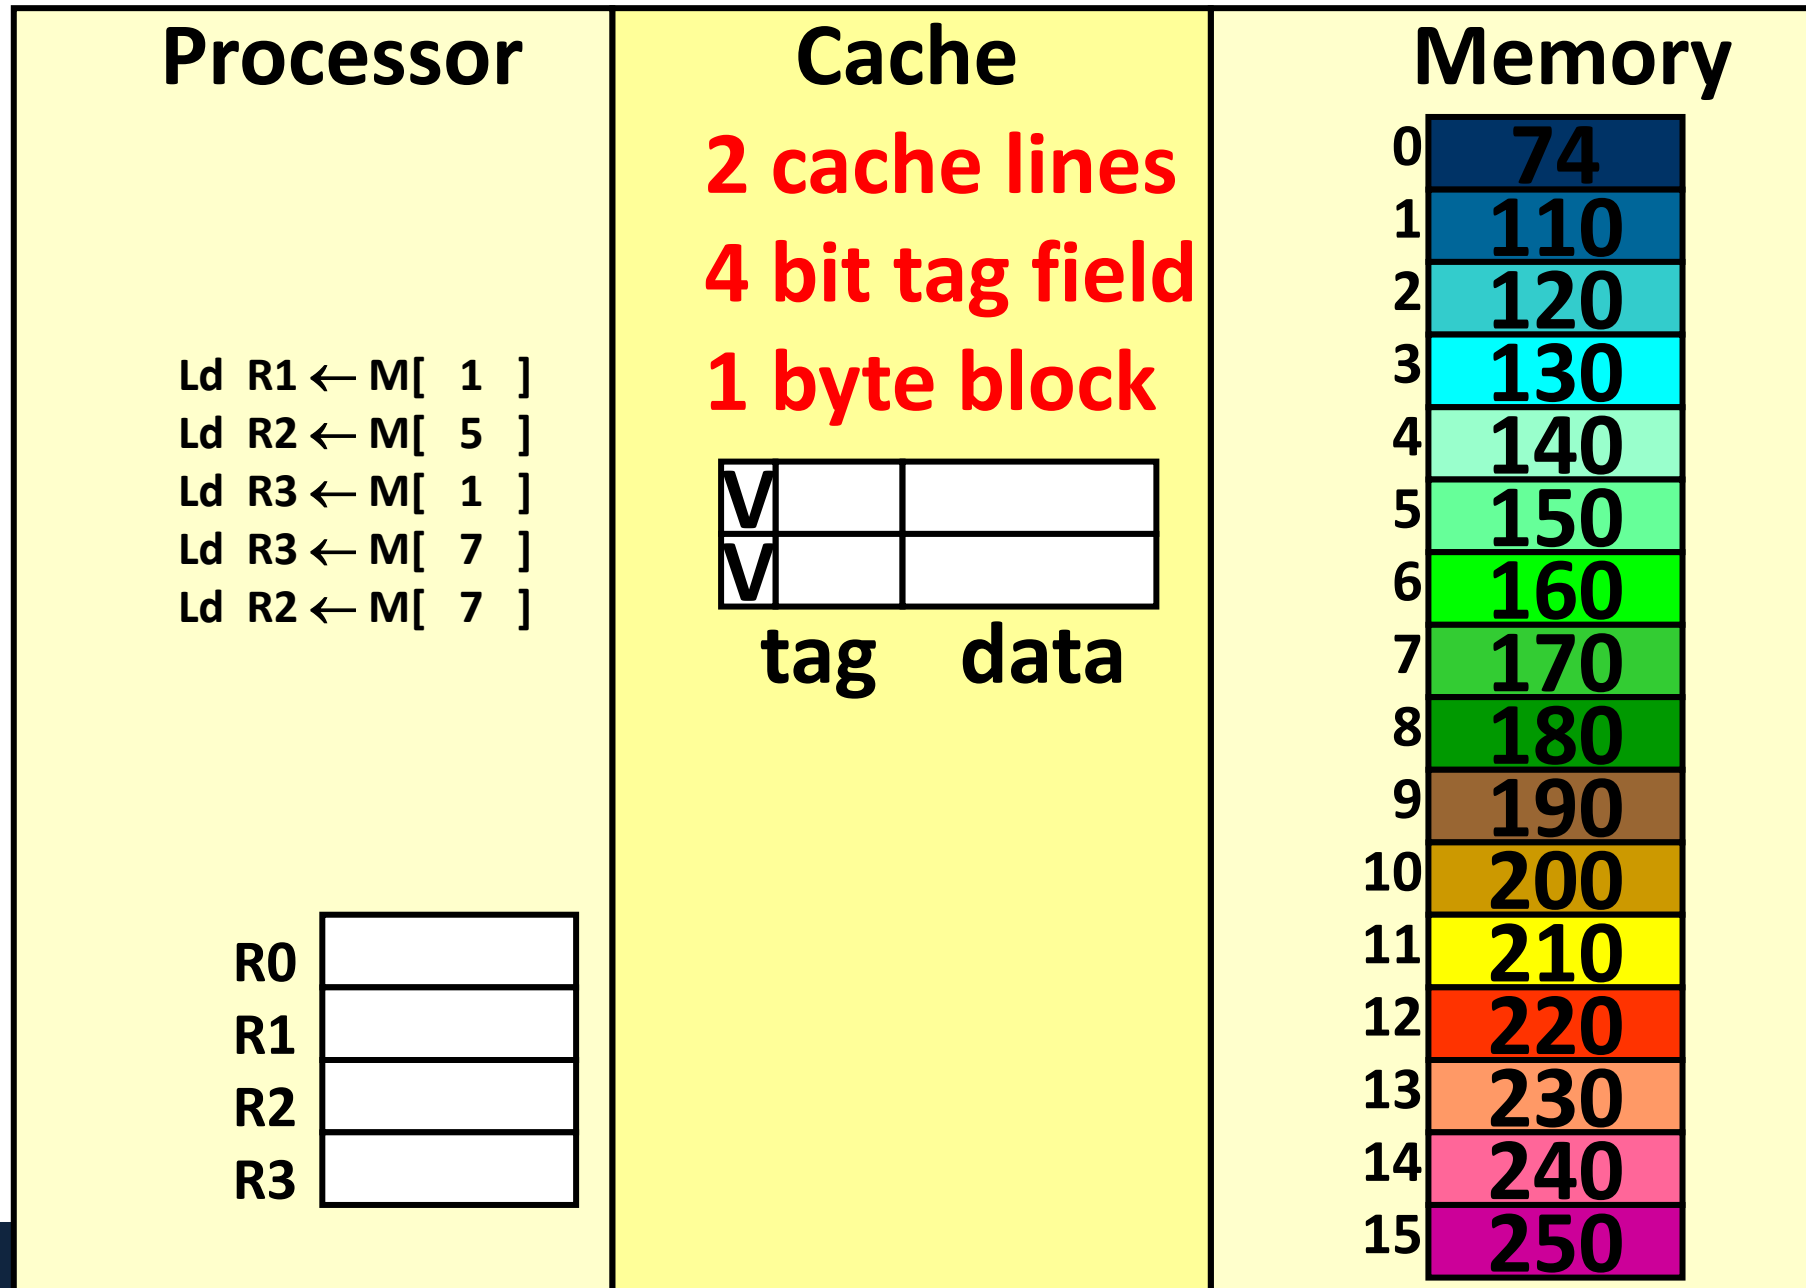
- Of course, we don't just access one memory location
- What if we access many memory locations, and cache isn't large enough to hold all of them?
- How do we choose what to keep in the cache?
  - How does hardware predict what's most likely to be needed soon?
- Answer: **locality**
  - Temporal locality
  - Spatial locality

# Temporal Locality

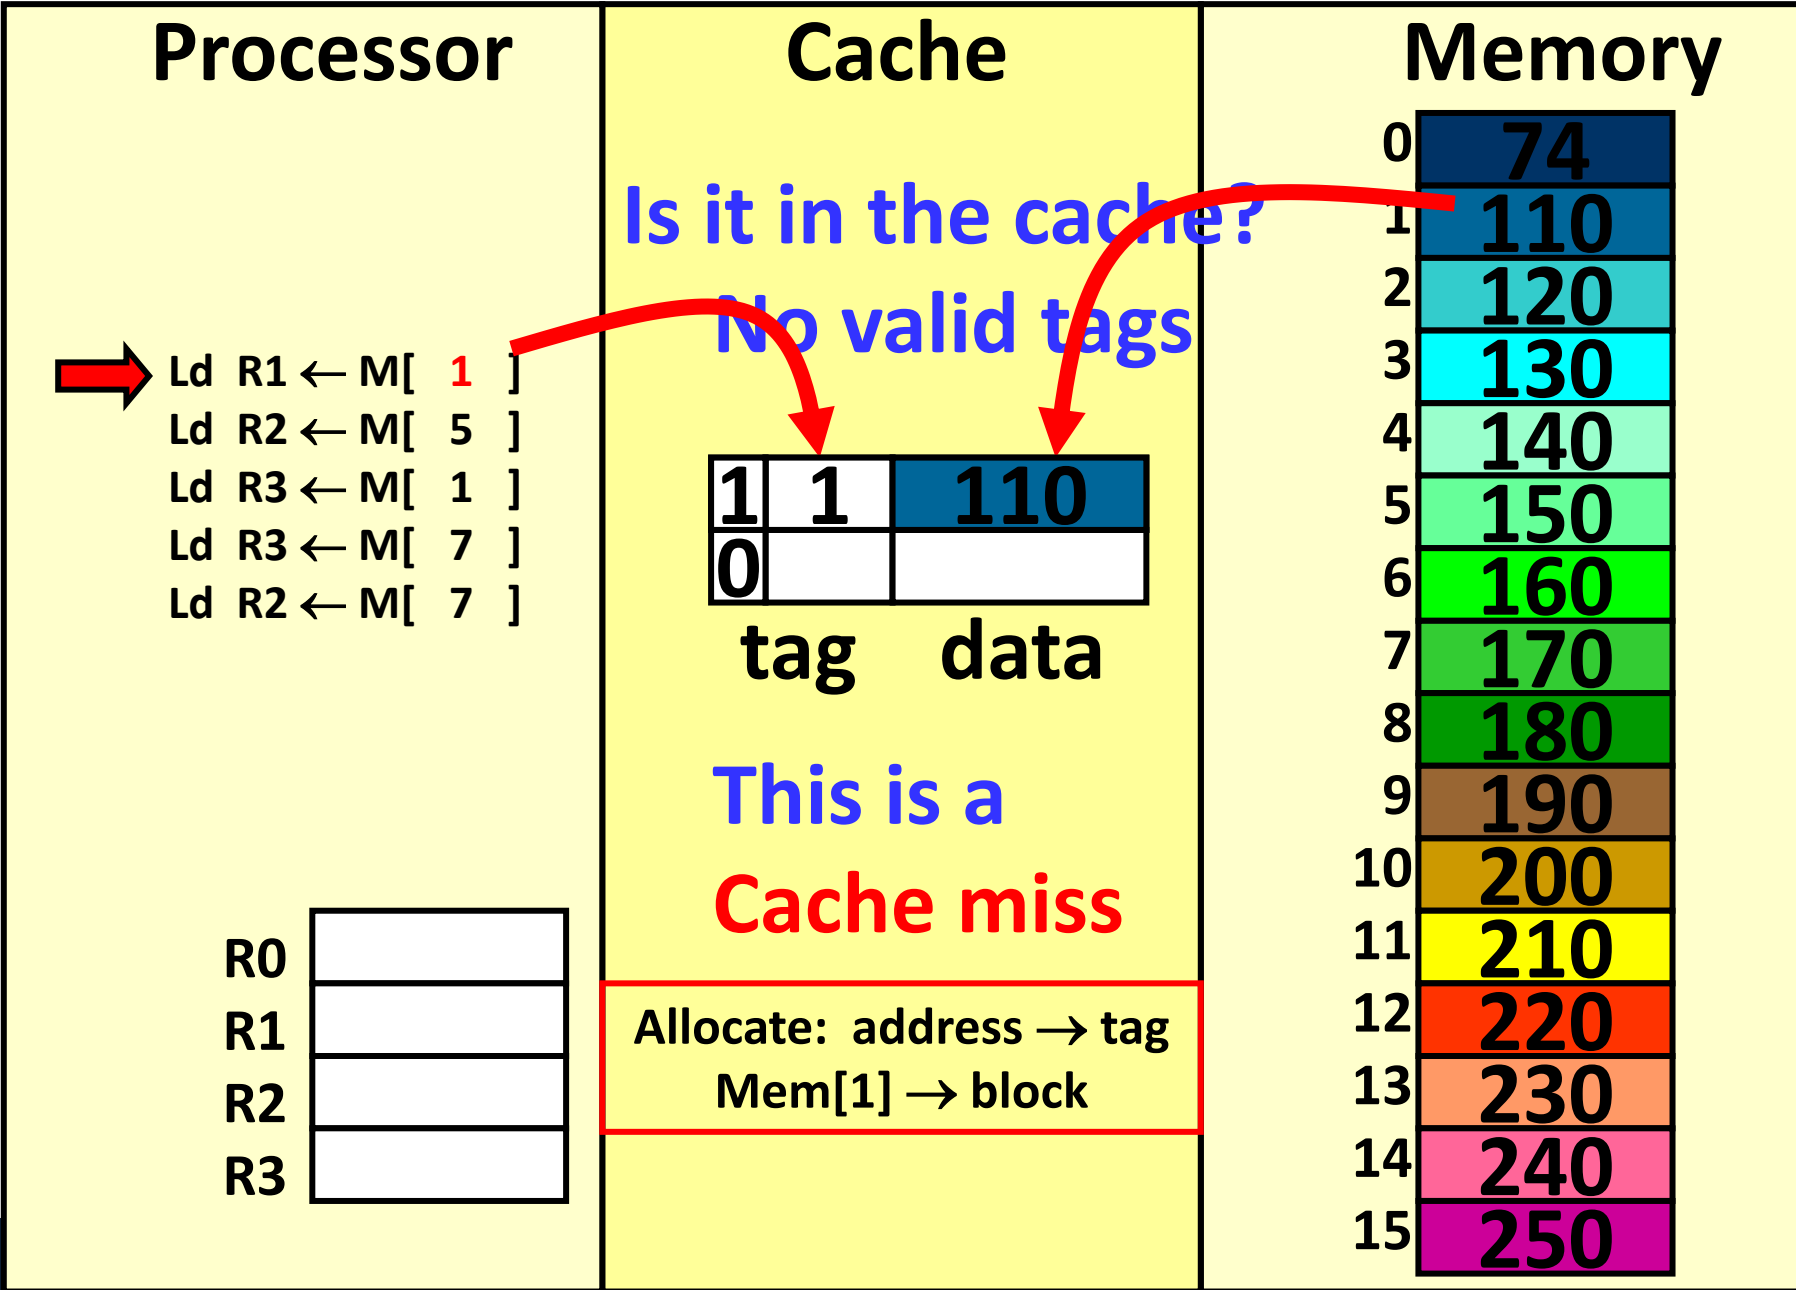
- Temporal locality: if a given memory location is referenced now, it will probably be used again in the near future
  - Why? Take a look at code you've written. You tend to use a variable multiple times
  - Corollary: if you haven't used a variable in a while, you probably won't need it very soon either
- Hardware should take advantage of this by:
  - Placing items we just accessed in the cache
  - When we need to evict something, evict whatever data was **least recently used (LRU)**



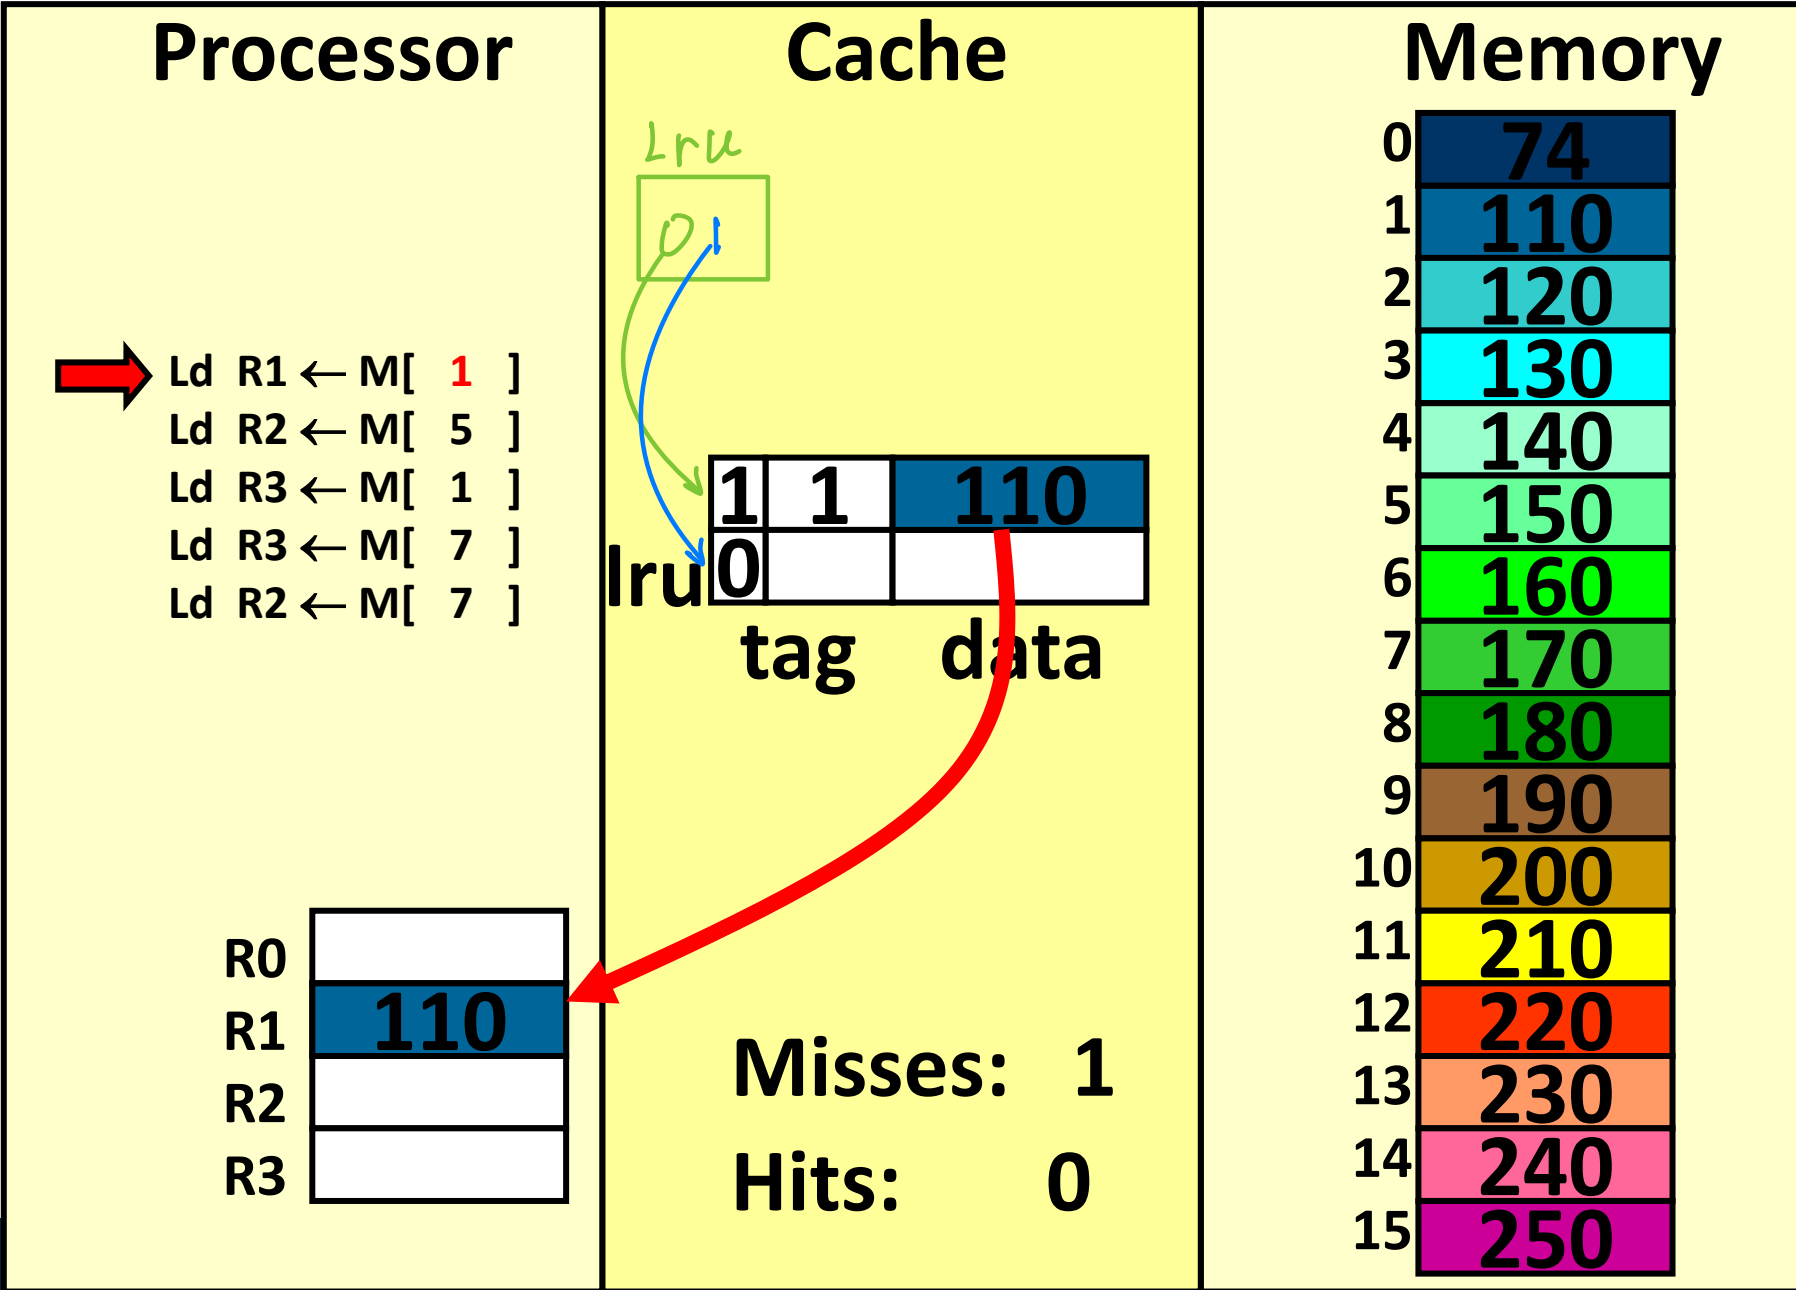
# A Very Simple Memory System



# A Very Simple Memory System

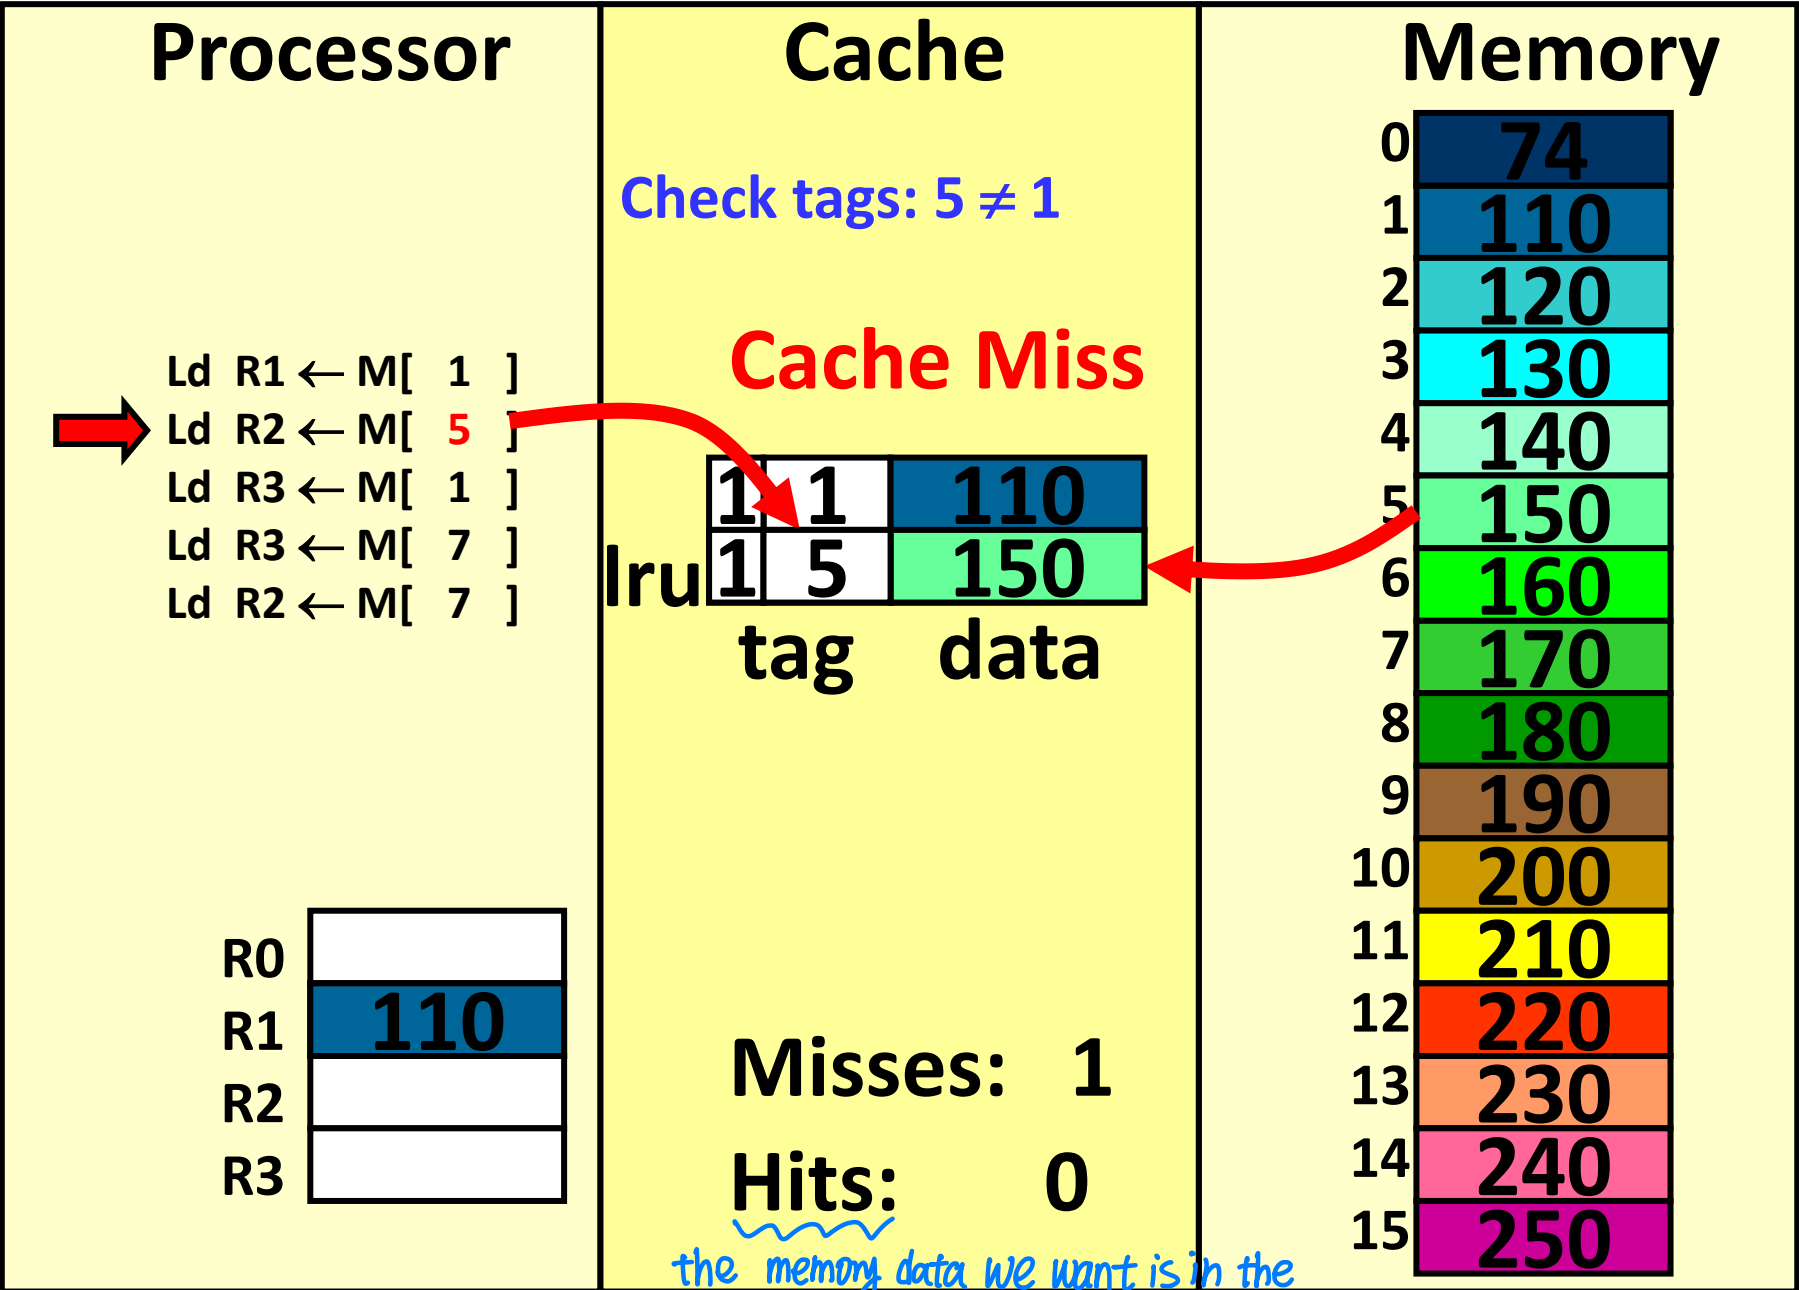


# A Very Simple Memory System

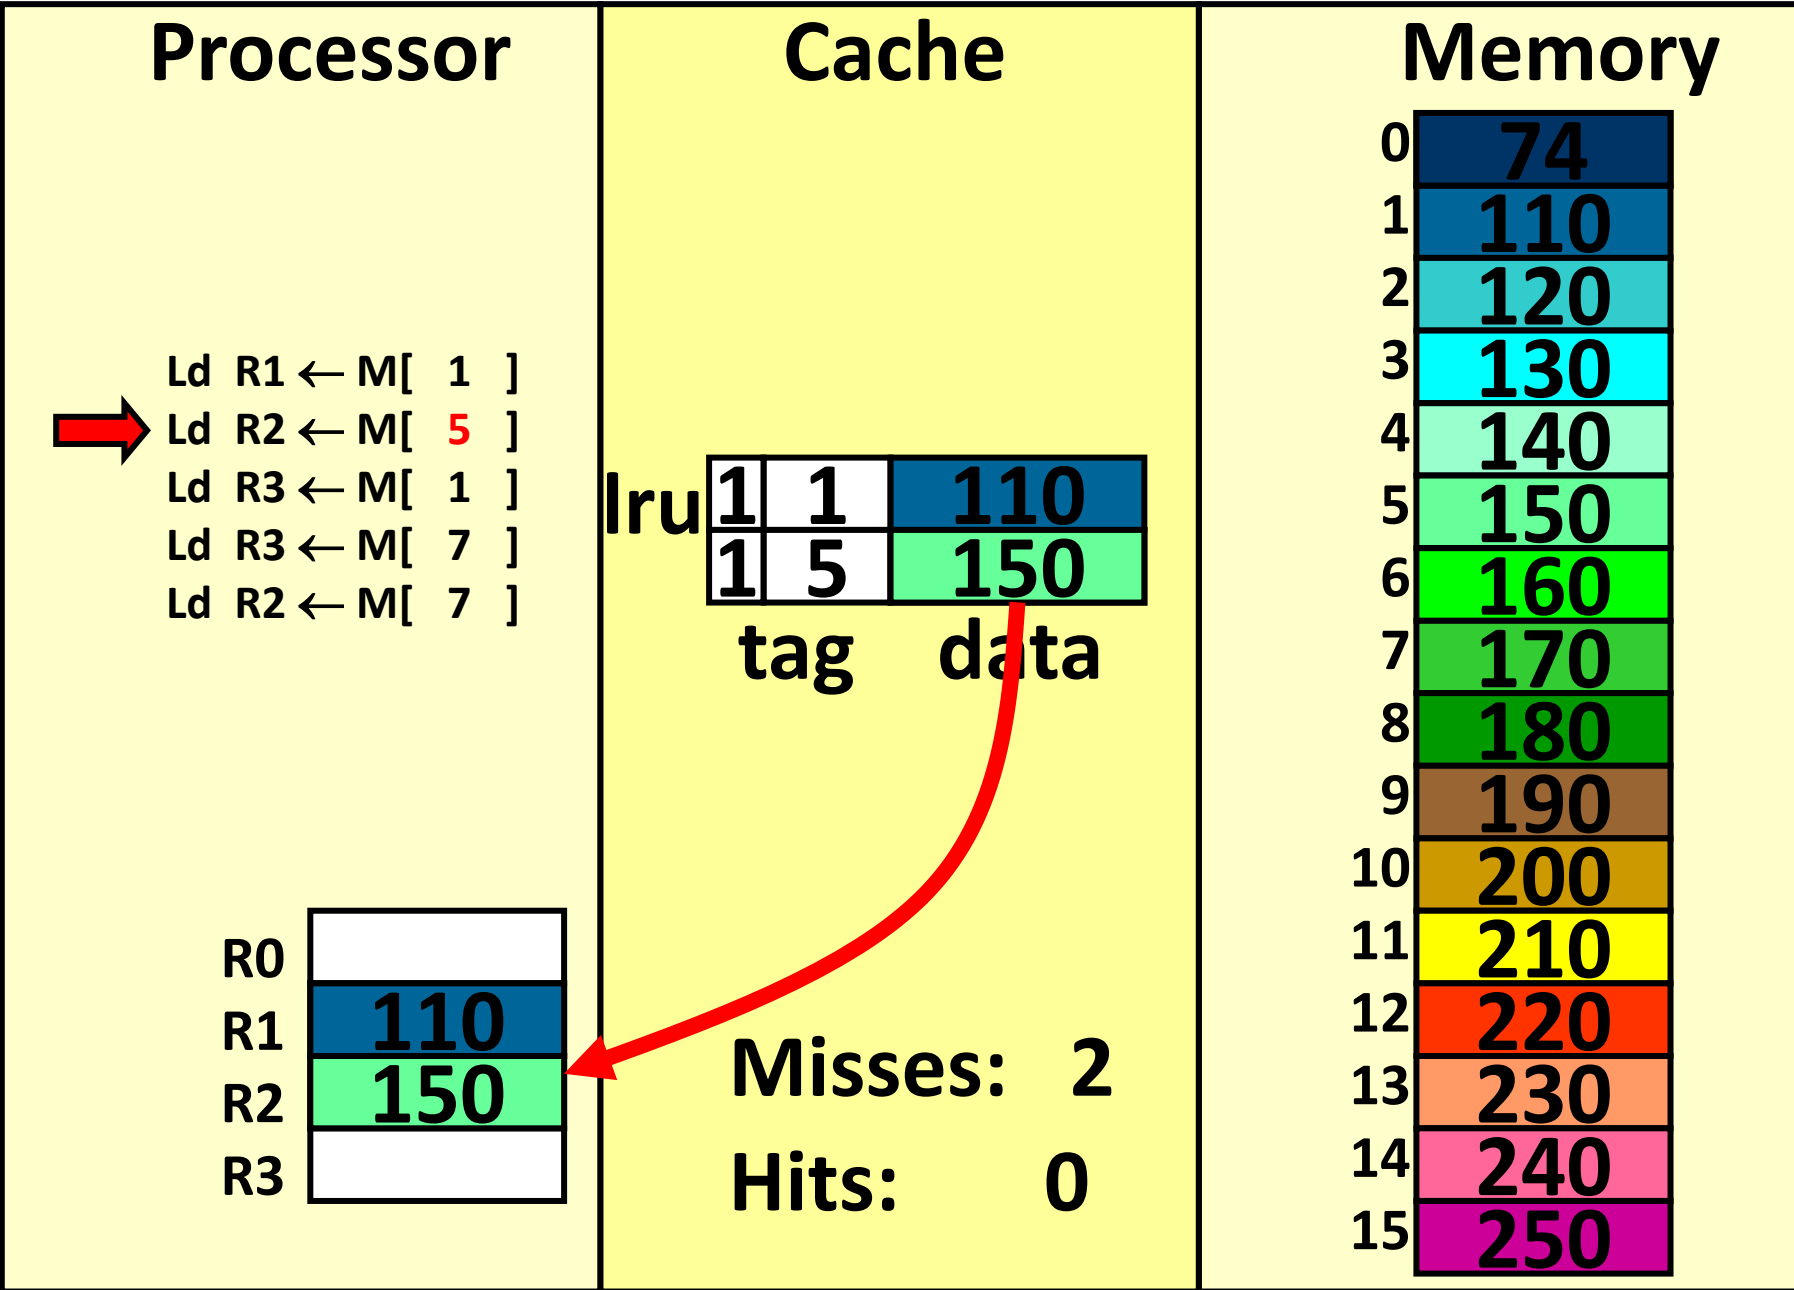


# A Very Simple Memory System

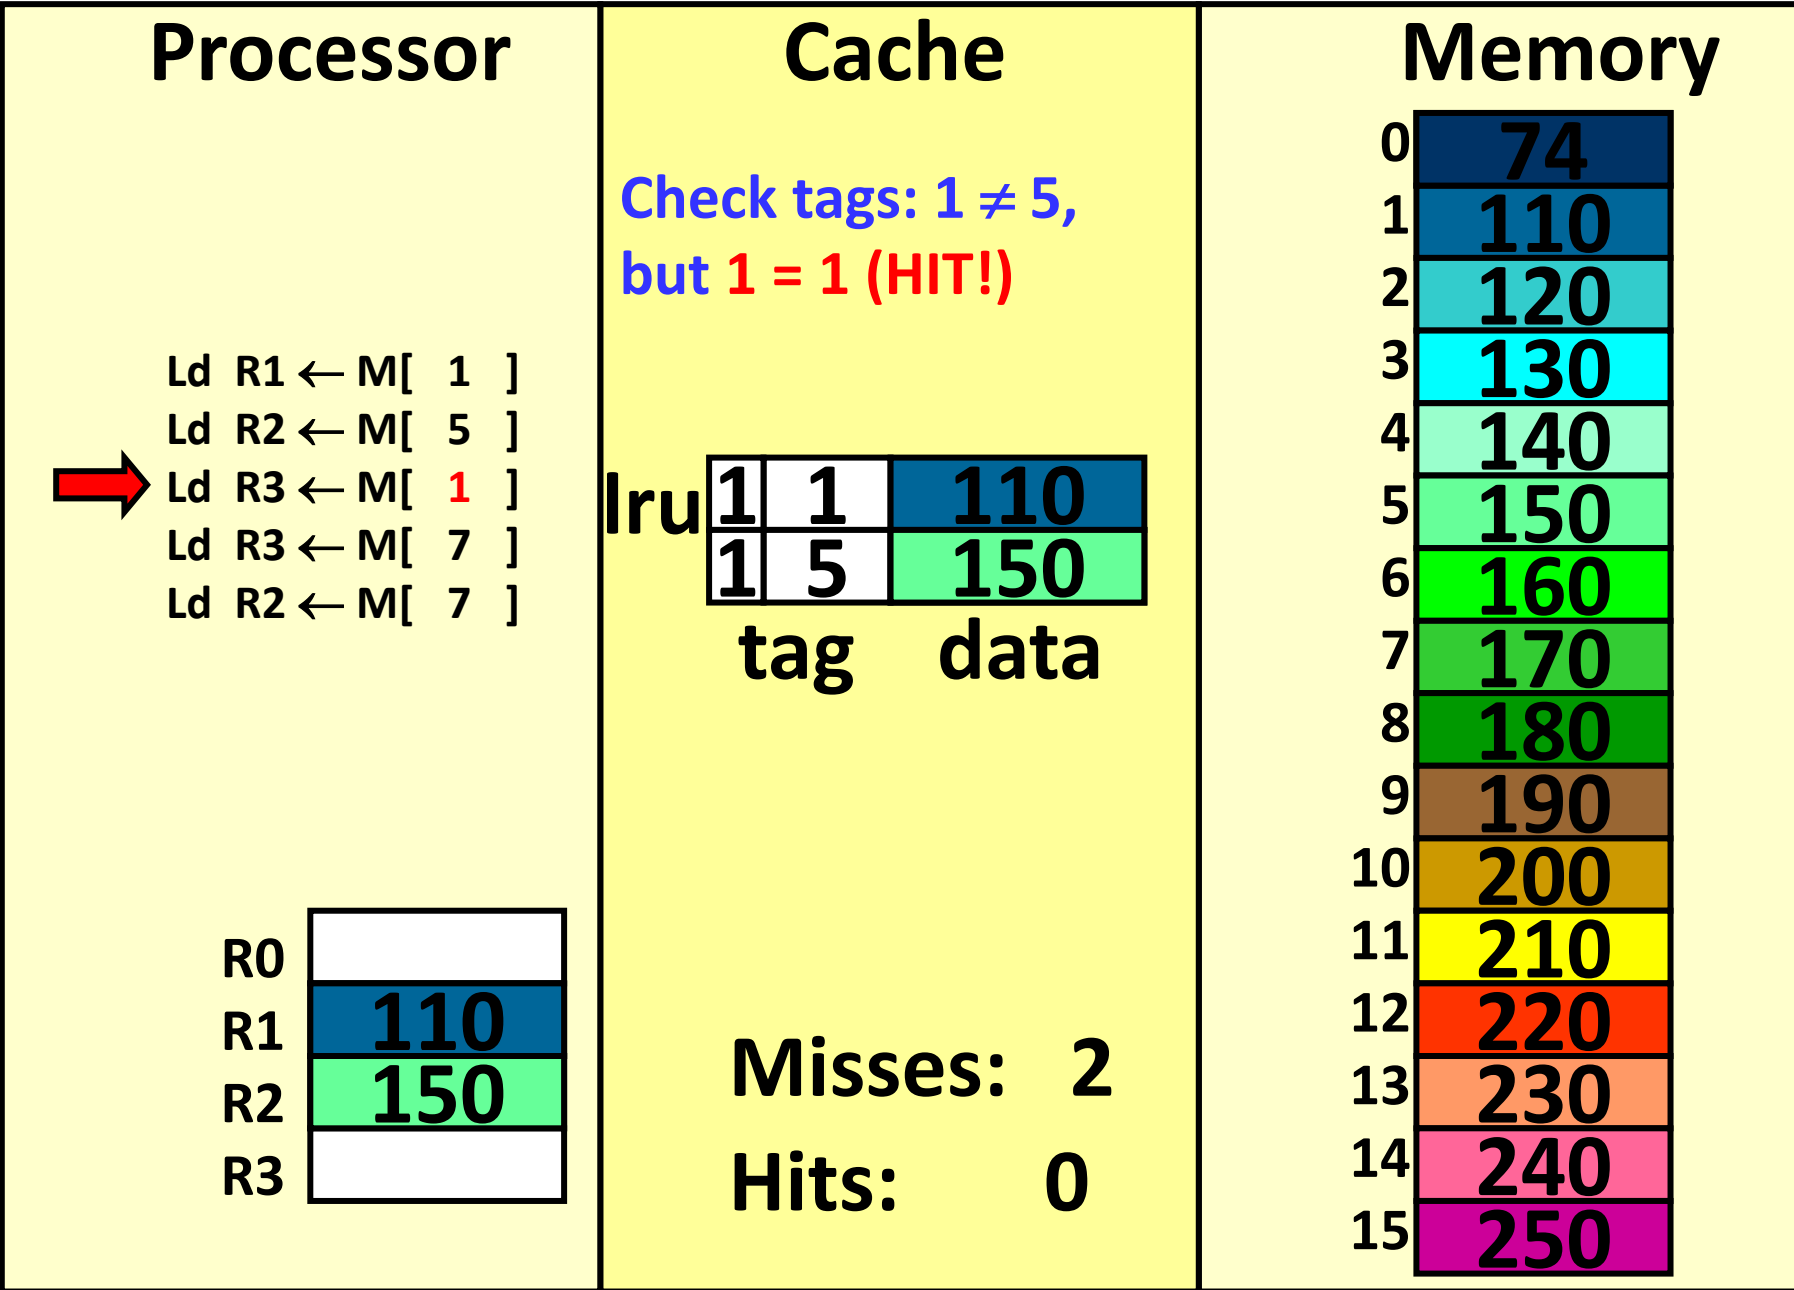
Tag comparison uses hardware called "content-addressable memory (CAM)"



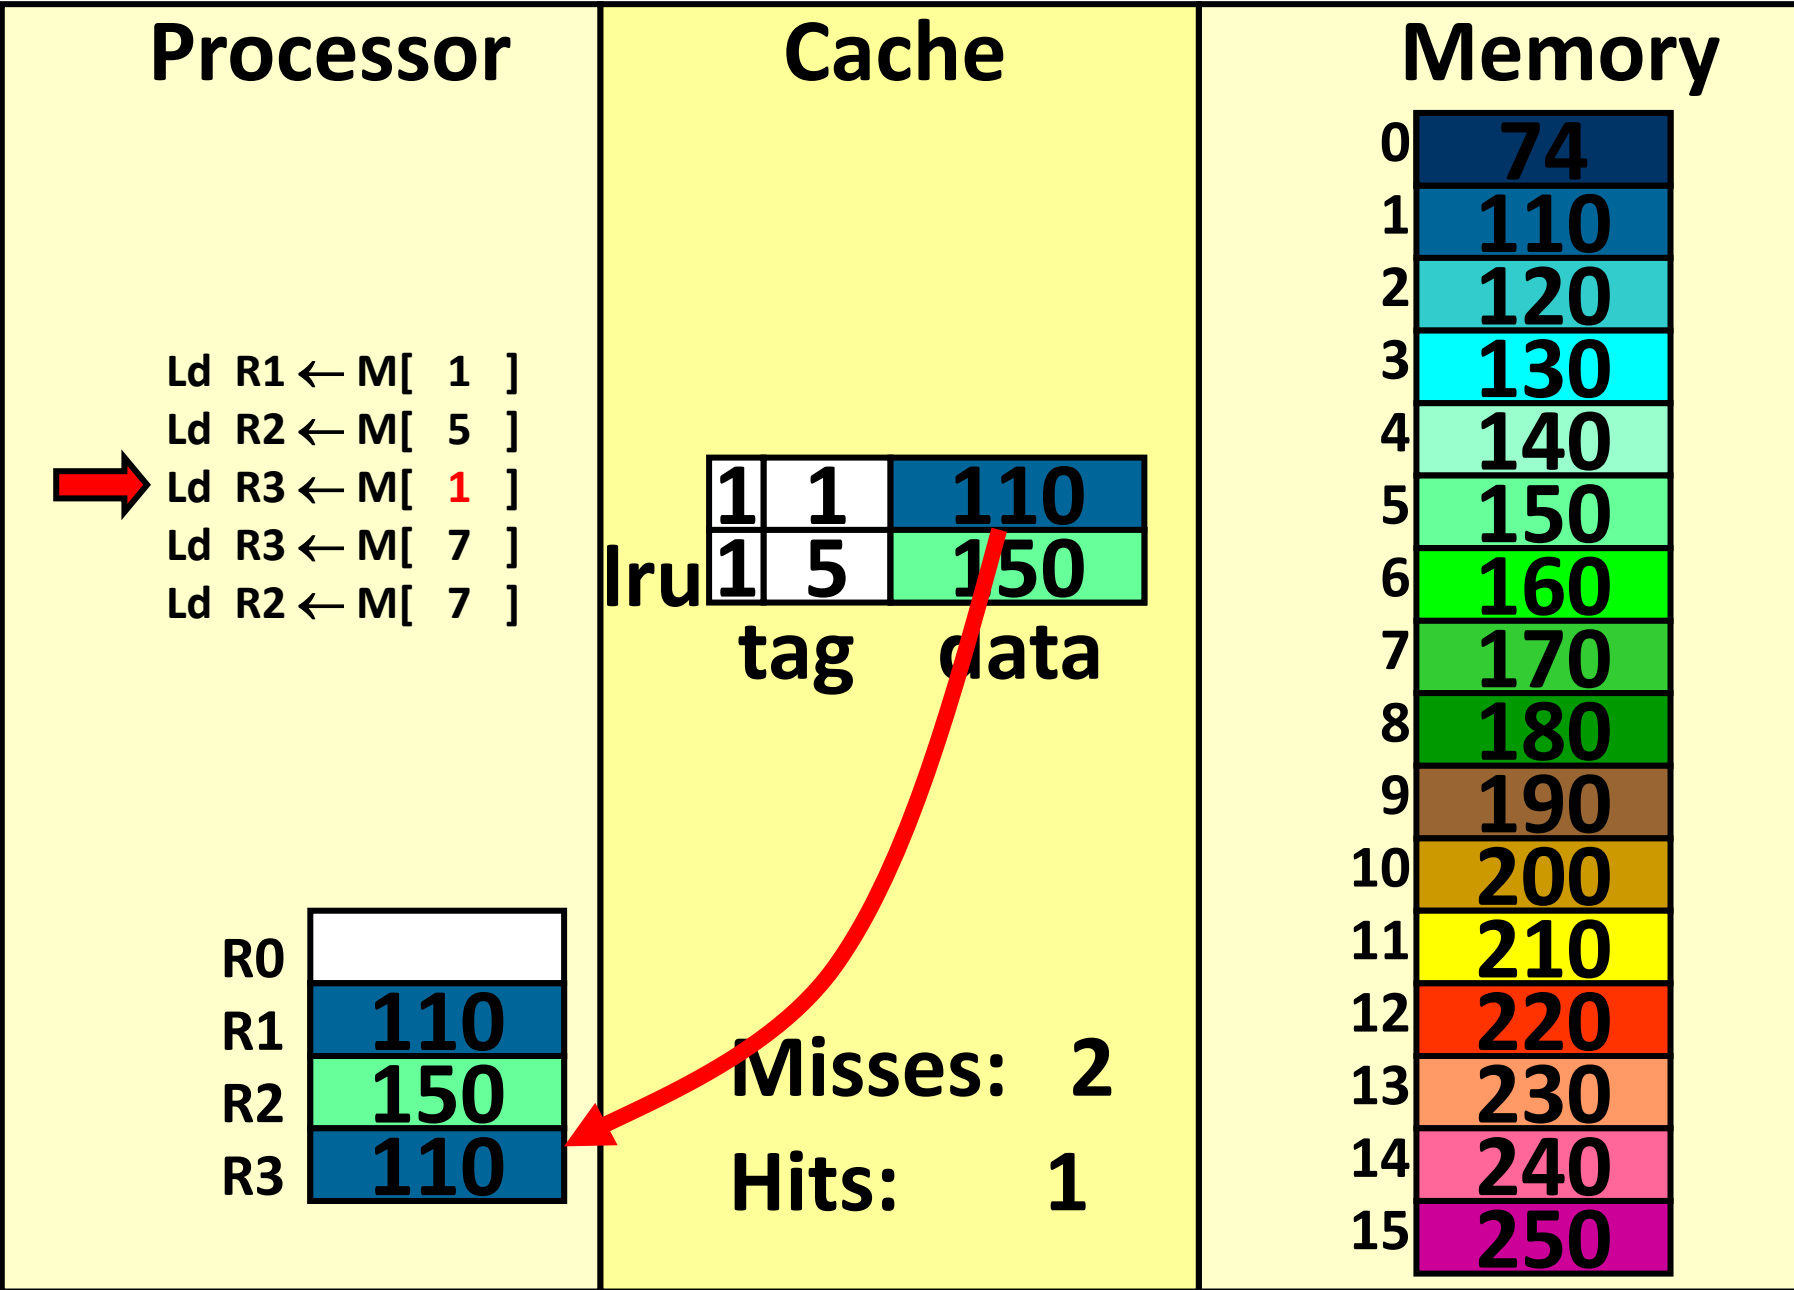
# A Very Simple Memory System



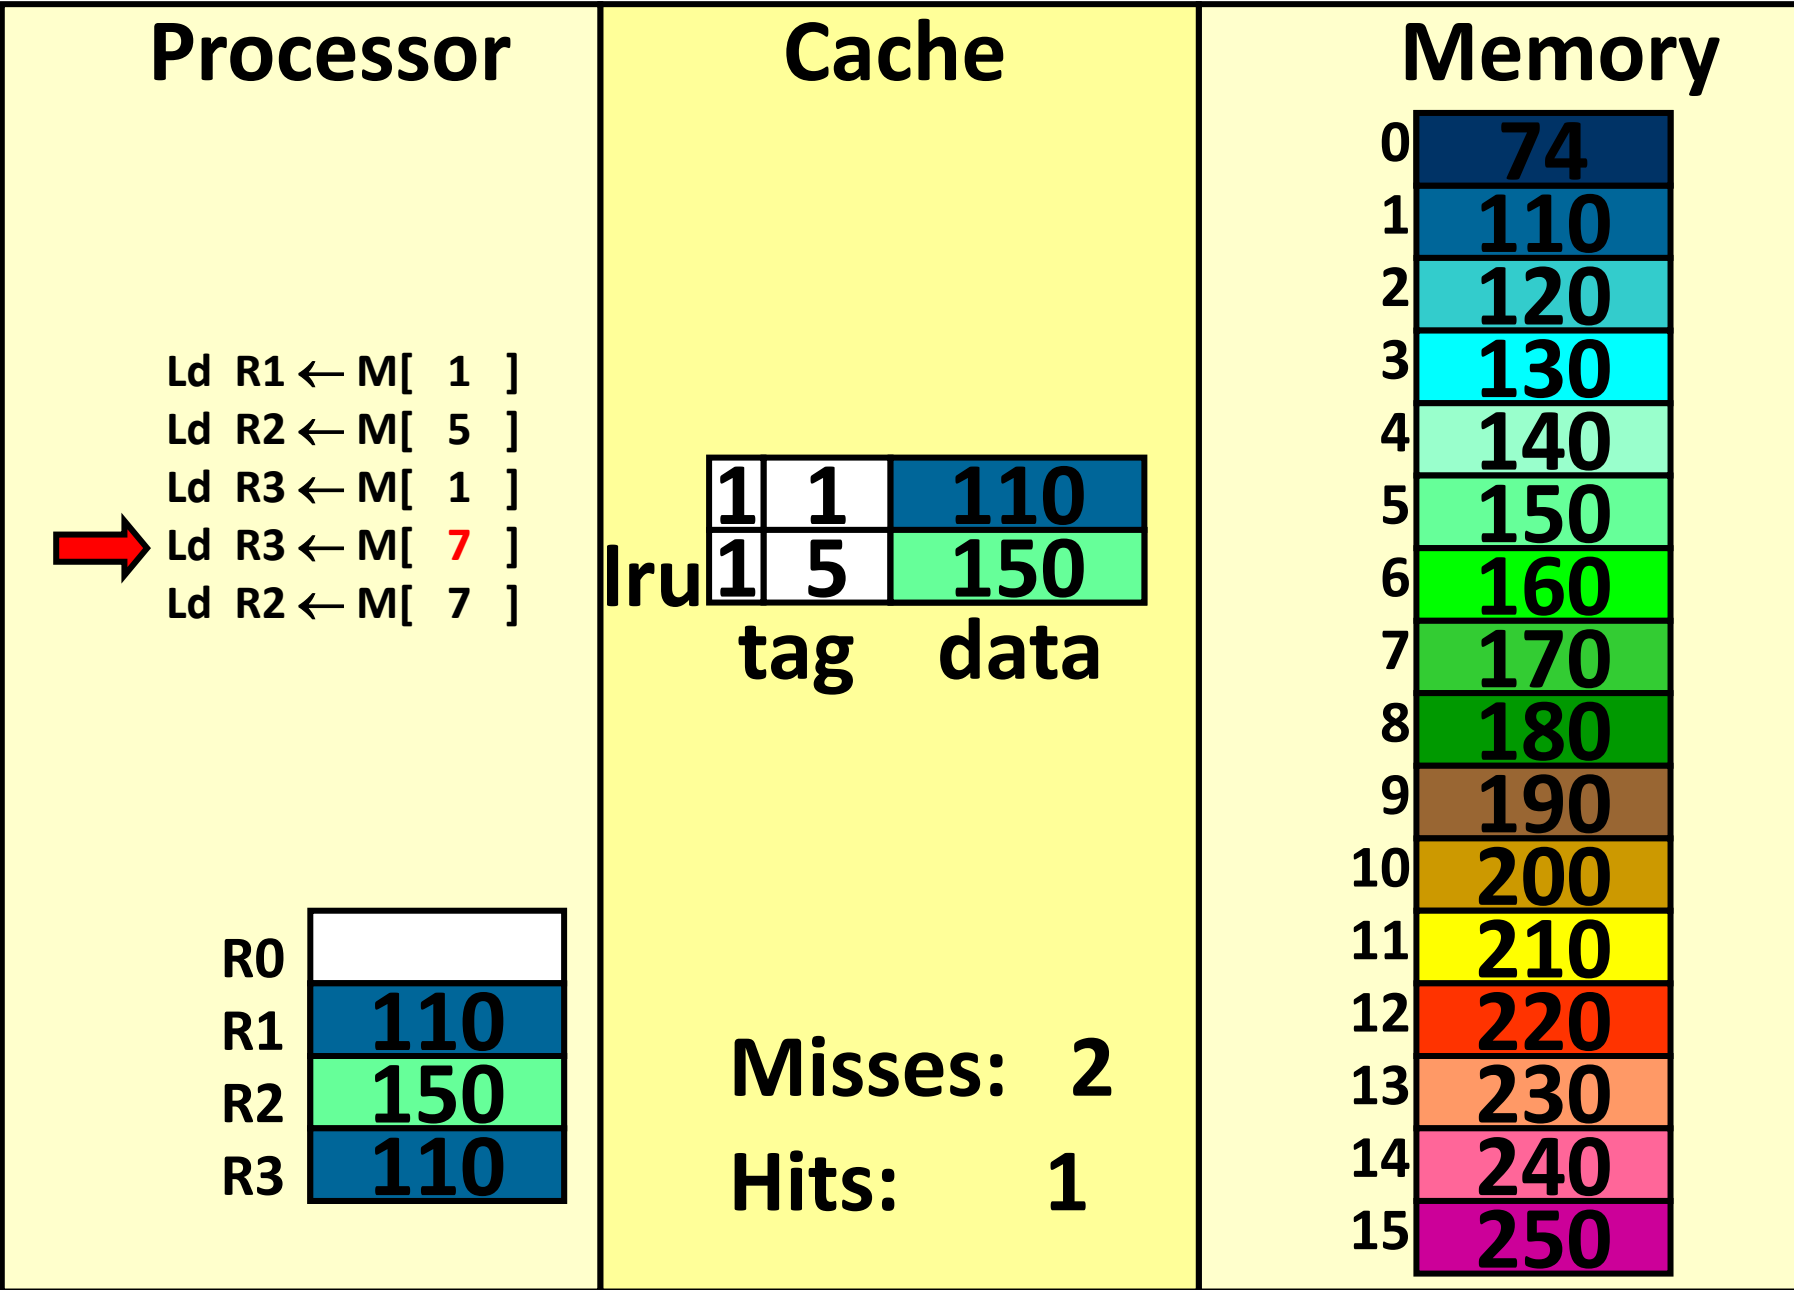
# A Very Simple Memory System



# A Very Simple Memory System

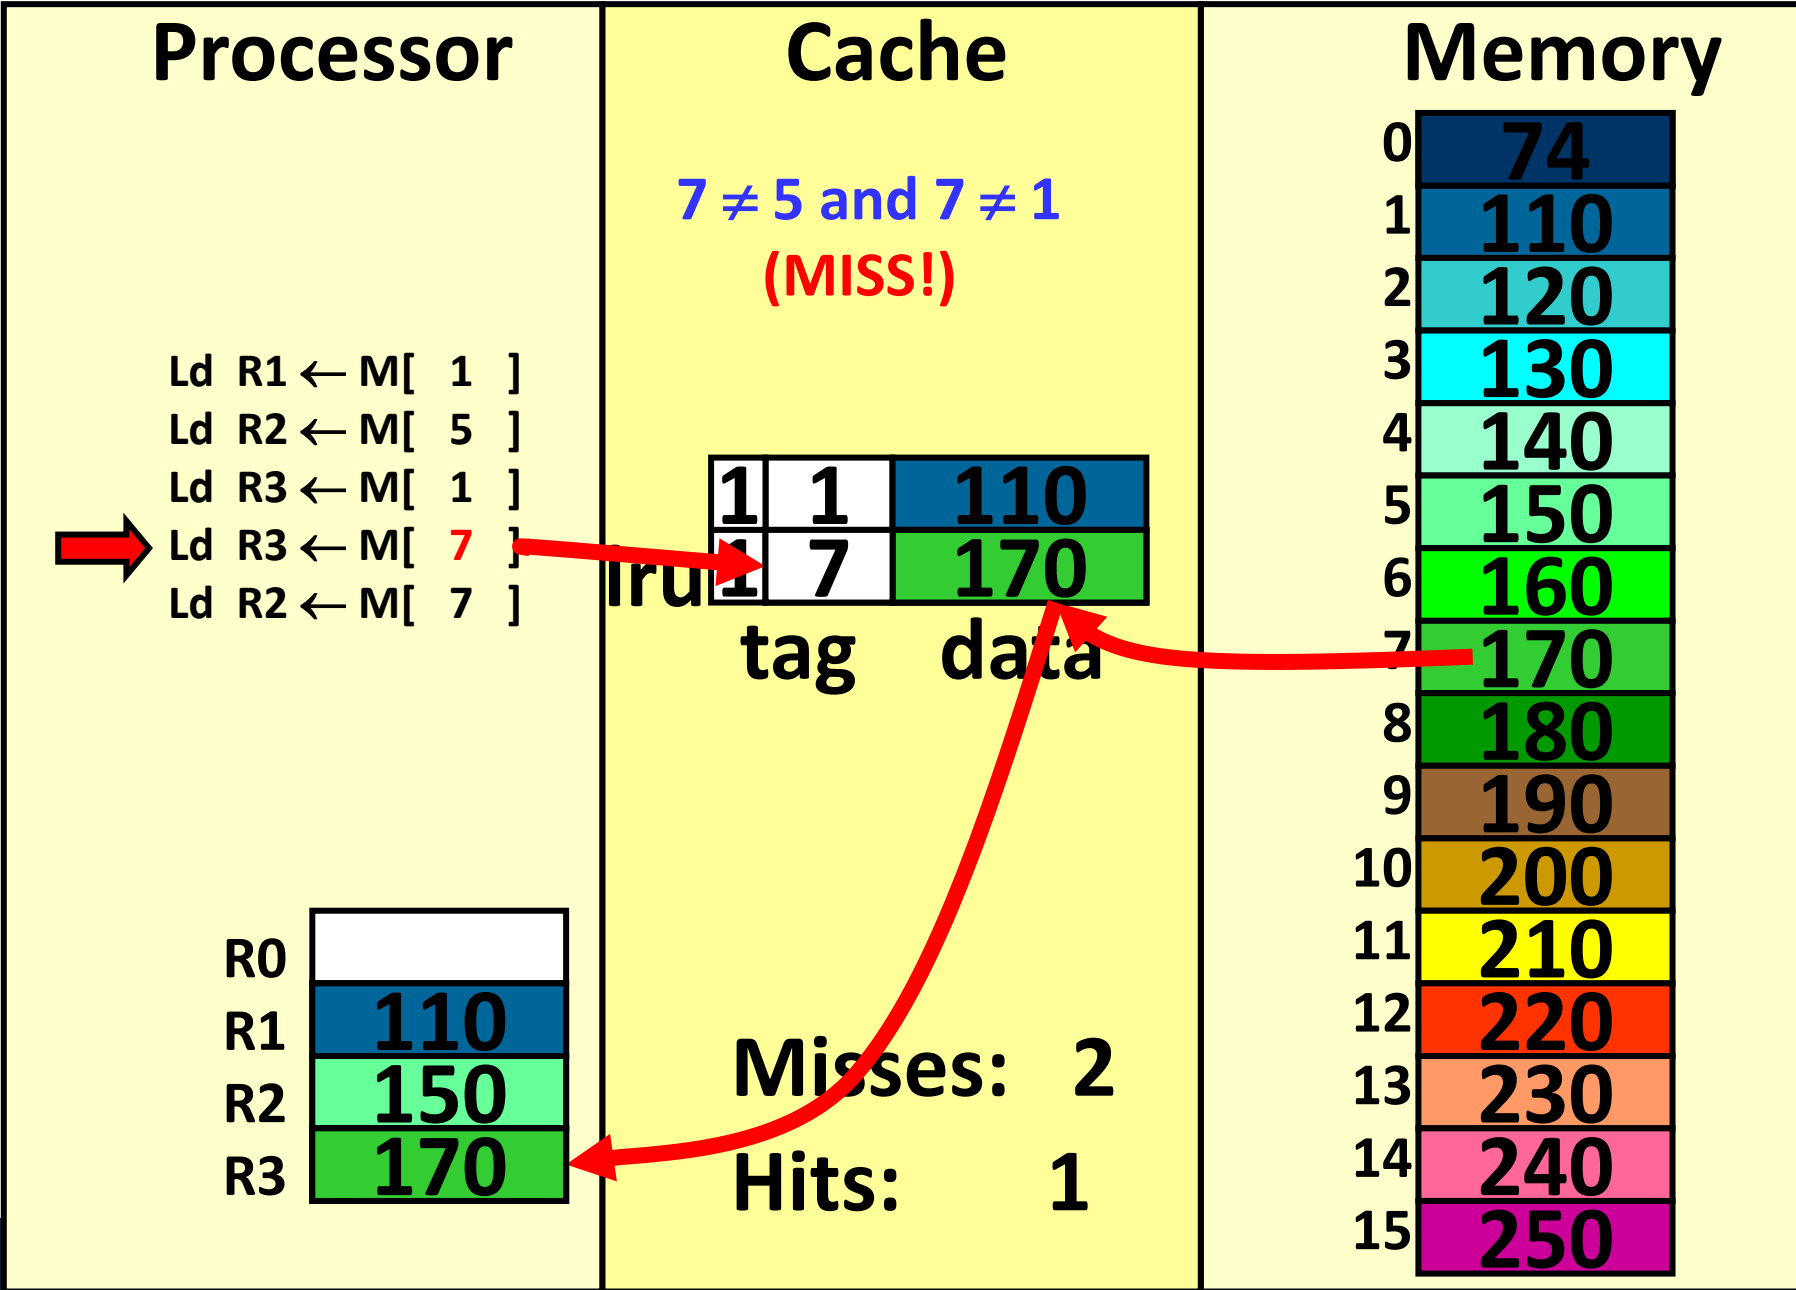


# A Very Simple Memory System

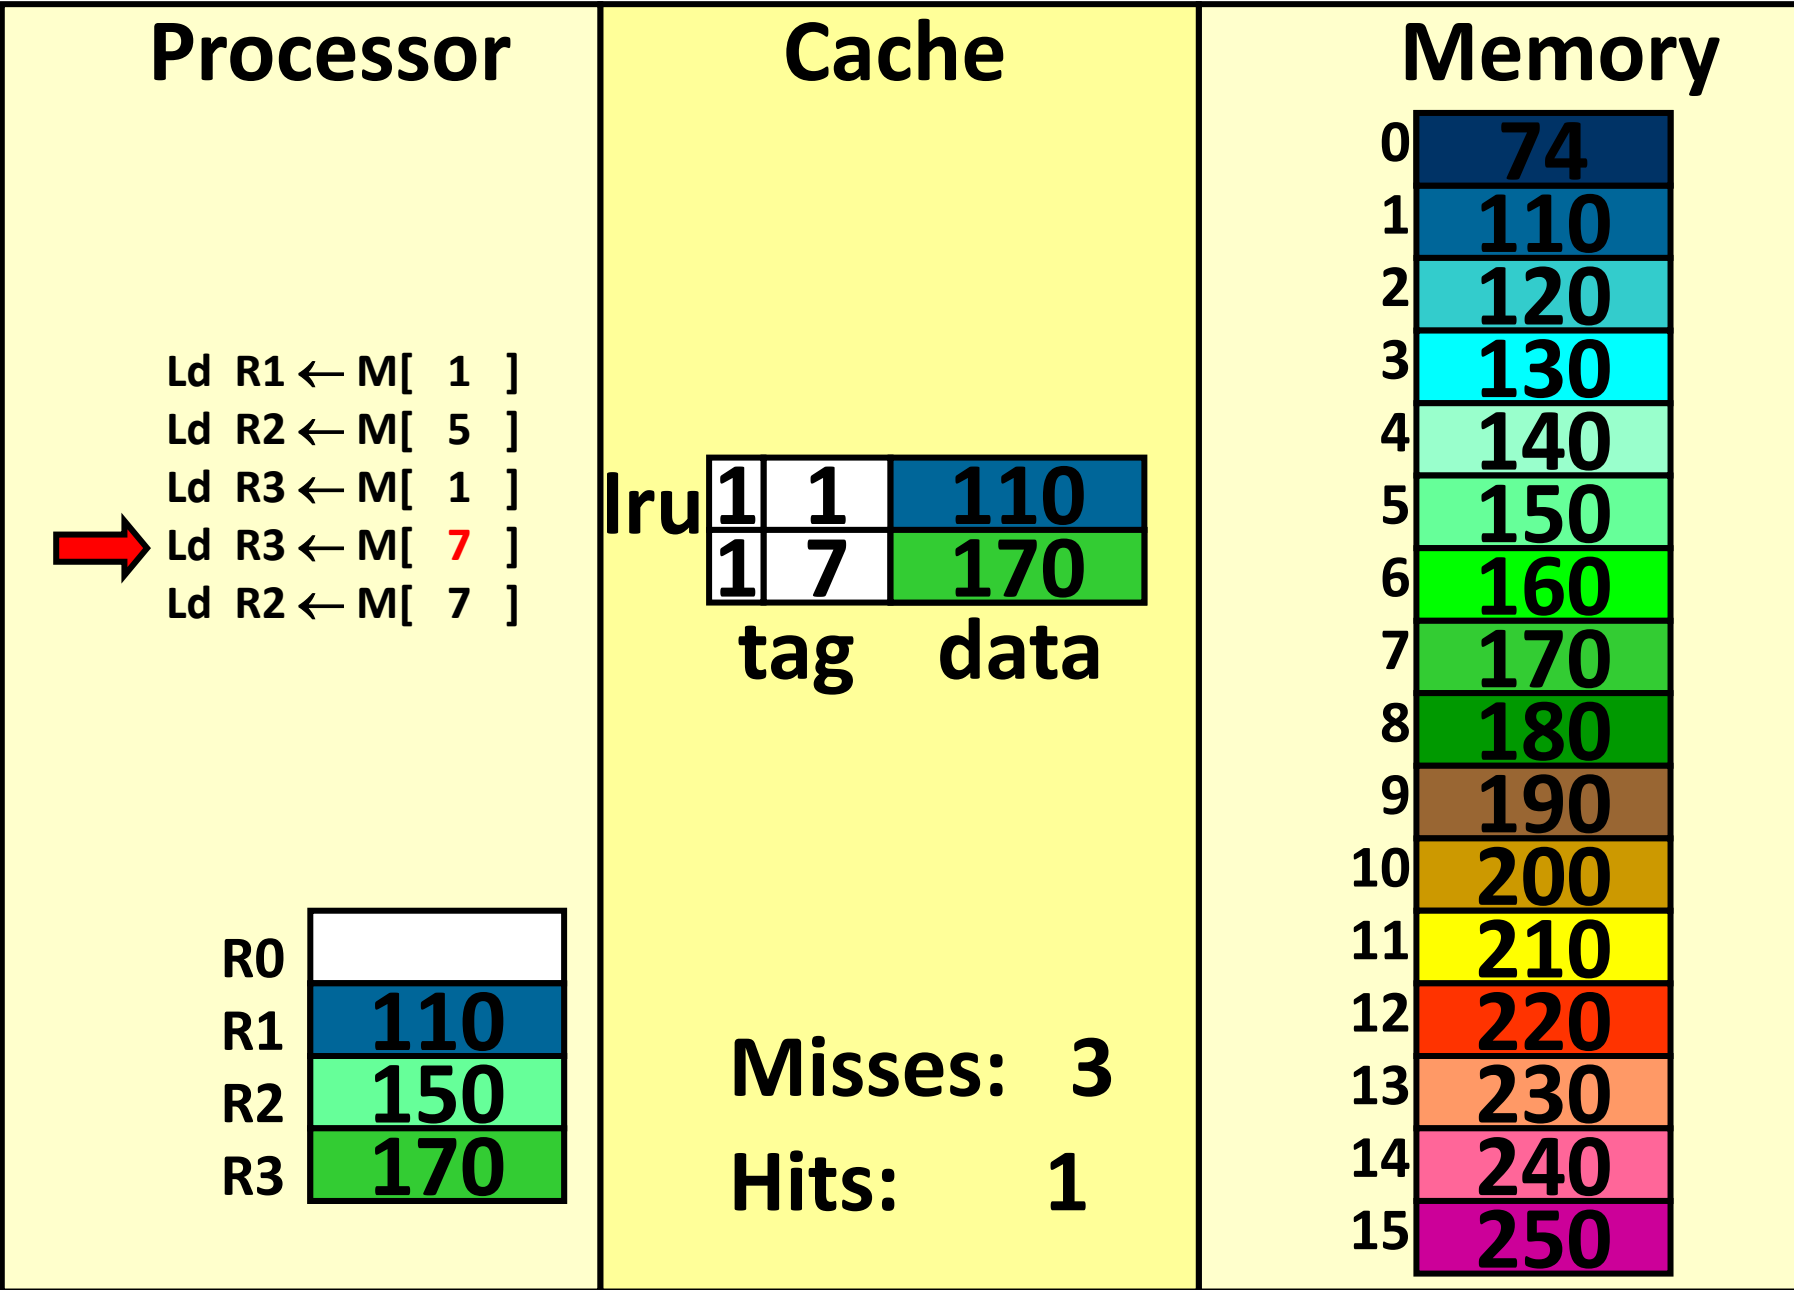




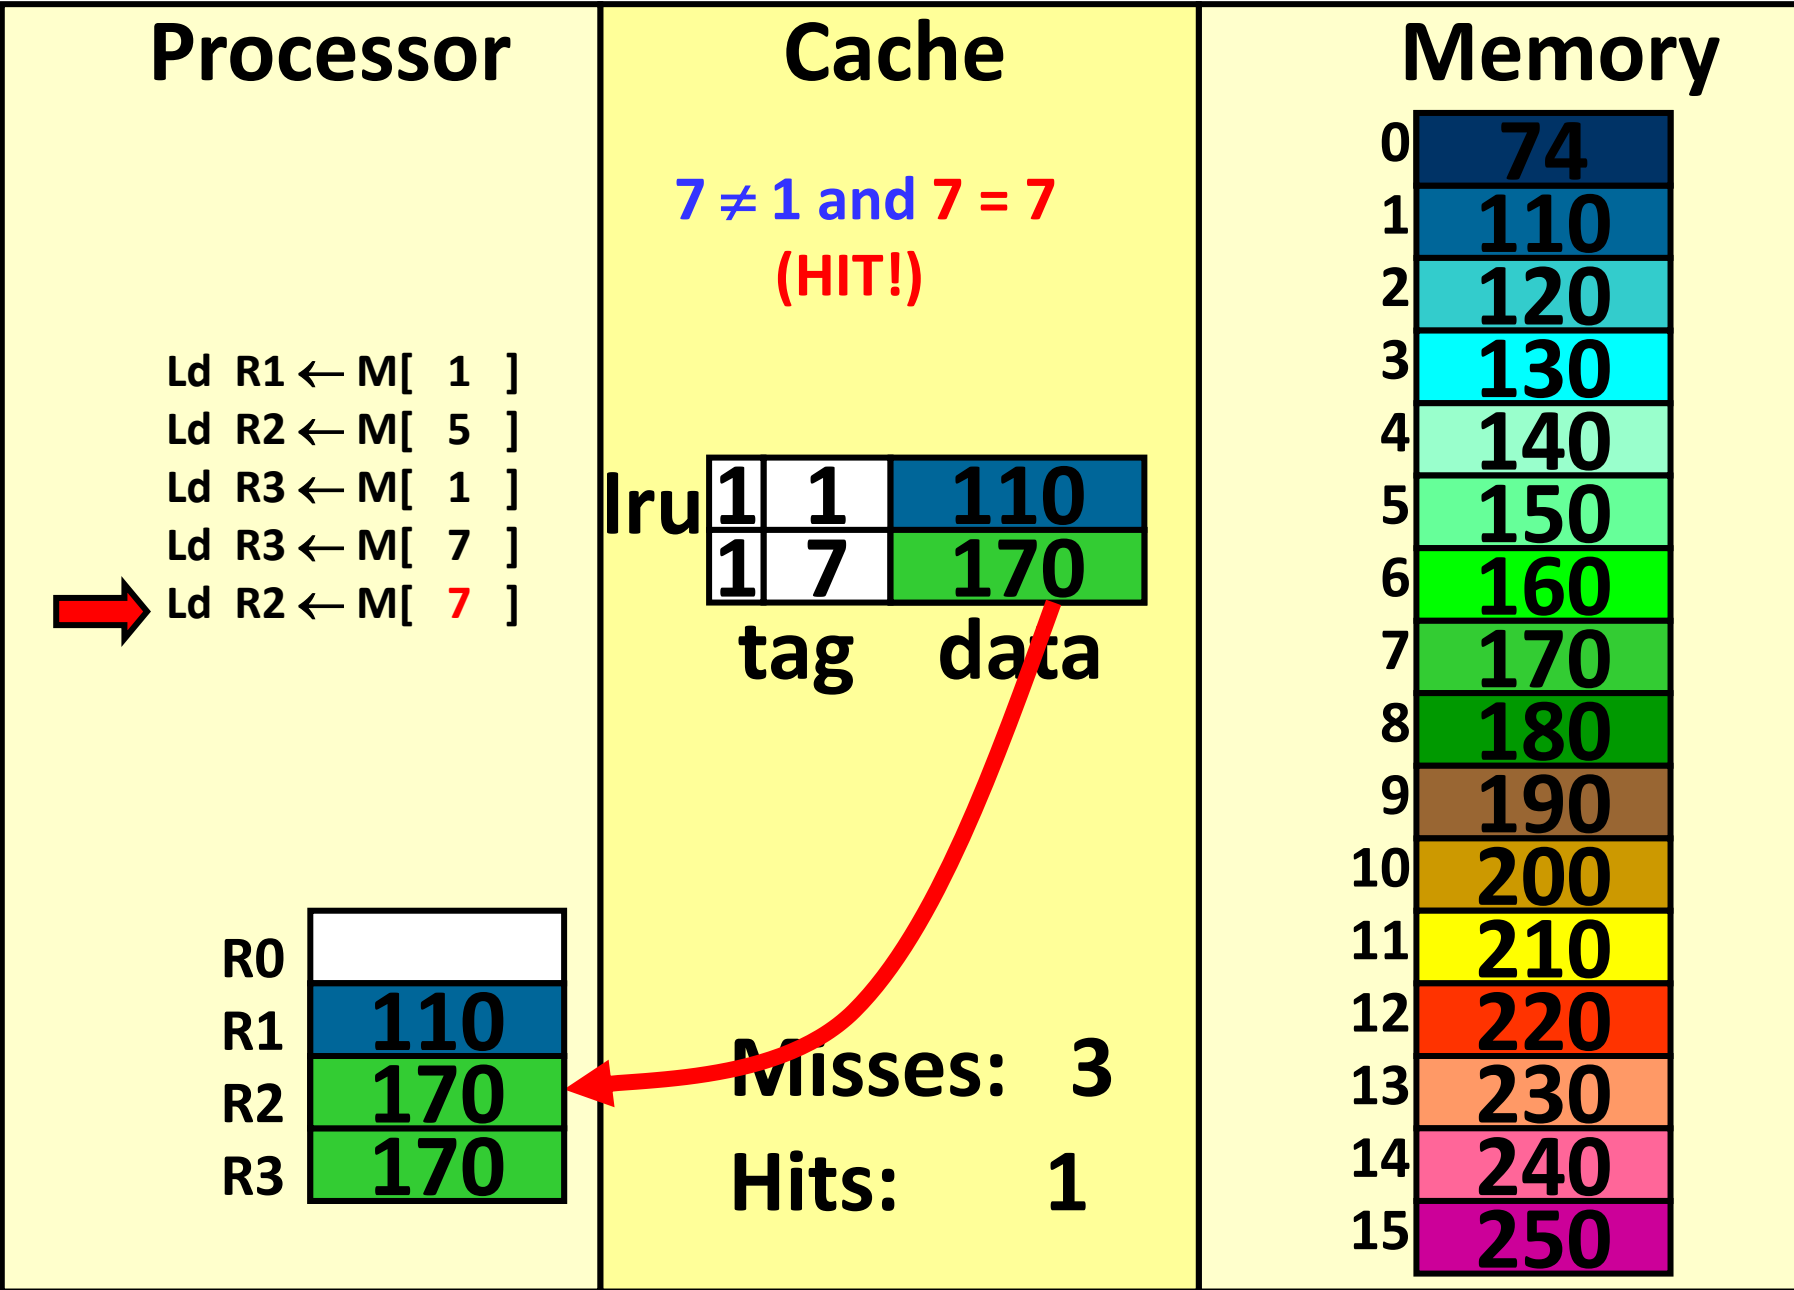
# A Very Simple Memory System



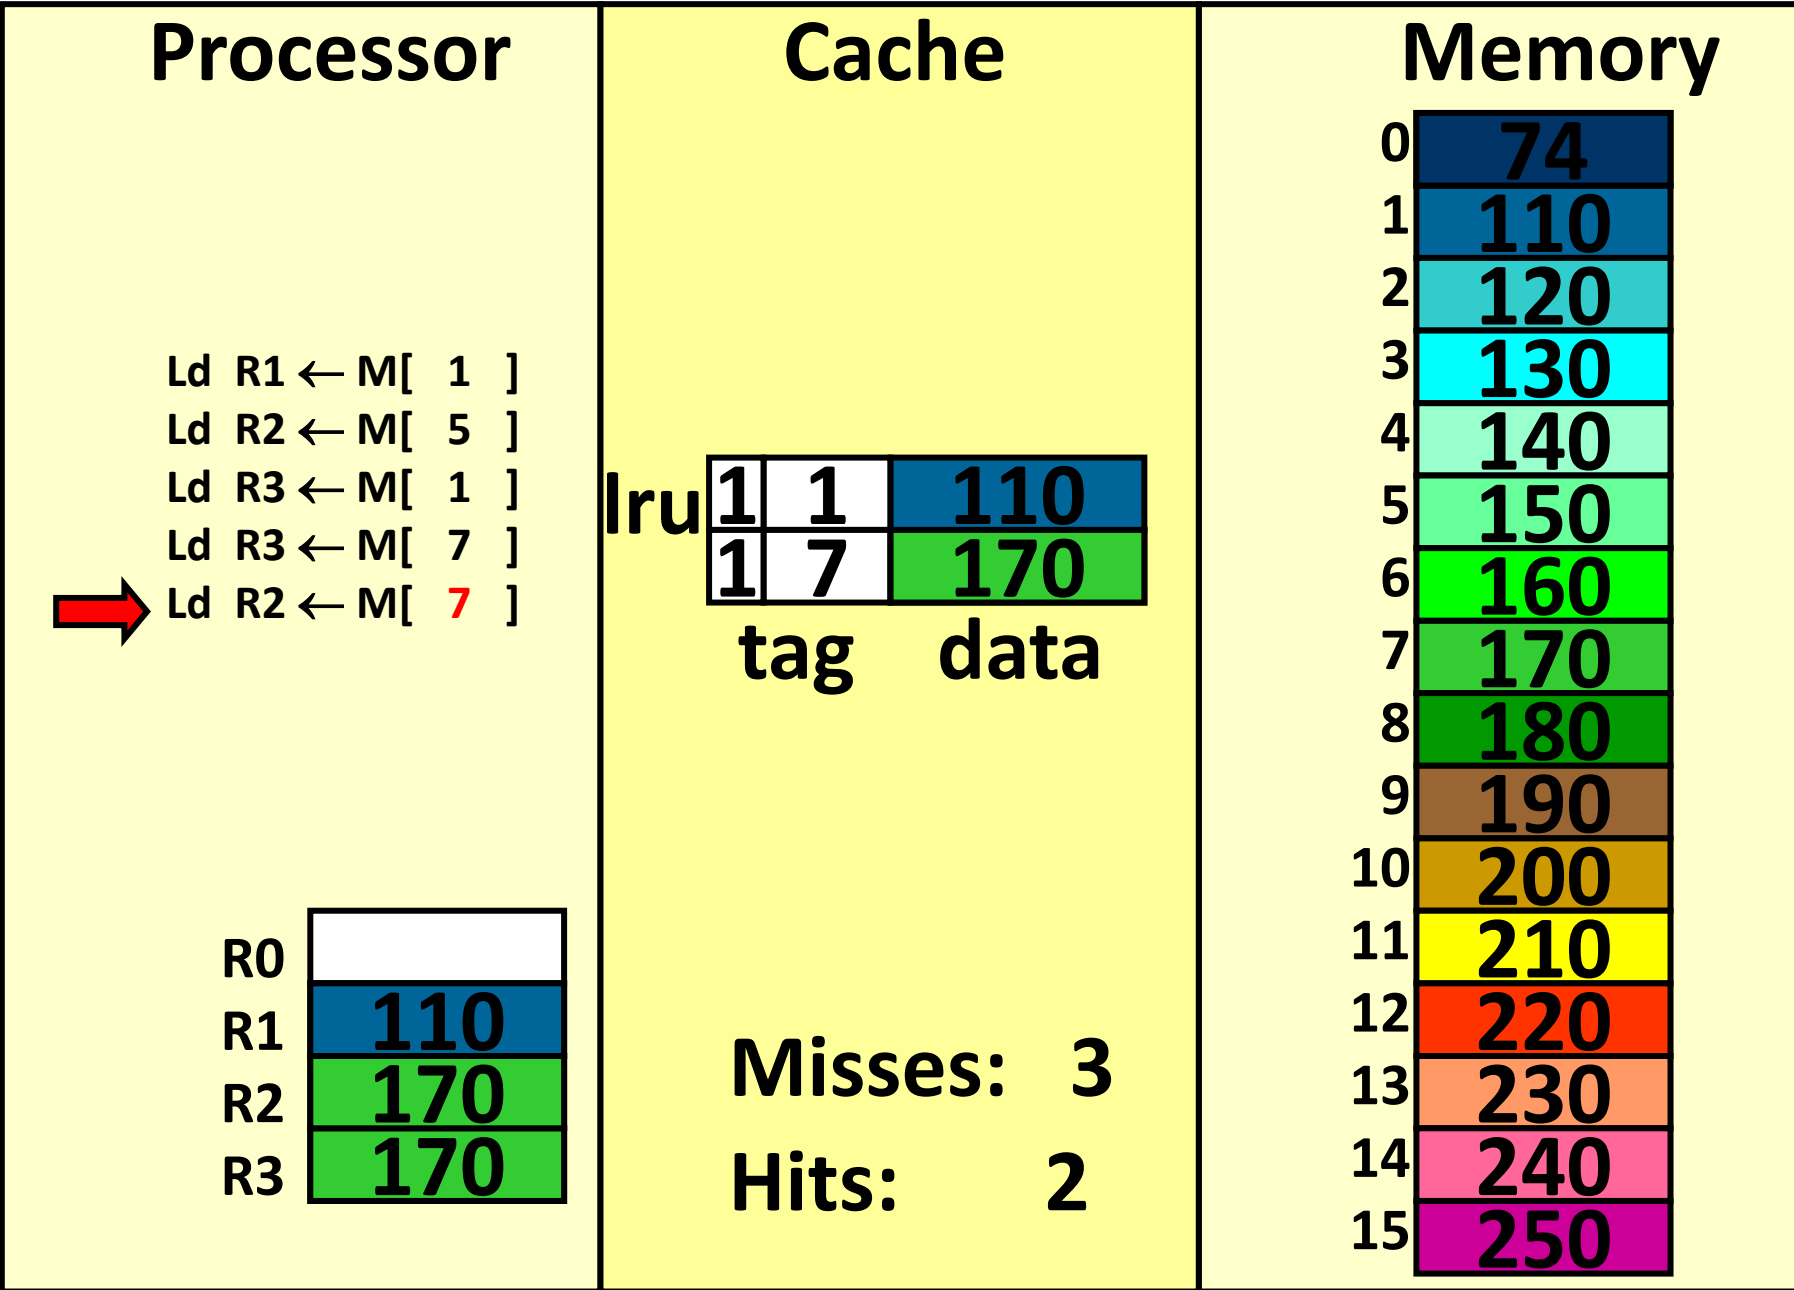
# A Very Simple Memory System



# A Very Simple Memory System



# A Very Simple Memory System



# Agenda

- Memory Types
- Memory Hierarchy and Cache Principles
- Cache example
- **How to improve cache**

# Definitions

- **Hit:** when data for a memory access is found in the cache
- **Miss:** when data for a memory access is not found in the cache
- **Hit/Miss rate:** percentage of memory accesses that hit/miss in the cache

# Example Problem

- Assume the following:
  - Cache has 1 cycle access time
  - If data is not found in cache, main memory is then accessed instead
  - Main memory has 100 cycle access time
- If we have a 90% hit rate in the cache, what is the average memory latency?

$$0.9 \times 1 + 0.1 \times 101 = 11$$

$$1 \times 1 + 0.1 \times 100 = 11$$

$$1 + 0.1 * ( 100 ) = 11$$

# Calculating Average Access Latency

- Average Latency:
  - $\text{cache latency} + (\text{memory latency} \cdot \text{miss rate})$
- Average latency for our example:
  - $1 \text{ cycle} + \left(15 \cdot \frac{3}{5}\right)$
  - $= 10 \text{ cycles per reference}$
- To improve latency, either:
  - Improve memory access latency, or
  - Improve cache access latency, or
  - Improve cache hit rate



# Next time

- Making our caches bigger and better