

# EECS 370

## Multi-Level Page Tables

Poll  
What's your favorite Thanksgiving dish?

- a) Turkey
- b) Stuffing
- c) Cranberry sauce
- d) Sweet, delicious Buckeye tears



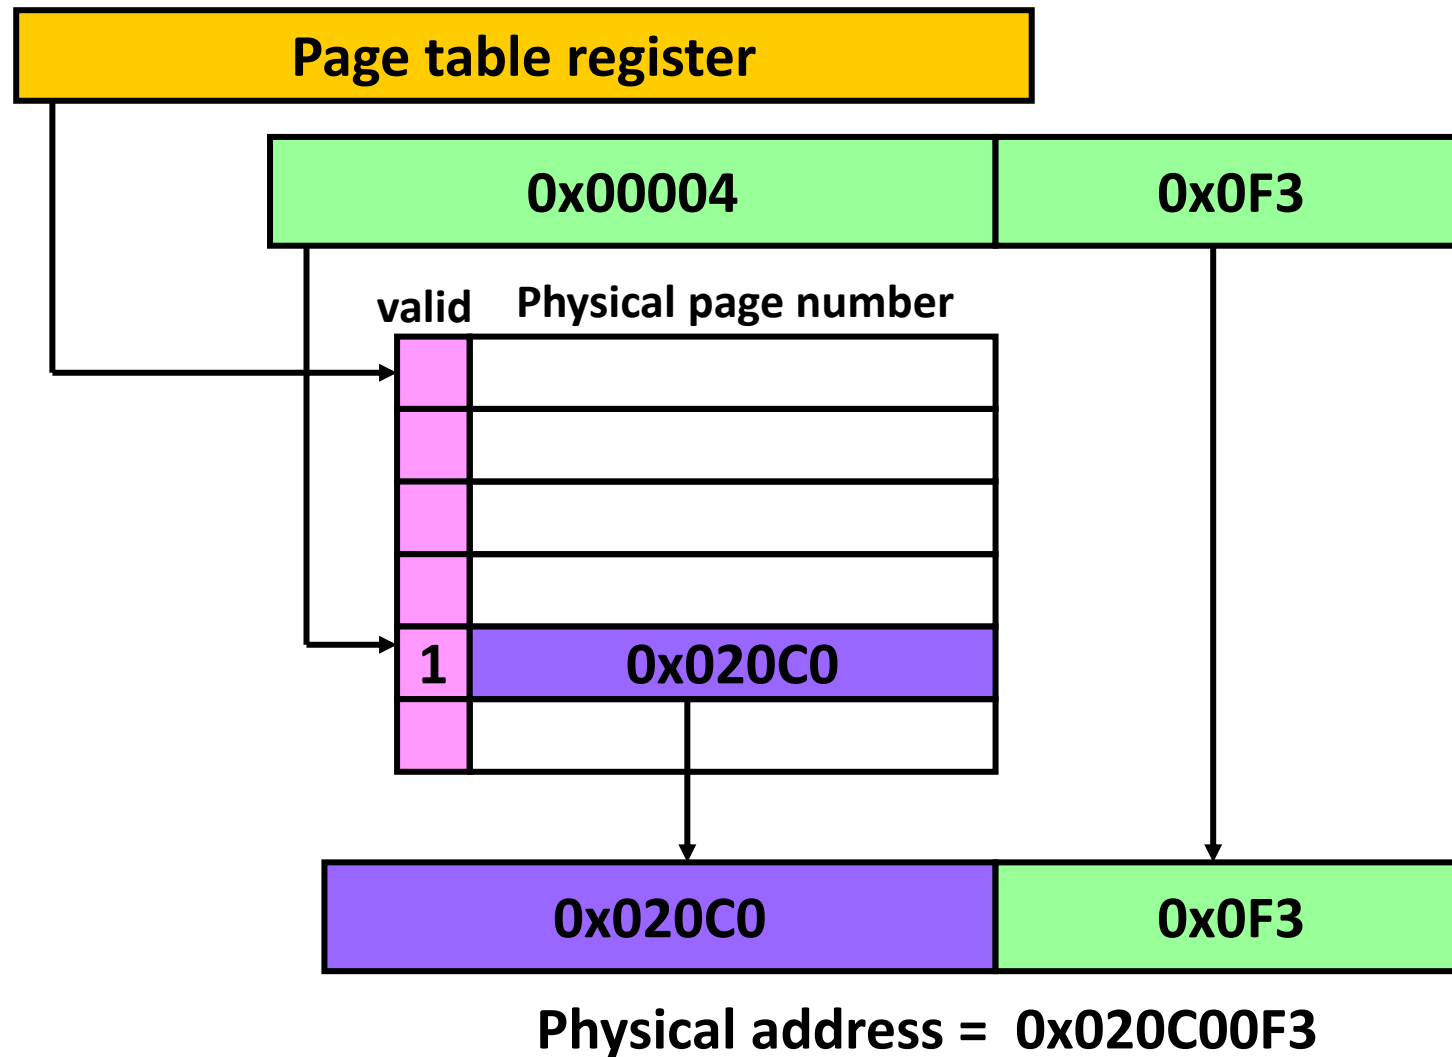
# Announcements

- P4
  - Last project!
  - Due Thur (11/30)
- HW 4
  - Last homework!
  - Due Mon (12/4)
- Lecture Thursday
  - Last lecture!
  - Review on Tuesday
- Lab Fr/M
  - Last lab!
  - Don't forget to submit lab 11 by Wed
- Final exam
  - ...Last exam!
  - Tue (12/12) @ 10:30 am



# Reminder: Page tables

Virtual address = 0x000040F3



# Agenda

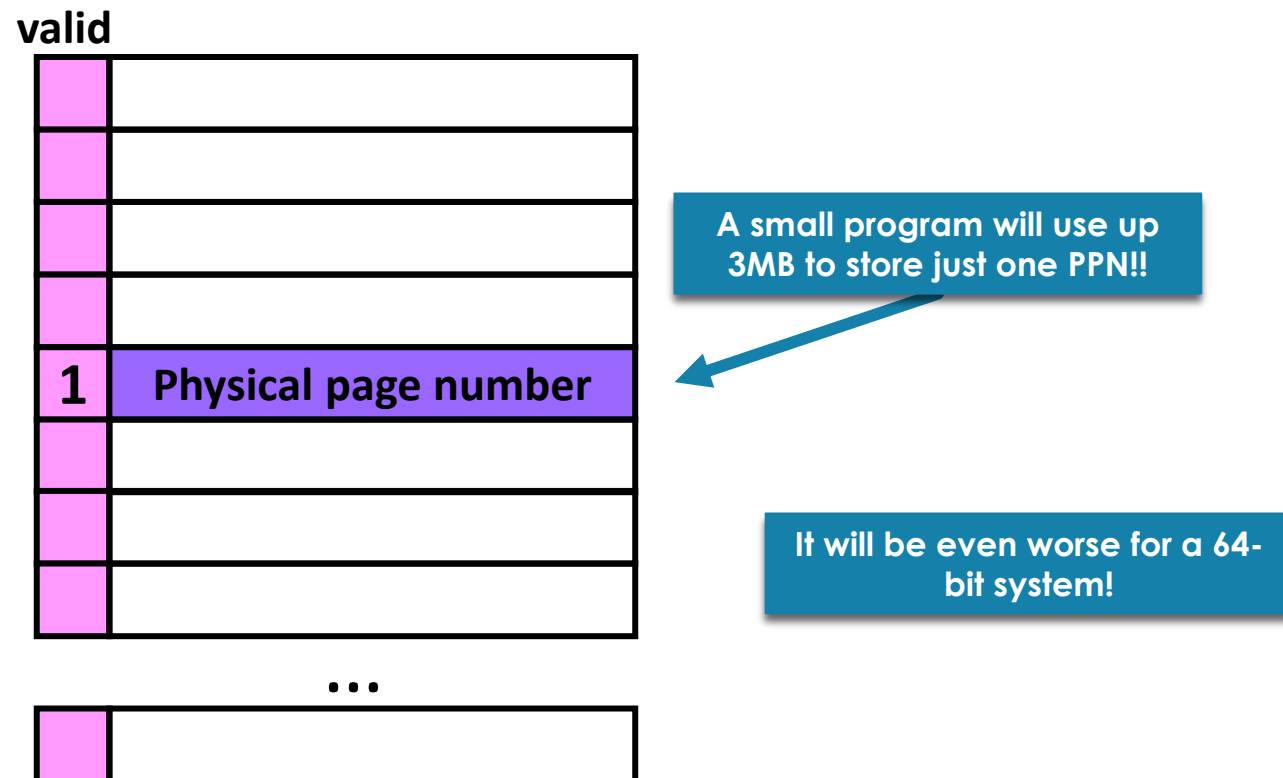
- **Motivation for Multi-level Page Tables**
- Example architecture
- Class Problem: 32bit Intel x86
- Class Problem: Multi-Level VM Design
- VM Miscellanea

# Size of the page table

- How big is a page table entry?
  - For 32-bit virtual address:
    - If the machine can support  $1\text{GB} = 2^{30}$  bytes of physical memory and we use pages of size  $4\text{KB} = 2^{12}$ ,
    - then the physical page number is  $30-12 = 18$  bits.  
Plus another **valid bit** + other useful stuff (**read only, dirty, etc.**)
    - Let say about 3 bytes.
- How many entries in the page table?
  - 1 entry per virtual page
  - ARM virtual address is 32 bits – 12 bit page offset = 20
  - Total number of virtual pages =  $2^{20}$
- Total size of page table = Number of virtual pages
  - \* Size of each page table entry
  - =  $2^{20} \times 3$  bytes  $\sim 3$  MB

# How can you organize the page table?

1. Single-level page table occupies continuous region in physical memory
  - Previous example always takes 3MB regardless of how much virtual memory is used



# How can you organize the page table?

## 2. Option 2: Use a multi-level page table

- 1<sup>st</sup> level page table (much smaller!) holds addresses 2<sup>nd</sup> level page tables
  - 2<sup>nd</sup> level page tables hold translation info, or 3<sup>rd</sup> level page tables if we wanna go deeper
  - Only allocate space for 2<sup>nd</sup> level page tables that are used

valid	PPN
1	0x1
1	0x2
1	0x3

Single-level: Tons of wasted space!

valid	2 <sup>nd</sup> level page table
1	0x1000

valid	PPN
1	0x1
1	0x2
1	0x3

# Multi-Level Page Table

- Only allocate second (and later) page tables when needed
- Program starts: everything is invalid, only first level is allocated

Single Level

valid	2 <sup>nd</sup> level page table
0	
0	
0	
0	

Multi-level: Size is proportional to amount of memory used

- As we access more, second level page tables are allocated

Multi Level



valid	2 <sup>nd</sup> level page table
1	0x1000
1	0x3500

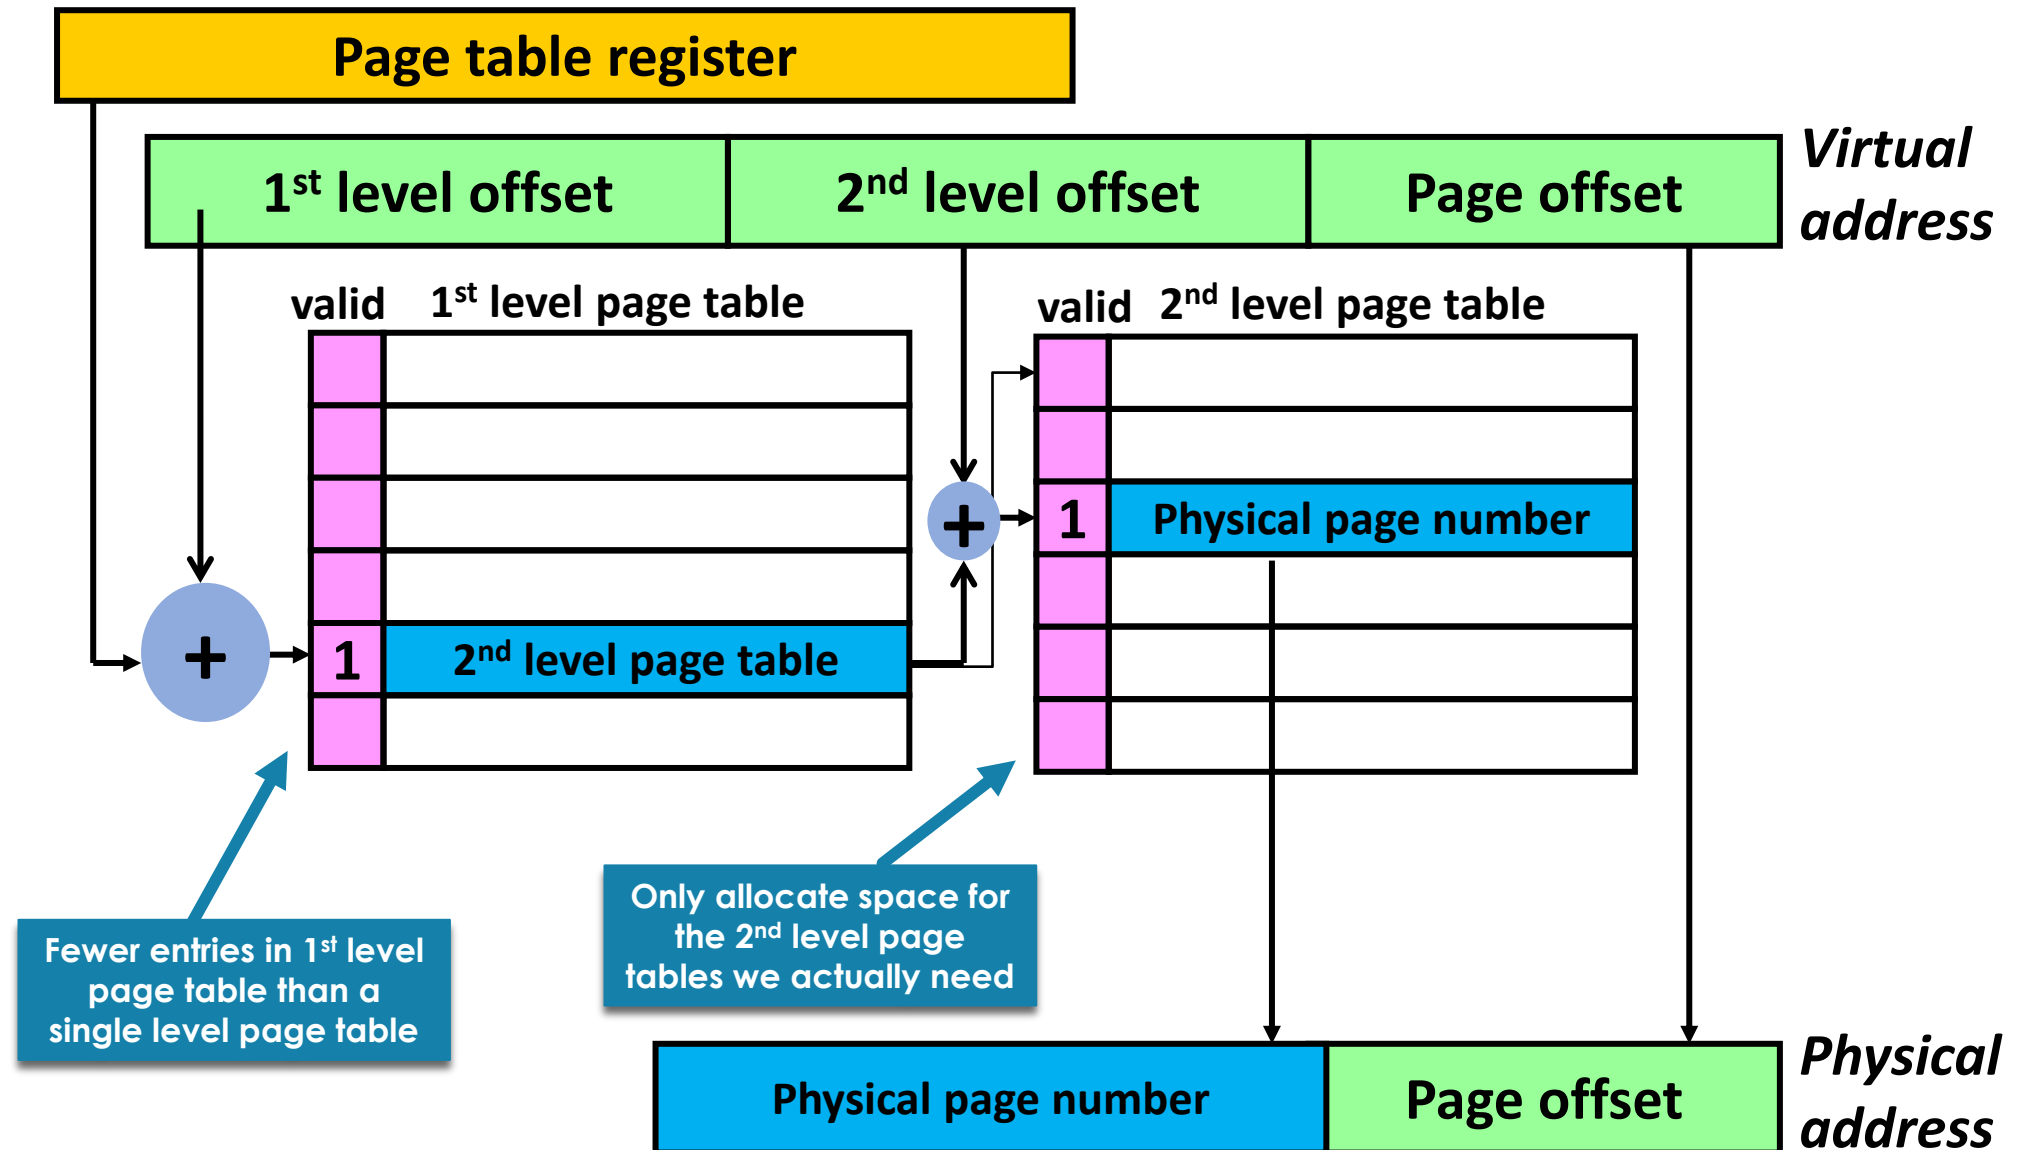
valid	PPN
1	0x1
1	0x6
1	0x2
1	0x1f

valid	PPN
1	0x9a
1	0x3
0	
1	0xff

Common case: most programs use small portion of virtual memory space




# Hierarchical page table



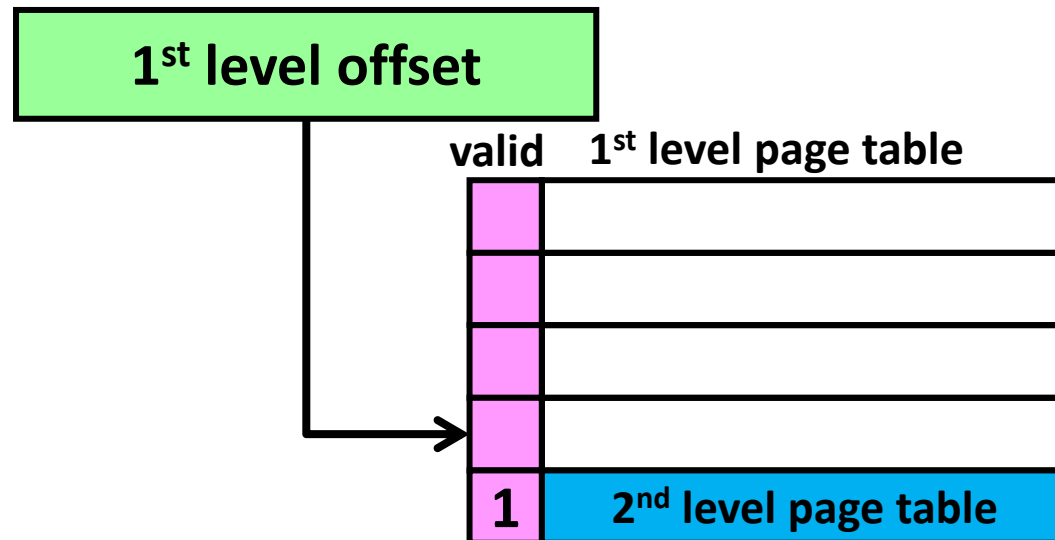
# Agenda

- Motivation for Multi-level Page Tables
- **Example architecture**
- Class Problem: 32bit Intel x86
- Class Problem: Multi-Level VM Design
- VM Miscellanea


# Hierarchical page table – 32bit Intel x86



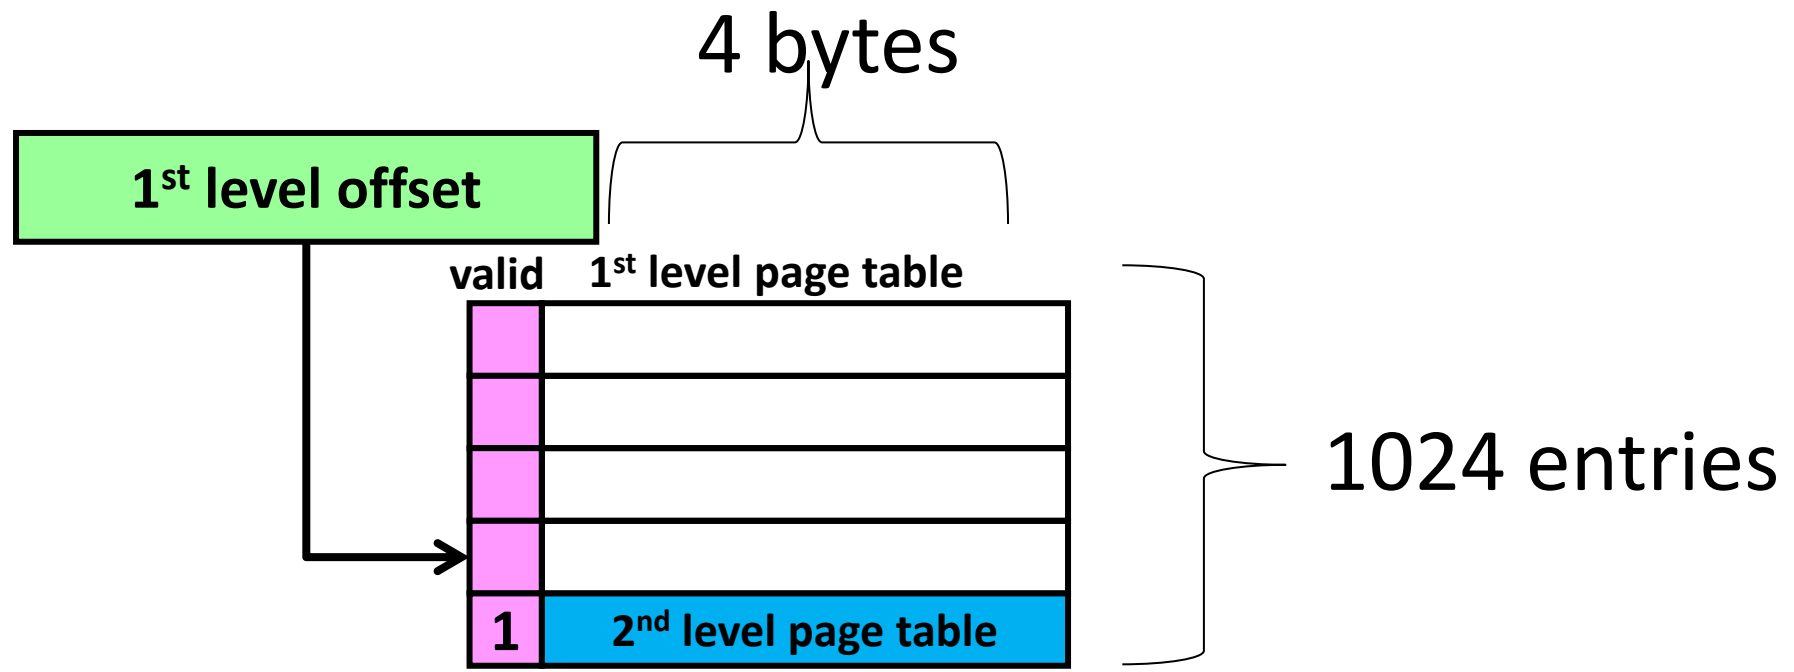
- How many bits in the virtual 1<sup>st</sup> level offset field? 10
- How many bits in the virtual 2<sup>nd</sup> level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1<sup>st</sup> level page table?  $2^{10}=1024$



# Hierarchical page table – 32bit Intel x86

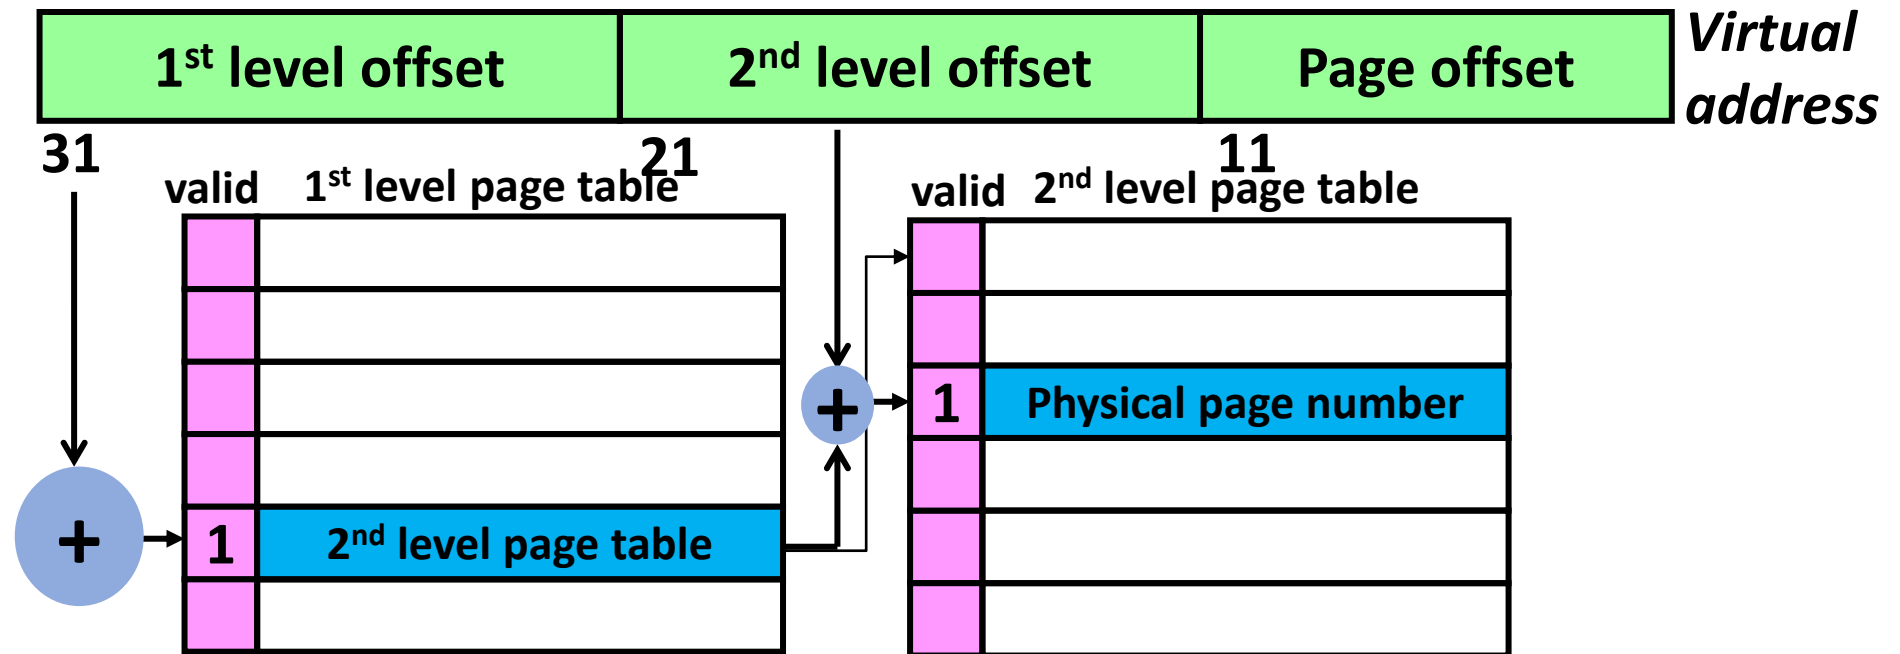


- Let's say physical address size + overhead bits is 4 bytes per entry
- Total size of 1<sup>st</sup> level page table
  - 4 bytes \* 1024 entries = 4 KB




# Hierarchical page table

- How many entries in the 2<sup>nd</sup> level of the page table?
  - $2^{10} = 1024$
- How many bytes for each VPN in a 2<sup>nd</sup> level table?
  - Let's round up to 4 bytes



# Hierarchical page table – 32bit Intel x86



- How many bits in the virtual 1<sup>st</sup> level offset field? 10
- How many bits in the virtual 2<sup>nd</sup> level offset field? 10
- How many bits in the page offset? 12
- How many entries in the 1<sup>st</sup> level page table?  $2^{10}=1024$
- How many bytes for each entry in the 1<sup>st</sup> level page table? 4
- How many entries in the 2<sup>nd</sup> level of the page table?  $2^{10}=1024$
- How many bytes for each entry in a 2<sup>nd</sup> level table?  $\sim 4$
- **What is the total size of the page table?**  
 (here  $n$  is number of valid entries in the 1<sup>st</sup> level page table)  **$4K+n*4K$**

# Agenda

- Motivation for Multi-level Page Tables
- Example architecture
- **Class Problem: 32bit Intel x86**
- Class Problem: Multi-Level VM Design
- VM Miscellanea

# Class Problem (32 bit x86)

1st      2nd  
4kB + n · 4kB

- What is the least amount of memory that could be used? When would this happen? *4kB . Only happen when we start the program and haven't even fetch*
- What is the most memory that could be used? When would this *the first instruction* happen?
- How much memory is used for this memory access pattern:  
0x00000ABC  
0x00000ABD  
0x10000ABC  
0x20000ABC
- How much memory if we used a single-level page table with 4KB pages? Assume entries are rounded to the nearest word (4B)



# Class Problem (32 bit x86)

- What is the least amount of memory that could be used? When would this happen?
  - 4KB for 1<sup>st</sup> level page table. Occurs when no memory has been accessed (before program runs)

- What is the most memory that could be used? When would this happen?



- 4KB for 1<sup>st</sup> level page table  
+ 1024\*4KB for all possible 2<sup>nd</sup> level page tables  
= 4100KB (which is slightly greater than 4096KB)

worse case. } 1024 active pages  
1024 away from each

But the spatial locality will avoid this.

All the second level page table.

- Occurs when program uses all virtual pages ( $= 2^{20}$  pages)


In the worse case, take slightly more memory than just take 1 level memory.

# Class Problem (32 bit x86)

if I have 4kB pages page offset is  $2^{12}$   
12 bits

- How much memory is used for this memory access pattern:

virtual page number	page offset	
0x00000	ABC	// Page fault
0x00000	ABD	
0x10000	ABC	// Page fault
0x20000	ABC	// Page fault



- 4KB for 1<sup>st</sup> level page table + 3\*4KB for each 2L page table  
= 16 KB

# Agenda

- Motivation for Multi-level Page Tables
- Example architecture
- Class Problem: 32bit Intel x86
- **Class Problem: Multi-Level VM Design**
- VM Miscellanea

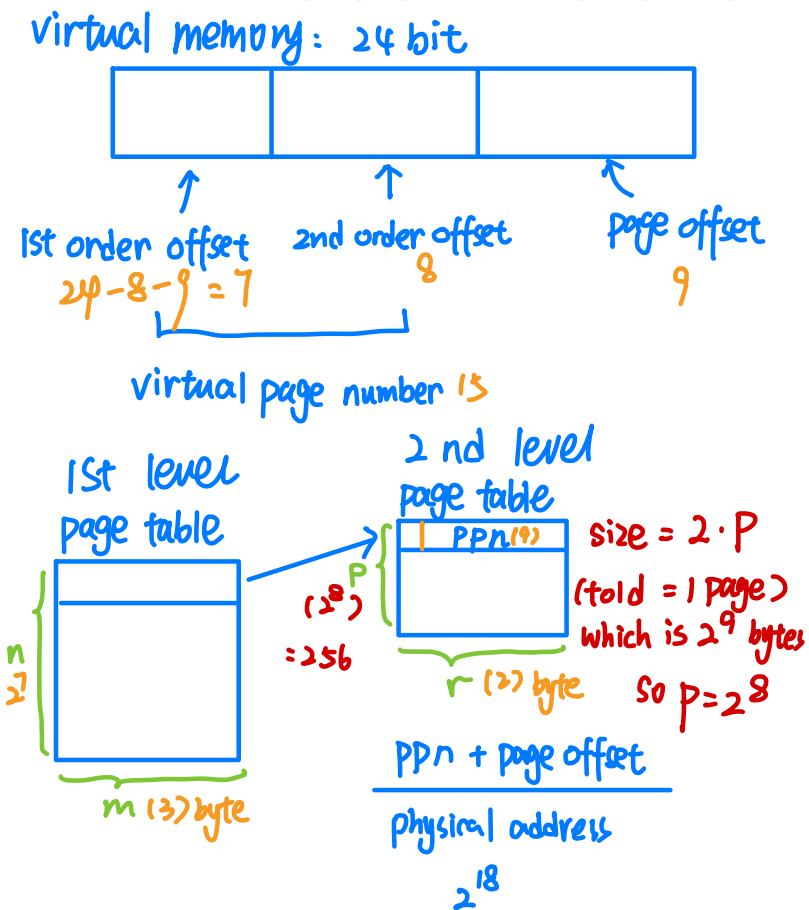
# Class Problem – Multi-level VM

- Design a two-level virtual memory system of a byte addressable processor with **24-bit long addresses**. No cache in the system.  
*this is virtual address.*  
*Are they mean virtual or physical?*
- **256Kbytes of memory** installed, and no additional memory can be added.  
*2<sup>18</sup>*  
*this is physical address.*  
*We don't install virtual memory.*
- **Virtual memory page: 512 Bytes**. Each page table entry must be an integer number of bytes, and must be the smallest size required to fit the physical page number + 1 bit to mark valid-entry.
- We want each second-level page table to fit exactly in one memory page, and 1<sup>st</sup> level page table entries are 3 bytes each (a memory address pointer to a 2L page table).
- Compute:
  - Number of entries in each 2<sup>nd</sup> level page table;
  - Number of virtual address bits used to index the 2<sup>nd</sup> level page table;
  - Number of virtual address bits used to index the 1<sup>st</sup> level page table;
  - Size of the 1<sup>st</sup> level page table.

# Class Problem – Multi-level VM

- Design a two-level virtual memory system of a byte addressable processor with **24-bit long addresses**. No cache in the system. **256Kbytes of memory** installed, and no additional memory can be added.
  - **Virtual memory page: 512 Bytes**. Each page table entry must be an integer number of bytes, and must be the smallest size required to fit the physical page number + 1 bit to mark valid-entry
  - We want each second-level page table to fit exactly in one memory page, and 1<sup>st</sup> level page table entries are 3 bytes each (a memory address pointer to a 2L page table).
- Compute:
  - Number of entries in each 2<sup>nd</sup> level page table;
  - Number of virtual address bits used to index the 2<sup>nd</sup> level page table;
  - Number of virtual address bits used to index the 1<sup>st</sup> level page table;
  - Size of the 1<sup>st</sup> level page table.

# Class Problem – Multi-level VM



- Since the physical address is  $2^{18}$   
and the virtual memory page = physical memory page =  $2^9$   
So physical page number in 2nd level page table is  $2^{18} \div 2^9 = 2^9$
- Also need to plus 1 bit valid-entry.  
so  $r = 1 + 9 = 10$ .
- Since each table entry must have integer number of bytes.  
So it should have 2 bytes  $\rightarrow$  16 bits.
- The total number of virtual pages.:  $2^{15}$  total entries for my second level page table.

# Class Problem – Multi-level VM

Page Offset: 9 bits (512B page size)

**Physical address = 18b (256KB Mem size)**

**Physical page number = 18b (256KB mem size) – 9b (offset)**

Physical page number = 9b	Page offset = 9b
---------------------------	------------------

2<sup>nd</sup> level page table entry size: 9b (physical page number) + 1b = ~ 2 bytes

2<sup>nd</sup> level page table **fits exactly in 1 page**

#entries in 2<sup>nd</sup> level page table is 512 bytes / 2 bytes = 256

#entries in 2<sup>nd</sup> level page table = 256 → Virtual page bits = 8b

Virtual 1<sup>st</sup> level page bits = 24 – 8 – 9 = 7b

1<sup>st</sup> level page table size =  $2^7 * 3 \text{ bytes} = 384\text{B}$

1 <sup>st</sup> level = 7b	2 <sup>nd</sup> level = 8b	Page offset = 9b
----------------------------	----------------------------	------------------

**Virtual address = 24b**

# Agenda

- Motivation for Multi-level Page Tables
- Example architecture
- Class Problem: 32bit Intel x86
- Class Problem: Multi-Level VM Design
- **VM Miscellanea**



# Page Replacement Strategies

- Page table indirection enables a fully associative mapping between virtual and physical pages.
- How do we implement LRU in OS?
  - True LRU is expensive, but LRU is a heuristic anyway, so approximating LRU is fine
  - Keep a “***accessed***” ***bit per page***, cleared occasionally by the operating system. Then pick any “unaccessed” page to evict

# Other VM Translation Functions

- Page data location
  - Physical memory, disk, uninitialized data
- Access permissions
  - Read only pages for instructions
  - **This is how your system detects segmentation faults**
- Gathering access information
  - Identifying dirty pages by tracking stores
  - Identifying accesses to help determine LRU candidate

# Next time

- Speeding up virtual memory