


EECS 370 - Lecture 5

C to Assembly



Announcements

- HW 1
 - Posted, due next-next Monday
- P1a
 - Due Thursday
- Labs
 - Lab 2 due Wednesday
 - Lab 3 meets Fr/M

Warm-Up Exercise

- Write ARM assembly code for the following C expression:
 - (assume an int is 4 bytes and that struct elements are stored contiguously)

```
struct { 4int a; 1unsigned char b, c; } y;  
y.a = y.b + y.c;
```

- Assume that a pointer to **y** is in X1.

```
LDURB x2, [x1, #4]      STURW x4, [x1, #0]
```

```
LDURB x3, [x1, #5]
```

```
add x4, x2, x3.
```


Warm-Up Exercise

- Write ARM assembly code for the following C expression:
 - (assume an int is 4 bytes and that struct elements are stored contiguously)

\uparrow
struct { int a; unsigned char b, c; } y;
 $\frac{y.a}{x_2} = \frac{y.b}{x_3} + \frac{y.c}{x_4}$

x1: y

- Assume that a pointer to **y** is in X1.

LDURB x3 [x1, #4] SDURW x2 [x1, #0]
LDURB x4 [x1, #5]
add x2 x3 x4

Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- **Lecture 5 : Converting C to assembly – basic blocks**
- Lecture 6 : Converting C to assembly – functions
- Lecture 7 : Translation software; libraries, memory layout



Agenda

- **Memory alignment**
 - Aligning Structs
- Control flow instructions
 - C-code examples
- Extra Problems

Calculating Load/Store Addresses for Variables

Datatype	size (bytes)
char	1
short	2
int	4
double	8

```
    2
short  a[100];
    1
char   b;
    4
int    c;
    8
double d;
    2
short  e;
struct {
    char f;
    int  g[1];
    char h;
} i;
```

- *Problem:* Assume data memory starts at address 100, calculate the total amount of memory needed

$$a = 2 \text{ bytes} * 100 = 200$$

$$b = 1 \text{ byte}$$

$$c = 4 \text{ bytes}$$

$$d = 8 \text{ bytes}$$

$$e = 2 \text{ bytes}$$

$$i = 1 + 4 + 1 = 6 \text{ bytes}$$

total = 221, right or wrong?

Memory layout of variables

- Compilers don't like variables placed in memory arbitrarily
- As we'll see later in the course, memory is divided into fixed sized **chunks**
 - When we load from a particular chunk, we really read the whole chunk
 - Usually an integer number of words (32 bits)
- If we read a single char (1 byte), it doesn't matter where it's placed

0x1000	0x1001	0x1002	0x1003
'a'	'b'	'c'	'd'

ldurb [x0, 0x1002]

- Reads [0x1000-0x1003], then throws away all but 0x1002, **fine**

Memory layout of variables

- BUT, if we read a 32-bit integer word, and that word starts at 0x1002:

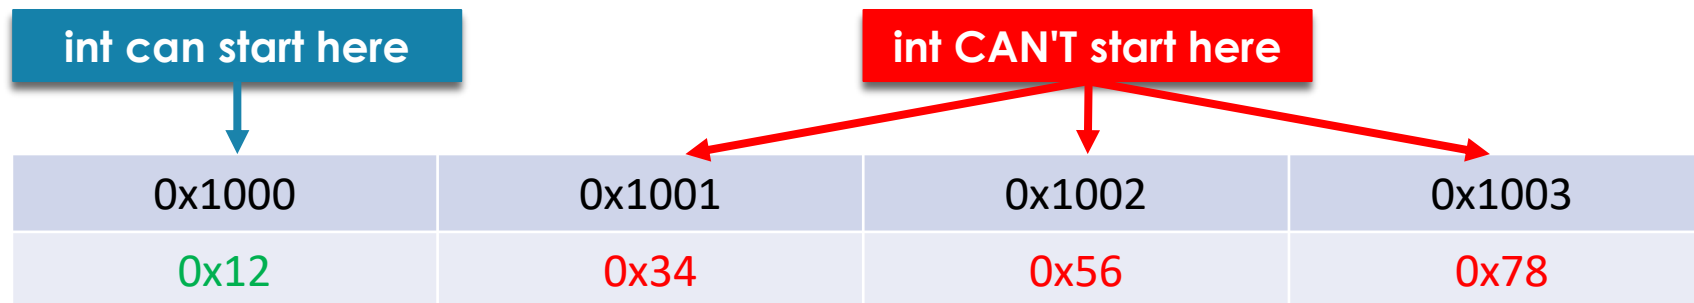
0x1000	0x1001	0x1002	0x1003
0xFF	0xFF	0x12	0x34
0x1004	0x1005	0x1006	0x1007
0x56	0x78	0xFF	0xFF

- First we need to read [0x1000-0x1003], throw away 0x1000 and 0x1001, **then** read [0x1004-0x1007]
- Need to read from memory twice! Slow! Complicated! **Bad!**


Solution: Memory Alignment

Poll: Where can chars start?

- Most modern ISAs require that data be aligned
 - An N-byte variable must start at an address A, such that $(A\%N) == 0$
- For example, starting address of a 32 bit **int** must be divisible by 4



- Starting address of a 16 bit **short** must be divisible by 2



Golden Rule of Alignment

You can't break into smaller size

- Every (primitive) object starts at an address divisible by its size
- "Padding" is placed in between objects if needed

```
char  c;  
short s;  
int   i;
```

0x1000	0x1001	0x1002	0x1003	0x1004	0x1005	0x1006	0x1007
[c]	[padding]	[s]		[i]			

- But what about non-primitive data types?
 - Arrays? Treat as independent objects
 - Structs? Trickier...

Agenda

- Memory alignment
 - **Aligning Structs**
- Control flow instructions
 - C-code examples
- Extra Problems

Problem with Structs

- If we align each element of a struct according to the Golden Rule, we can still run into issues
 - E.g.: An array of structs

```
char c; 1000

struct {
  char c; 1001
  int i; 1004-1007
} s[2]; 1001-1007
```

Amount of padding
is different across
different instances

1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	s[0].c	[pad]	[pad]	s[0].i				s[1].c	[pad]	[pad]	[pad]	s[1].i			

- Why is this bad?
- It makes "for" loops very difficult to write!
 - Offsets need to be different on each iteration

Structure Alignment

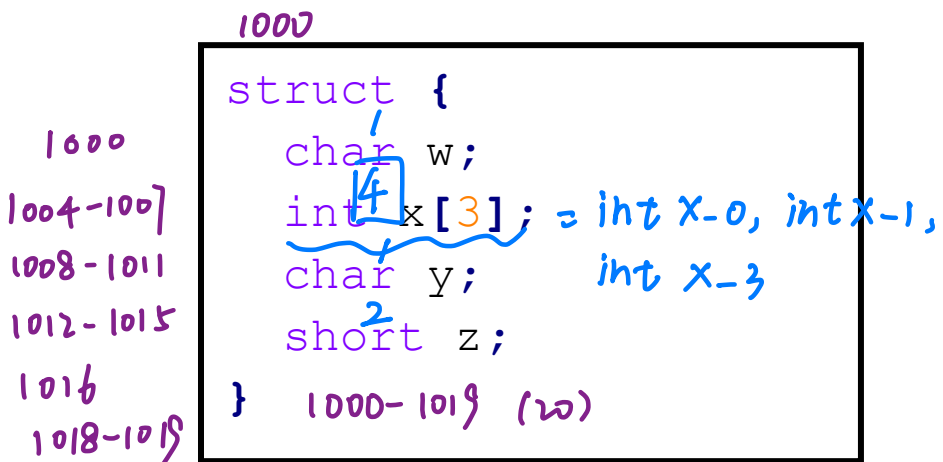
- Solution: in addition to laying out each field according to Golden Rule...
 - Identify largest (primitive) field
 - Starting address of overall struct is aligned based on the largest field
 - Padded in the back so total size is a multiple of the largest primitive

```
char c;  
  
struct {  
    char c;  
    int i;  
} s[2];
```

1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	100A	100B	100C	100D	100E	100F
c	[pad]	[pad]	[pad]	s[0].c	[pad]	[pad]	[pad]	s[0].i				s[1].c	[pad]	[pad]	[pad]

Guaranteed to lay
out each instance
identically

Structure Example



B

Poll: What boundary should this struct be aligned to?

- a) 1 byte
- ☒ b) 4 bytes
- c) 12 bytes
- d) 2 bytes
- e) 19 bytes

- Assume struct starts at location 1000,
 - char w → 1000
 - x[0] → 1004-1007, x[1] → 1008 – 1011, x[2] → 1012 – 1015
 - char y → 1016
 - short z → 1018 – 1019

Total size = 20 bytes!

Datatype	size (bytes)
char	1
short	2
int	4
double	8

Calculating Load/Store Addresses for Variables

8 B08
38

²short a[100];
¹char b;
⁴int c;
⁸double d;
²short e;
struct {
 ¹char f;
 ⁴int g[1];
 ¹char h;
}
i;

→ [100 299]
→ [300 300]
→ [304 307]
→ [312 319]
→ [320 321]
→ [324 325]
→ [328 331]
→ [332 332]
→ [324 335]

We need to add padding to ensure that the total size of this struct is divisible by four.

- *Problem:* Assume data memory starts at address 100, calculate the total amount of memory needed

a = 200 bytes (100-299)
 b = 1 byte (300-300)
 c = 4 bytes (304-307)
 d = 8 bytes (312-319)
 e = 2 bytes (320-321)
 struct: largest field is 4 bytes, start at 324
 f = 1 byte (324-324) *g: 324 - 327*
 g = 4 bytes (328-331) *f: 328 - 328*
 h = 1 byte (332-332) *h: 329 - 329*
 i = 12 bytes (324-335) *i: [324 - 331]*

236 bytes total!! (compared to 221, originally)

$6+2 = 8+1 = 9 \quad 12-9 = 3$



Data Layout – Why?

- Does gcc (or another compiler) reorder variables in memory to avoid padding?
- No, C99 forbids this
 - Memory is laid out in order of declaration for structs
- The programmer (i.e., you) are expected to manage data layout of variables for your program and structs.
- Two optimal strategies:
 - Order fields in struct by datatype size, smallest first
 - Or by largest first

Agenda

- Memory alignment
 - Aligning Structs
- **Control flow instructions**
 - C-code examples
- Extra Problems

ARM/LEGv8 Sequencing Instructions

- Sequencing instructions change the flow of instructions that are executed
 - This is achieved by modifying the program counter (PC)
- Unconditional branches are the most straightforward they ALWAYS change the PC and thus “jump” to another instruction out of the usual sequence
- Conditional branches

If (condition_test) goto target_address

condition_test examines the four flags from the processor status word (SPSR)

target_address is a 19 bit signed word displacement on current PC

LEGv8 Conditional Instructions

- Two varieties of conditional branches
 1. One type compares a register to see if it is equal to zero.
 2. Another type checks the condition codes set in the status register.

PC + offset : In LEGv8, we don't need PC + 1 + offset.

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + <u>100</u> <i>25x4</i>	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Let's look at the first type: CBZ and CBNZ
 - CBZ: Conditional Branch if Zero
 - CBNZ: Conditional Branch if Not Zero

LEGv8 Conditional Instructions

- CBZ/CBNZ: test a register against zero and branch to a PC relative address
 - The relative address is a *19 bit signed integer*—the number of instructions. Recall **instructions are 32 bits of 4 bytes**

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

- Example: CBNZ X3, Again
 - If X3 doesn't equal 0, then branch to label "Again"
 - "Again" is an offset from the PC of the current instruction (CBNZ)
 - Why does "25" in the above table result in PC + 100?

LEGv8 Conditional Instructions

- Example: What would the offset or displacement be if there were two instructions between ADDI and CBNZ?

0 Again: ADDI X3, X3, #-1
1 -----
2 -----
3 CBNZ X3, Again

$PC = PC + \text{offset}$
 $3 - 0 = 3$
 $\text{offset} = 3 \times 4 = 12$
 $PC' = 0 \quad PC = 3$
 $\text{offset} = 3$



-12

Poll: What's the offset?

- a) -16
- ~~b) -12~~
- c) -4
- ☒ d) -3
- e) 0

*We will go back 12 bytes,
But when we code the offset,
We only care how many instructions should
I go forward or backward.*

LEGv8 Conditional Instructions

- Motivation:
 - Some types of branches makes sense to check if a certain value is zero or not
 - while(a)
 - But not all:
 - if(a > b)
 - if(a == b)
 - Using an extra **program status register** to check for various conditions allows for a greater breadth of branching behavior

LEGv8 Conditional Instructions Using FLAGS

0 or 1 one bit

- FLAGS: NZVC record the results of (arithmetic) operations
Negative, Zero, oVerflow, Carry—not present in LC2K
- We explicitly set them using the “set” modification to ADD/SUB etc.
- Example: ADDS causes the 4 flag bits to be set according as the outcome is negative, zero, overflows, or generates a carry

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	ADD X1, X2, X3	$X1 = X2 + X3$	Three register operands
	subtract	SUB X1, X2, X3	$X1 = X2 - X3$	Three register operands
	add immediate	ADDI X1, X2, 20	$X1 = X2 + 20$	Used to add constants
	subtract immediate	SUBI X1, X2, 20	$X1 = X2 - 20$	Used to subtract constants
	add and set flags	ADDS X1, X2, X3	$X1 = X2 + X3$	Add, set condition codes
	subtract and set flags	SUBS X1, X2, X3	$X1 = X2 - X3$	Subtract, set condition codes
	add immediate and set flags	ADDIS X1, X2, 20	$X1 = X2 + 20$	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	$X1 = X2 - 20$	Subtract constant, set condition codes



ARM Condition Codes Determine Direction of Branch

- In LEGv8 only ADDS / SUBS / ADDIS / SUBIS / CMP /CMPI set the condition codes FLAGS or condition codes in PSR—the program status register
- Four primary condition codes evaluated:
 - N – set if the result is **negative** (i.e., bit 63 is non-zero)
 - Z – set if the result is **zero** (i.e., all 64 bits are zero)
 - ~~• C – set if last addition/subtraction had a **carry**/borrow out of bit 63~~
 - ~~• V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produce a positive result)~~
- Don't worry about the C and V for this class

ARM Condition Codes Determine Direction of Branch--continued

	Encoding	Name (& alias)	Meaning (integer)	Flags
→	0000	EQ	Equal	Z==1
→	0001	NE	Not equal	Z==0
	0010	HS (CS)	Unsigned higher or same (Carry set)	C==1
	0011	LO (CC)	Unsigned lower (Carry clear)	C==0
	0100	MI	Minus (negative)	N==1
	0101	PL	Plus (positive or zero)	N==0
	0110	VS	Overflow set	V==1
	0111	VC	Overflow clear	V==0
	1000	HI	Unsigned higher	C==1 && Z==0
	1001	LS	Unsigned lower or same	!(C==1 && Z==0)
→	1010	GE	Signed greater than or equal	N==V
→	1011	LT	Signed less than	N!=V
→	1100	GT	Signed greater than	Z==0 && N==V
→	1101	LE	Signed less than or equal	!(Z==0 && N==V)
→	1110	AL	Always	Any
	1111	NV [†]		

Need to know
the 7 with the
red arrows

```
CMP X1, X2
B.LE Label1
```

For this example,
we branch if X1 is
≥ to X2

Conditional Branches: How to use

pseudo instruction

- CMP instruction lets you compare two registers.
 - Could also use SUBS etc. *and then throw away the result (eg: to X31 Zero register)*
 - That could save you an instruction. *But then it will update the status register flags*
- B.cond lets you branch based on that comparison.

- Example:

```
CMP    X1, X2  
B.GT Label1
```

check condition

if $X_1 > X_2$, then branch to Label1.

- Branches to Label1 if X1 is greater than X2.

→ We don't have to specify any register, we can still tell the previous comparison $X_1 > X_2$, and then branch based of that.

The advantage: (split this across multiple instructions).

since we don't specify any other registers, so we can make the label pretty long. So can branch far away

Agenda

- Memory alignment
 - Aligning Structs
- Control flow instructions
 - **C-code examples**
- Extra Problems

Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

if condition is true, then go to the next line, if not jump to the other position.

However, Pc in default is go to the next line.

```
int x, y;
if (x == y)
    x++;
else
    y++;
// ...
```

invert

cmp x1 x2

b.ne else

addi x1 x1 #1

b end

or b.eq end

else addi x2 x2 #1

end

← unconditional branch.

no matter what, go to end label.

Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;
if (x == y)
    x++;
else
    y++;
// ...
```

Handwritten notes:

- x_1 points to `x`, x_2 points to `y`
- `cmp x1, x2`
- `B.NE else`
- `if ADDI x1, x1, #1`
- `else ADDI x2, x2, #1`

Branch—Example

- Convert the following C code into LEGv8 assembly (assume x is in X1, y in X2):

```
int x, y;  
if (x == y)  
    x++;  
else  
    y++;  
// ...
```

Using Labels

```
CMP X1, X2  
B.NE L1  
ADD X1, X1, #1  
B L2  
L1: ADD X2, X2, #1  
L2: ...
```

Note that conditions in assembly are often the inverse of the "if" condition. Why?

Without Labels

```
CMP X1, X2  
B.NE 3  
ADD X1, X1, #1  
B 2  
ADD X2, X2, #1
```

Assemblers must deal with labels and assign displacements

Loop—Example

// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

of branch instructions
= $3 \cdot 10 + 1 = 31$

a.k.a. while-do template

	MOV	X1, XZR	} initialize i and sum = 0
	MOV	X2, XZR	
Loop1:	CMPI	X1, #10	
	B.EQ	endLoop	
	LSL	X6, X1, #3	
	LDUR	X5, [X6, #100]	
	CMPI	X5, #0	
	B.LT	endif	
	ADD	X2, X2, X5	
endif:	ADDI	X1, X1, #1	
	B	Loop1	
endLoop:			

Loop—Example

// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;
for (i=0; i < 10; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

Handwritten annotations:
- $\leftarrow X_1$ above `i=0`
- $\leftarrow X_2$ above `sum = 0`
- $\leftarrow X_2$ above `sum += a[i]`
- $\leftarrow X_3$ above `a[i]`
- $\leftarrow X_3$ address above `a[i]`

of branch instructions
= $3 \cdot 10 + 1 = 31$

a.k.a. while-do template

```
MOV X1, XZR
MOV X2, XZR
for CMPI X2, #10
    B.GE endfor
    LSL X3, X1, #3
    LDUR X4, X3, #100
    CMPI X4, #0
    B.LT endif
    ADD X1, X1, X4
    ADDI X2, X2, #1
endif
B.for
endfor
```

Agenda

- Memory alignment
 - Aligning Structs
- Control flow instructions
 - C-code examples
- **Extra Problems**

Extra Example: Do-while Loop

// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

of branch instructions
= $2 * 10 = 20$

a.k.a. do-while template

	MOV	X1, XZR
	MOV	X2, XZR
Loop1:	LSL	X6, X1, #3
	LDUR	X5, [X6, #100]
	CMPI	X5, #0
	B.LT	endif
	ADD	X2, X2, X5
endif:	ADDI	X1, X1, #1
	CMPI	X1, #10
	B.LT	Loop1
endLoop:		

Extra Example: Do-while Loop

// assume all variables are long long integers (64 bits or 8 bytes)
// i is in X1, start of a is at address 100, sum is in X2

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

of branch instructions
= $2 \times 10 = 20$

a.k.a. do-while template



Extra Problem – For Your Reference

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in X1  
// struct {int val; struct node *next;} node;  
// struct node *ptr;
```

```
if ((ptr != NULL) && (ptr->val > 0))  
    ptr->val++;
```

Extra Problem

- Write the ARM assembly code to implement the following C code:

```
// assume ptr is in X1
// struct {int val; struct node *next;} node;
// struct node *ptr;
```

```
if ((ptr != NULL) && (ptr->val > 0))
    ptr->val++;
```

```
cmp r1, #0
beq Endif
ldursw r2, [r1, #0]
cmp r2, #0
b.le Endif
add r2, r2, #1
str r2, [r1, #0]
Endif : ....
```

Extra Class Problem

- How much memory is required for the following data, assuming that the data starts at address 200 and is a 32 bit address space?

```
int a;  
struct {double b, char c, int d} e;  
char* f;  
short g[20];
```

Poll: How much memory?

- a) $x < 40$ bytes
- b) $40 < x < 50$ bytes
- c) $50 < x < 60$ bytes
- d) $60 < x$ bytes



Next Time

- More C-to-Assembly
 - Function calls
- Lingering questions / feedback? Post it to Slido