

**Poll: Who is responsible for
erasing a whiteboard in a public
space?**

- a) The person who is done using it
- b) The person who is about to use it

EECS 370 - Lecture 6

Function Calls



Announcements

- P1
 - Part a due tonight @ 11:55 via Autograder
 - Project 1 s + m due next Thu



- HW 1
 - Due next-next Monday
- Lab 3 meets Fr/M
- Get exam conflicts and SSD accommodations sent to us **in the next week**
 - Forms listed on the website

Instruction Set Architecture (ISA) Design Lectures

- Lecture 2: ISA - storage types, binary and addressing modes
- Lecture 3 : LC2K
- Lecture 4 : ARM
- Lecture 5 : Converting C to assembly – basic blocks
- **Lecture 6 : Converting C to assembly – functions**
- Lecture 7 : Translation software; libraries, memory layout



Agenda

- **Using branches more generally**
- Function calls and the call stack
- Assigning variables to memory locations
- Saving registers
- Caller/callee example

Branching far away

- Underlying philosophy of ISA design: **make the common case fast**
- Most branches target nearby instructions
 - Displacement of 19 bits is usually enough
- BUT what if we need to branch really far away (more than 2^{19} words)?

CBZ X15, FarLabel

- The assembler is smart enough to replace that with

CBNZ X15, L1

B FarLabel

L1:

← replace

- The simple branch instruction (B) has a 26 bit offset which spans about 64 million instructions!
- In LC2K, we can do a similar thing by using JALR instead of BEQ

Unconditional Branching Instructions

Unconditional branch	branch	B	2500	go to PC + 10000	Branch to target address; PC-relative
	branch to register	BR	X30	go to X30	For switch, procedure return
	branch with link	BL	2500	X30 = PC + 4; PC + 10000	For procedure call PC-relative

Handwritten notes:

- We don't specify X30 here → So X30 is implicitly been used.* (pointing to X30 in the BL row)
- number of instruction* (under 2500 in the BL row)
- each instruction is 4 bytes* (under PC + 4 in the BL row)
- this hold the memory address* (under PC + 10000 in the BL row)

- There are three types of unconditional branches in the LEGv8 ISA.

- The first **(B)** is the PC relative branch with the 26 bit offset from the last slide.
- The second **(BR)** jumps to the address contained in a register (X30 above)
- The third **(BL)** is like our PC relative branch but it does something else.
 - It sets X30 (always) to be the current PC+4 before it branches.

Look inside the register which hold the address that we want to branch to . (64 bits)

- Why is BL storing PC+4 into a register?

Branch with Link (BL)

- Branch with Link is the branch instruction used to call functions
 - Functions need to know where they were called from so they can return.
 - In particular they will need to return to right after the function call
 - Can use “BR X30” *saving where you are, when you return from the function, you know where you go.*

- Say that we execute the instruction BL #200 when at PC 1000. *means 200 instruction*

- What address will be branched to? *1800*

- What value is stored in X30? *X30 = 1004*

- How is that value in X30 useful?

we will finally branch back to X30

1000 : BL #200 $\xrightarrow{\text{go to}}$ 1000 + 4x200 = 1800

Poll

Agenda

- Using branches more generally
- **Function calls and the call stack**
- Assigning variables to memory locations
- Saving registers
- Caller/callee example

Converting function calls to assembly code

C: factorial(5);

- Need to pass parameters to the called function—factorial
- Need to save return address of caller so we can get back
- Need to save register values (why?)
- Need to jump to factorial



Execute instructions for factorial()
Jump to return address

The diagram consists of two black arrows. The first arrow originates from the bullet point 'Need to jump to factorial' and points to the text 'Execute instructions for factorial()'. The second arrow originates from the text 'Jump to return address' and points to the bullet point 'Need to get return value (if used)'.

- Need to get return value (if used)
- Restore register values

Task 1: Passing parameters

- Where should you put all of the parameters?
 - Registers?
 - Fast access but few in number and wrong size for some objects
 - Memory?
 - Good general solution but slow
- ARMv8 solution—and the usual answer:
 - Both
 - Put the first few parameters in registers (if they fit) (X0 – X7)
 - Put the rest in memory on the call stack— **important concept**

Call stack

- ARM conventions (and most other processors) allocate a region of memory for the “call” stack
 - This memory is used to manage all the storage requirements to simulate function call semantics
 - Parameters (that were not passed through registers)
 - Local variables
 - Temporary storage (when you run out of registers and need somewhere to save a value)
 - Return address ← *Before we call other function, we need to store the return address for this level function.*
 - Etc.
- Sections of memory on the call stack [**stack frames**] are allocated when you make a function call, and de-allocated when you return from a function

The stack grows as functions are called

FUNCTION CALLS

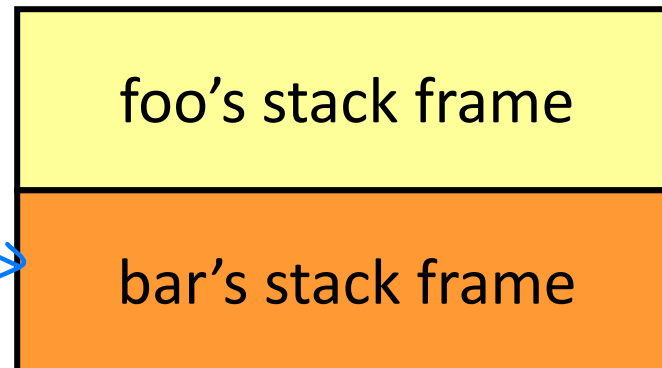
```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

inside foo

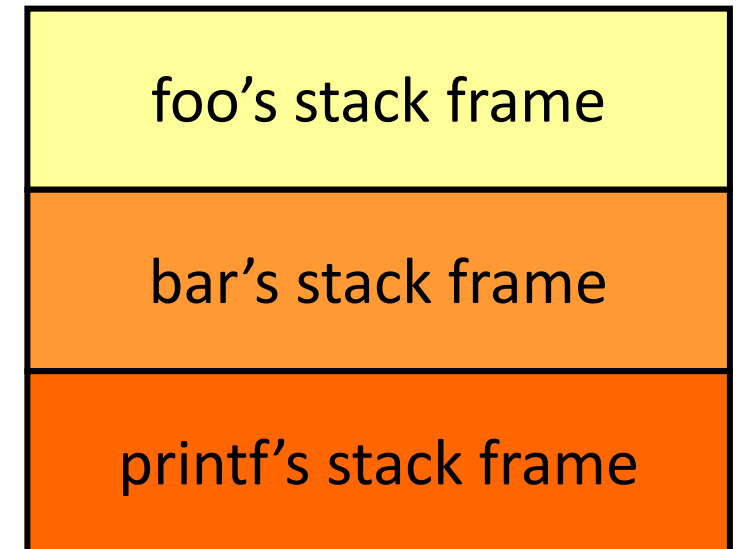


```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

foo calls bar



bar calls printf



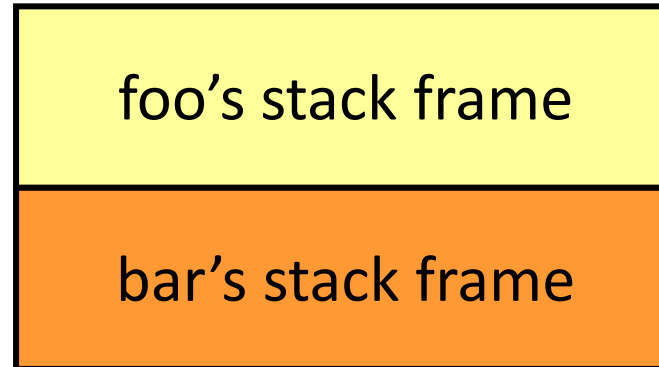
The stack shrinks as functions return

FUNCTION CALLS


```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}
```

```
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

printf returns



bar returns



Stack frame contents

FUNCTION CALLS

foo was called by main

foo's stack frame

```
void foo()  
{  
    int x, y[2];  
    bar(x);  
}  
  
void bar(int x)  
{  
    int a[3];  
    printf();  
}
```

return addr to main
x
y[0]
y[1]
<u>spilled registers in foo</u>

*as we doing internal computation
if they don't fit inside of our
own registers .*

Stack frame contents (2)

```
void foo()
{
    int x, y[2];
    bar(x);
}

void bar(int x)
{
    int a[3];
    printf();
}
```

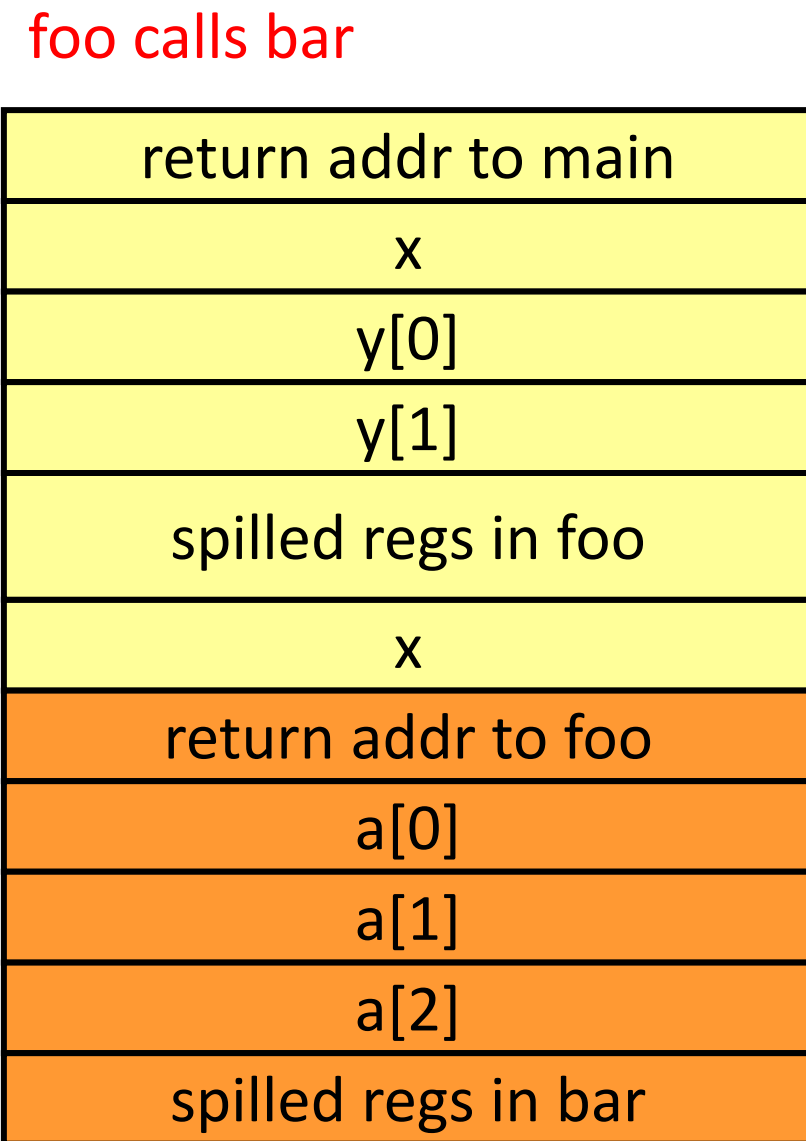
foo would write x to this location to pass the parameter to bar, And bar can read from that location to see what that parameter is.

*two copy of x:
① foo has its local copy
② bar also has it.*

Spill data—not enough room in x0-x7 for params and also caller and callee saves

foo's frame

bar's frame




Agenda

- Using branches more generally
- Function calls and the call stack
- **Assigning variables to memory locations**
- Saving registers
- Caller/callee example

Review: Where do the variables go?

Assigning variables to memory spaces

FUNCTION CALLS



Assigning variables to memory spaces

FUNCTION CALLS

```
int w;  
void foo(int x)  
{  
    static int y[4];  
    char* p;  
    p = malloc(10);  
    //...  
    printf("%s\n", p);  
}
```

w goes in static, as it's a global

x goes on the stack, as it's a parameter

bigger

y goes in static, 1 copy of this!!

p goes on the stack


allocate 10 bytes on heap, ptr
set to the address

string goes in static, pointer
to string on stack, p goes on
stack

The addresses of local variables
will be different depending on
where we are in the call stack

So we can't hard code the address

smaller



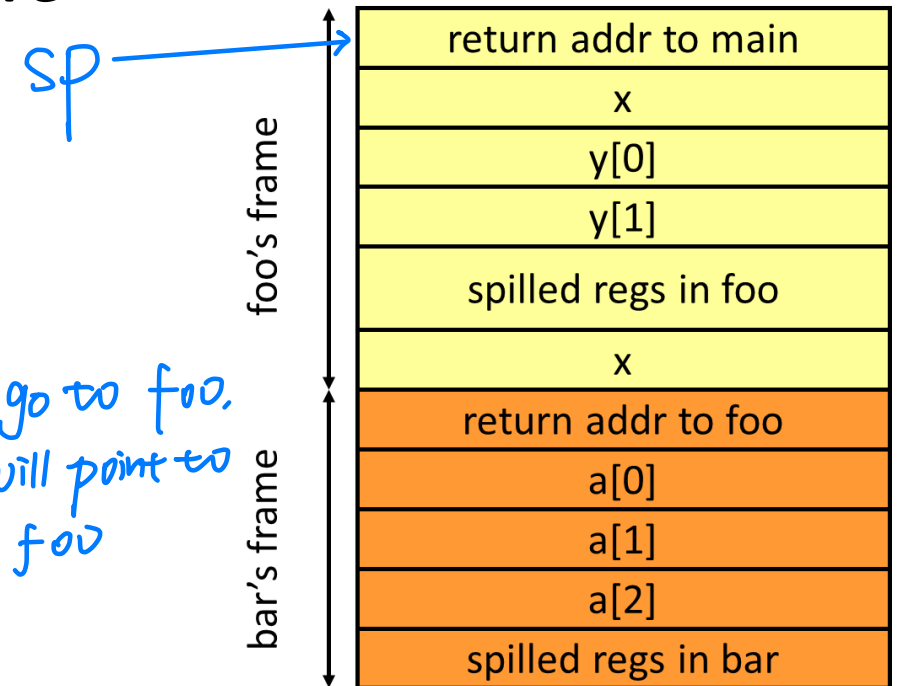
Accessing Local Variables

- Stack pointer (SP):
 - register that keeps track of current top of stack
- Compiler (or assembly writer) knows relative offsets of objects in stack
- Can access using lw/sw offsets

specific register
↓
SW x30 [SP, #0]
SW x1 [SP, #4]
SW x2 [SP, #8]

*All the address is related to sp
when you express them*

*when we go to foo,
the sp will point to
top of foo*



Agenda

- Using branches more generally
- Function calls and the call stack
- Assigning variables to memory locations
- **Saving registers**
- Caller/callee example

What about registers?

- Higher level languages (like C/C++) provide many abstractions that don't exist at the assembly level
- E.g. in C, each function has its own local variables
 - Even if different function have local variables with the same name, they are independent and guaranteed not to interfere with each other!

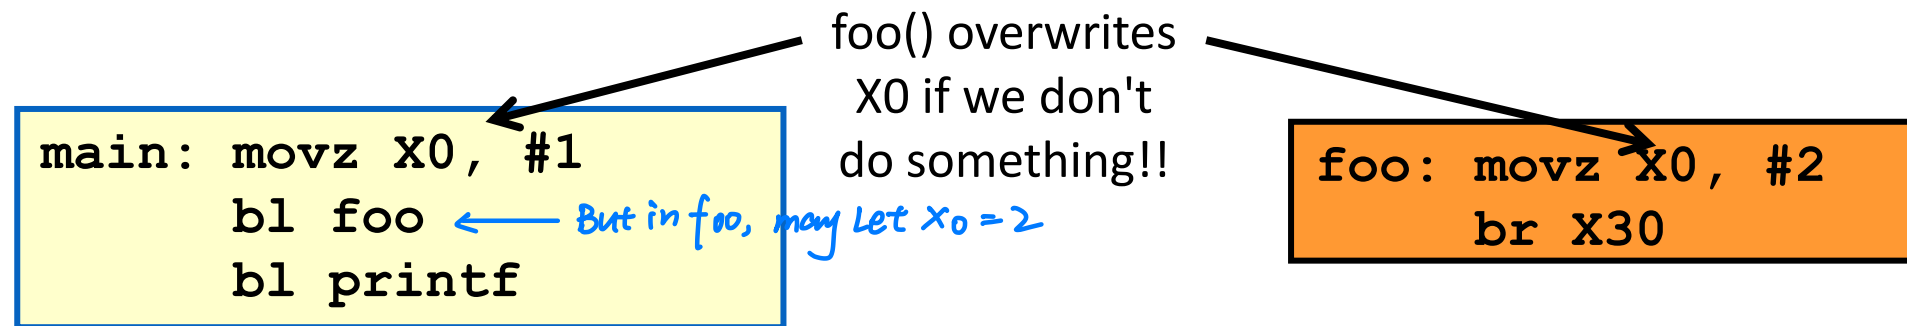
```
void foo() {  
    int a=1;  
    bar();  
    printf(a);  
}
```

Still prints "1"..
these don't
interfere

```
void bar() {  
    int a=2;  
    return;  
}
```

What about registers?

- But in assembly, all functions share a small set (e.g. 32) of registers
 - Called functions will overwrite registers needed by calling functions



- "Someone" needs to save/restore values when a function is called to ensure this doesn't happen

Two Possible Solutions

- Either the **called** function **saves** register values before it overwrites them and **restores** them before the function returns (**callee** saved)...

```
main: movz X0, #1
      bl  foo
      bl  printf
```

Store X0 out to Stack

```
foo: stur X0, [stack]
     movz X0, #2
     ldur X0, [stack]
     br  X30 load back to X0
```

- Or the **calling** function **saves** register values before the function call and **restores** them after the function call (**caller** saved)...

```
main: movz X0, #1
      stur X0, [stack]
      bl  foo
      ldur X0, [stack]
      bl  printf
```

```
foo: movz X0, #2
     br  X30
```

Another example

Original C Code

```
void foo() {
    int a,b,c,d;

    a = 5; b = 6;
    c = a+1; d=c-1;
```

```
    bar();
```

```
    d = a+d;
    return();
}
```

only reading two of them after bar has returned

No need to
save r2/r3.
Why?

Additions for Caller-save

```
void foo() {
    int a,b,c,d;

    a = 5; b = 6;
    c = a+1; d=c-1;
    save r1 to stack
    save r4 to stack
    bar();
    restore r1
    restore r4
    d = a+d;
    return();
}
```

Handwritten: r1 points to a, r4 points to d

Assume bar() will
overwrite registers
holding a,d

Additions for Callee-save

```
void foo() {
    int a,b,c,d;
    save r1
    save r2
    save r3
    save r4
    a = 5; b = 6;
    c = a+1; d=c-1;
    bar();
    d = a+d;
    restore r1
    restore r2
    restore r3
    restore r4
    return();
}
```

bar() will save a,b, but
now foo() must save
main's variables

“caller-save” vs. “callee-save”

- Caller-save

- What if bar() doesn't use r1/r4?
- No harm done, but wasted work

```
void foo(){  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    save r1 to stack  
    save r4 to stack  
    bar();  
    restore r1  
    restore r4  
    d = a+d;  
    return();  
}
```


- Callee-save

- What if main() doesn't use r1-r4?
- No harm done, but wasted work

```
void foo(){  
    int a,b,c,d;  
    save r1  
    save r2  
    save r3  
    save r4  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
    bar();  
    d = a+d;  
    restore r1  
    restore r2  
    restore r3  
    restore r4  
    return();  
}
```

Another helpful visual

CALLER-CALLEE



Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- **Caller-saved**
 - Only needs saving if value is “live” across function call
 - **Live** = contains a useful value: Assign value before function call, use that value after the function call
 - In a leaf function (a function that calls no other function), caller saves can be used without saving/restoring

a, d are live

b, c are NOT
live

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

Saving/Restoring Optimizations

CALLER-CALLEE

- Where can we avoid loads/stores?
- Callee-saved
 - Only needs saving at beginning of function and restoring at end of function
 - Only save/restore it if function overwrites the register

Only use r1-
r4

No need to
save other
registers

```
void foo() {  
    int a,b,c,d;  
  
    a = 5; b = 6;  
    c = a+1; d=c-1;  
  
    bar();  
  
    d = a+d;  
    return();  
}
```

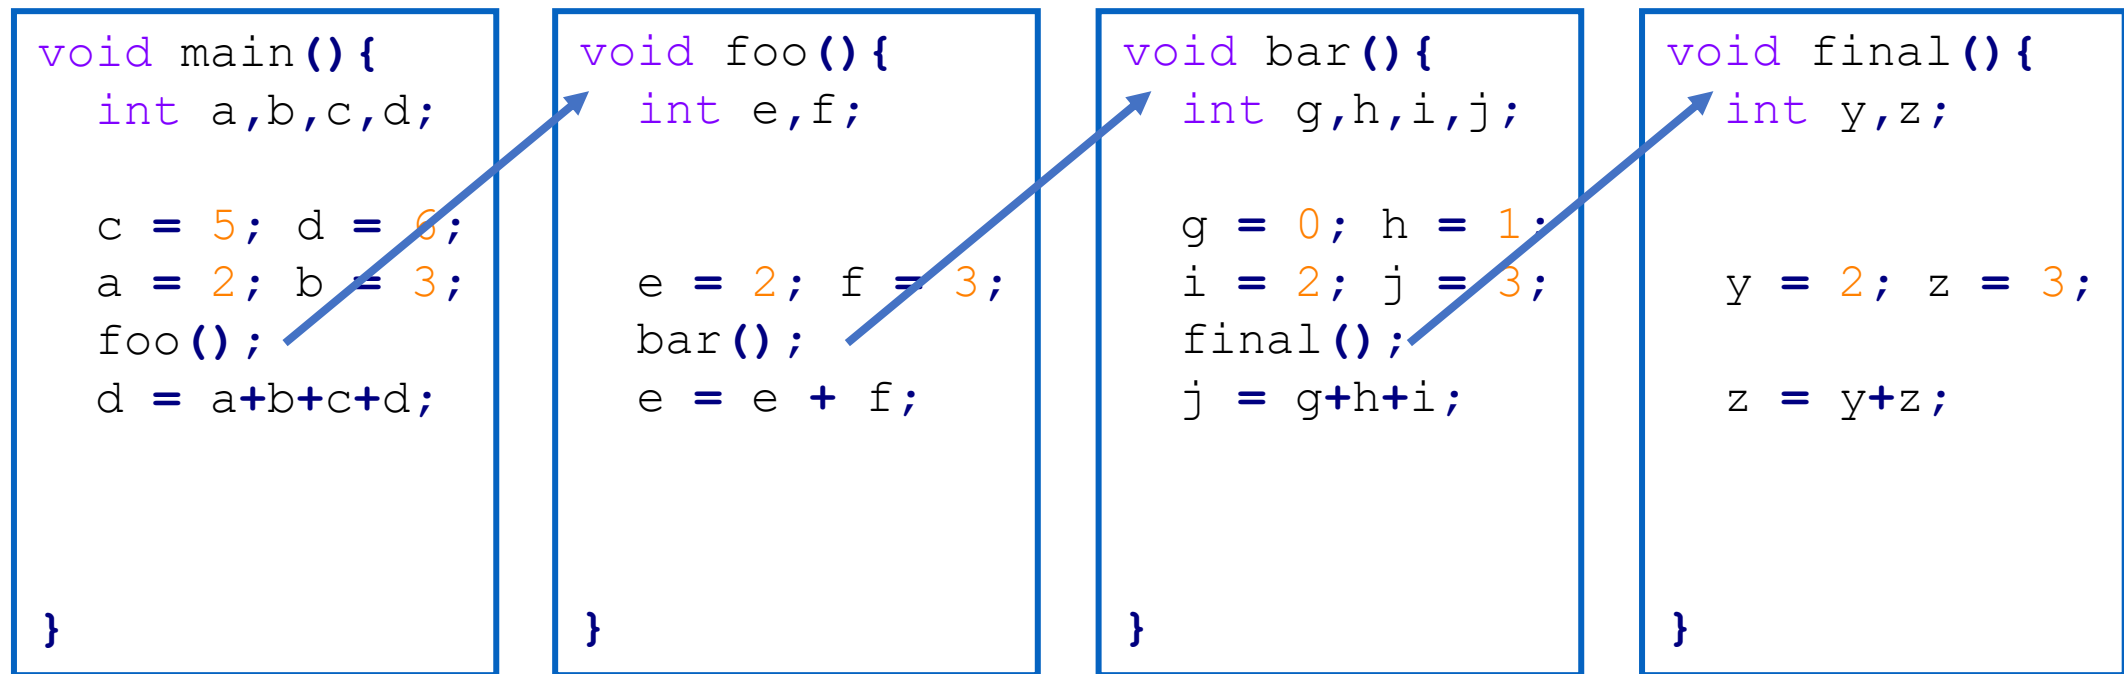
Agenda

- Using branches more generally
- Function calls and the call stack
- Assigning variables to memory locations
- Saving registers
- **Caller/callee example**

Caller versus Callee

- Which is better??
- Let's look at some examples...
- Simplifying assumptions:
 - A function can be invoked by many different call sites in different functions.
 - Assume no inter-procedural analysis (hard problem)
 - A function has no knowledge about which registers are used in either its caller or callee
 - Assume `main()` is not invoked by another function
- Implication
 - Any register allocation optimization is done using function local information

Caller-saved vs. callee saved – Multiple function case



Note: assume main does not have to save any callee registers

Caller-saved vs. callee saved – Multiple function case

- Questions:

1. How many registers need to be saved/restored if we use a **caller-save** convention?
2. How many registers need to be saved/restored if we use a **callee-save** convention?
3. How many registers need to be saved/restored if we use a mix of **caller-save** and **callee-save**?

Question 1: Caller-save

```
void main() {  
    int a,b,c,d;  
    c = 5; d = 6;  
    a = 2; b = 3;  
    [4 STUR]  
    foo();  
    [4 LDUR]  
    d = a+b+c+d;  
    ↑  
}
```

```
void foo() {  
    int e,f;  
  
    e = 2; f = 3;  
    [2 STUR]  
    bar();  
    [2 LDUR]  
    e = e + f;  
    2  
}
```

```
void bar() {  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    [3 STUR]  
    final();  
    [3 LDUR]  
    j = g+h+i;  
    3  
}
```

```
void final() {  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
}
```

Total: 9 STUR / 9 LDUR

Question 2: Callee-save

Poll: How many ld/st pairs are needed?

```
void main() {  
    int a,b,c,d;  
  
    c = 5; d = 6;  
    a = 2; b = 3;  
    foo();  
    d = a+b+c+d;  
}
```

```
void foo() {  
    [2 STUR]  
    int e,f;  
  
    e = 2; f = 3;  
    bar();  
    e = e + f;  
  
    [2 LDUR]  
}
```

```
void bar() {  
    [4 STUR]  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    final();  
    j = g+h+i;  
  
    [4 LDUR]  
}
```

```
void final() {  
    [2 STUR]  
    int y,z;  
  
    y = 2; z = 3;  
    z = y+z;  
  
    [2 LDUR]  
}
```

Total: 8 STUR / 8 LDUR

Is one better?

- **Caller-save** works best when we don't have many live values across function call
- **Callee-save** works best when we don't use many registers overall
- We probably see functions of both kinds across an entire program
- Solution:
 - Use both!
 - E.g. if we have 6 registers, use some (say r0-r2) as **caller-save** and others (say r3-r5) as **callee-save**
 - Now each function can optimize for each situation to reduce saving/restoring

Poll: What's the optimal number of caller/callee registers for final?

Question 3: Mixed 3 caller / 3 callee

- For main, ideally put all variables in callee-save registers
- But we only 3 are available
- One variable needs to go in caller-saved register

```
void main() {  
    int a,b,c,d;  
    c = 5; d = 6;  
    a = 2; b = 3;  
    [1 STUR]  
    foo();  
    [1 LDUR]  
    d = a+b+c+d;  
}
```

1 caller r
3 callee r

```
void foo() {  
    [2 STUR]  
    int e,f;  
  
    e = 2; f = 3;  
    bar();  
    e = e + f;  
  
    [2 LDUR]  
}
```

2 callee r
(2 caller would
be equivalent)

```
void bar() {  
    [3 STUR]  
    int g,h,i,j;  
    g = 0; h = 1;  
    i = 2; j = 3;  
    final();  
    j = g+h+i;  
  
    [3 LDUR]  
}
```

1 caller r
3 callee r

```
void final() {  
    int y,z;  
  
    y = 2; z = 3;  
  
    z = y+z;  
}
```

2 caller r

Total: 6 STUR / 6 LDUR

LEGv8 ABI- Application Binary Interface

- The ABI is an agreement about how to use the various registers
- Not enforced by hardware, just a convention by programmers / compilers
- If you won't your code to work with other functions / libraries, **follow these**
- Some register conventions in ARMv8
 - X30 is the **link register** – used to hold return address
 - X28 is **stack pointer** – holds address of top of stack
 - X19-X27 are **callee-saved** – function must save these before writing to them
 - X0-15 are **caller-saved** –function must save live values before call
 - X0-X7 used for **arguments** (memory used if more space is needed)
 - X0 used for **return value**

Next Time

- Finish Up Function Calls
- Talks about linking – the final puzzle piece of software