



L5 – Data Layout and Control Flow

EECS 370 – Introduction to Computer Organization – Winter 2022

Outline

- Data layout
 - Structs
- Control Flow
 - Branch instructions

L5_1 Assembly – Data Layout

EECS 370 – Introduction to Computer Organization – Winter 2022

Learning Objectives

- Understand mapping of C-code data structures (`struct`) to data layout in memory (e.g., stack)

Reminders

- HW1 due Monday 1/24
- P1a due Thursday 1/27
- P1.s, P1.m due next Thursday 2/3
- HW2 due Monday 2/7
- Midterm Thursday 3/10 @ 7–9pm

Resources

- Video reviews of many topics!
 - <https://www.youtube.com/channel/UCVHvjtdIZI40TWd5-64hTLQ/videos>

The screenshot shows the course resources page for EECS 370. On the left is a dark sidebar with white text containing links to 'EECS 370', 'Calendar', 'Lecture', 'Discussion', 'Assignments', 'Exams', 'Admin Requests', 'Schedule', 'Course Resources', 'Staff', and 'Syllabus'. The main content area has a light gray background and features a title 'Course Resources' at the top. Below it are four large rectangular boxes arranged horizontally. The first box is titled 'Review Content' and contains two items: 'EECS 370 Youtube Channel' and 'Binary, Hex, and 2's complement Review Sheet'. The second box is titled 'Simulators' and contains two items: 'Cache Simulator' and 'Pipeline Simulator'. The third box is titled 'Reference Material' and contains three items: 'Green LEGv8 Cheat Sheet', 'C for C++ users by Ian Cooke', and 'Symbol Table and Relocation Table for EECS 370'. The fourth box is titled 'GDB Content' and contains two items: 'GDB Tutorial' and 'GDB Reference Card'. At the bottom of the page, there is a section titled 'Staff'.

Converting C to assembly – Example #2

ARM ISA
REVIEW!

Write ARM assembly code for the following C expression (assume an int is 4 bytes, unsigned char is 1 byte)

Register to variable mapping

X1→pointer to y

C-code instructions

a is an integer, and an integer is 4 bytes
struct { int a; unsigned char b, c; } y;
y.a = y.b + y.c; only one byte

ARM LEGv8

```
LDURB X2, [X1, #4] // load y.b  
LDURB X3, [X1, #5] (4+1) // load y.c  
ADD X4, X2, X3 // calculate y.b + y.c  
STURW X4, [X1, #0] // store y.a
```

since a is at the beginning of this struct

See supplemental video for detailed explanation

How do you determine offsets for struct sub-fields?
THIS lecture will detail

Calculating Load/Store Addresses

Problem: Calculate the total amount of memory needed for the struct instance x

Assume data memory starts at address 1000

Datatype	size (bytes)
short	2
char	1
int	4
double	8

C-code

```
short a[100];
char b;
int c;
double d;
short e;
struct {
    char f;
    int g[1];
    char h;
} x;
```

*{ compound variable
(class & struct).*

Calculating Load/Store Addresses

Data
Layout

Problem: Calculate the total amount of memory needed for the struct instance x

Assume data memory starts at address 1000

Datatype	size (bytes)
short	2
char	1
int	4
double	8

C-code

```
short a[100];
char b;
int c;
double d;
short e;
struct {
    char f;
    int g[1];
    char h;
} x;
```

Calculating Load/Store Addresses

Problem: Calculate the total amount of memory needed for the struct instance x
 Assume data memory starts at address 1000

Datatype	size (bytes)
short	2
char	1
int	4
double	8

C-code

```
short a[100];
char b;
int c;
double d;
short e;
struct {
    char f;
    int g[1];
    char h;
} x;
```

Solution???

```
a = 2 bytes * 100 = 200
b = 1 byte
c = 4 bytes
d = 8 bytes
e = 2 bytes
x = 1 + 4 + 1 = 6 bytes
```

Total: 221 bytes???

Correct or incorrect?

Because we have alignment restriction

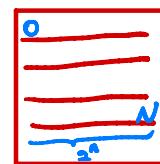
对齐

tell you where a variable can start in a memory

Memory layout of variables

- Most modern ISAs require that data be aligned
- What do we mean by alignment in this context?
 - An N-byte variable must start at an address A, such that $(A \% N) == 0$
- “Golden” rule – Address of a variable is aligned based on the size of the variable
 - **char** is byte aligned (any address is fine)
 - **short** is half-word (H) aligned (LSBit of address must be 0) 
 - **int** is word aligned (W) (2 LSBits of addr must be 0) 
- This greatly simplifies hardware needed for loads and stores
 - Otherwise, multiple memory accesses need to be used to access one piece of data

The reason why the Golden rule



The reason why 2^n is because we you want to know where to end, you need use size of variable divide size of row. —> shifting

rows of rows of bits

Structure Alignment

- Each field is laid out in the order it is declared using the Golden Rule for alignment
- Identify largest (primitive) field
 - A • Starting address of overall struct is aligned based on the largest field
 - B • Size of overall struct is a multiple of the largest field
 - Reason for this is so we can have an array of structs
 - Guarantees that each instance of struct is aligned the same way

this is because : We could have an array of struct.

A | B

Structure Alignment - Example

Data
Layout

Problem: What boundary should this struct be aligned to?
What is the total size of the struct?

C-code

```
struct {  
    char w;  
    int x[3];  
    char y;  
    short z;  
} s;
```

Datatype	size (bytes)
short	2
char	1
int	4
double	8

Assume struct starts at address
1000, what is the data layout
of the struct?

Structure Alignment - Example

Data
Layout

Problem: What boundary should this struct be aligned to?
What is the total size of the struct?

C-code

```
struct {  
    char w;  
    int x[3];  
    char y;  
    short z;  
} s;
```

Datatype	size (bytes)
short	2
char	1
int	4
double	8

Assume struct starts at address 1000, what is the data layout of the struct?

Structure Alignment - Example

Problem: What boundary should this struct be aligned to?
What is the total size of the struct?

C-code

```
struct {  
    char w;  
    int x[3];  
    char y;  
    short z;  
} s;
```

Largest field is 4 bytes (int), therefore:

- struct size will be multiple of 4 **B**
- struct starting address is word aligned, since a word is 4 bytes **A**

Assume struct starts at address 1000, what is the data layout of the struct?

Datatype	size (bytes)
short	2
char	1
int	4
double	8

Structure Alignment - Example

Problem: What boundary should this struct be aligned to?
What is the total size of the struct?

C-code
<pre>struct { char w; int x[3]; char y; short z; } s;</pre>

But we can try to change #3
the order of parameters
to save the memory.

Largest field is 4 bytes (int), therefore:

- struct size will be multiple of 4
- struct starting address is word aligned, since a word is 4 bytes

Assume struct starts at address 1000, what is the data layout of the struct?

char w -> 1000
 // padding 1001-1003 unused storage
 x[0] -> 1004-1007 since $1004 \bmod 4 = 0$
 x[1] -> 1008-1011
 x[2] -> 1012-1015
 char y -> 1016
 // padding 1017
 short z -> 1018-1019 since $1018 \bmod 2 = 0$

start: 1000

end: 1019

Total size = 20 bytes

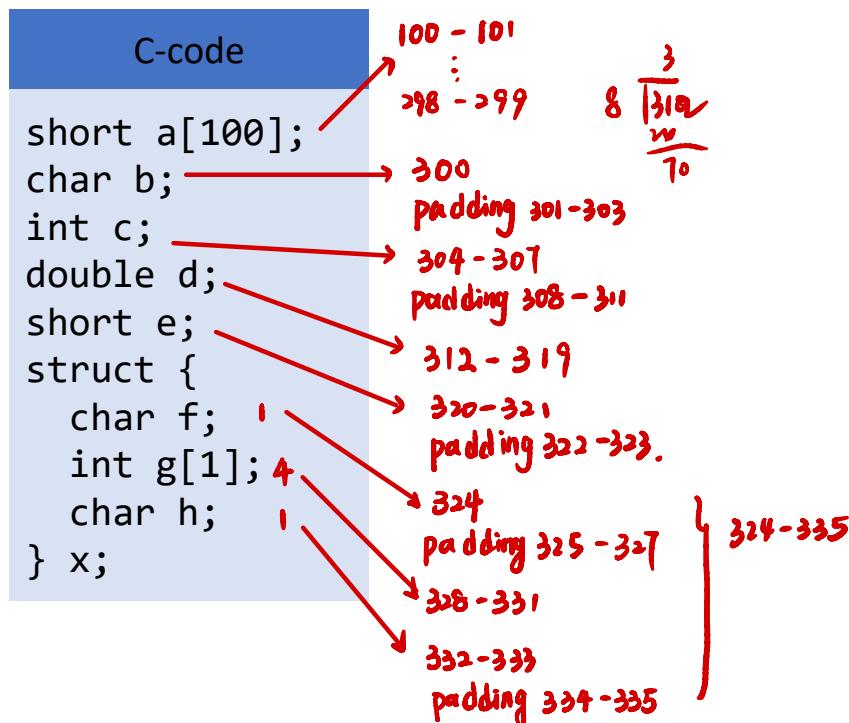
Datatype	size (bytes)
short	2
char	1
int	4
double	8

Why padding?
 “Golden” rule – Address of a variable is aligned based on the size of the variable

Calculating Load/Store Addresses – 2nd Try

Problem: Calculate the total amount of memory needed for the declarations.

Assume data memory starts at address 100



Datatype	size (bytes)
short	2
char	1
int	4
double	8

An N-byte variable must start at an address A, such that
 $(A \% N) == 0$

Calculating Load/Store Addresses – 2nd Try

Data
Layout

Problem: Calculate the total amount of memory needed for the declarations.

Assume data memory starts at address 100

C-code
short a[100]; char b; int c; double d; short e; struct { char f; int g[1]; char h; } x;

Datatype	size (bytes)
short	2
char	1
int	4
double	8

An N-byte variable must start at an address A, such that $(A \% N) == 0$

Calculating Load/Store Addresses – 2nd Try

Data Layout

Problem: Calculate the total amount of memory needed for the declarations.

Assume data memory starts at address 100

C-code	C-code	Bytes	Start	End	Notes
short a[100];	short a[100];	200	100	299	
char b;	char b;	1	300	300	
		3	301	303	padding
int c;	int c;	4	304	307	
		4	308	311	padding
double d;	double d;	8	312	319	
short e;	short e;	2	320	321	
struct {		2	322	323	padding
char f;	struct {	12	324	335	largest field: 4 bytes
int g[1];	char f;	1	324	324	
char h;		3	325	327	padding
}	int g[1];	4	328	331	
x;	char h;	1	332	332	
		3	333	335	padding
	}	12	324	335	
	x;				

Datatype	size (bytes)
short	2
char	1
int	4
double	8

An N-byte variable must start at an address A, such that
 $(A \% N) == 0$

Total size: 236 bytes

Pause

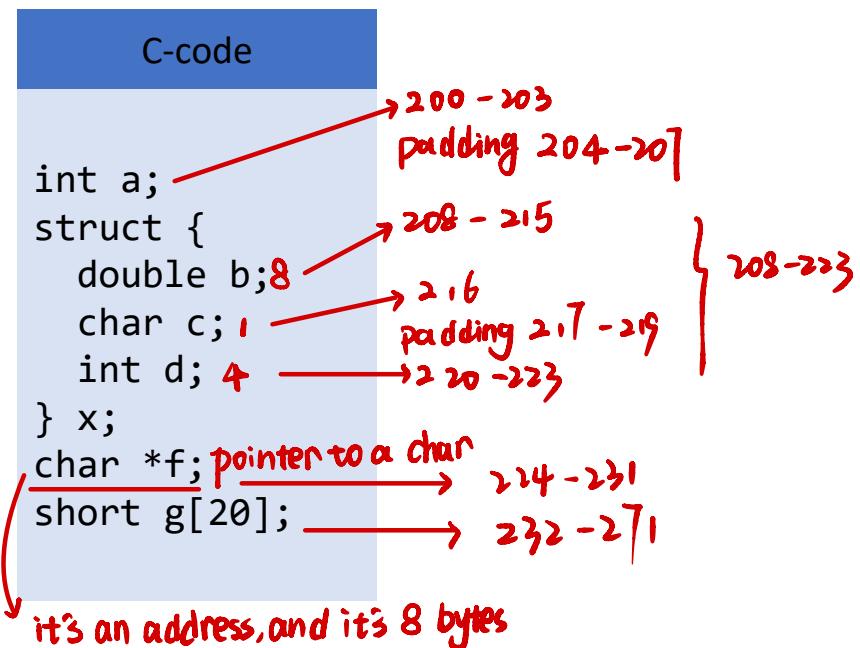
The next example is a review of Lecture 5 worksheet 1.

Pause, complete the worksheet, then proceed.

Example 2

Problem: Calculate the total amount of memory needed for the declarations.

Assume data memory starts at address 200



Datatype	size (bytes)
short	2
char	1
int	4
double	8
address	8

Example 2

Problem: Calculate the total amount of memory needed for the declarations.

Assume data memory starts at address **200**

C-code	C-code	Bytes	Start	End	Notes
int a;	int a;	4	200	203	
struct {		4	204	207	padding
double b;	struct {	16	208	223	largest field: 8 bytes
char c;	double b;	8	208	215	
int d;	char c;	1	216	216	
} x;		3	217	219	padding
char *f;	int d;	4	220	223	
short g[20];	} x;	16	208	223	
	char *f;	8	224	231	
	short g[20];	40	232	271	
	TOTAL:	72	200	271	

Datatype	size (bytes)
short	2
char	1
int	4
double	8
address	8

Data Layout – Why?

- Does gcc (or another compiler) reorder variables in memory to avoid padding?
- No, a compiler will not optimize data layout to remove padding.
- C99 standard prohibits this
 - Memory is laid out in order of declaration for structs.
- gcc has implemented an option for this, then later removed it
- The programmer (i.e., you) are expected to manage data layout of variables for your program and structs.
- For a start: order fields in struct by datatype size, smallest first

L5_2 Assembly – Control Flow

EECS 370 – Introduction to Computer Organization – Winter 2022

Learning Objectives

- Recognize the set of branching instructions for ARM ISA and be able to describe the operations and operands for instructions
 - LEGv8 subset
- Understand mapping of complex C-code branching instructions into corresponding assembly code instructions

ARM/LEGv8 Sequencing Instructions

ARM
Branching

- **Sequencing instructions** change the flow of instructions that are executed
 - This is achieved by modifying the program counter (PC)
- ① • **Unconditional branches** are the most straightforward they ALWAYS change the PC and thus “jump” to another instruction out of the usual sequence *Eg: End of the “for loop.”*
- ② • **Conditional branches** *since every instruction is 4 bytes, So it go to target-address*
 - if (*condition_test*) goto *target_address* *or go to PC + 4.*
 - *condition_test* examines the four flags from the processor status word (SPSR)
 - *target_address* is the 19-bit signed word displacement from current PC

LEGv8 Conditional Instructions

- Two varieties of conditional branches
 1. One type compares a register to see if it is equal to zero.
 2. Another type checks the condition codes set in the status register.

3 conditional
branches

Conditional branch	① compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100	Equal 0 test; PC-relative branch
	② compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	③ branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

compare instruction & branch (conditional)

- Let us look at the first type: CBZ and CBNZ
 - CBZ – Conditional Branch if Zero
 - CBNZ – Conditional Branch if Not Zero

pack compare & branch into same
instruction
The reason is: They only compare with zero

"or gate" → compare with 0.

LEGv8 Compare and Branch Instructions

ARM
Branching

- CBZ/CBNZ: test a register against zero and branch to a PC relative address
 - The relative address is a 19-bit signed integer—the number of instructions. Recall instructions are 32 bits of 4 bytes
- But we will also learn how to jump further.*

ARM LEGv8	Description
CBNZ X3, foo	<ul style="list-style-type: none">• if X3 does not equal 0, then branch to label foo
CBNZ X3, <u>25</u>	<ul style="list-style-type: none">• 25 is an offset from the PC of the current instruction (CBNZ)

jump ahead 25 instructions

*PC + displacement * 4 = target.*

LEGv8 Compare and Branch Instructions

- CBZ/CBNZ: test a register against zero and branch to a PC relative address
 - The relative address is a 19-bit signed integer—the number of instructions. Recall instructions are 32 bits of 4 bytes

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if (X1 == 0) go to PC + 100 (25x4)	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if (X1 != 0) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

Why does 25 in the table result in PC + 100?

Offset is # of instructions (words)

Conditional Branch Offset Example

Problem: Calculate the numerical offset for the CBNZ instruction.

ARM LEGv8

```
loop:    ADDI X3, X3, #1
          SUBI X4, X4, #1
          ADD  X5, X3, + X4
          CBNZ X5, loop
          ↓ is 0      -3   pc + displacement * 4 = target.
                           -12
```

$\text{is } 0$

-3

$\text{pc} + \text{displacement} * 4 = \text{target.}$

-12

$3 : 0011$

$-3 : 1100+1 = 1101$

Conditional Branch Offset Example

Problem: Calculate the numerical offset for the CBNZ instruction.

ARM LEGv8

```
loop:    ADDI X3, X3, #1
          SUBI X4, X4, #1
          ADD  X5, X3, X4
          CBNZ X5, loop
```

Conditional Branch Offset Example

Problem: Calculate the numerical offset for the CBNZ instruction.

ARM LEGv8

```
loop:    ADDI X3, X3, #1
          SUBI X4, X4, #1
          ADD  X5, X3, X4
          CBNZ X5, loop
```

Answer: -3

Offset field: 19-bit, signed

111 1111 1111 1111 1101

The assembler will calculate the offset

If any instructions are added or removed while writing code, using a label saves from recalculating the offset

Conditional Branch Offset Example

How is the branch target address calculated?

ARM LEGv8

CBNZ X5, #-3

Conditional Branch Offset Example

How is the branch target address calculated?

ARM LEGv8

CBNZ X5, #-3

Conditional Branch Offset Example

How is the branch target address calculated?

ARM LEGv8

CBNZ X5, #-3

1. Offset field: 19 bits (-3 decimal)

111 1111 1111 1111 1101 -3

1. Append two zeros

1 1111 1111 1111 1111 0100 -12

1. Sign extend to 64 bits

since the PC is 64-bit

1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 1111 0100

1. Add offset to PC of CBNZ instruction

LEGv8 Compare and Branch Instructions

- B.cond: lets you branch based on the flags set by CMP (ADDS, etc.)

Conditional branch	compare and branch on equal 0	CBZ X1, 25	if ($X1 == 0$) go to PC + 100	Equal 0 test; PC-relative branch
	compare and branch on not equal 0	CBNZ X1, 25	if ($X1 != 0$) go to PC + 100	Not equal 0 test; PC-relative branch
	branch conditionally	B.cond 25	if (condition true) go to PC + 100	Test condition codes; if true, branch

LEGv8 Conditional Instructions Using FLAGS

ARM
Branching

- FLAGS: **NZVC** record the results of (arithmetic) operations
Negative, Zero, oVerflow, Carry—not present in LC-2K
- We explicitly set them using the “Set” modification to ADD/SUB etc.

Category	Instruction	Example	Meaning	Comments
Arithmetic	add	ADD X1, X2, X3	X1 = X2 + X3	Three register operands
	subtract	SUB X1, X2, X3	X1 = X2 - X3	Three register operands
	add immediate	ADDI X1, X2, 20	X1 = X2 + 20	Used to add constants
	subtract immediate	SUBI X1, X2, 20	X1 = X2 - 20	Used to subtract constants
	add and set flags	ADDS X1, X2, X3	X1 = X2 + X3	Add, set condition codes
	subtract and set flags	SUBS X1, X2, X3	X1 = X2 - X3	Subtract, set condition codes
	add immediate and set flags	ADDIS X1, X2, 20	X1 = X2 + 20	Add constant, set condition codes
	subtract immediate and set flags	SUBIS X1, X2, 20	X1 = X2 - 20	Subtract constant, set condition codes

CMP X₁, X₂ → X₁-X₂ and set

ARM LEGv8	Description
ADDS X1, X2, X3	Causes the 4 flag bits to be set accordingly as the outcome is negative, zero, overflows, or generates a carry bit

LEGv8 Condition Codes

- In LEGv8 only ADDS / SUBS / ADDIS / SUBIS / CMP / CMPI set the condition codes FLAGS or condition codes in PSR—the program status register
- Four primary condition codes evaluated:
 - N – set if the result is **negative** (i.e., bit 63 is non-zero)
 - Z – set if the result is **zero** (i.e., all 64 bits are zero)
 - C – set if last addition/subtraction had a **carry**/borrow out of bit 63
 - V – set if the last addition/subtraction produced an **overflow** (e.g., two negative numbers added together produce a positive result)
- Do not worry about the C and V bits *per se*. They are important but are tricky to understand.
 - Instead we will just be using branches based on these results for signed numbers which is a lot easier to deal with.

Conditional Branches

- CMP instruction lets you compare two registers, set NZVC flags
 - Could also use ADDS etc.
- B.*condition* lets you branch based on the flags set by CMP (ADDS, etc.)

when CMP, it will compare every test : \leq \neq $>$... , and the branch will tell the CMP, which test do we care about

ARM LEGv8	Description	care about
CMP X1, $>$ X2	Branches to label1 if value in register X1 greater than value in register X2	
B.GT label1		

if $X1 \leq X2$, then we go to Next instruction

- What is the set of conditions for B.*condition* ?

LEGv8 Conditional Instructions

All These test will be done once you do cmp instruction.

Encoding	Name (& alias)	Meaning (integer)	Flags
0000	EQ	Equal $=$	$Z==1$
0001	NE	Not equal \neq	$Z==0$
0010	HS (CS)	Unsigned higher or same (Carry set)	$C==1$
0011	LO (CC)	Unsigned lower (Carry clear)	$C==0$
0100	MI	Minus (negative)	$N==1$
0101	PL	Plus (positive or zero)	$N==0$
0110	VS	Overflow set	$V==1$
0111	VC	Overflow clear	$V==0$
1000	HI	Unsigned higher	$C==1 \& Z==0$
1001	LS	Unsigned lower or same	$! (C==1 \& Z==0)$
1010	GE	Signed greater than or equal	$N==V$
1011	LT	Signed less than	$N \neq V$
1100	GT	Signed greater than	$Z==0 \& N==V$
1101	LE	Signed less than or equal	$! (Z==0 \& N==V)$
1110	AL		
1111	NV ¹	Always	Any

After you do CMP, the flag will hold the value, and before you do the branch you can't use flags.

ARM
Branching

ARM LEGv8

CMP X1, > X2

B.GT label1

Branching if x_1 is greater than x_2

CMP X3, = X4

B.EQ label2

CMP X5, <= X6

B.LE label3

You need to know the
7 with red arrows

Branching Far Away

- The underlying philosophy of ISA design and microarchitecture in general is to **make the common case fast**
- In the case of branches, you are commonly going to branch to other instructions nearby.
 - In ARMv8, the encoding for the displacement of conditional branches is 19 bits.
 - Having a displacement of 19 bits is usually enough
- BUT what if we need jump to a target (Label) that we cannot get to with a 19-bit displacement from the current PC?
CBZ X15, FarLabel
*So we will change CBZ to CBNZ.
since conditional Branch can only have 19-bit offset*
- The assembler is smart enough to replace that with
- The simple branch instruction (B) has a 26-bit offset which spans about 64 million instructions!

```
CBNZ X15, L1
B    FarLabel
L1: unconditional
```

Unconditional Branching Instructions

Unconditional branch	branch (26-bit) if not big enough → branch to register that a register value and through to the PC (any address)	B 2500	go to PC + 10000	Branch to target address; PC-relative
	branch to register that a register value and through to the PC (any address)	BR X30	go to X30	For switch, procedure return
	branch with link	BL 2500	X30 = PC + 4; PC + 10000	For procedure call PC-relative

- There are three types of unconditional branches in the LEGv8 ISA.
 - The first (**B**) is the PC relative branch with the 26-bit offset from the last slide.
 - The second (**BR**) jumps to the address contained in a register (X30 above)
 - The third (**BL**) is like our PC relative branch but it does something else.
 - It sets X30 (always) to be the current PC+4 before it branches.
 - Why?
 - Function calls – return to next instruction

L5_3 C-to-Assembly Examples

EECS 370 – Introduction to Computer Organization – Winter 2022

Learning Objectives

- Translate C-code *statements* to ARM assembly code
 - Break down complex C-code branching instructions into a series of assembly operations
 - Map conditions in C to comparison and branch instructions in assembly

Branching - Example

Problem: Convert the C code to LEGv8 assembly: 1: using labels, 2. without labels

Register to variable mapping	C-code instructions
X1→x X2→y	int x, y; // update x,y if (x == y) T : x++; else F : y++;
	<i>mutually exclusive. You can't run two code together.</i>

Branching - Example

Problem: Convert the C code to LEGv8 assembly: 1: using labels, 2. without labels

Register to
variable mapping

X1→x
X2→y

C-code
instructions

```
int x, y;  
// update x,y  
if (x == y)  
    x++;  
else  
    y++;
```

Branching - Example

ARM
Branching

Problem: Convert the C code to LEGv8 assembly: 1: using labels, 2. without labels

Register to variable mapping
X1→x
X2→y

C-code instructions

```
int x, y;
// update x,y
if (x == y)
    x++;
else
    y++;
```

ARM LEGv8 – w/labels

```
CMP X1, X2
B.NE lbl1
ADDI X1, X1, #1
B lbl12
lbl1: ADDI X2, X2, #1
lbl12:
```

- this is Because :
- For the code:
if the condition is true,
then go to the next line,
if not, go to the target (else)
 - For the compare Branch.
if the condition is true,
then you go to the target
if not, go to the next line.

unconditional Branch at the end of true code
is because , they are mutually exclusive

Branching - Example

Problem: Convert the C code to LEGv8 assembly: 1: using labels, 2. without labels

Register to variable mapping	C-code instructions	ARM LEGv8 – w/o labels
X1→x X2→y	int x, y; // update x,y if (x == y) x++; else y++;	CMP X1, X2 B.NE ③ ADDI X1, X1, #1 B ② lbl1: ADDI X2, X2, #1 lbl2:

Loop - Example

Problem: Convert the C code to LEGv8 assembly (assume no registers initialized)

Register to variable mapping

X1→i
X2→sum
X4→#10
X5→a[i]
X6→i*8

a is array of long long integers (64 bits, 8 bytes)
Start of a at address 100,
sum starts at address 96

C-code instructions

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

Loop - Example

Problem: Convert the C code to LEGv8 assembly (assume no registers initialized)

Register to variable mapping

X1→i
X2→sum
X4→#10
X5→a[i]
X6→i*8

a is array of long long integers (64 bits, 8 bytes)
Start of a at address 100,
sum starts at address 96

C-code instructions

```
sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}
```

Loop - Example

Problem: Convert the C code to LEGv8 assembly (assume no registers initialized)

Register to variable mapping
X1→i
X2→sum
X4→#10
X5→a[i]
X6→i*8
a is array of long long integers (64 bits, 8 bytes)
Start of a at address 100,
sum starts at address 96

C-code instructions

```

sum = 0;
for (i=0 ; i < 10 ; i++) {
    if (a[i] >= 0) {
        sum += a[i];
    }
}

```

ARM LEGv8

```

MOV X1, XZR
MOV X2, XZR
MOVZ X4, #10
Loop1: CMP X1, X4
        B.GE >=
LSL X6, X1, #3
LDUR X5, [X6, #100]
CMPI X5, #0
        B.LT <
endif: ADD X2, X2, X5
       STURW X2, [XZR, #96]
endif: ADDI X1, X1, #1
       B Loop1
endLoop:

```

loop test

Loop - Example

ARM
Branching

Alternate Solution: do-while

Register to variable mapping

X1→i
X2→sum
X4→#10
X5→a[i]
X6→i*8

a is array of long
long integers (64
bits, 8 bytes)
Start of a at
address 100,
sum starts at
address 96

C-code instructions

```
sum = 0;  
for (i=0 ; i < 10 ; i++) {  
    if (a[i] >= 0) {  
        sum += a[i];  
    }  
}
```

ARM LEGv8

```
MOV    X1, XZR  
MOV    X2, XZR  
MOVZ   X4, #10  
Loop1: LSL   X6, X1, #3  
        LDUR  X5, [X6, #100]  
        CMPI  X5, #0  
        B.LT  endif  
        ADD   X2, X2, X5  
        STUR  X2, [XZR, #96]  
endIf: ADDI  X1, X1, #1  
        CMP   X1, X4  
        B.LT  Loop1 ← with a conditional  
endLoop: test at the end of the loop. Branch.
```