

Descriptors

EECS 442 – Jeong Joon Park
Winter 2024, University of Michigan

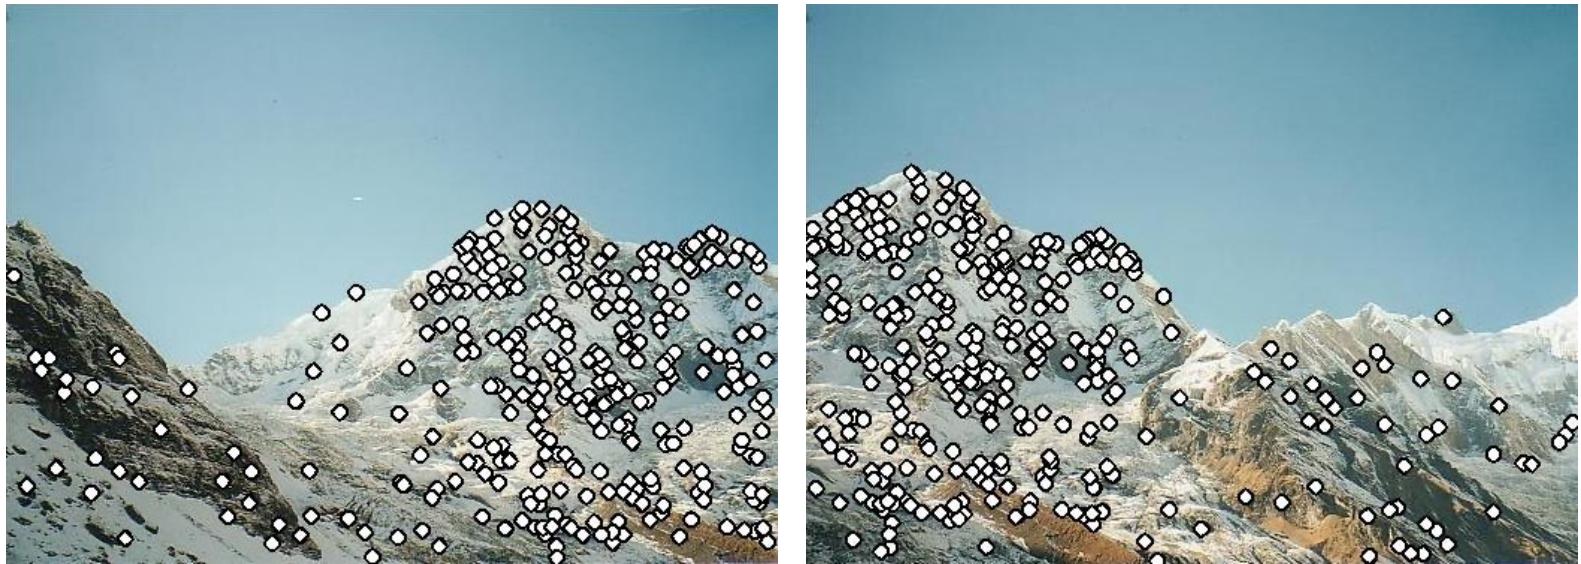
Administrivia

- Questions to the Instructor (me)
 - Best used by asking conceptual questions
 - Detailed HW-related questions: TAs. Why?
 - TAs make, edit HWs, prepare solutions, etc.

Administrivia

- Project
- We will post details on Wednesday
- Will provide sample topics

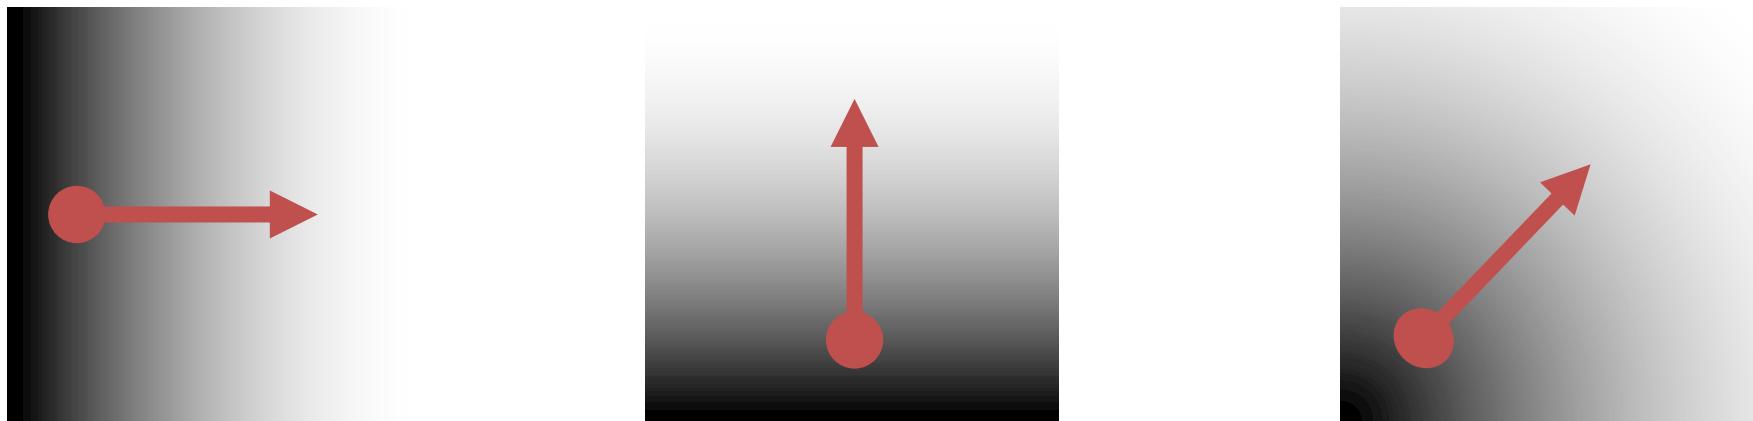
Recap: Motivation



1: find corners+features

Last Time – Gradients

Image gradients – treat image like function of x, y – gives edges, corners, etc.



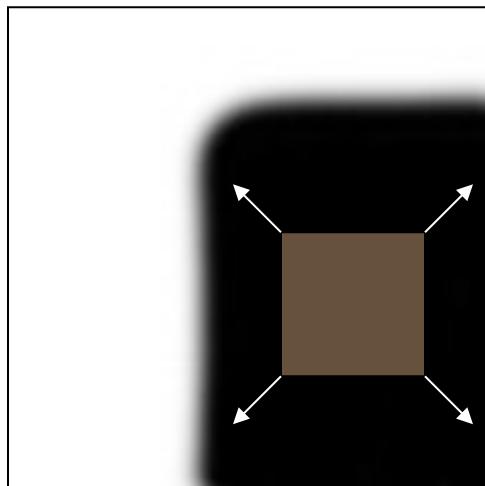
$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$

$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

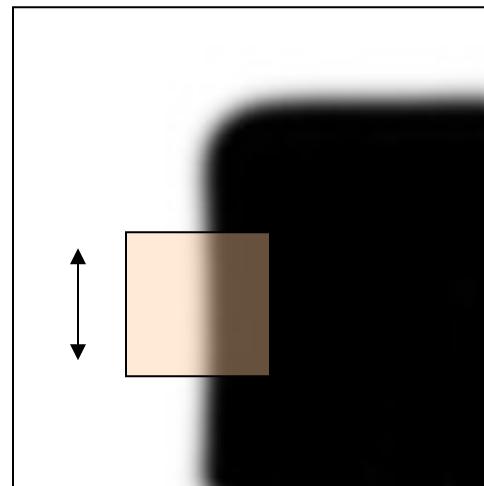
$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Last Time – Corner Detection

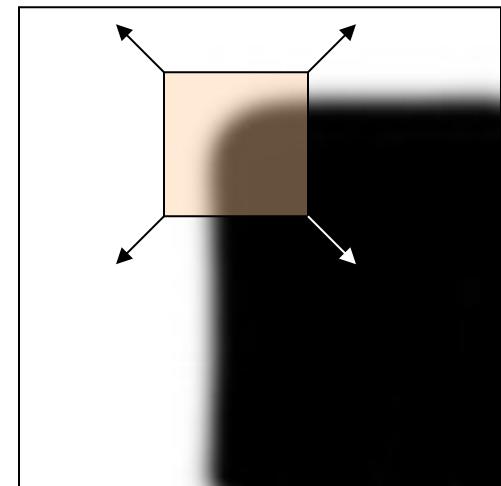
Can localize the location, or any shift → big intensity change.



“flat” region:
no change in
all directions



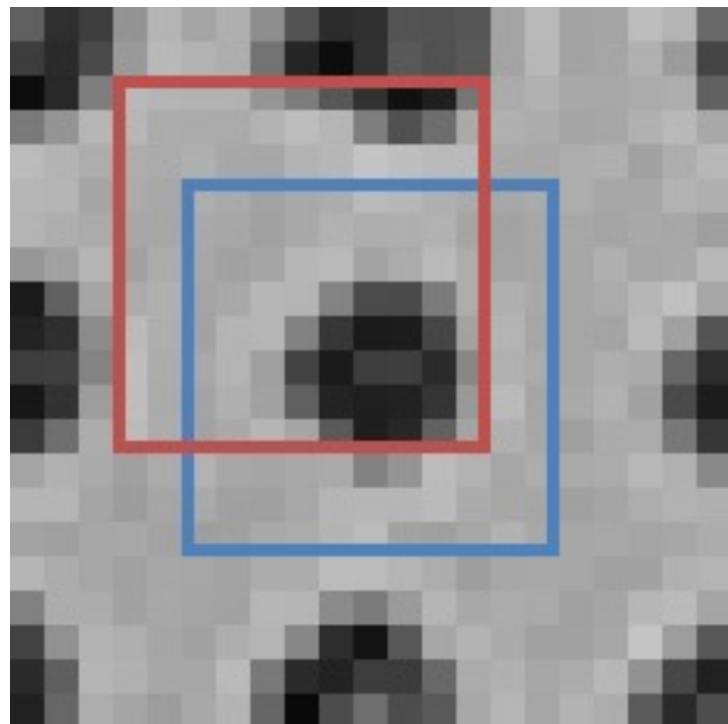
“edge”:
no change
along the edge
direction



“corner”:
significant
change in all
directions

Last Time – Corner Detection

Zoom-In at x, y



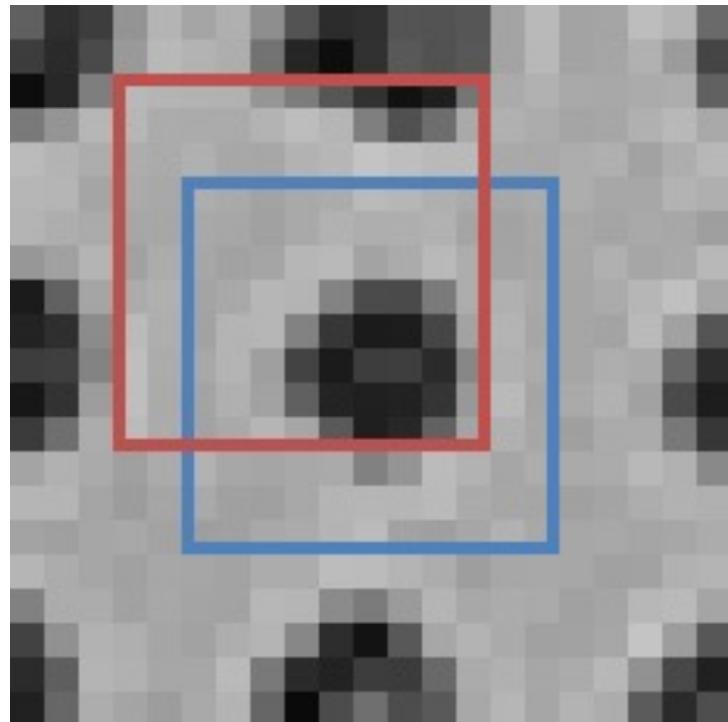
Window with and w/o Offset

“Window”
At $x+u, y+v$
Here: $u=-2, v=-3$

“Window”
At x, y

Last Time – Corner Detection

Zoom-In at x,y



Error (Sum Sq) for u,v offset

$$E(u, v) = \sum_{(x,y) \in W} (I[x + u, y + v] - I[x, y])^2$$

$$\left(\text{[4x4 window with red border]} - \text{[4x4 window with blue border]} \right)^2$$

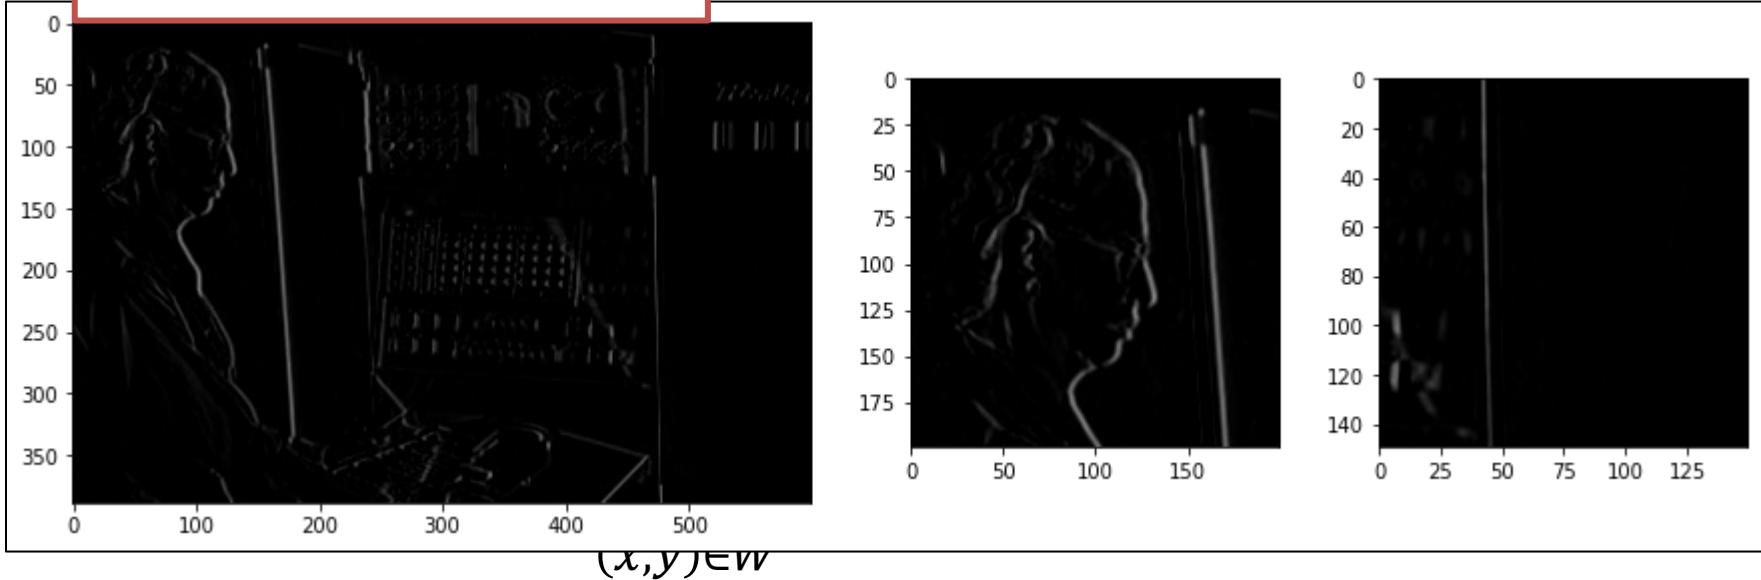
Formalizing Corner Detection

By linearizing image, we can approximate $E(u,v)$ with quadratic function of u and v

$$\begin{aligned} E(u, v) &\approx \sum_{(x,y) \in W} (I_x^2 u^2 + 2I_x I_y u v + I_y^2 v^2) \\ &= [u, v] \mathbf{M} [u, v]^T \end{aligned}$$

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix}$$

$I_x = x$ derivative

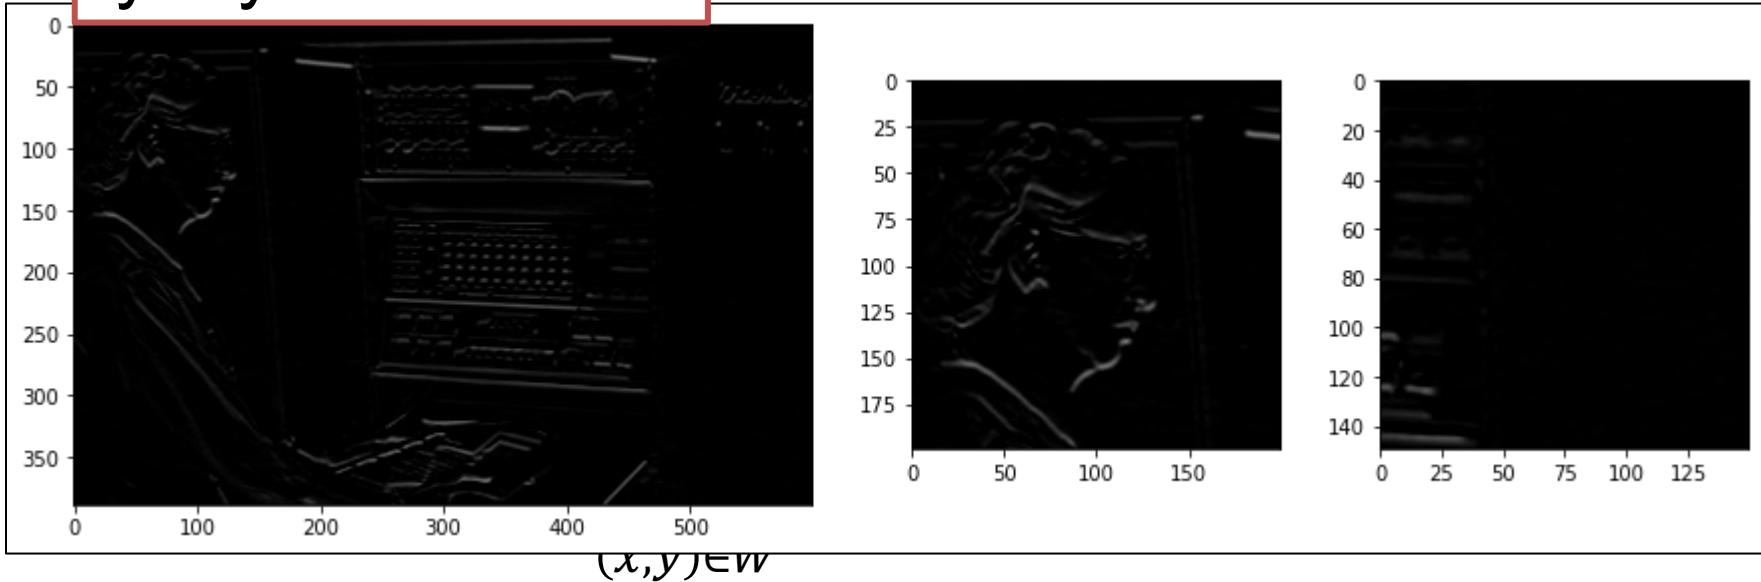


$(x, y) \in W$

$$= [u, v] M [u, v]^T$$

$$M = \begin{bmatrix} \sum_{x, y \in W} I_x^2 & \sum_{x, y \in W} I_x I_y \\ \sum_{x, y \in W} I_x I_y & \sum_{x, y \in W} I_y^2 \end{bmatrix}$$

$I_y = y$ derivative



$(x, y) \in W$

$$= [u, v] \mathbf{M} [u, v]^T$$

$$\mathbf{M} = \begin{bmatrix} \sum_{x, y \in W} I_x^2 & \sum_{x, y \in W} I_x I_y \\ \sum_{x, y \in W} I_x I_y & \sum_{x, y \in W} I_y^2 \end{bmatrix}$$

Sum goes over all the pixels in window W:

$$\sum_{x,y \in W} I_x^2 = \sum_{x,y \in W} (I_x[y, x])^2$$

i.e., sum of squares of x gradients in window

$$= [u, v] \mathbf{M} [u, v]^T$$

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix}$$

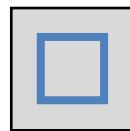
Intuitively what is M?

Pretend gradients are *either* vertical or horizontal

Obviously Wrong!

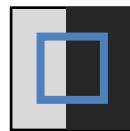
$$M = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} \approx \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$$

a,b both small: flat



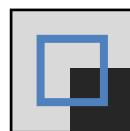
$$\begin{bmatrix} 0.1 & 0 \\ 0 & 0.1 \end{bmatrix}$$

One big,
other small:
edge



$$\begin{bmatrix} 50 & 0 \\ 0 & 0.1 \end{bmatrix} \text{ or } \begin{bmatrix} 0.1 & 0 \\ 0 & 50 \end{bmatrix}$$

a,b both big:
corner



$$\begin{bmatrix} 50 & 0 \\ 0 & 50 \end{bmatrix}$$

Intuitively what is M?

Pretend gradients are *either* vertical or horizontal at a pixel (so $I_x I_y = 0$)

$$M = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} \approx? \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$$

a,b both small:

flat

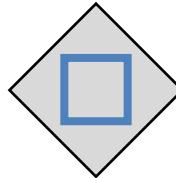
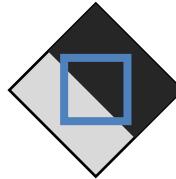


Image might be rotated by rotation θ !

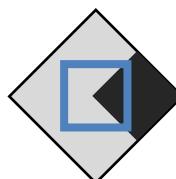
One big,
other small:

edge



a,b both big:

corner

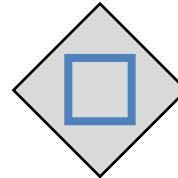


Intuitively what is M?

Pretend gradients are *either* vertical or horizontal at a pixel (so $I_x I_y = 0$)

$$M = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} = V^{-1} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} V$$

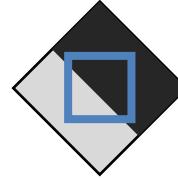
a,b both small: flat



If image rotated by rotation θ / matrix V

One big,
other small:

edge

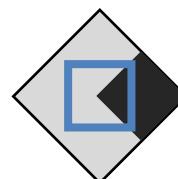


M will look like

$$V^{-1} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} V$$

a,b both big:

corner



So What Now?

Can calculate \mathbf{M} at pixel, by summing nearby gradients, but need access to a and b .

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} = \mathbf{V}^{-1} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} \mathbf{V}$$

Given \mathbf{M} , can decompose it into eigenvectors \mathbf{V} and eigenvalues λ_1, λ_2 with $\mathbf{M} = \mathbf{V}^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} \mathbf{V}$.

Really slow. Why?

So What Now?

Can calculate M at pixel, by summing nearby gradients, but need access to a and b .

$$M = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} = V^{-1} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} V$$

Instead: compute quantity R from M

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

Easy fast formula
for 2x2

Fast – sum the diagonal

Empirical value,
usually 0.04-0.06

The tl;dr

TL;DR: Taylor expansion for error $E(u,v)$. All terms in equation are sums of image gradients and in \mathbf{M}

Should know

Can compute at each pixel

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} =$$

$I_x = I_x$ at point (x,y) , $I_y = I_y$ at point (x,y)

Optional

Directions

Amounts

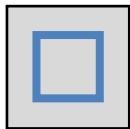
$$V^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} V$$

Putting It Together

$$R = \det(\mathbf{M}) - \alpha \operatorname{trace}(\mathbf{M})^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

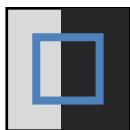
det, trace are *fast*

flat



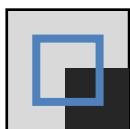
$$\lambda_1, \lambda_2 \approx 0$$

edge

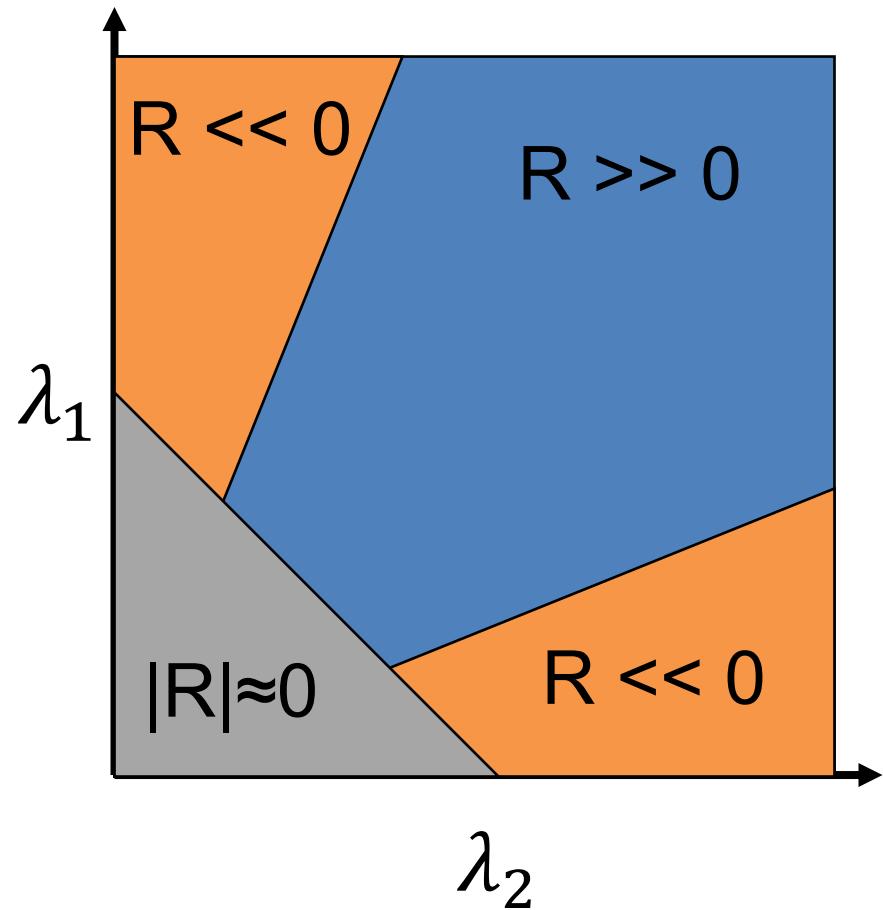


$$\begin{aligned}\lambda_1 &\gg \lambda_2 \gg 0 \\ \lambda_2 &\gg \lambda_1 \gg 0\end{aligned}$$

corner



$$\lambda_1 \approx \lambda_2 \gg 0$$



Remake of standard diagram from S. Lazebnik from original Harris paper.

In Practice

1. Compute partial derivatives I_x , I_y per pixel
2. Compute \mathbf{M} at each pixel, using Gaussian weighting w

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} w(x,y) I_x^2 & \sum_{x,y \in W} w(x,y) I_x I_y \\ \sum_{x,y \in W} w(x,y) I_x I_y & \sum_{x,y \in W} w(x,y) I_y^2 \end{bmatrix}$$

C.Harris and M.Stephens. [A Combined Corner and Edge Detector.](#)
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

In Practice

1. Compute partial derivatives I_x, I_y per pixel
2. Compute \mathbf{M} at each pixel, using Gaussian weighting w
3. Compute response function R

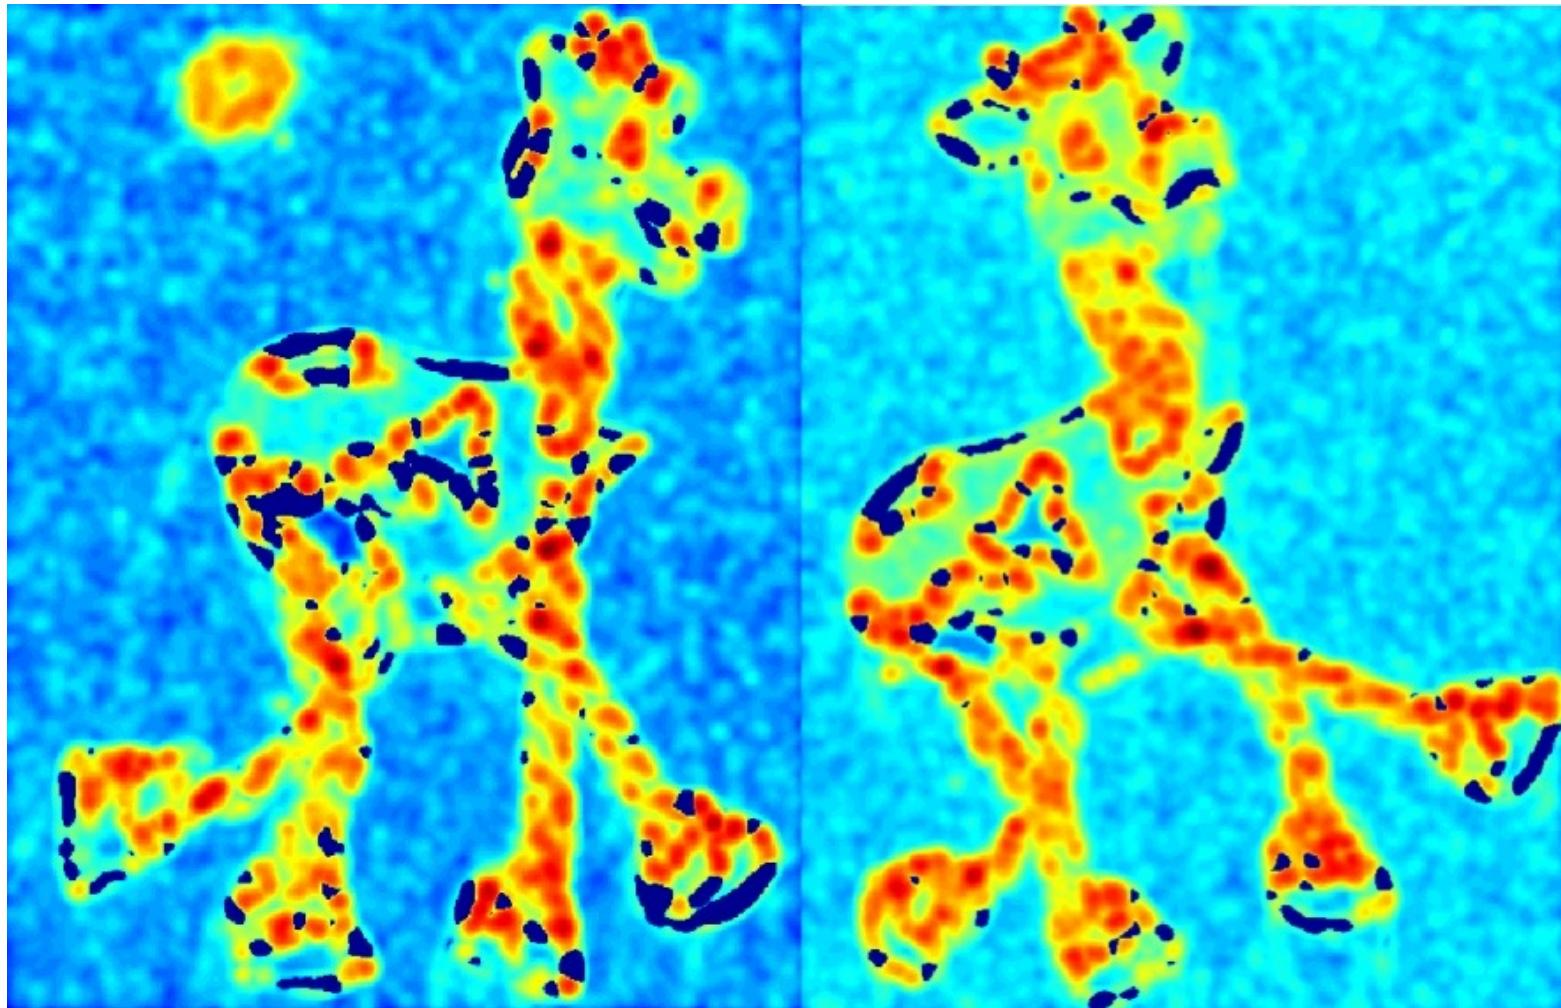
$$\begin{aligned} R &= \det(\mathbf{M}) - \alpha \operatorname{trace}(\mathbf{M})^2 \\ &= \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2 \end{aligned}$$

C.Harris and M.Stephens. ["A Combined Corner and Edge Detector."](#)
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Computing R



Computing R

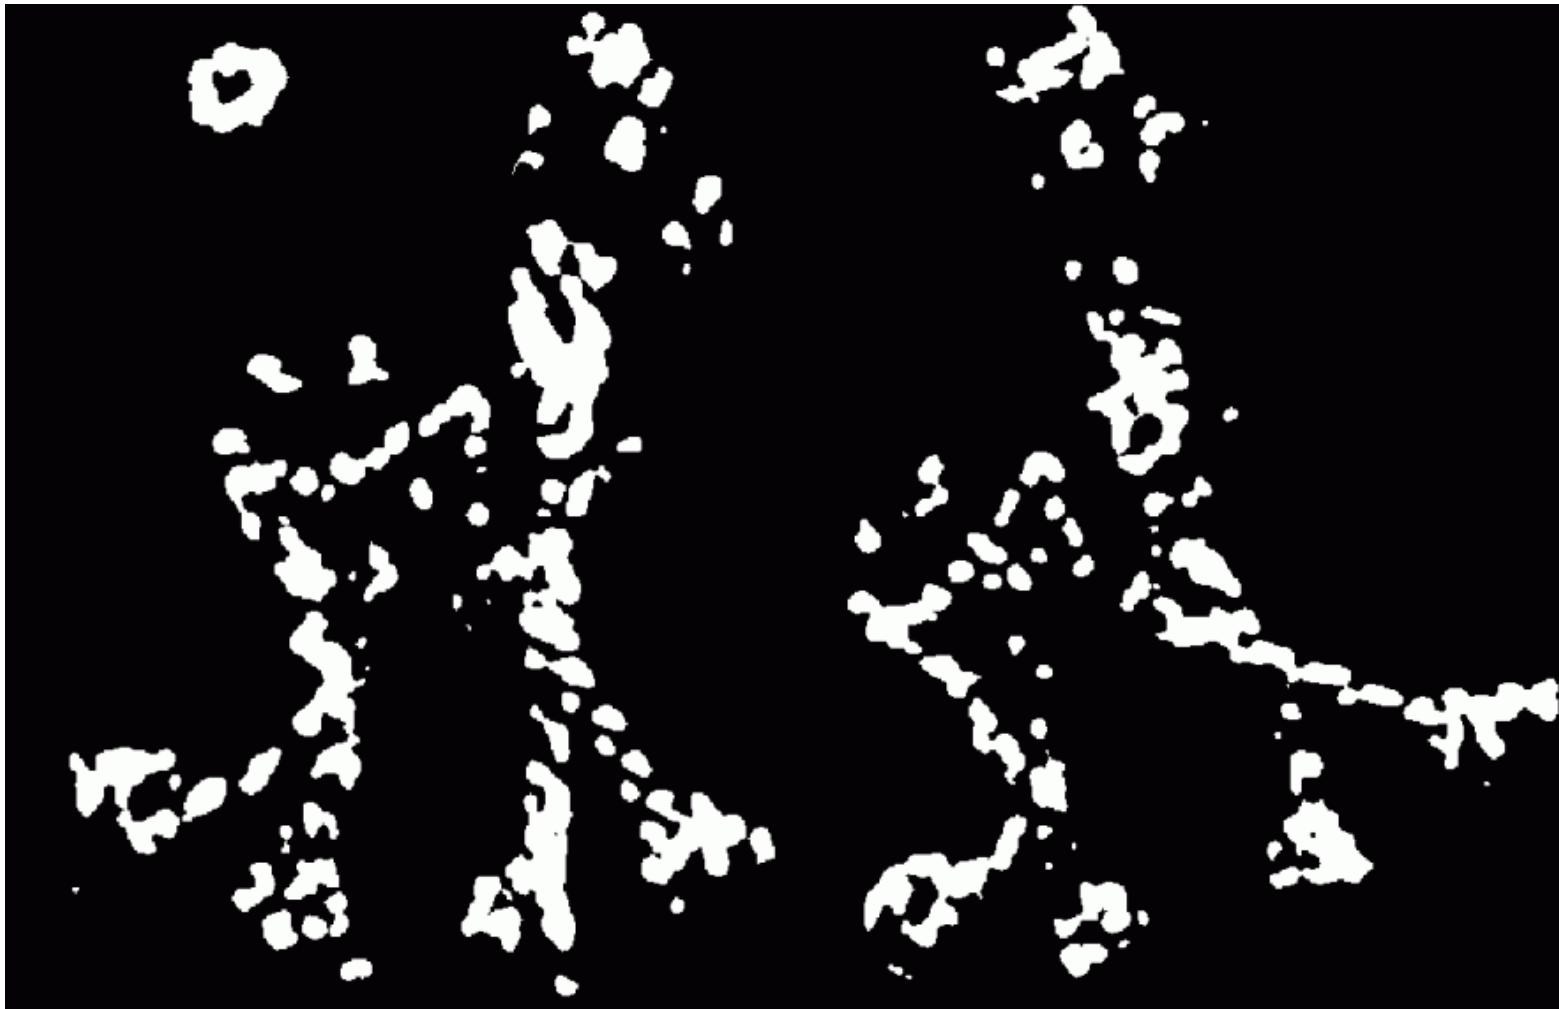


In Practice

1. Compute partial derivatives I_x , I_y per pixel
2. Compute \mathbf{M} at each pixel, using Gaussian weighting w
3. Compute response function R
4. Threshold R

C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)"
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Thresholded R

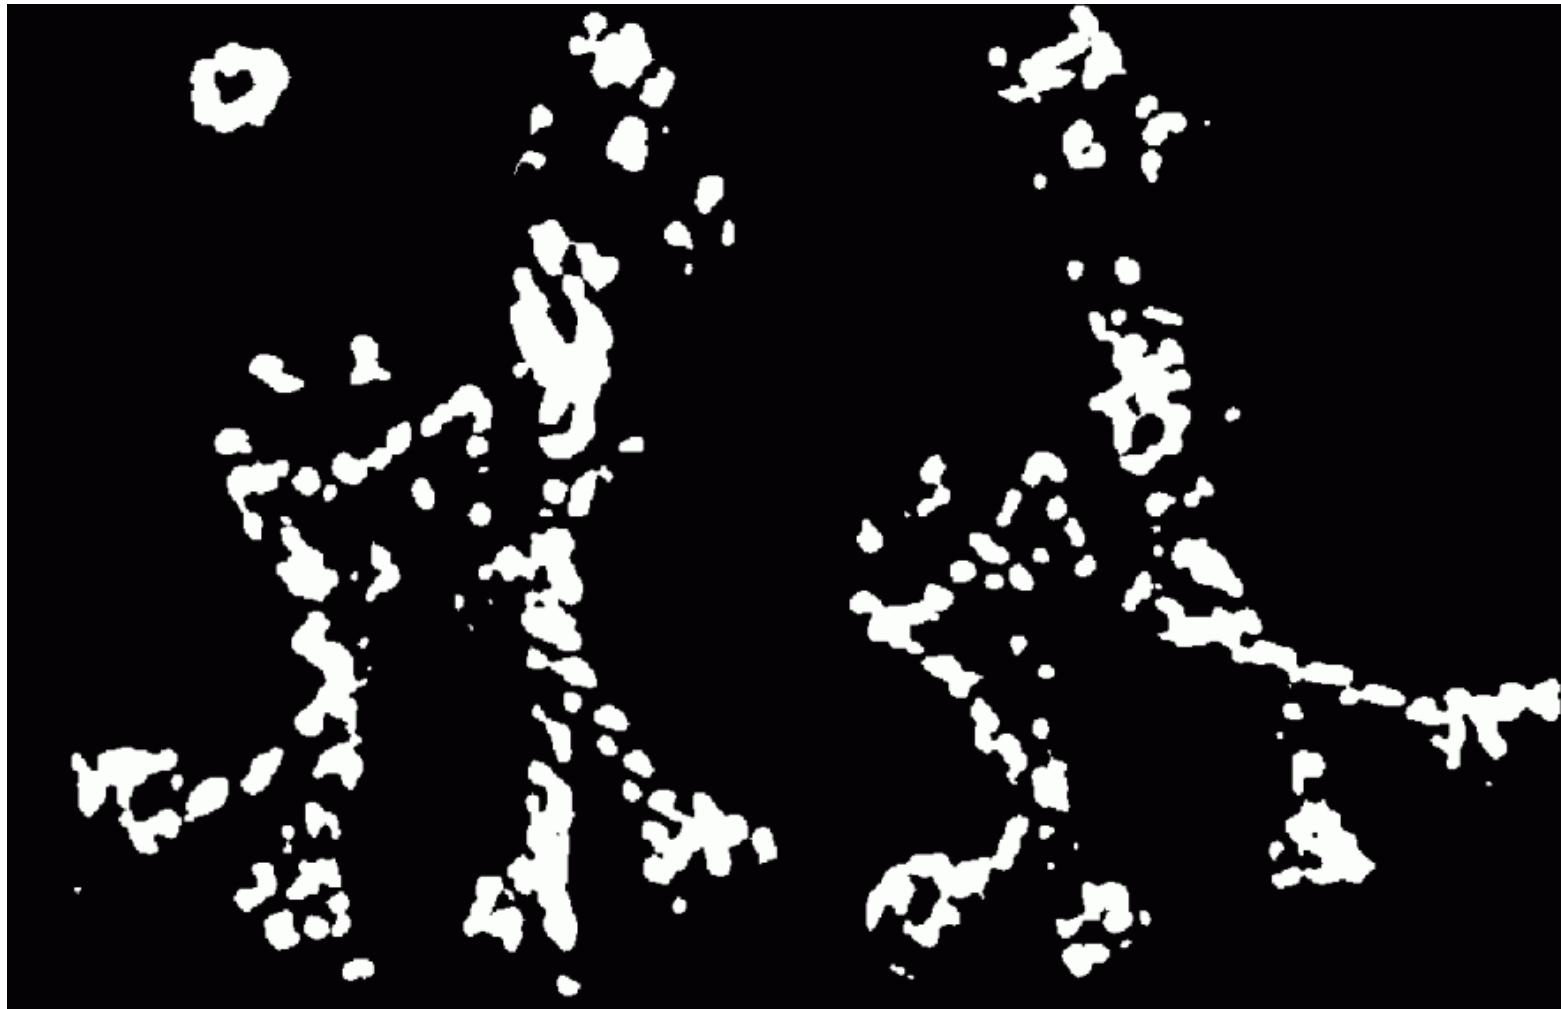


In Practice

1. Compute partial derivatives I_x , I_y per pixel
2. Compute \mathbf{M} at each pixel, using Gaussian weighting w
3. Compute response function R
4. Threshold R
5. Take only local maxima (called non-maxima suppression)

C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)"
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Thresholded



Final Results



Desirable Properties

If our detectors are repeatable, they should be:

- **Invariant** to some things: image is transformed and corners remain the same
- **Covariant/equivariant** with some things: image is transformed and corners transform with it.

Recall Motivating Problem

Images may be different in lighting and geometry

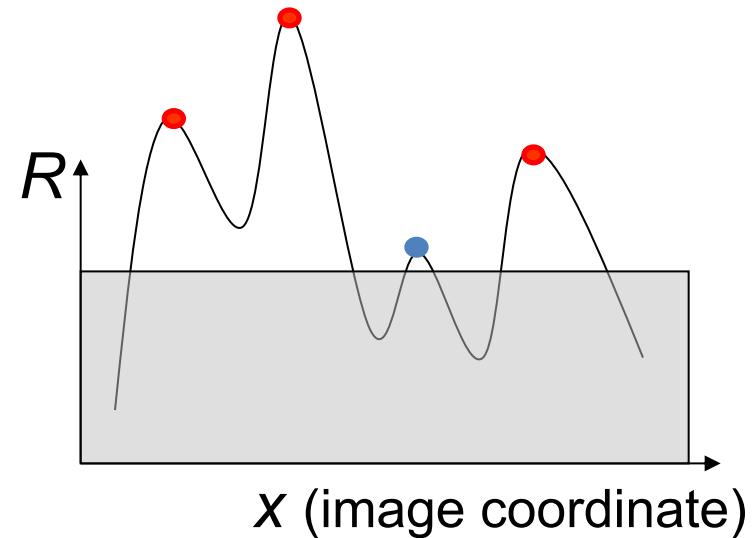
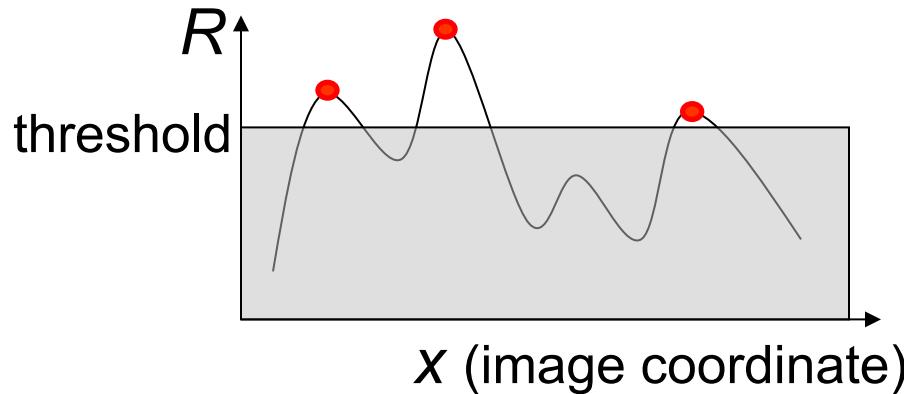


Affine Intensity Change

$$I_{new} = aI_{old} + b$$

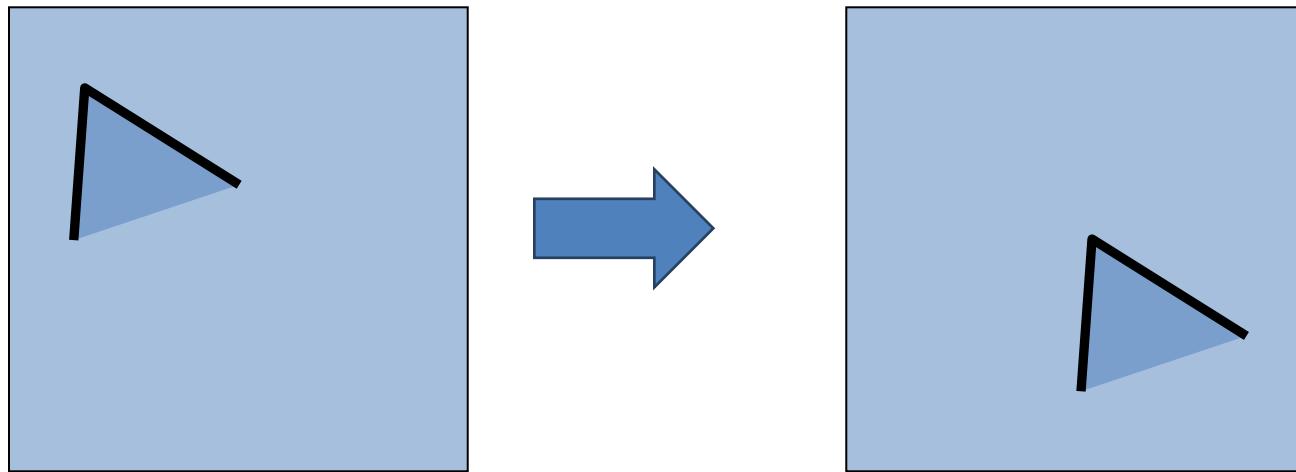
M only depends on derivatives, so b is irrelevant

But a scales derivatives and there's a threshold



Partially invariant to affine intensity changes

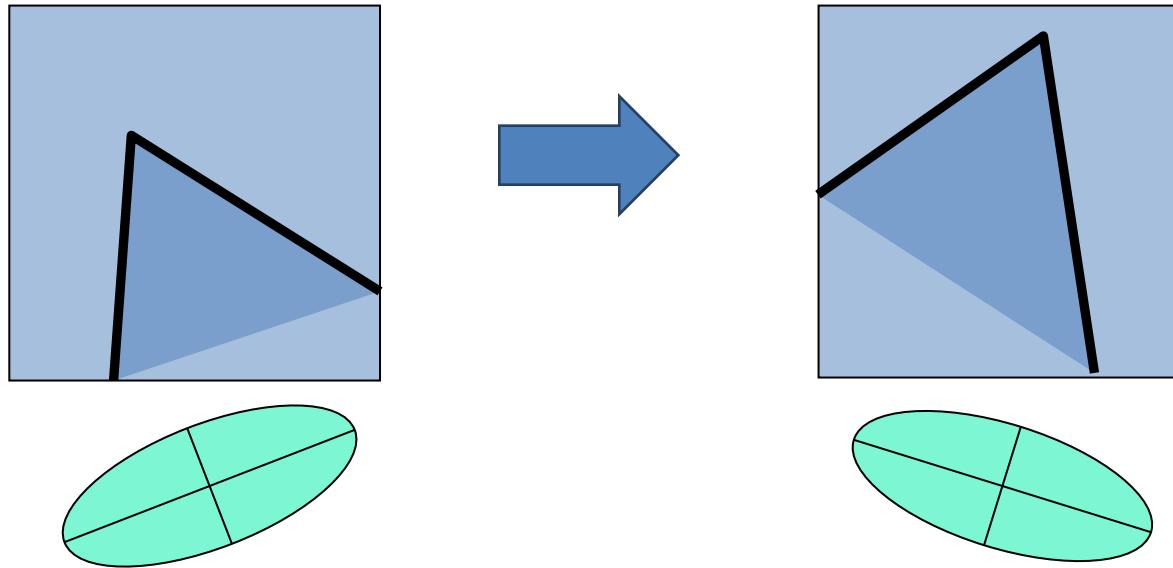
Image Translation



All done with convolution. Convolution is translation equivariant.

Equivariant with translation

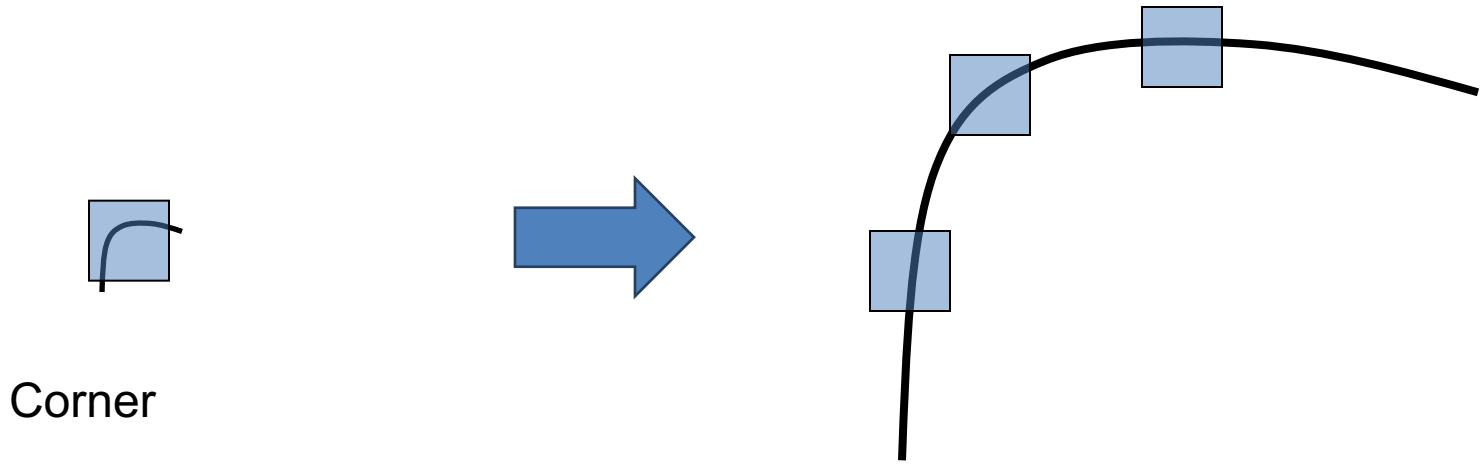
Image Rotation



Rotations just cause the corner rotation matrix to change. Eigenvalues remain the same.

Equivariant with rotation

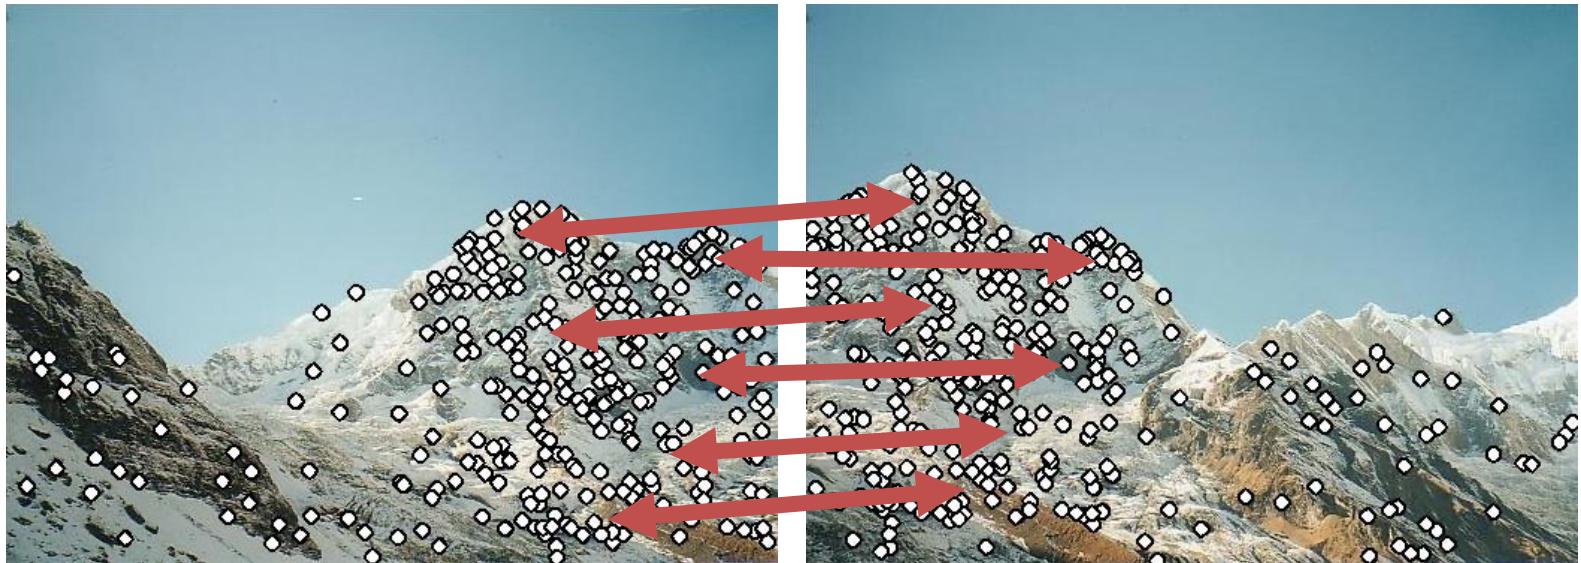
Image Scaling



Corner in one scale → Edges in another scale

Not equivariant with scaling
How do we fix this?

Recap: Motivation



- 1: find corners+features
- 2: match based on local image data

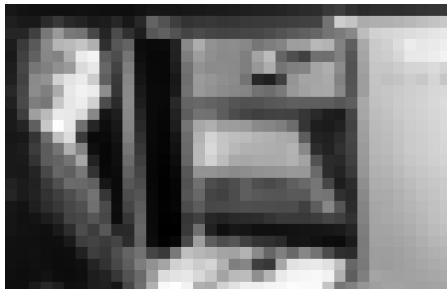
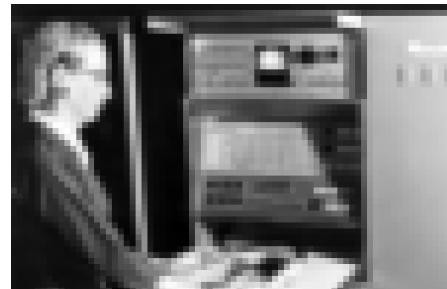
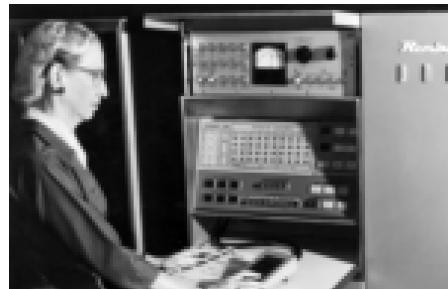
How?

Today

- Fixing scaling by making detectors in both location **and scale**
- Enabling matching between features by **describing regions**

Key Idea: Scale Space

Left to right: each image is half-sized
Upsampled with big pixels below

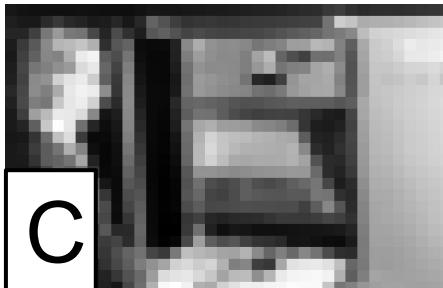
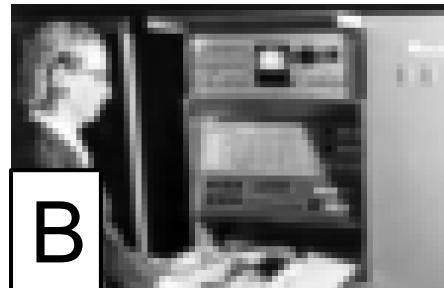


Note: I'm also slightly blurring to prevent aliasing (<https://en.wikipedia.org/wiki/Aliasing>)

Key Idea: Scale Space

Left to right: each image is half-sized

If I apply a $K \times K$ filter, how much of information of the original image does it see in each image?



Note: I'm also slightly blurring to prevent aliasing (<https://en.wikipedia.org/wiki/Aliasing>)

Solution to Scales

Try them all!

Harris Detection



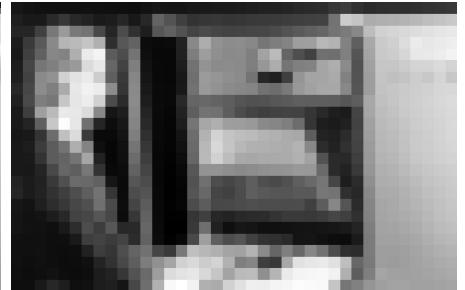
Harris Detection



Harris Detection



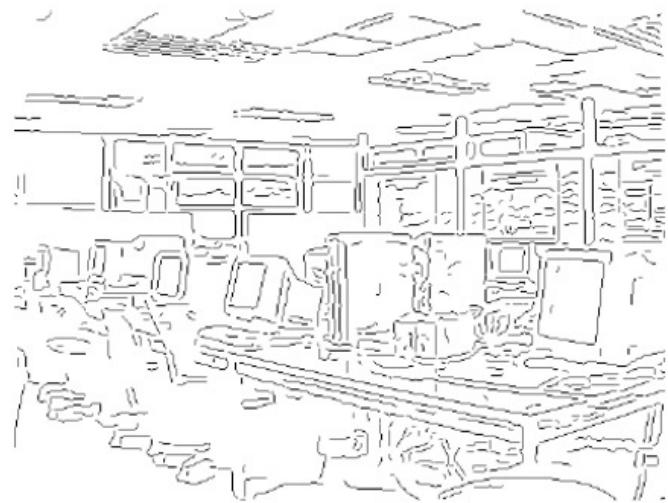
Harris Detection



See: Multi-Image Matching using Multi-Scale Oriented Patches, Brown et al. CVPR 2005

What about the Internal Structures?

Edges & Corners convey boundary information



What about interior texture of the object?

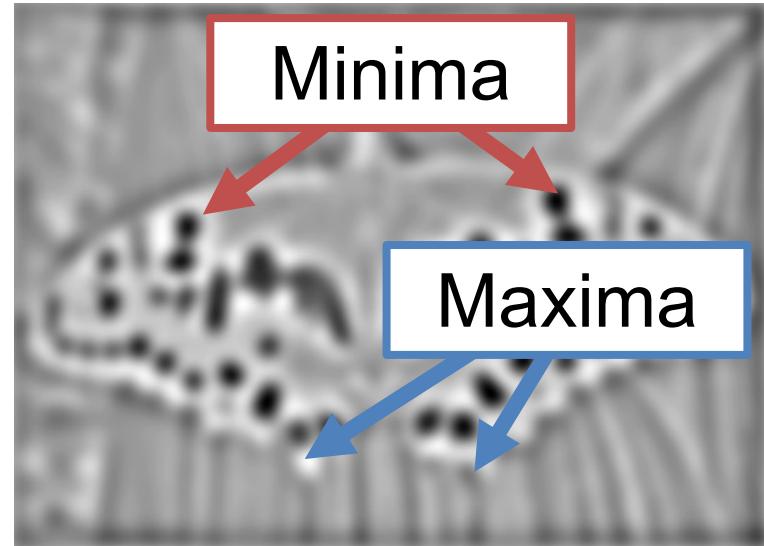


Blob Detection

Another detector: Detect and describe “Regions”

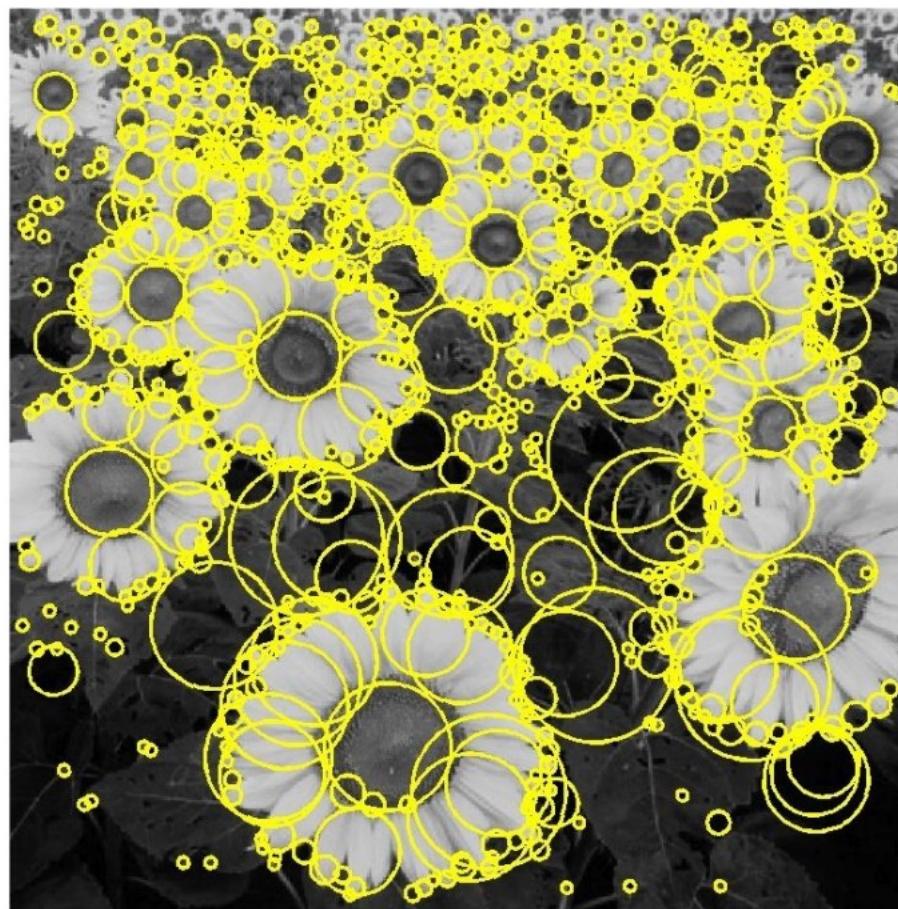


$$* \bullet =$$



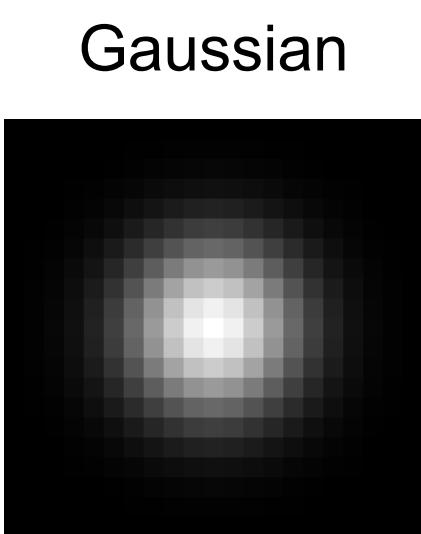
Find maxima ***and minima*** of blob filter response in
scale and space

Blob Detection with Scale Selection

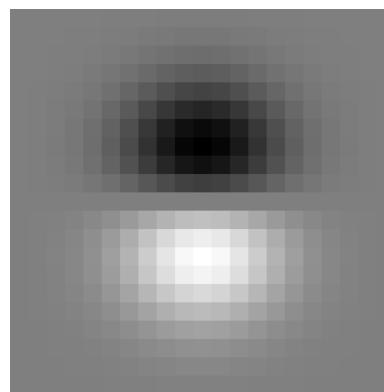


Slide credit: N. Snavely

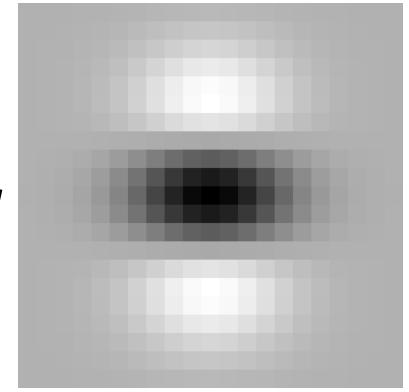
Gaussian Derivatives



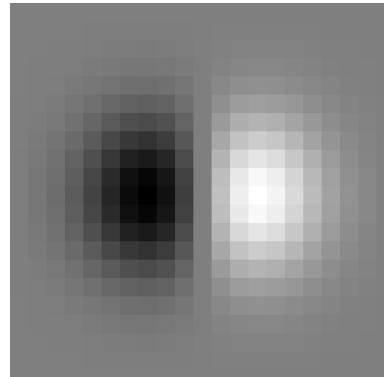
$$\frac{\partial}{\partial y} g$$



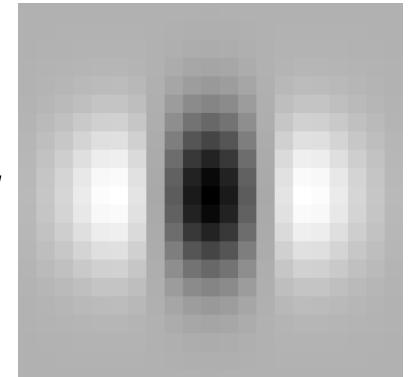
$$\frac{\partial^2}{\partial^2 y} g$$



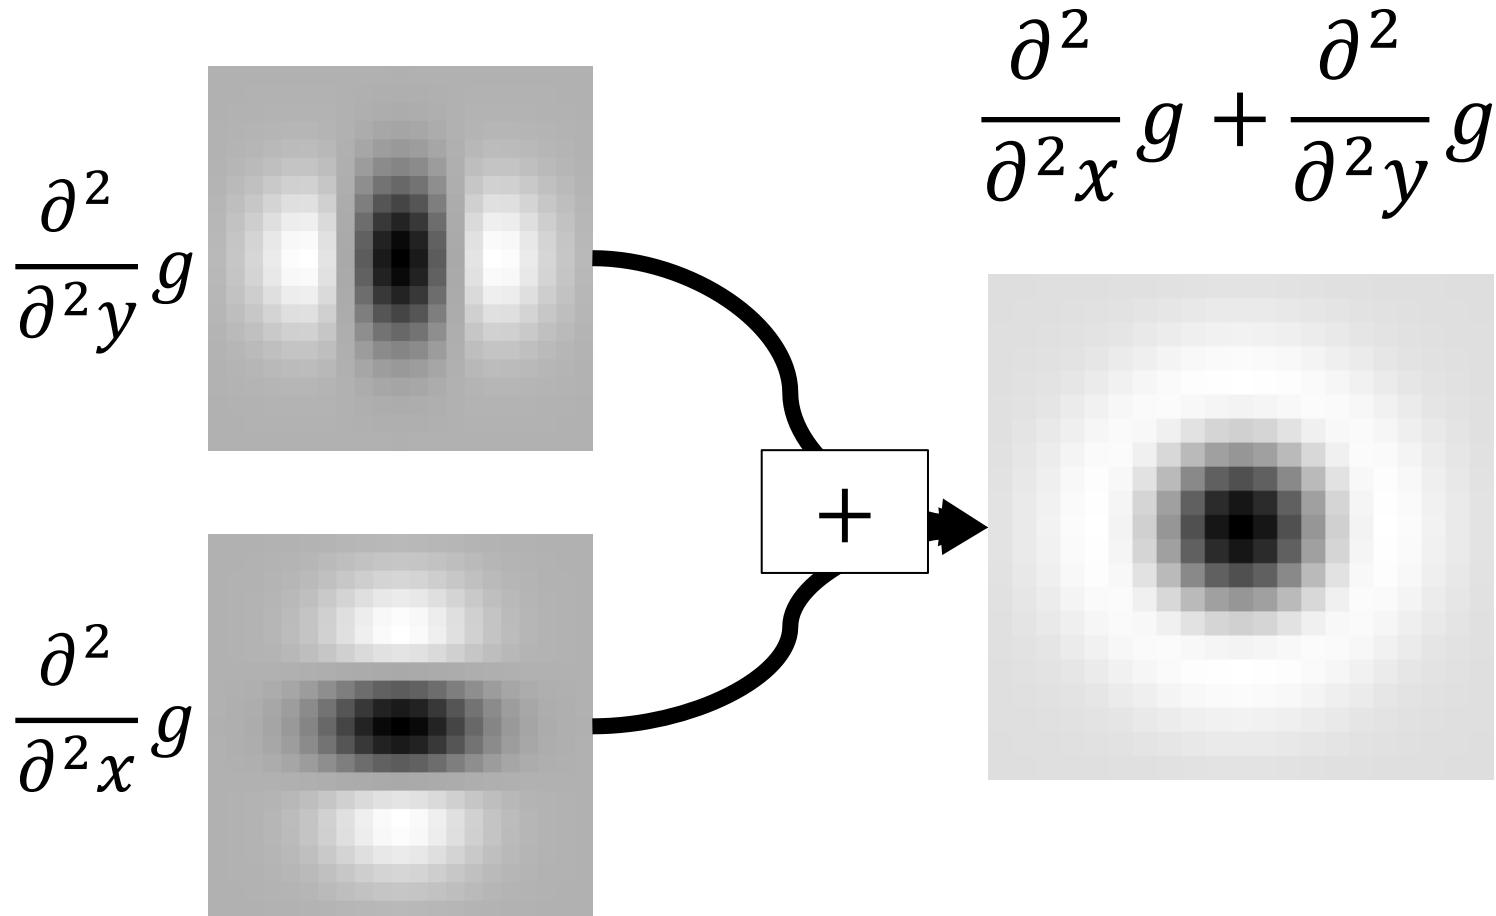
$$\frac{\partial}{\partial x} g$$



$$\frac{\partial^2}{\partial^2 x} g$$



Laplacian of Gaussian (LoG)

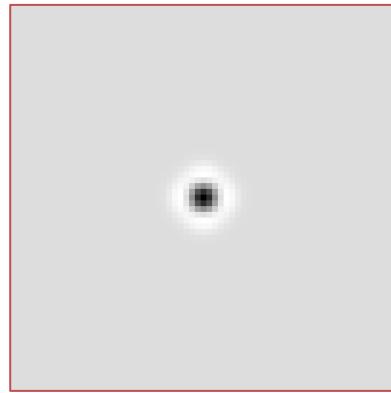


Slight detail: for technical reasons, you need to scale the Laplacian of Gaussian if you want to compare across sigmas.

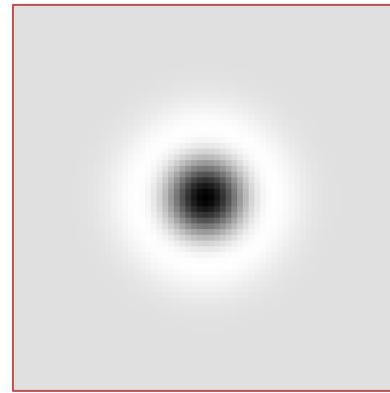
$$\nabla_{norm}^2 = \sigma^2 \left(\frac{\partial^2}{\partial x^2} g + \frac{\partial^2}{\partial y^2} g \right)$$

Size of Laplacian of Gaussian

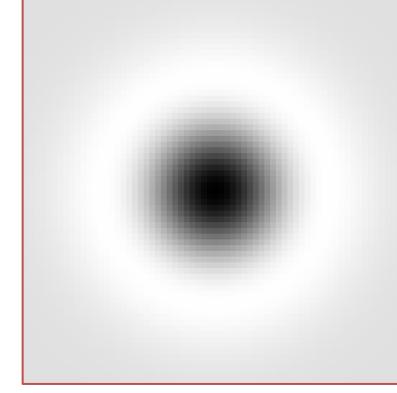
$\sigma = 2$



$\sigma = 6$



$\sigma = 10$

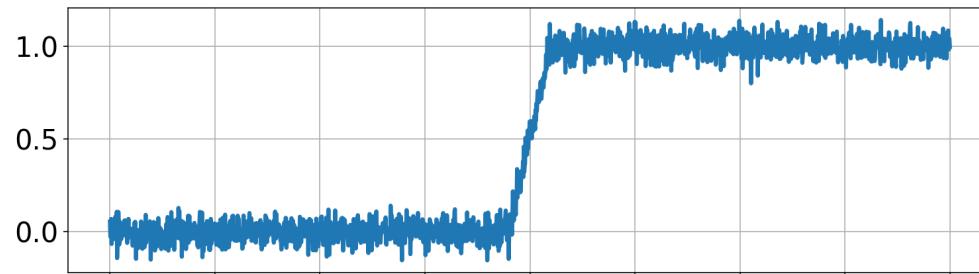


Slight detail: for technical reasons, you need to scale the Laplacian of Gaussian if you want to compare across sigmas.

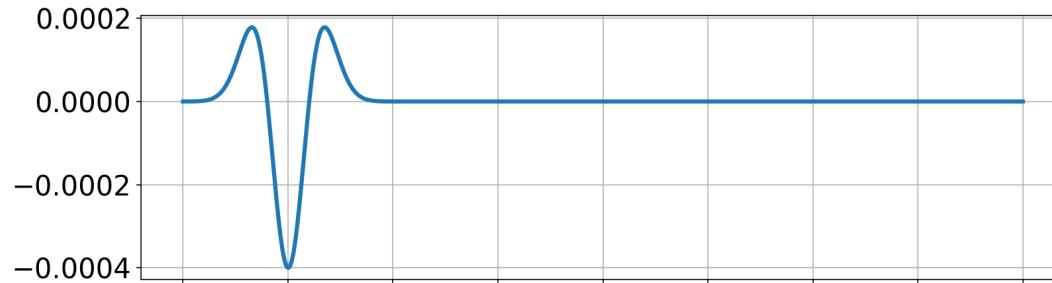
$$\nabla^2_{norm} = \sigma^2 \left(\frac{\partial^2}{\partial x^2} g + \frac{\partial^2}{\partial y^2} g \right)$$

Edge Detection with LoG

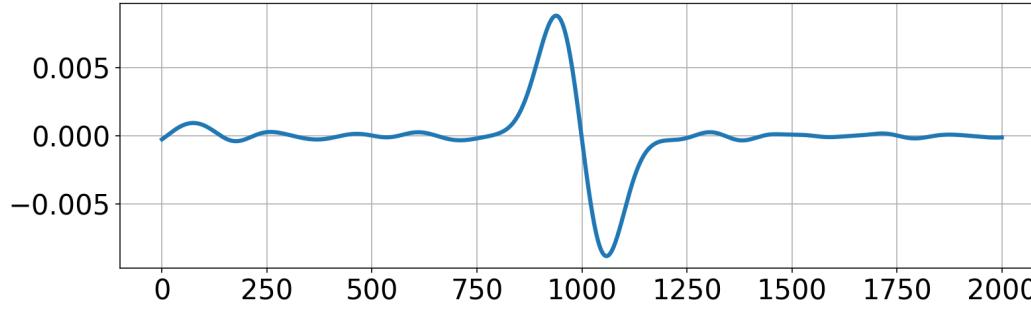
f



$\frac{\partial^2}{\partial^2 x} g$



$f * \frac{\partial^2}{\partial^2 x} g$

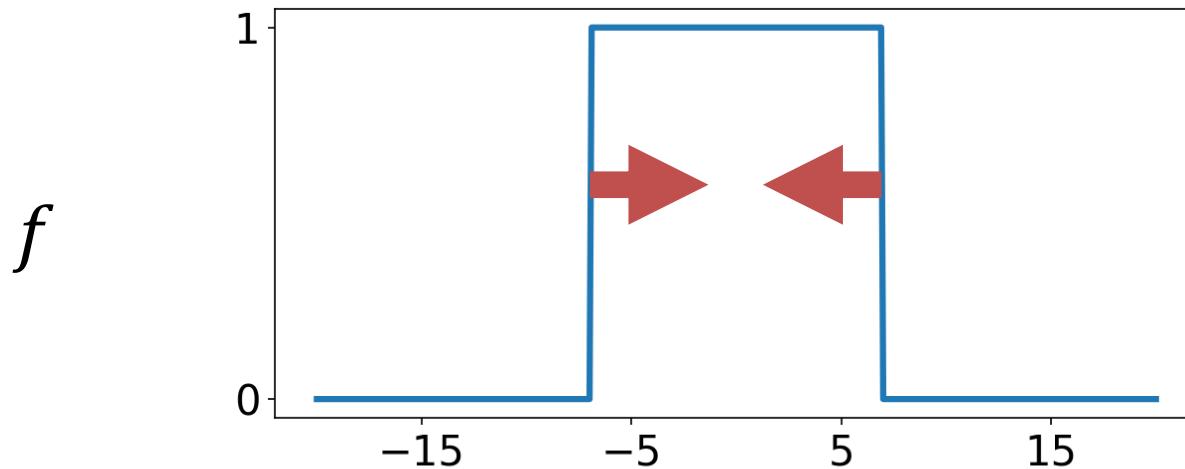
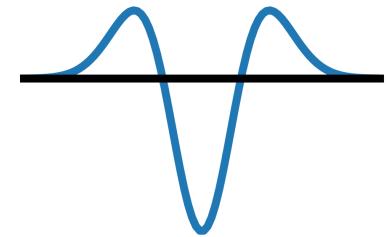


Edge

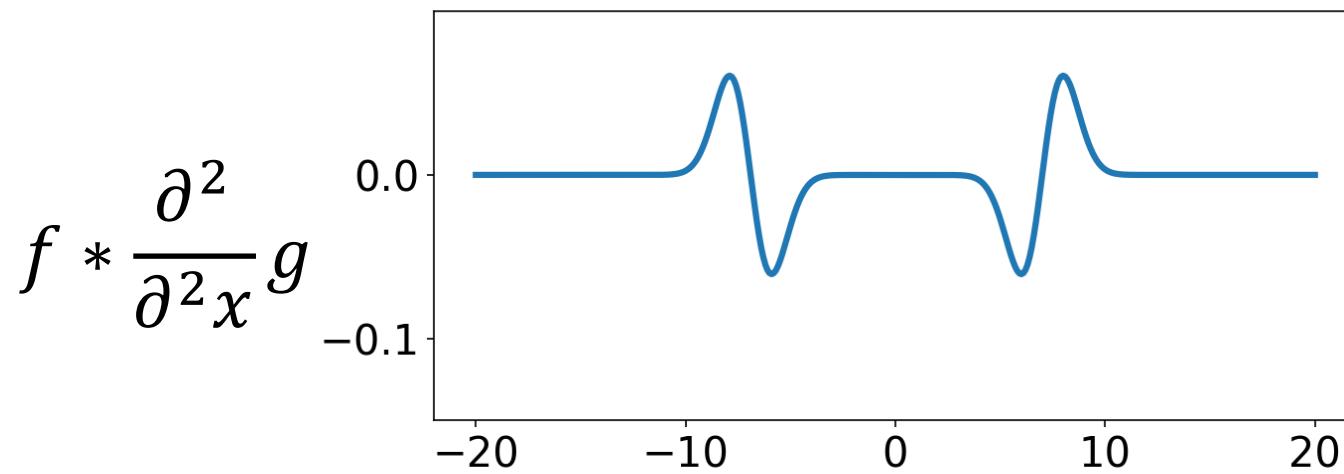
Laplacian
Of Gaussian
(LoG)

Edge =
Zero-crossing
Ripple

Blob Detection with LoG

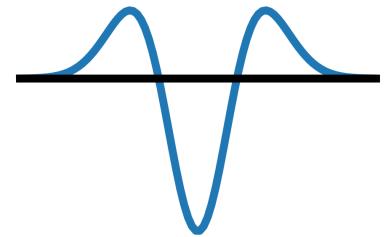


Edges



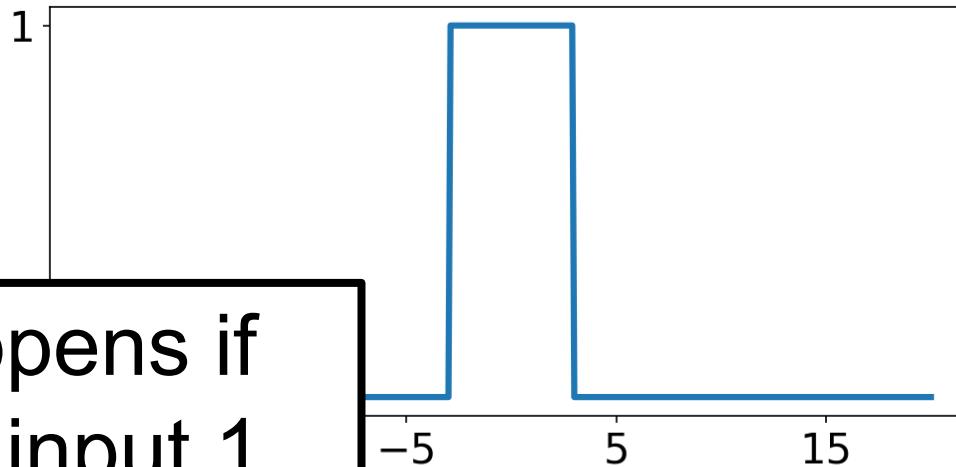
Blob * LoG =
Zero-crossings
Superposition
of Ripples

Blob Detection with LoG



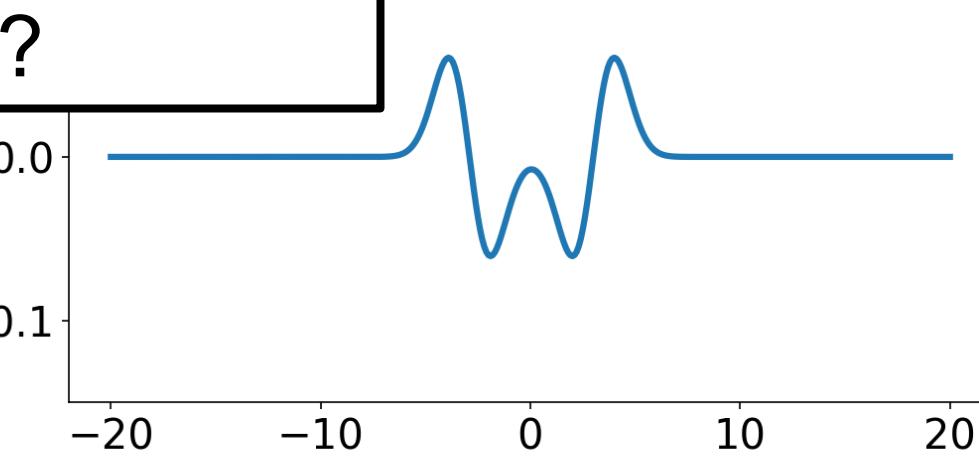
f

What happens if
we make input 1
unit wide?



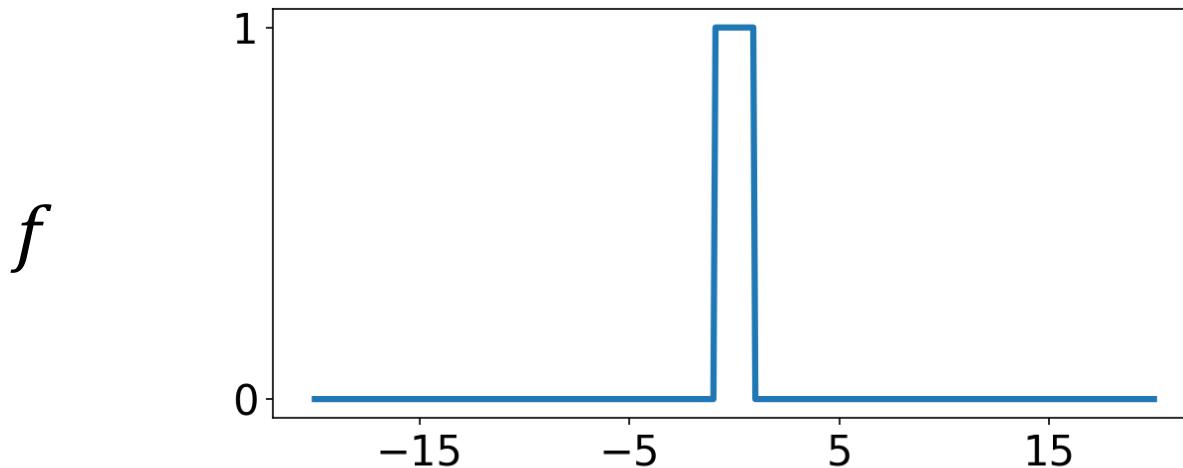
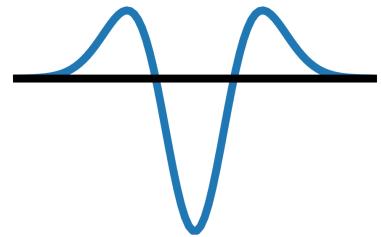
Edges

$f * \frac{\partial^2}{\partial^2 x} g$

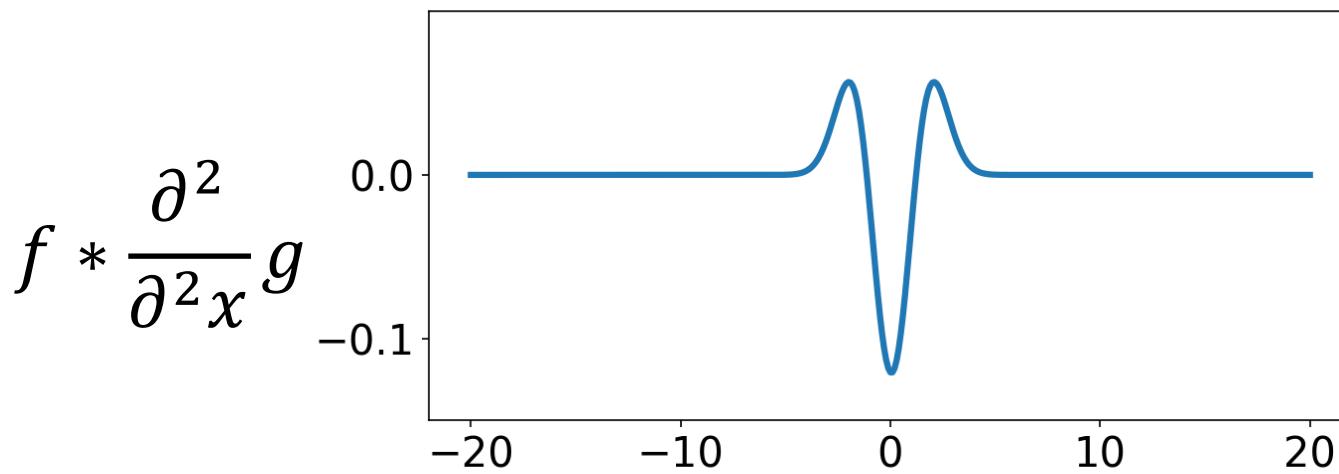


Edges *
LoG =
Zero-crossings
Superposition
of Ripples

Blob Detection with LoG



Edge

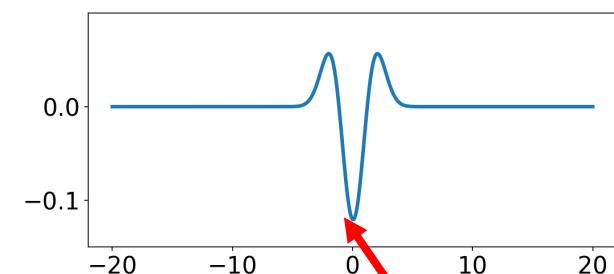
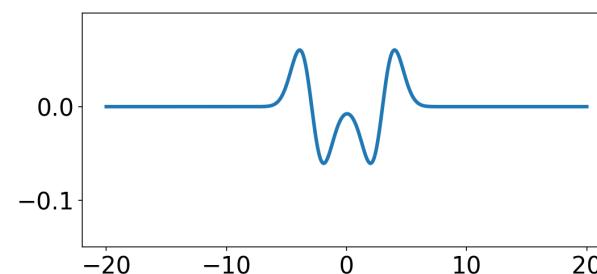
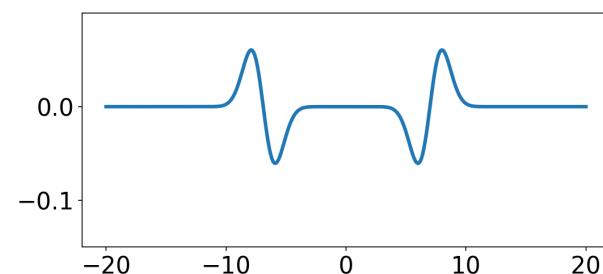
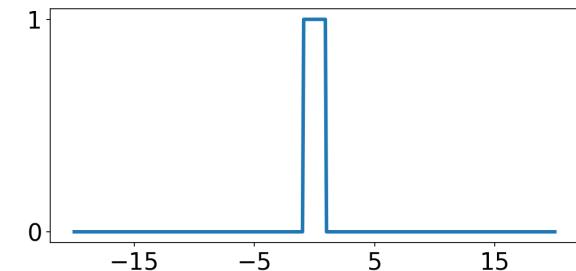
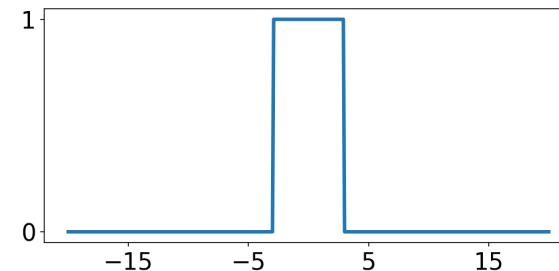
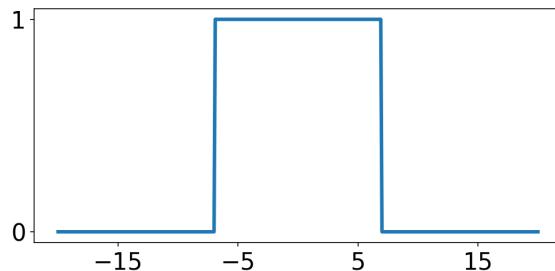


Edge *
LoG =
Zero-crossings
Constructive
Interference

Blob Detection with LoG

Each LoG kernel has a specific scale that responds highly

Original Signal

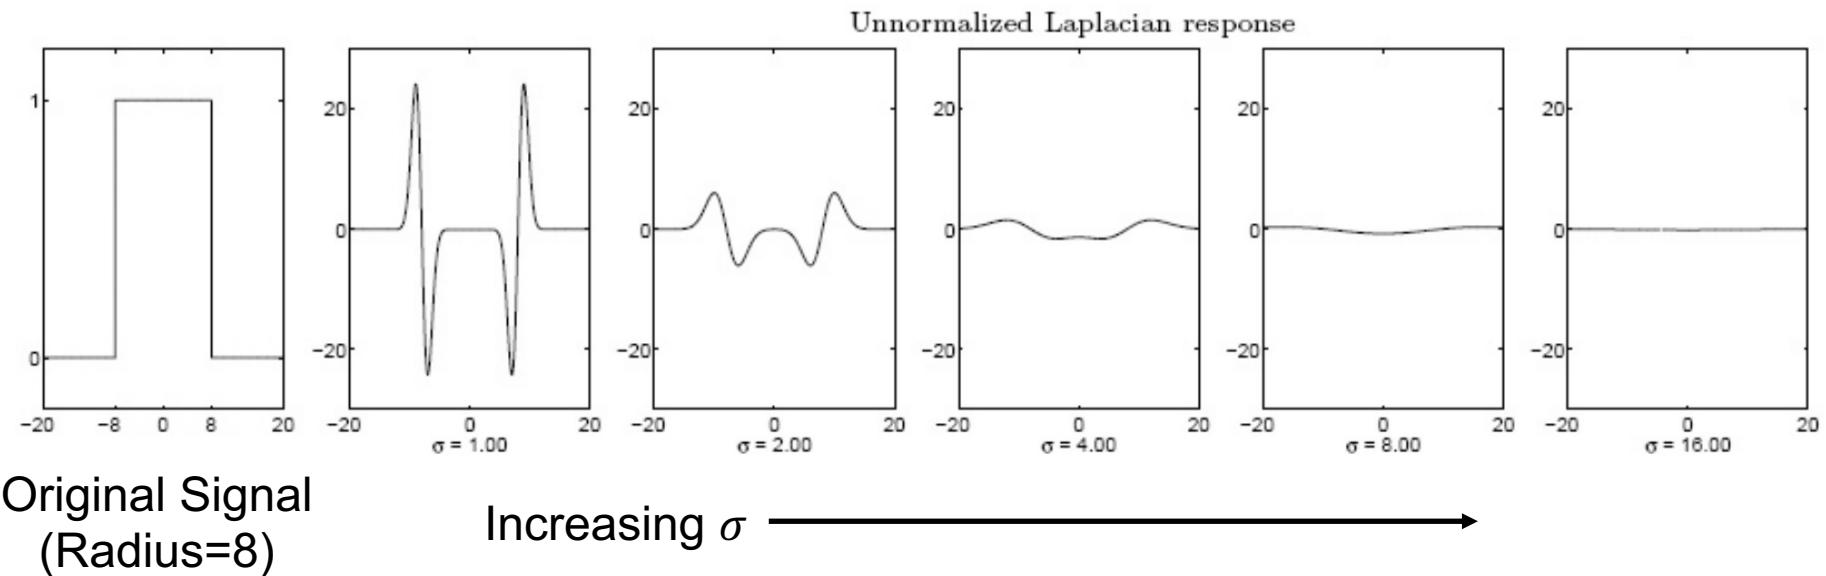


Convolved with Laplacian ($\sigma = 1$)

Maximum

Scale Selection

- Convolve with Laplacians at several scales and look for the maximum response
- However, Laplacian response decays as scale increases: energy dispersed widely



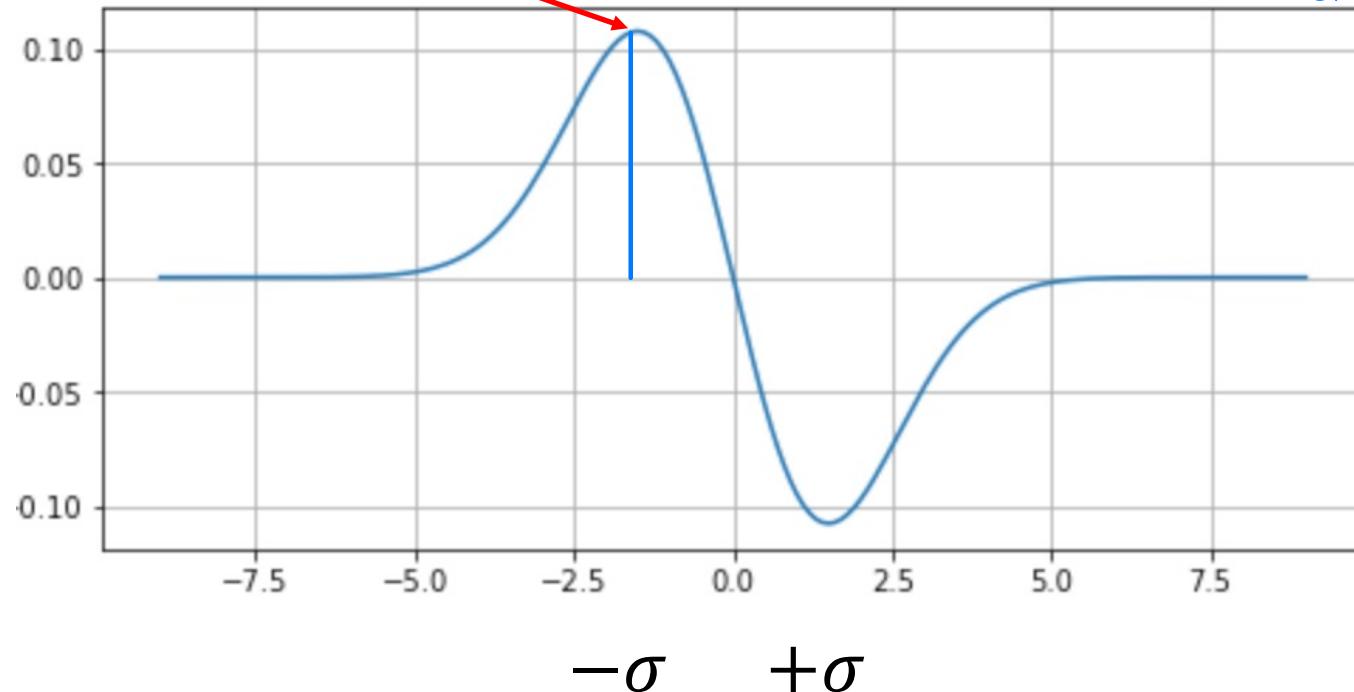
Scale Normalization

- The response of a derivative of Gaussian filter to an edge decreases as σ increases

Peak height $\propto \frac{1}{\sigma}$

1st order derivative of Gaussian Filter

$$\begin{aligned} G(x) &= \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2}{2\sigma^2}} \\ G'(x) &= \frac{1}{\sqrt{2\pi}\sigma} \cdot e^{-\frac{x^2}{2\sigma^2}} \cdot \left(-\frac{x}{\sigma^2}\right) \\ &= \frac{-x}{\sqrt{2\pi}\sigma^3} e^{-\frac{x^2}{2\sigma^2}} \\ G''(x) & \end{aligned}$$

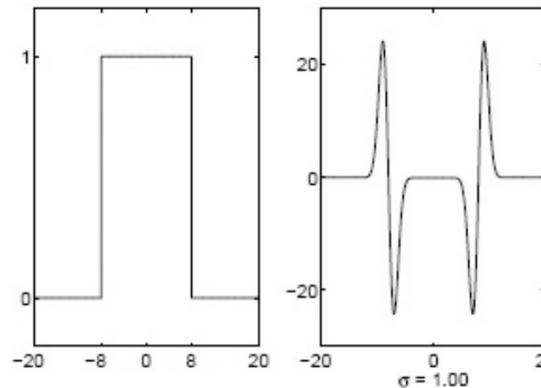


Scale Normalization

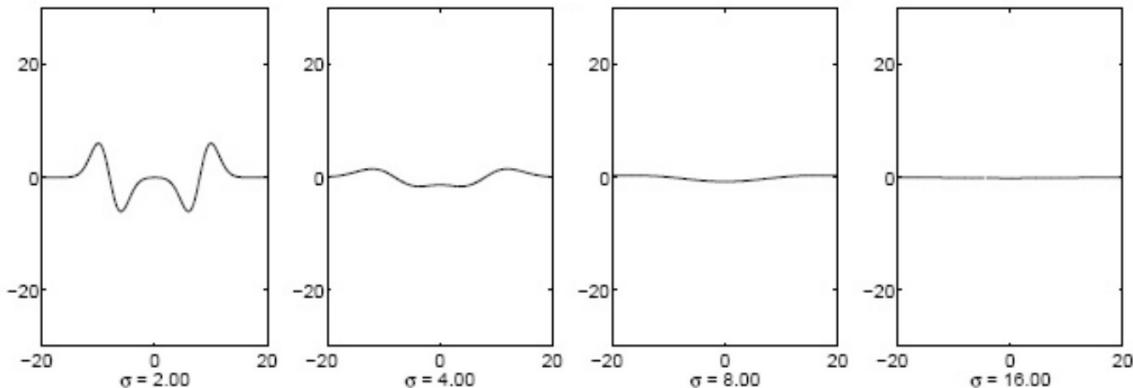
- The response of a derivative of Gaussian filter to a perfect step edge decreases as σ increases
- Peak height $\propto \frac{1}{\sigma}$
- To keep response the same (scale-invariant), must multiply Gaussian derivative by σ
- Laplacian is the second Gaussian derivative, so it must be multiplied by σ^2

Scale Normalization

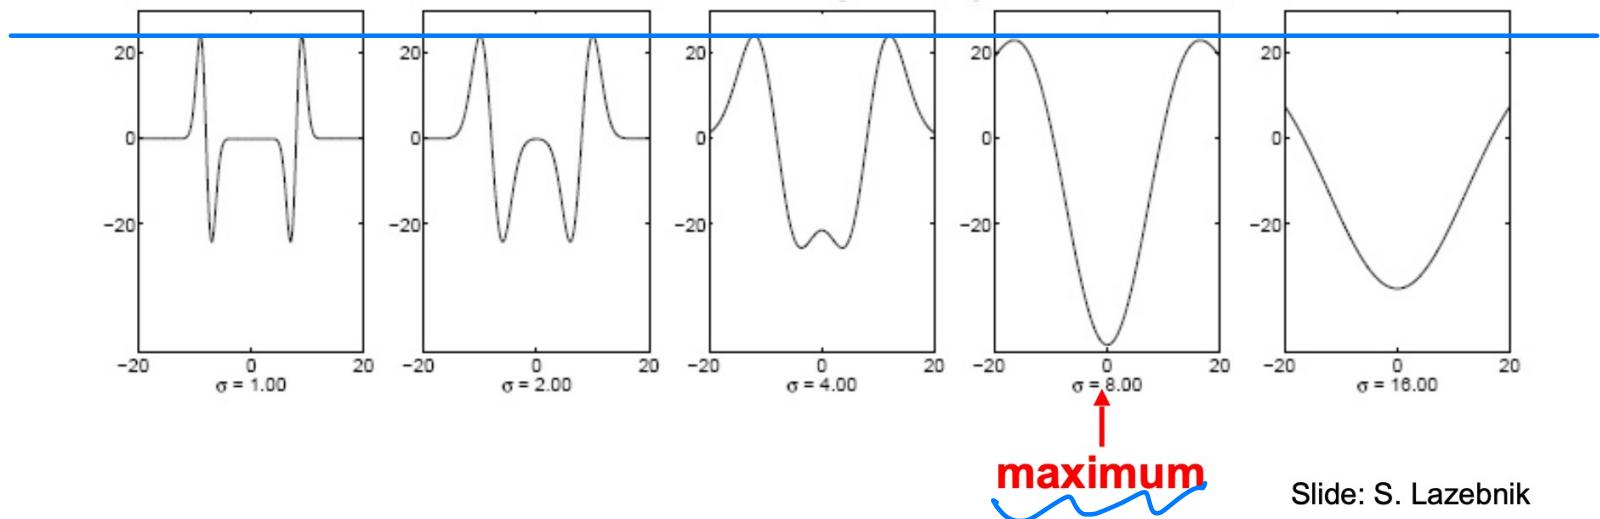
Original signal



Unnormalized Laplacian response



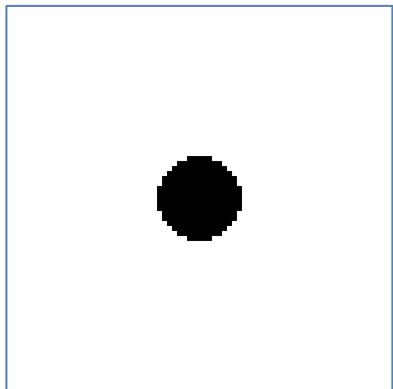
Scale-normalized Laplacian response



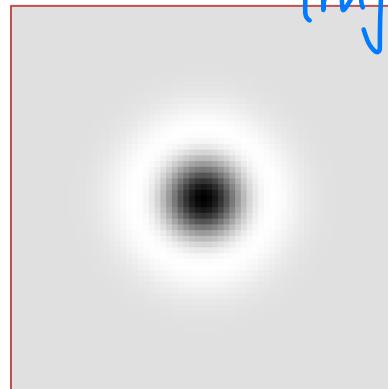
Scale Selection

Given binary circle and Laplacian filter of scale σ , we can compute the response as a function of the scale.

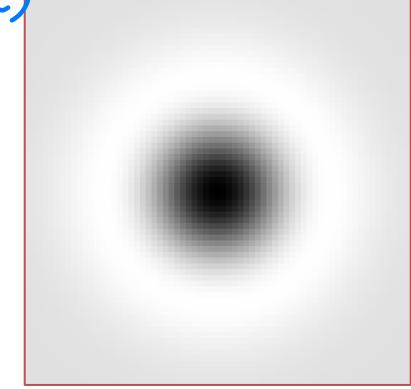
Image $\sigma = 2$
Radius: 8 *maximum height of this response*



$\sigma = 6$
R: 2.9 *(highest)*

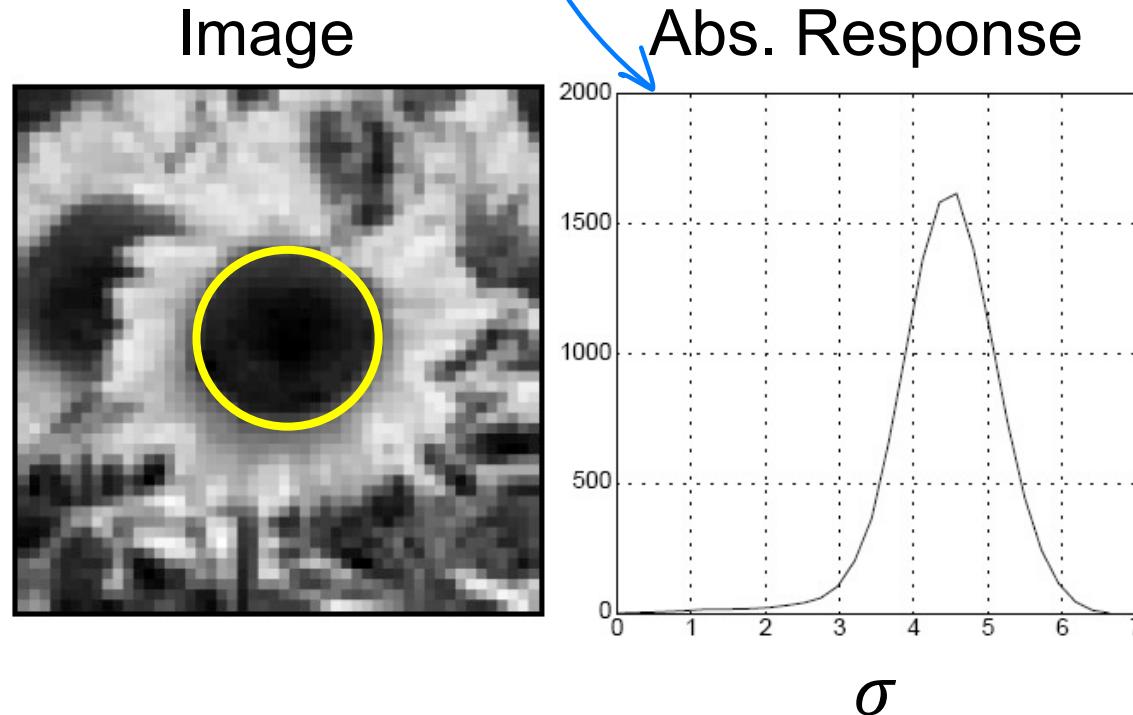


$\sigma = 10$
R: 1.8



Characteristic Scale

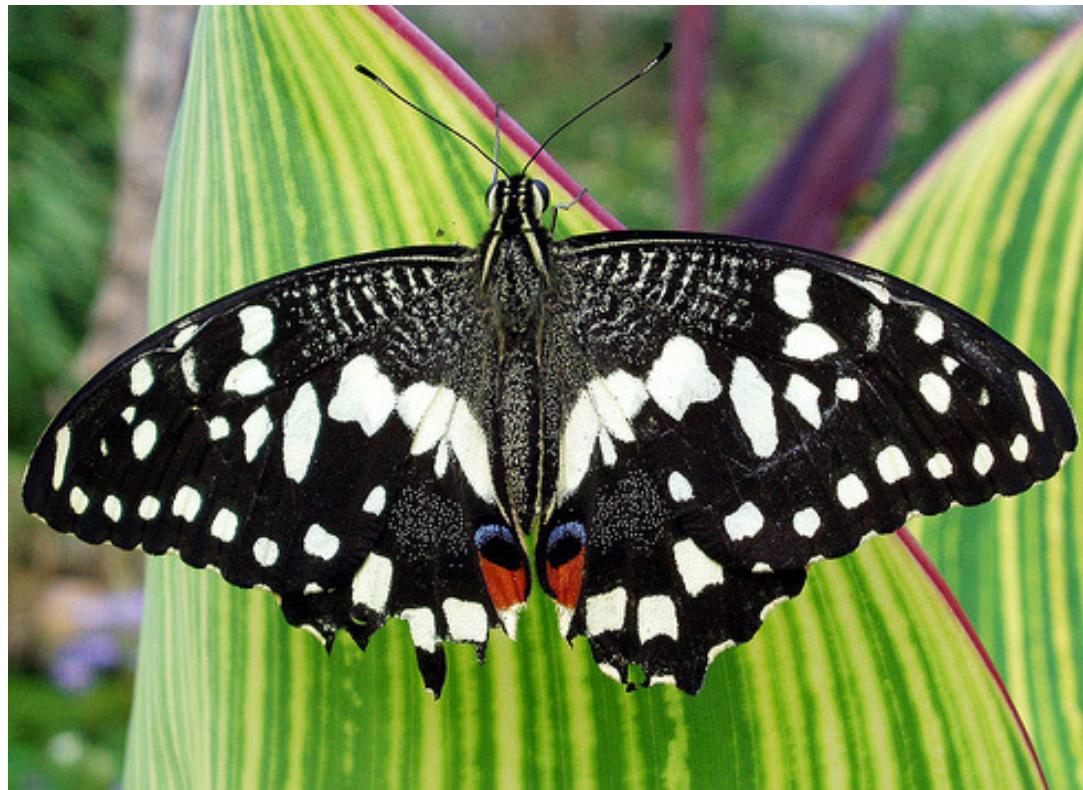
Characteristic scale of a blob is the scale that produces the maximum response



Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales

Scale-space blob detector: Example

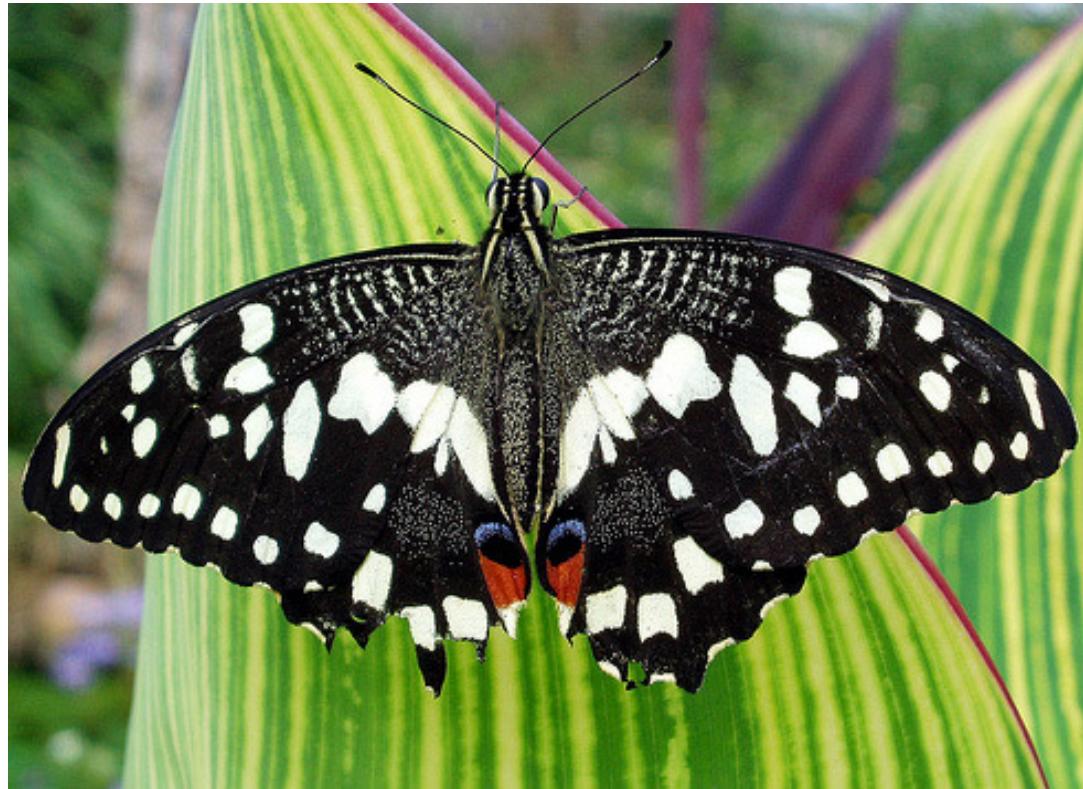


Scale-space blob detector: Example



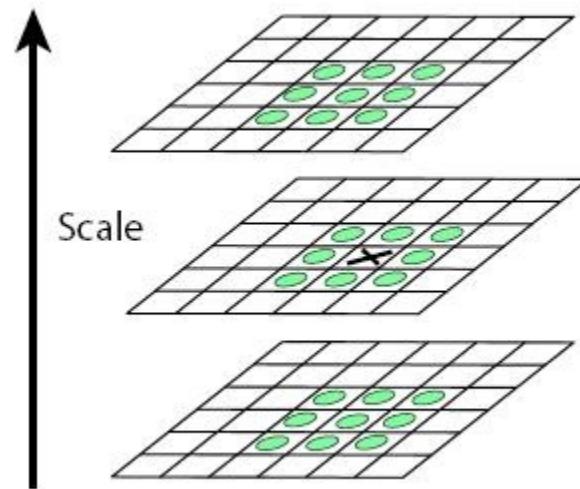
$\sigma = 11.9912$

Scale-space blob detector: Example



Scale-space blob detector

1. Convolve image with scale-normalized Laplacian at several scales
2. Find local maxima of squared Laplacian response in scale-space



Finding Spatial Maxima

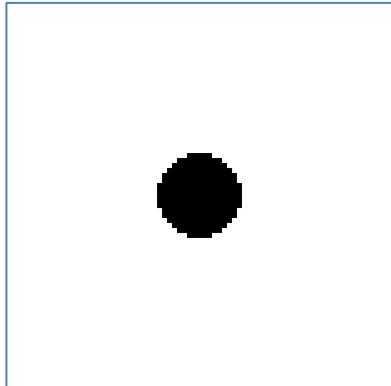
Point i,j is maxima (minima if you flip sign) in image I if it's bigger than all neighbors (8)

```
for y=range(i-1,i+1+1):
    for x in range(j-1,j+1+1):
        if y == i and x== j: continue
        #below has to be true
        I[y,x] < I[i,j]
```

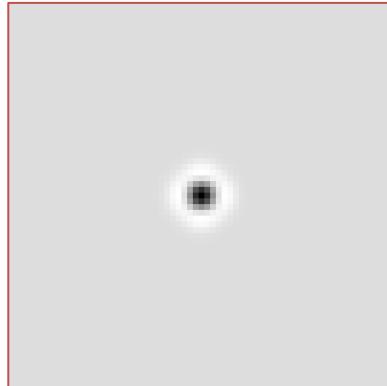
Scale Space

Blue lines are image-space neighbors (should be just one pixel over but that's impossible to draw)

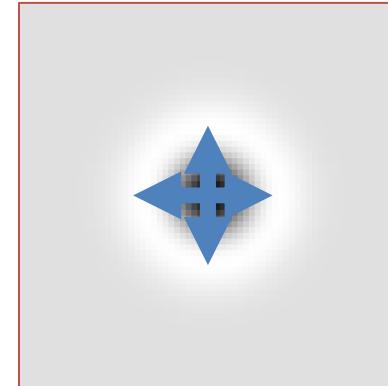
Image
Radius: 8



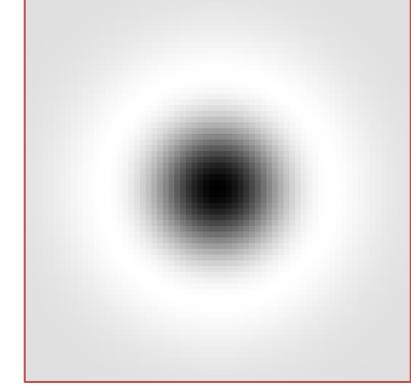
$\sigma = 2$
R: 0.02



$\sigma = 6$
R: 2.9



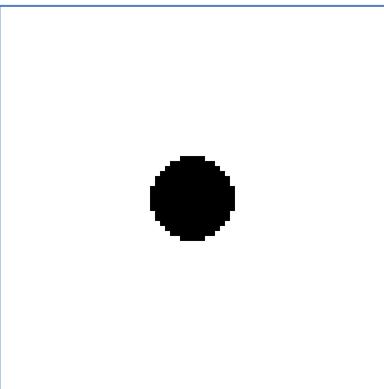
$\sigma = 10$
R: 1.8



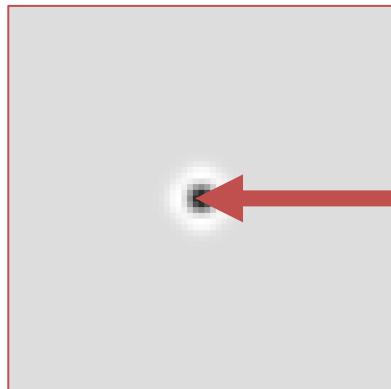
Scale Space

Red lines are the scale-space neighbors

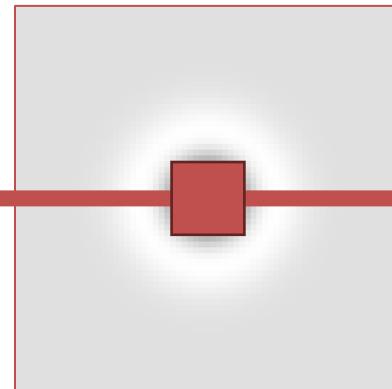
Image
Radius: 8



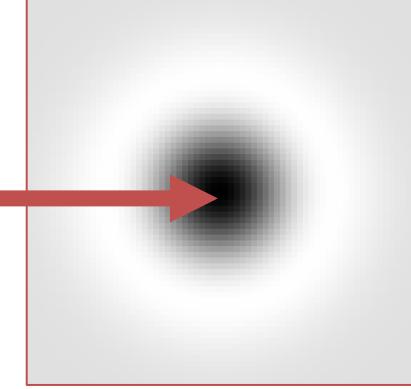
$\sigma = 2$
R: 0.02



$\sigma = 6$
R: 2.9



$\sigma = 10$
R: 1.8



Finding Scale-Space Maxima

Suppose $I[:, :, k]$ is image at scale k . Point i, j, k is maxima (minima if you flip sign) in image I if:

for $y = \text{range}(i-1, i+1+1)$:

 for $x = \text{range}(j-1, j+1+1)$:

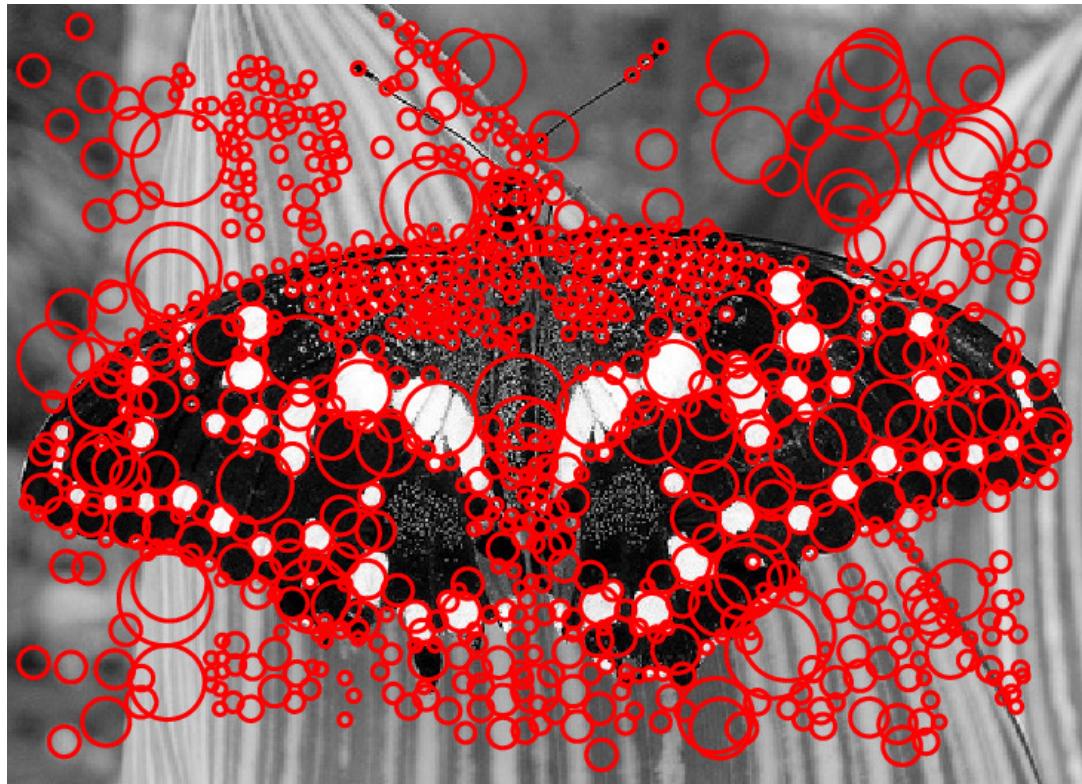
 for $c = \text{range}(k-1, k+1+1)$:

 if $y == i$ and $x == j$ and $c == k$:
 continue

#below has to be true

$I[y, x, c] < I[i, j, k]$

Scale-space blob detector: Example



Efficient implementation

- Approximating the Laplacian with a difference of Gaussians:

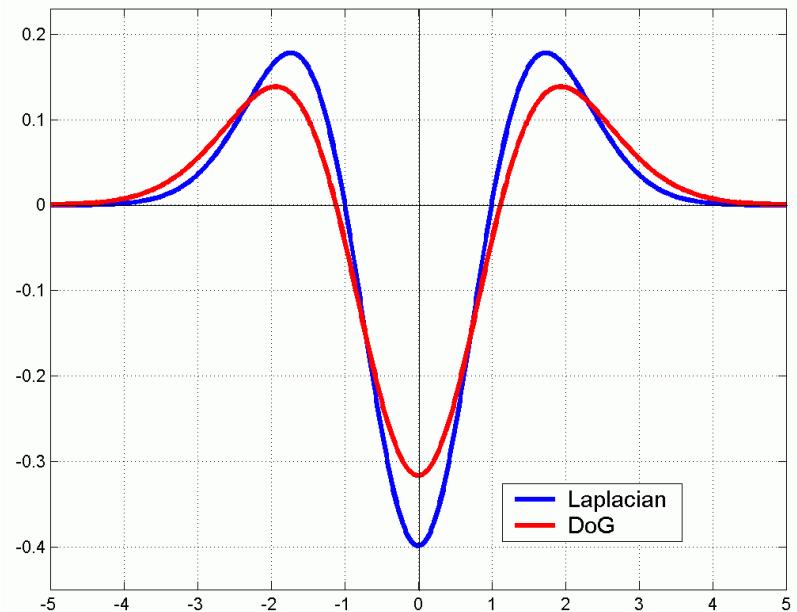
$$L \triangleq \sigma^2 [G_{xx}(x, y, \sigma) + G_{yy}(x, y, \sigma)]$$

(Laplacian)

$$DoG \triangleq G(x, y, k\sigma) - G(x, y, \sigma)$$

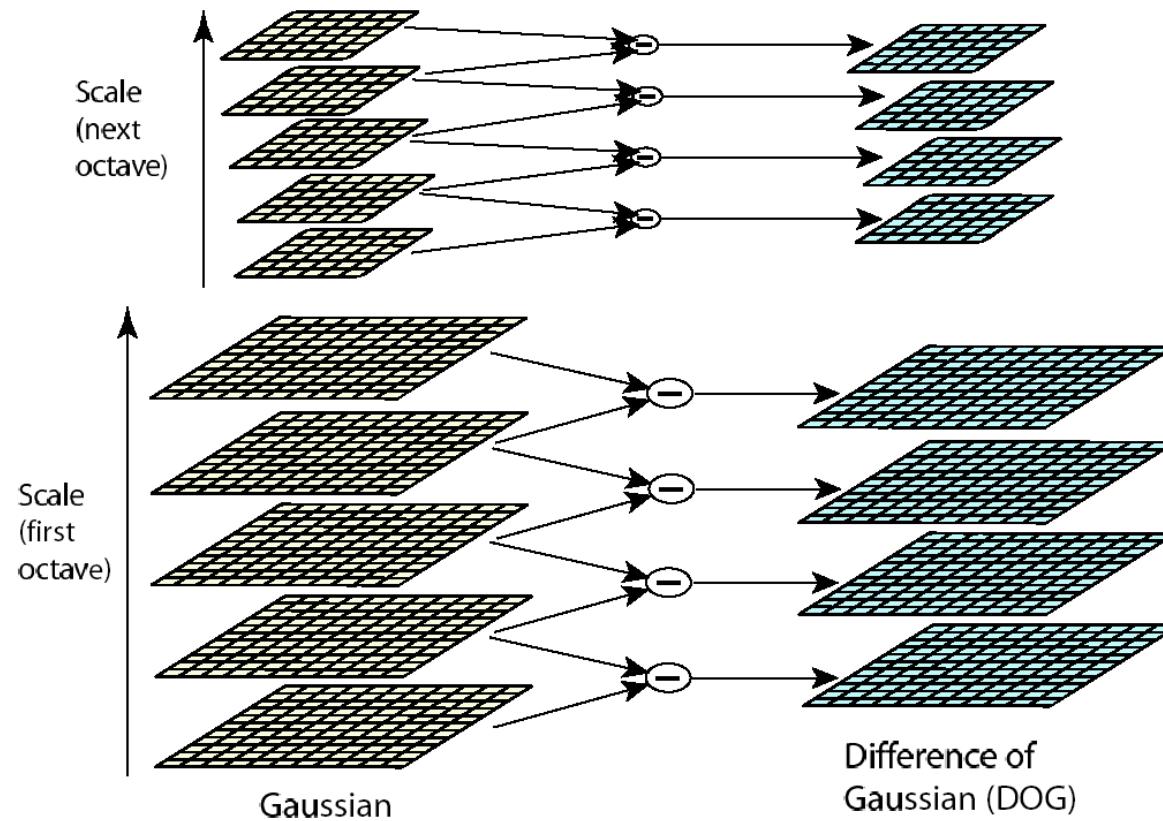
(Difference of Gaussians)

k is empirically determined



Efficient implementation

...



David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" IJCV 60 (2), pp. 91-110, 2004.

Slide credit: S. Lazebnik

Problem 1 Solved

- How do we deal with scales: try them all
- **Why is this efficient?**

Vast majority of effort is in the first and second scales

$$1 + \frac{1}{4} + \frac{1}{16} + \frac{1}{64} + \frac{1}{4^i} \dots = \frac{4}{3}$$

Problem 2 – Describing Features

Image – 40

1/2 size, rot. 45°
Lightened+40

Full
Image



100x100 crop
at Glasses



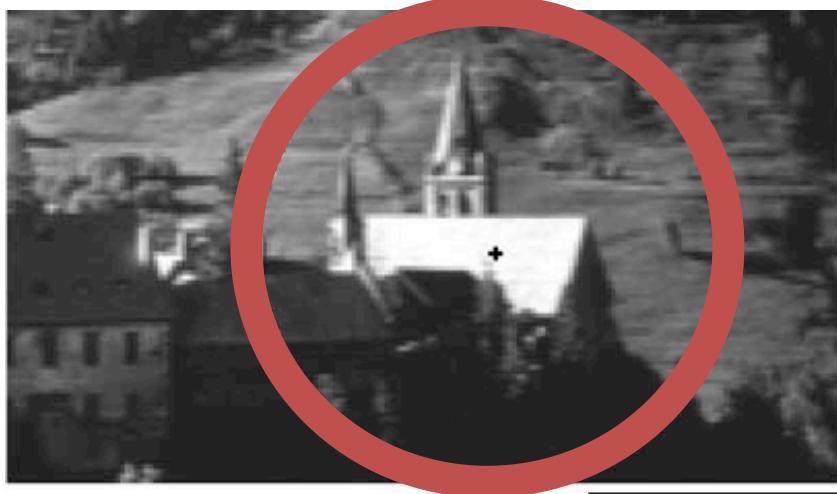
Problem 2 – Describing Features

Once we've found a corner/blobs, we can't just use the nearby patch. What about:

1. Scale?
2. Rotation?
3. Additive light?

Handling Scale

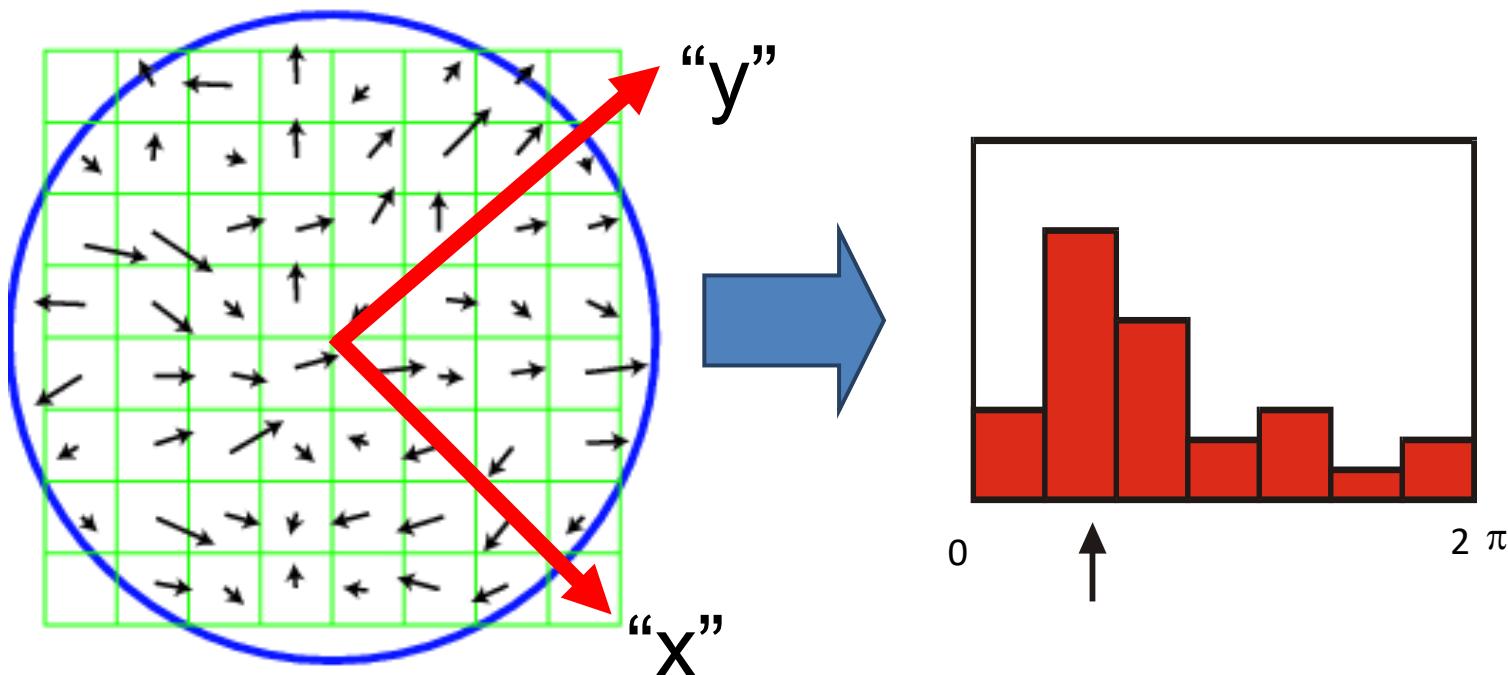
Given characteristic scale (maximum Laplacian response), we can just rescale image



Handling Rotation

Create histogram of local gradient directions in the patch

Given window, can compute “dominant orientation”
and then rotate image



Rectangles at different sizes, sizes tell you the detection of of blond at certain size. So that is the patch that we are interested in. The rectangles are rotated, which tells you the domination of these blobs.

Scale and Rotation

SIFT features at characteristic scales and dominant orientations



Picture credit: S. Lazebnik. Paper: David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" *IJCV* 60 (2), pp. 91-110, 2004.

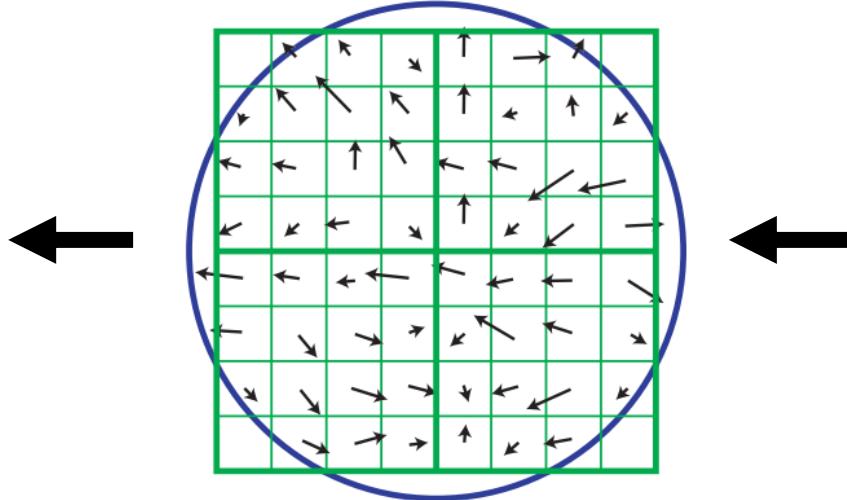
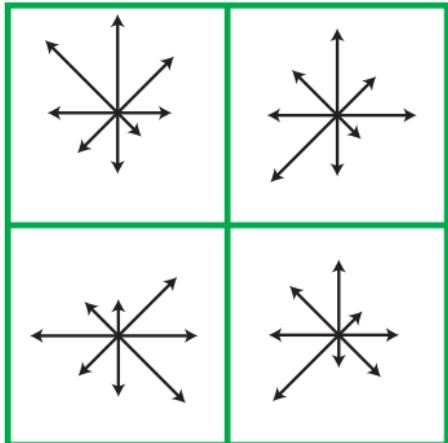
Scale and Rotation



Picture credit: S. Lazebnik. Paper: David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" *IJCV* 60 (2), pp. 91-110, 2004.

SIFT Descriptors

All these arrows tell you the gradient direction of each pixel.



1. Compute gradients
2. Build histogram (2x2 shown, 4x4 in practice)
3. (4x4) cells * 8 orientations = 128-dim descriptor

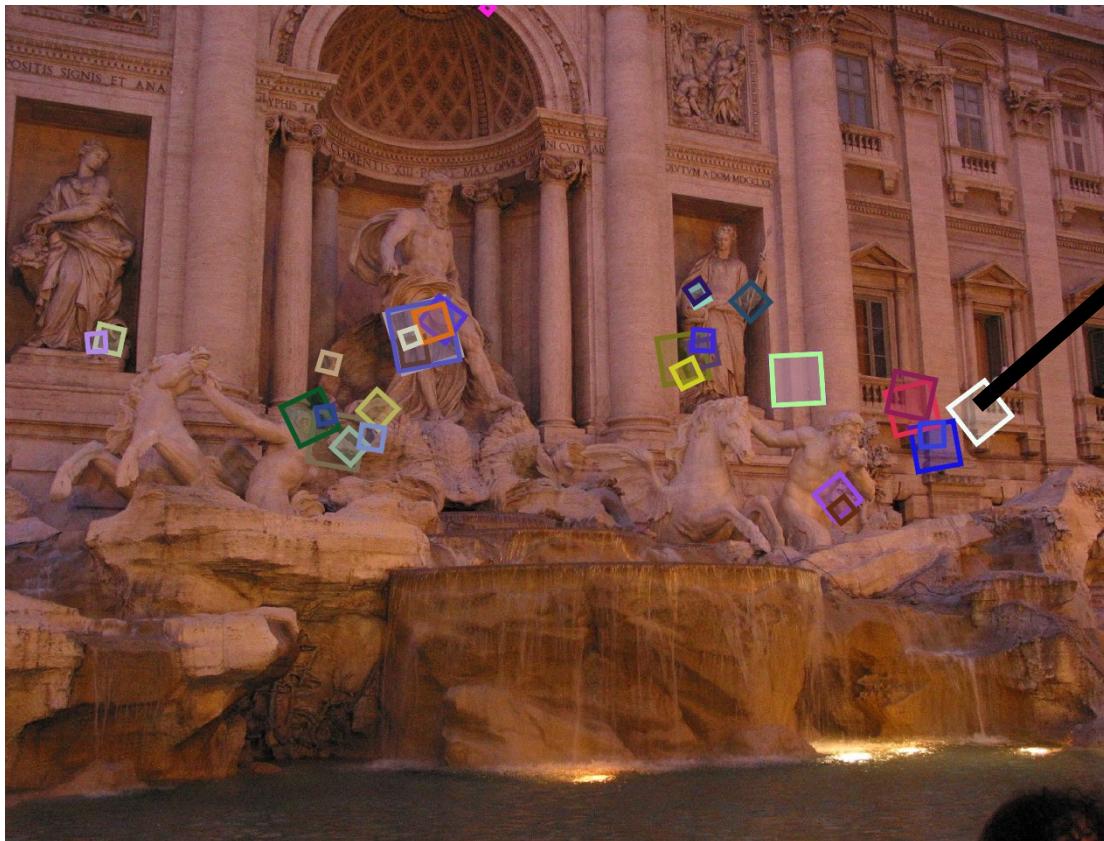
Gradients ignore global illumination changes

SIFT Descriptors

- In principle: build a histogram of the gradients
- In reality: quite complicated
 - Gaussian weighting: smooth response
 - Normalization: reduces illumination effects
 - Clamping
 - Tons of more stuff

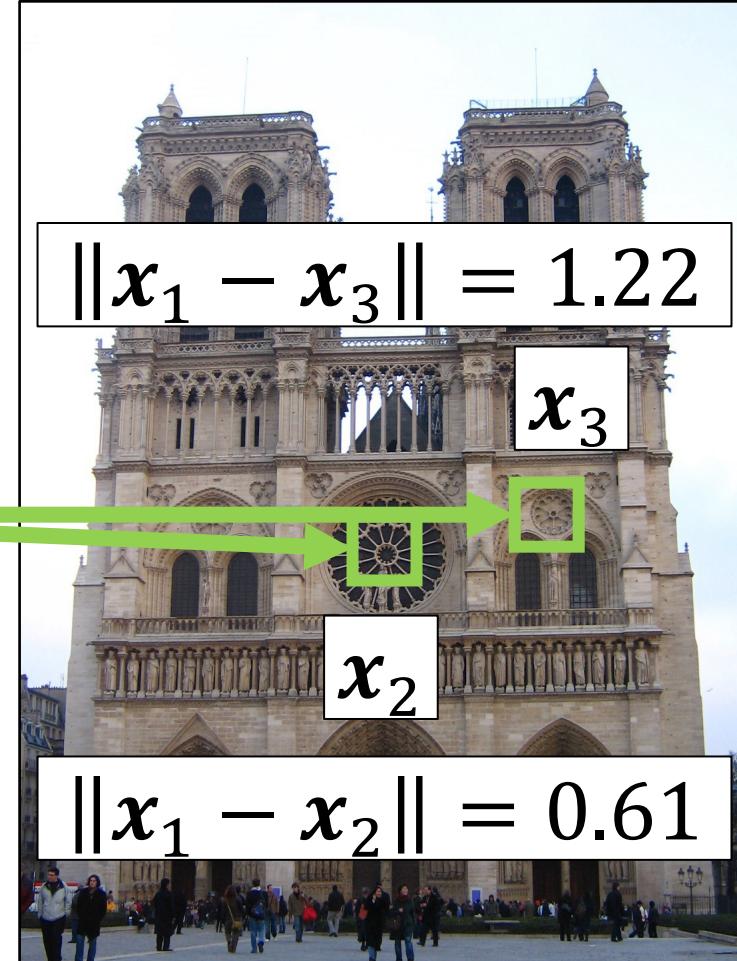
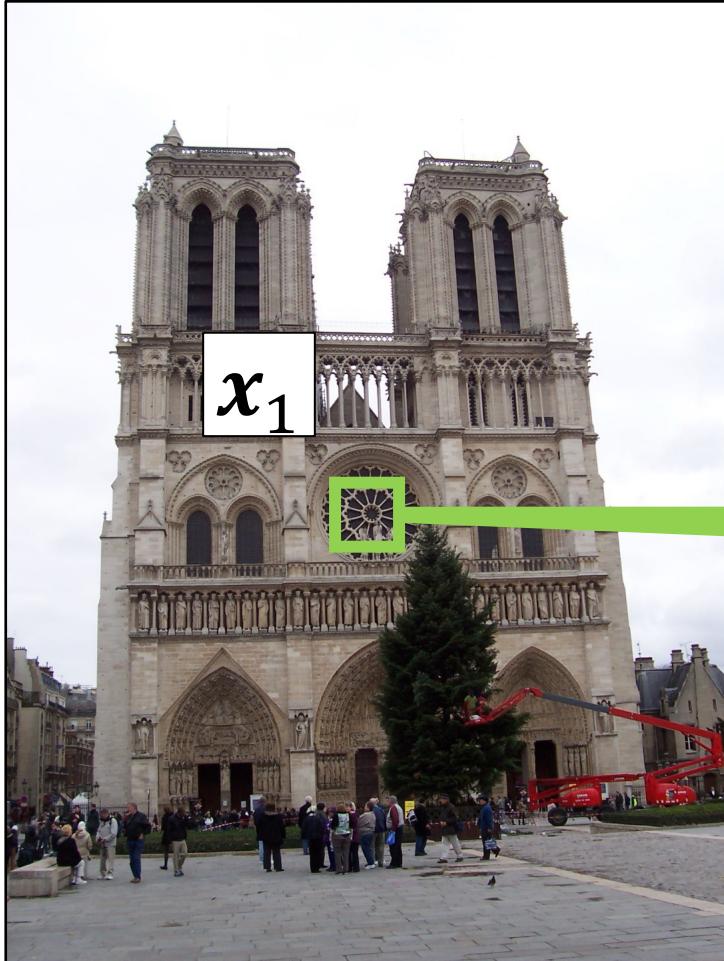
Feature Descriptors

Think of feature as some non-linear filter that maps each pixel to a 128-dimensional feature



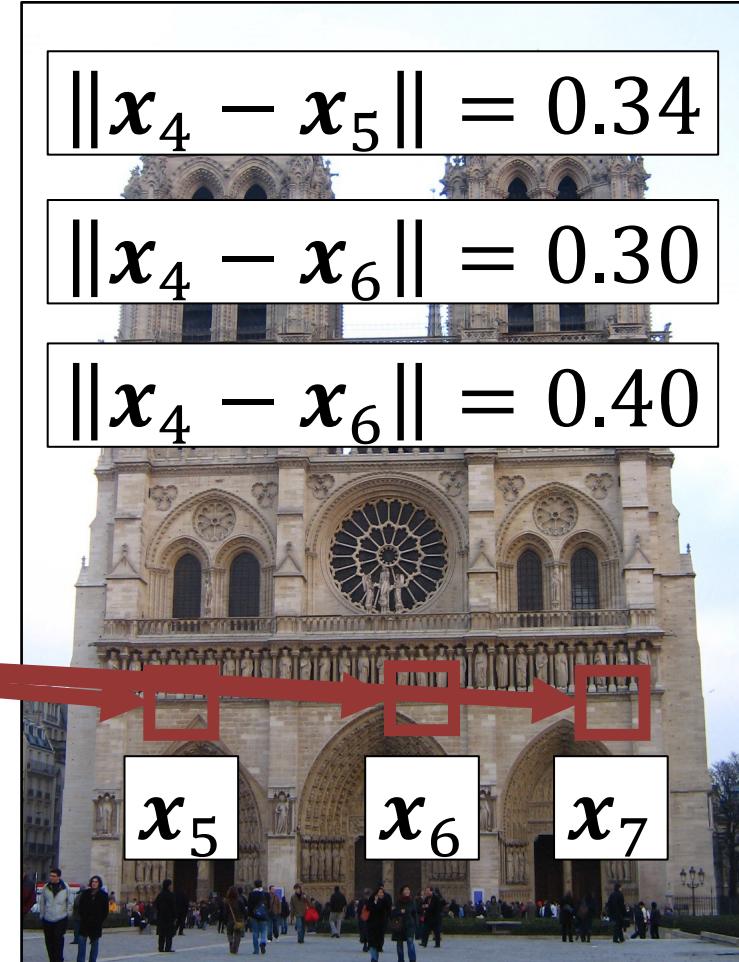
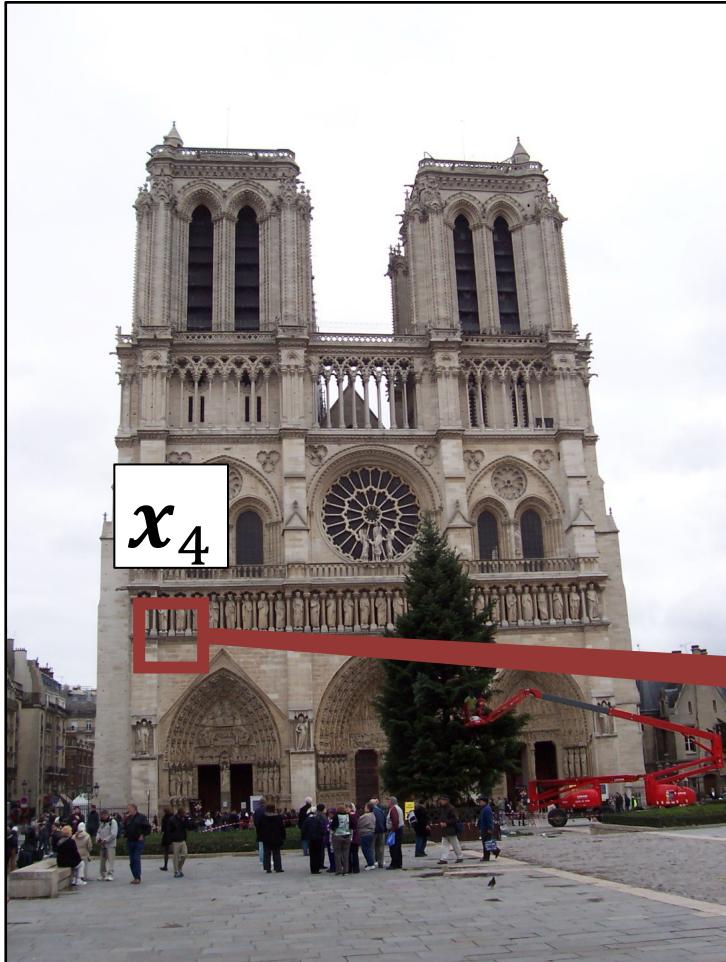
128D
vector \mathbf{x}

Instance Matching



Example credit: J. Hays

Instance Matching



Example credit: J. Hays

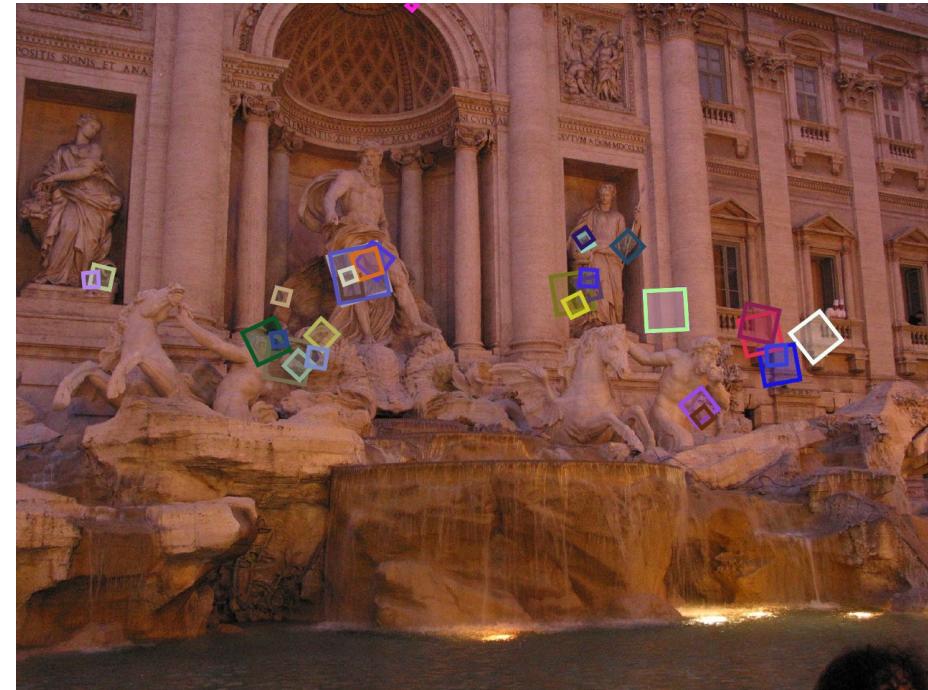
2nd Nearest Neighbor Trick

- Given a feature x_q , nearest neighbor to x is a good match, but distances can't be thresholded.
- Instead, find nearest neighbor (x_{1NN}) and second nearest neighbor (x_{2NN}). This ratio is a good test for matches:

$$r = \frac{\|x_q - x_{1NN}\|}{\|x_q - x_{2NN}\|}$$

Properties of SIFT

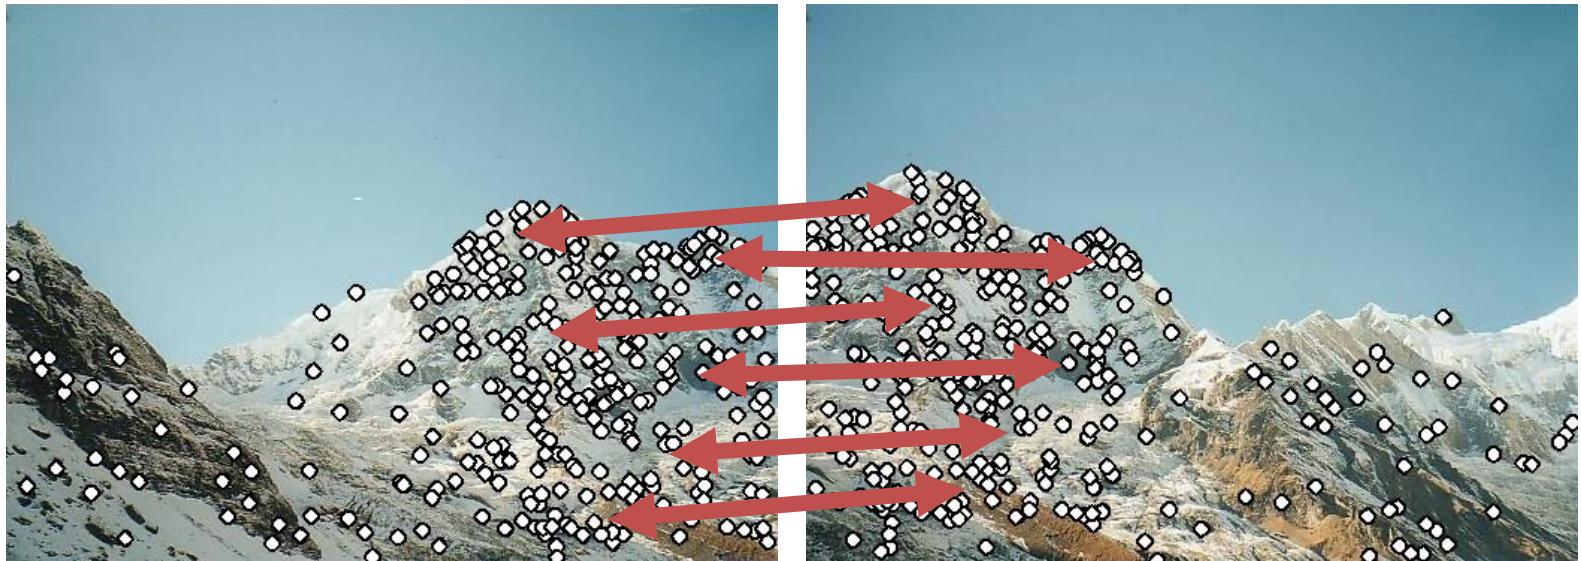
- Can handle: up to ~60 degree out-of-plane rotation, changes of illumination (night vs day)
- Fast, efficient, code available (but was patented)



Summing Up

1. Convolve image with scale-normalized Laplacian at several scales
2. Find local maxima of squared Laplacian response in scale-space
3. Compute gradients
4. Compute SIFT Descriptors (Histogram of Grad)
5. Match between features

So Far; What's Next?



- 1: find corners+features
- 2: match based on local image data
- 3: next time: compute transforms from matches

Extra Reading for the Curious

Aside: This Trick is Common

Given a 50x16 person detector, how do I detect:

- (a) 250x80 (b) 150x48 (c) 100x32 (d) 25x8 people?



Sample people from image



Aside: This Trick is Common

Detecting all the people
The red box is a fixed size



Sample people from image



Aside: This Trick is Common

Detecting all the people
The red box is a fixed size

Sample people from image



Aside: This Trick is Common

Detecting all the people
The red box is a fixed size

Sample people from image



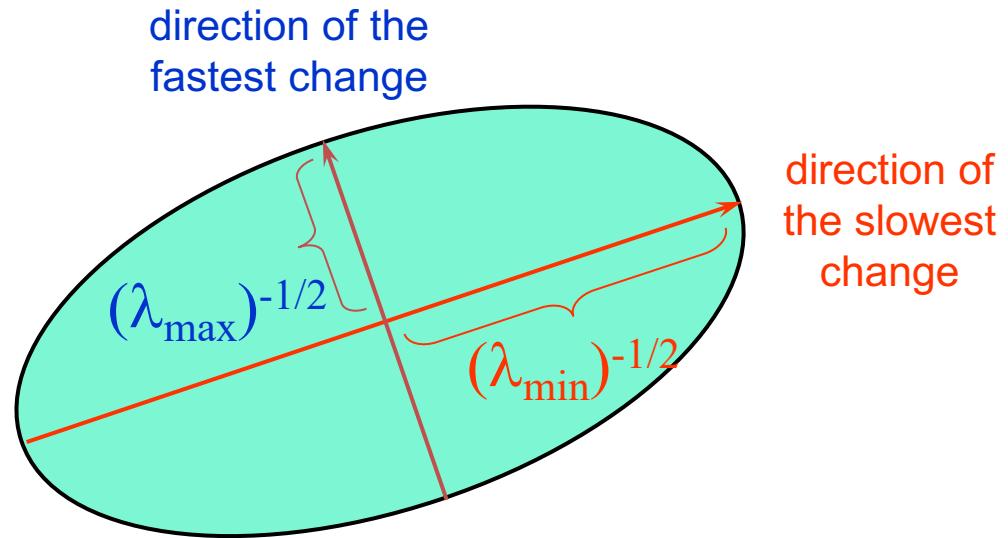
Affine adaptation

Consider the second moment matrix of the window containing the blob:

$$M \triangleq \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \triangleq R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

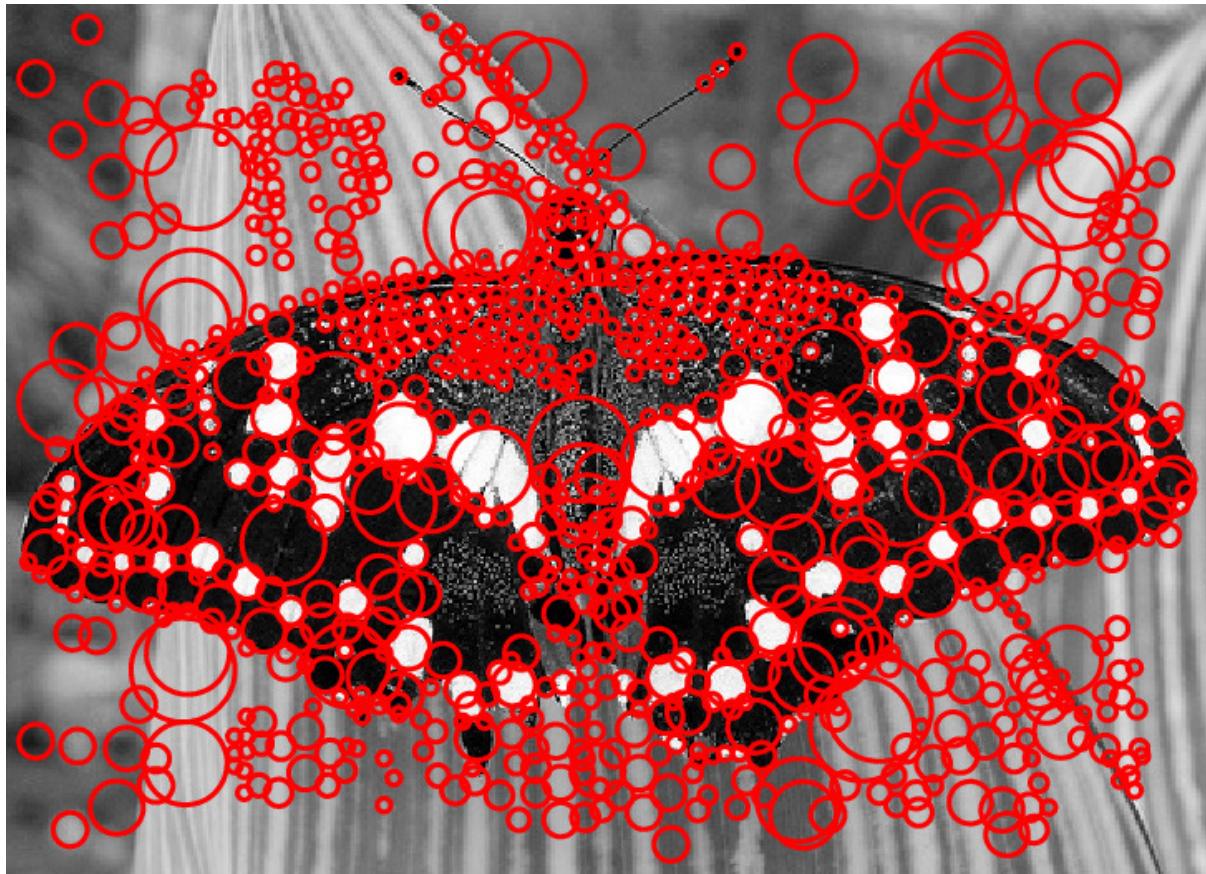
Recall:

$$[u \ v] M \begin{bmatrix} u \\ v \end{bmatrix} \triangleq \text{const}$$



This ellipse visualizes the “characteristic shape” of the window

Affine adaptation example



Scale-invariant regions (blobs)

Affine adaptation example



Affine-adapted blobs

2nd Nearest Neighbor Trick

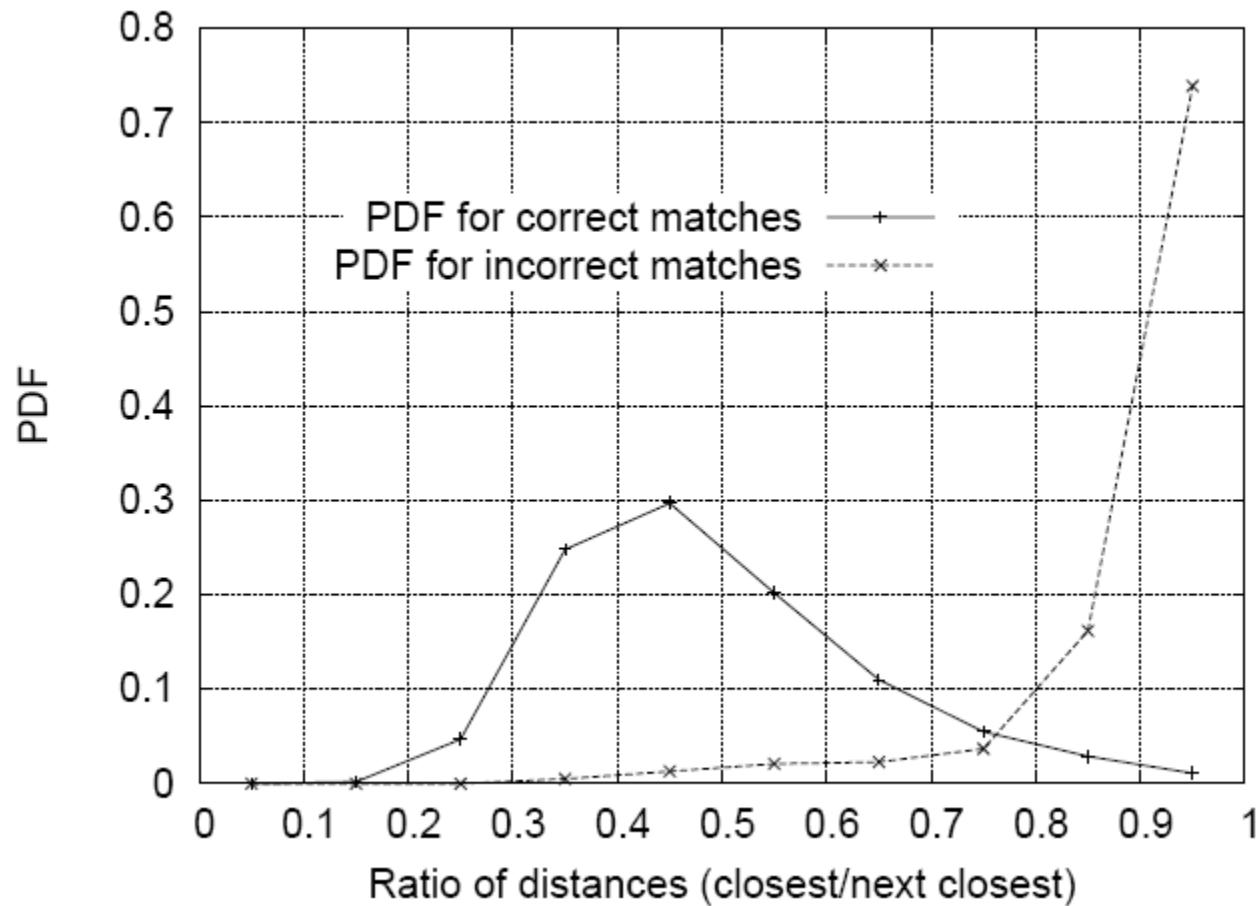


Figure from David G. Lowe. "[Distinctive image features from scale-invariant keypoints.](#)" *IJCV* 60 (2), pp. 91-110, 2004.