

Detectors and Descriptors

EECS 442 – Jeong Joon Park
Winter 2024, University of Michigan
<https://eecs442.github.io/>

Goal

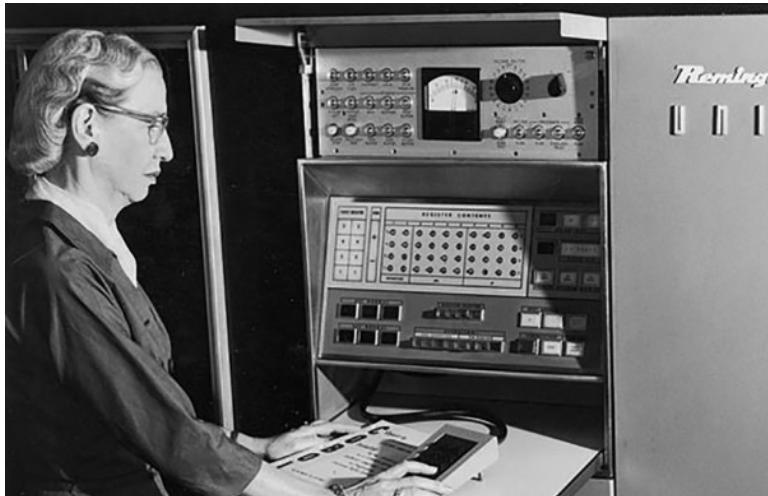
Extract Characteristics from this image

How big is this image as a vector?

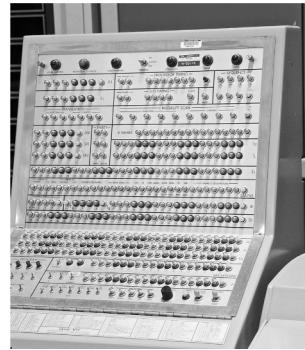
$389 \times 600 = 233,400$ dimensions (**big**)



Applications To Have In Mind



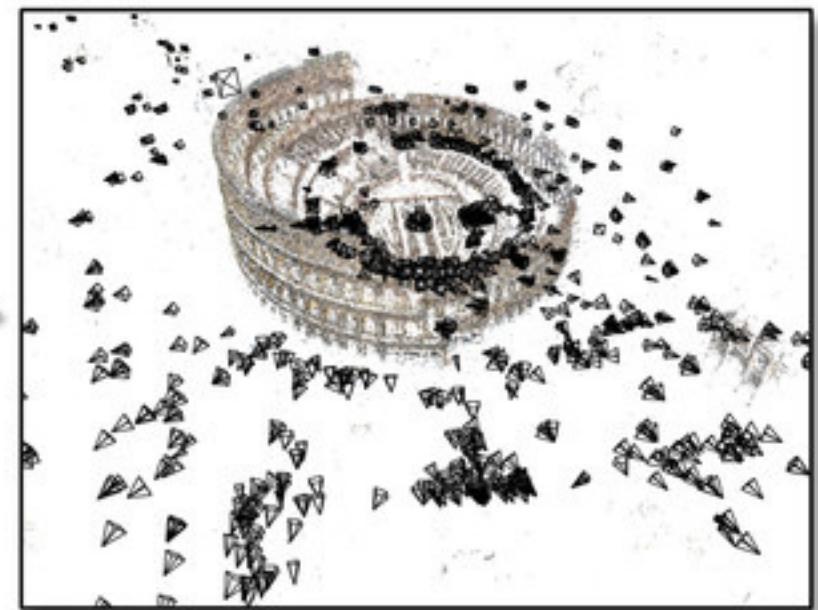
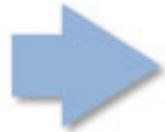
Part of the
same
photo?



Same
computer
from another
angle?

Applications To Have In Mind

Building a 3D Reconstruction Out Of Images



Applications To Have In Mind

(panorama)

Stitching photos taken at different angles



One Example

Given two images: how do you align them?



One Solution

```
for y in range(-ySearch,ySearch+1):  
    for x in range(-xSearch,xSearch+1):  
        #Touches all HxW pixels!  
        check_alignment_with_images()
```

One Motivating Example

Given these images: how do you align them?



These aren't off by a small 2D translation but instead by a 3D rotation + translation of the camera.

One Solution

```
for y in yRange:  
    for x in xRange:  
        for z in zRange:  
            for xRot in xRotVals:  
                for yRot in yRotVals:  
                    for zRot in zRotVals:  
                        #touches all HxW pixels!  
                        check_alignment_with_images()
```

This code should make you really unhappy

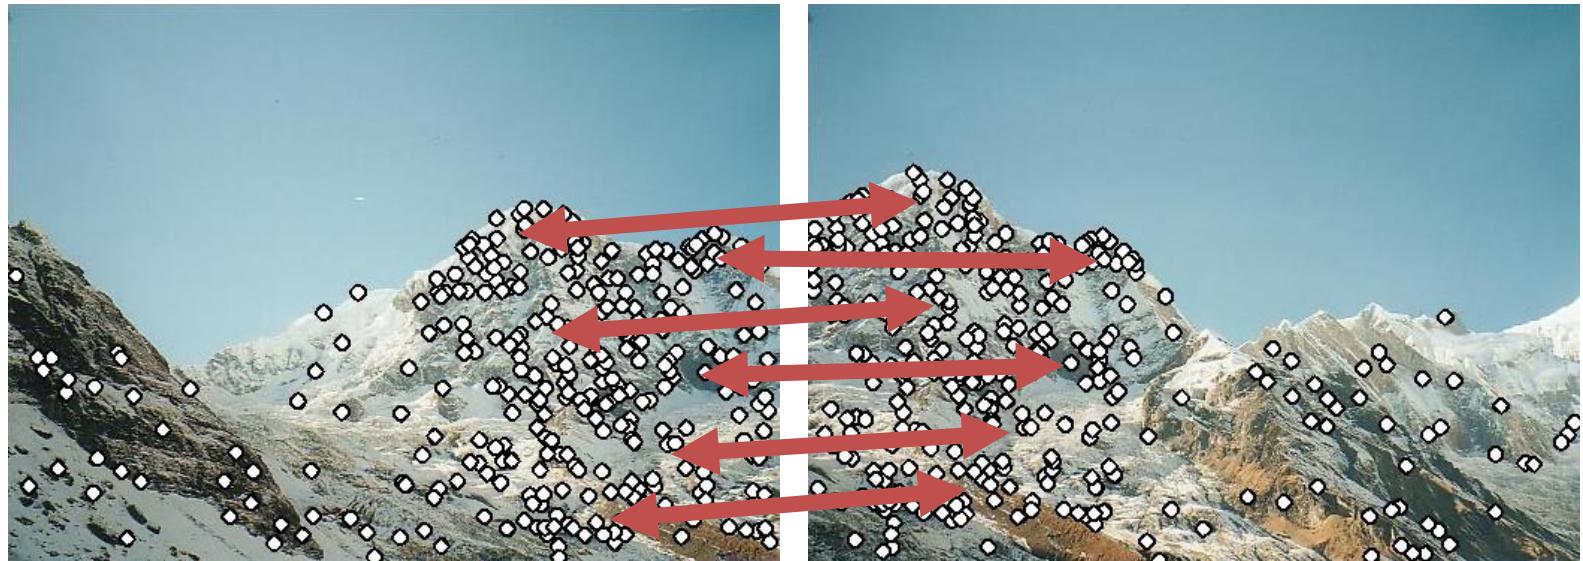
An Alternate Approach

Given these images: how would you align them?



An Alternate Approach

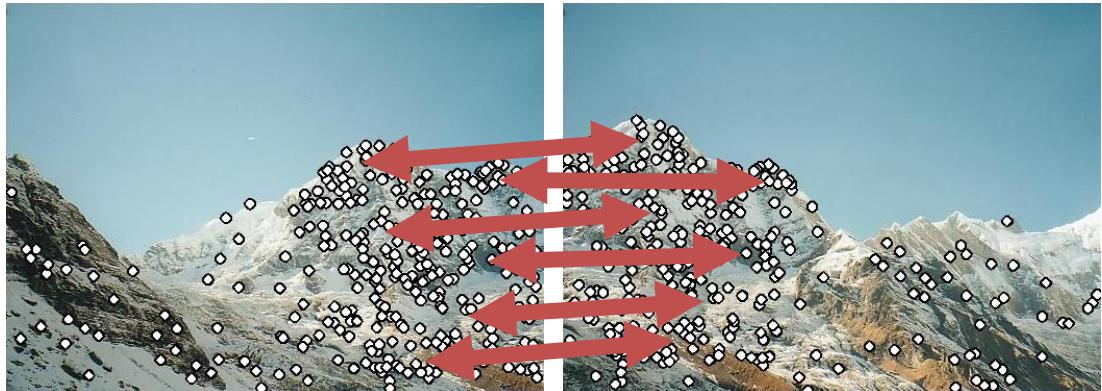
Finding and Matching



- 1: find corners+features
- 2: match based on local image data

What Now?

Given pairs
 p_1, p_2 of
correspondence,
how do I align?

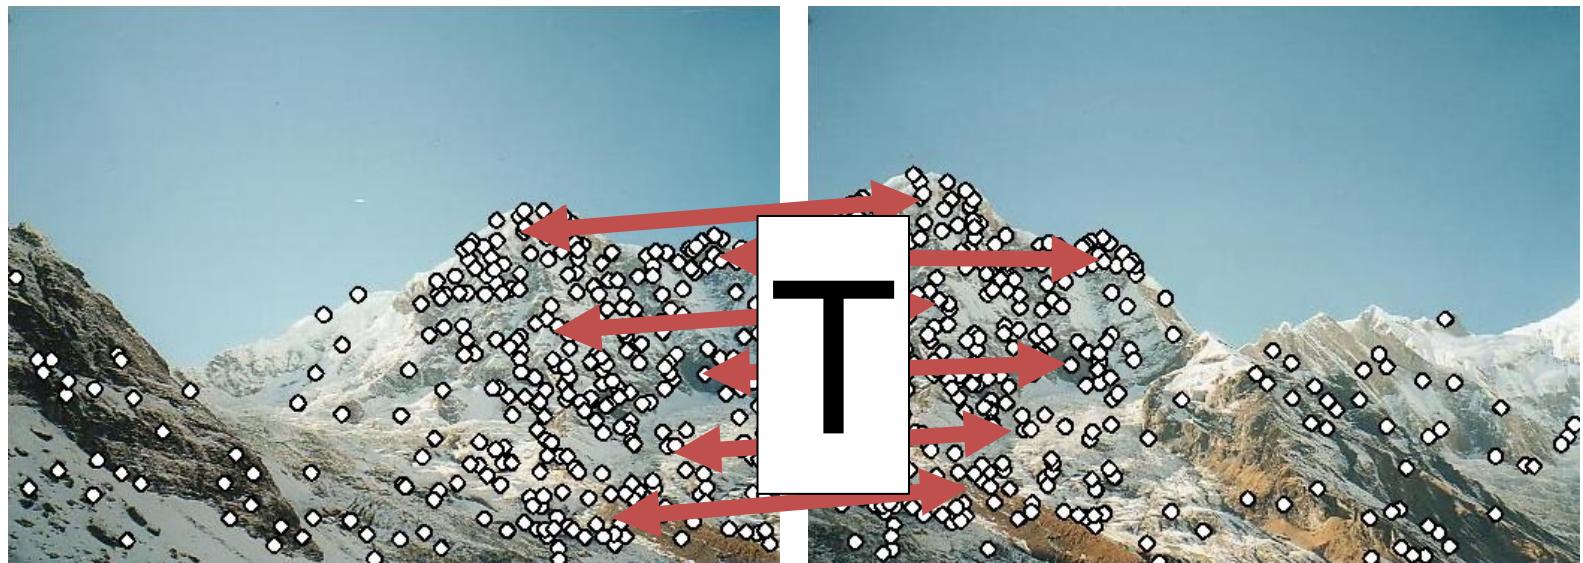


Consider translation-
only case.



An Alternate Approach

Solving for a Transformation



3: Solve for transformation T (e.g. such that $p_1 \equiv T p_2$) that fits the matches well

Note the homogeneous coordinates, you'll see them again.

An Alternate Approach

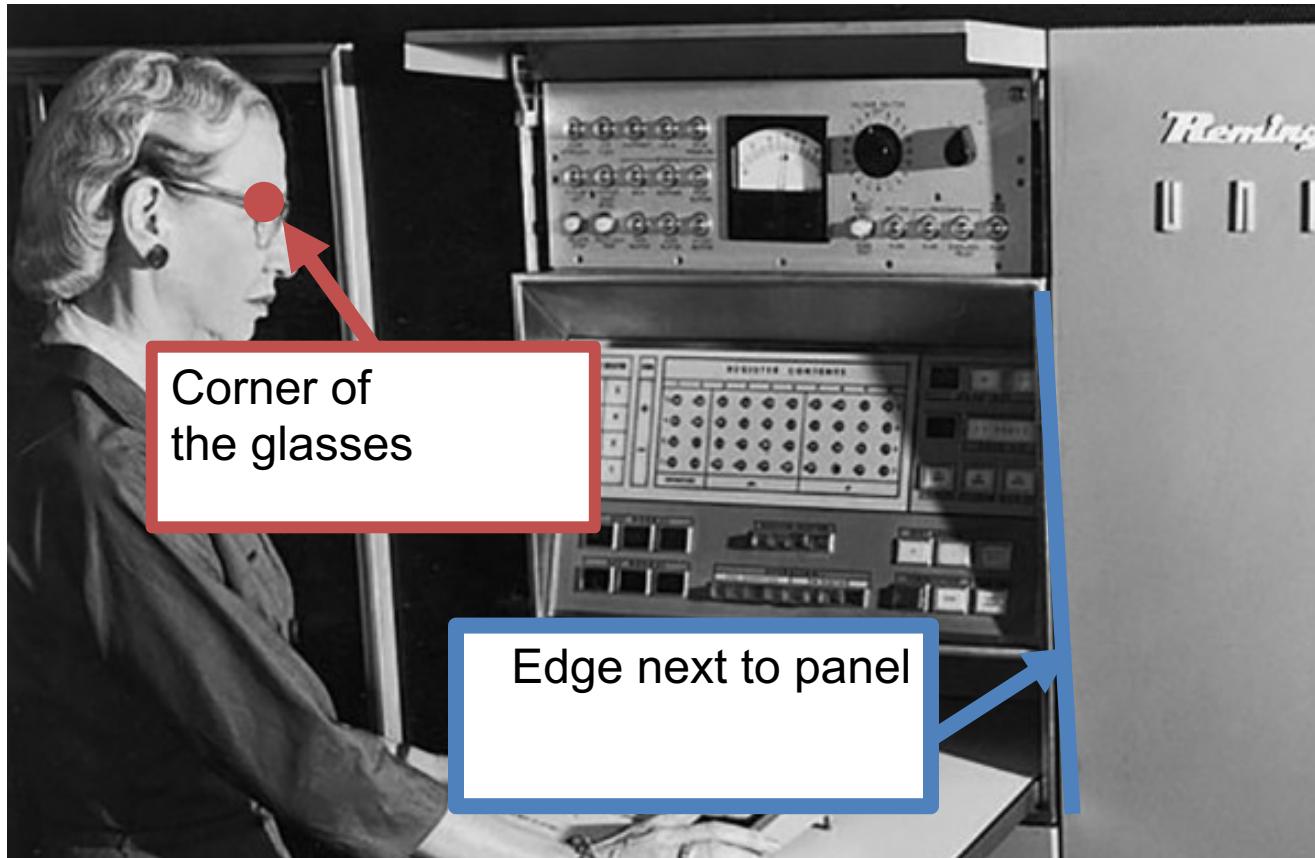
Blend Them Together



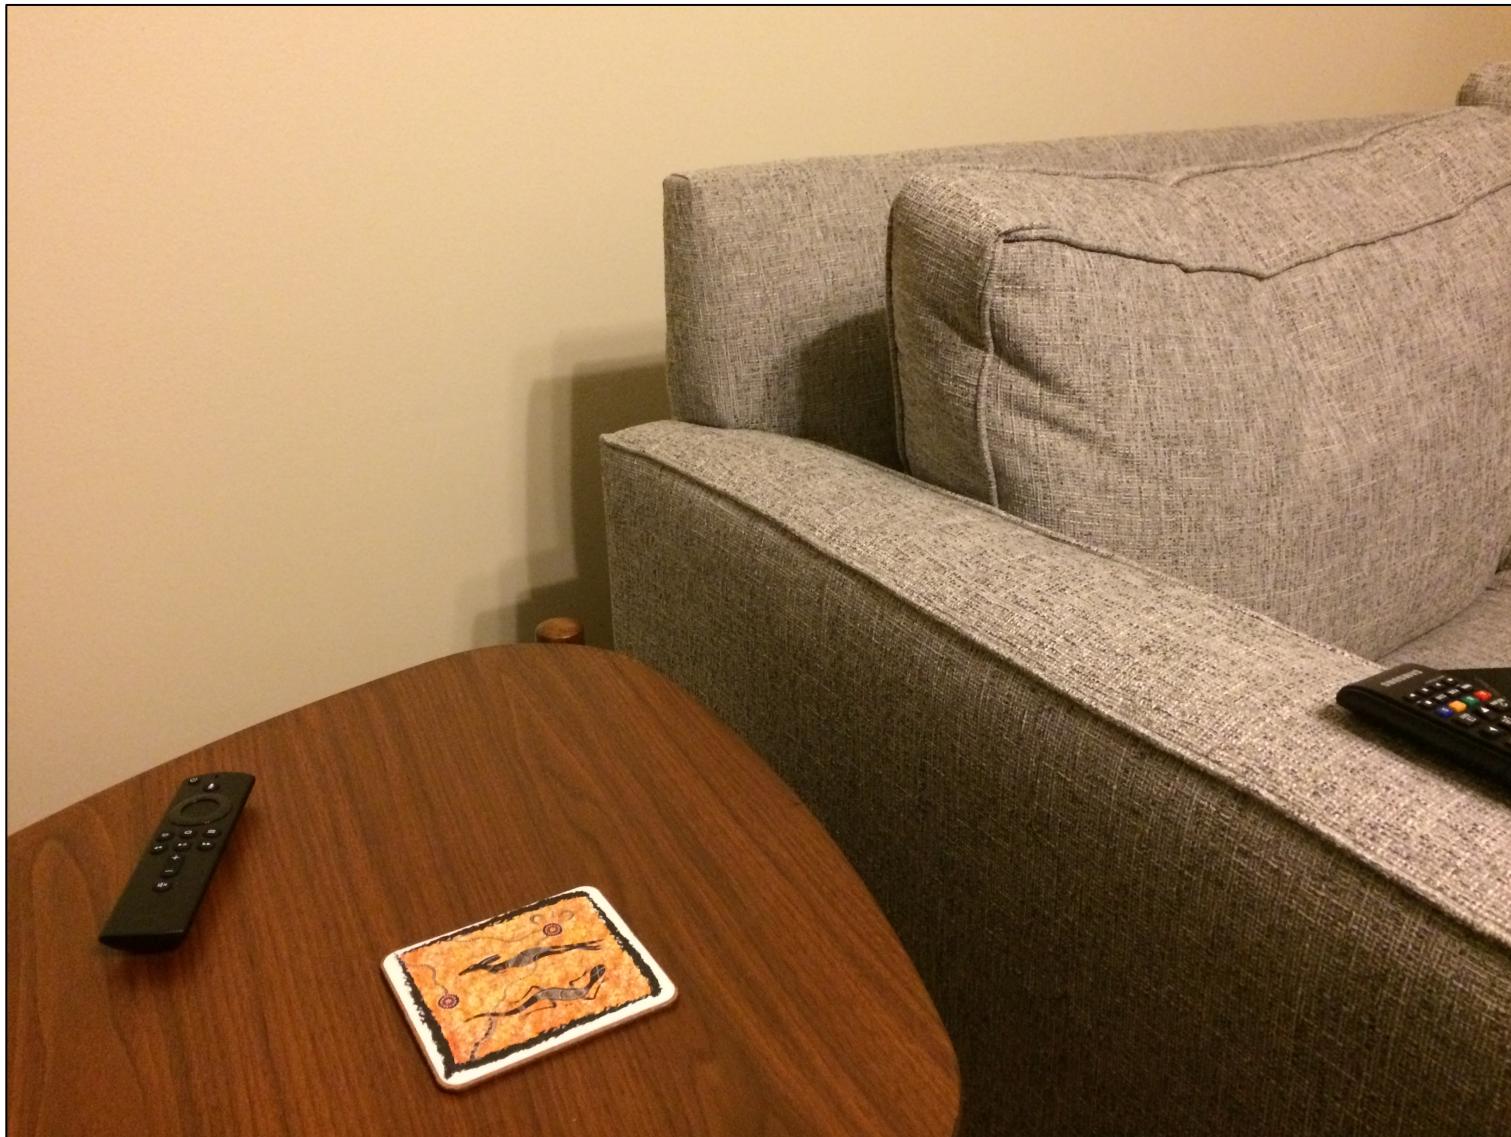
Key insight: we don't work with full image. We work with only parts of the image (features).

Today

Finding edges (part 1) and corners
(part 2) in images.



Where do Edges Come From?



Where do Edges Come From?

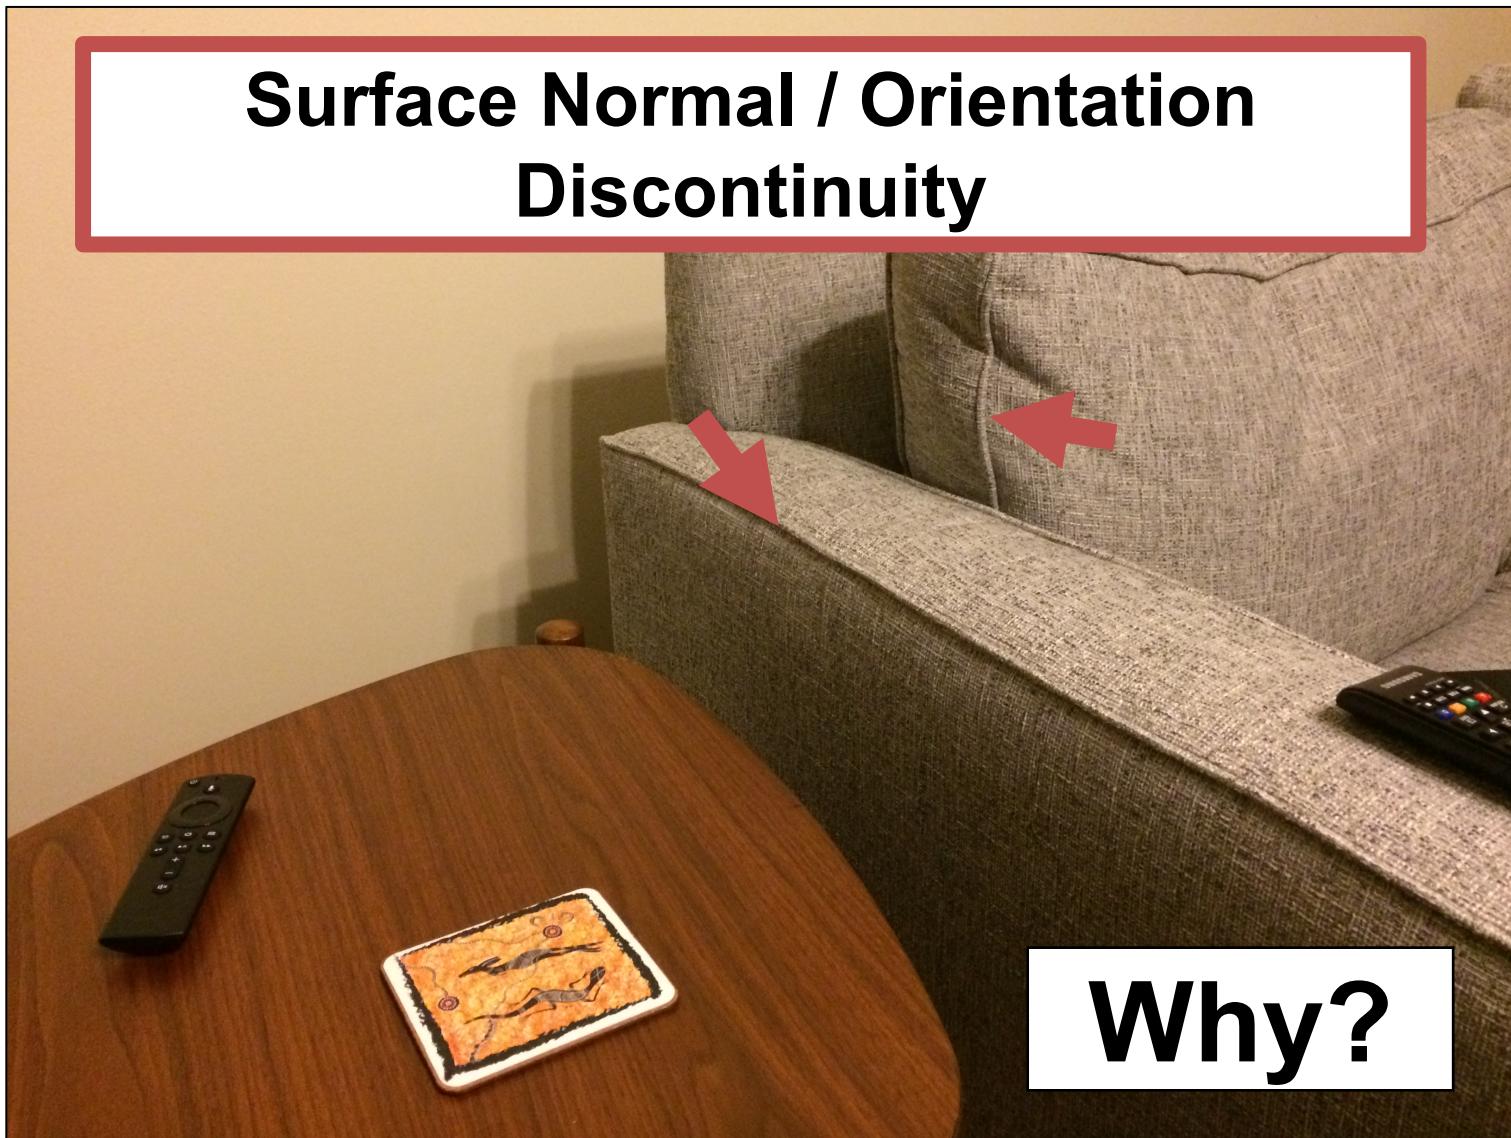
**Depth / Distance
Discontinuity**



Why?

Where do Edges Come From?

**Surface Normal / Orientation
Discontinuity**



Why?

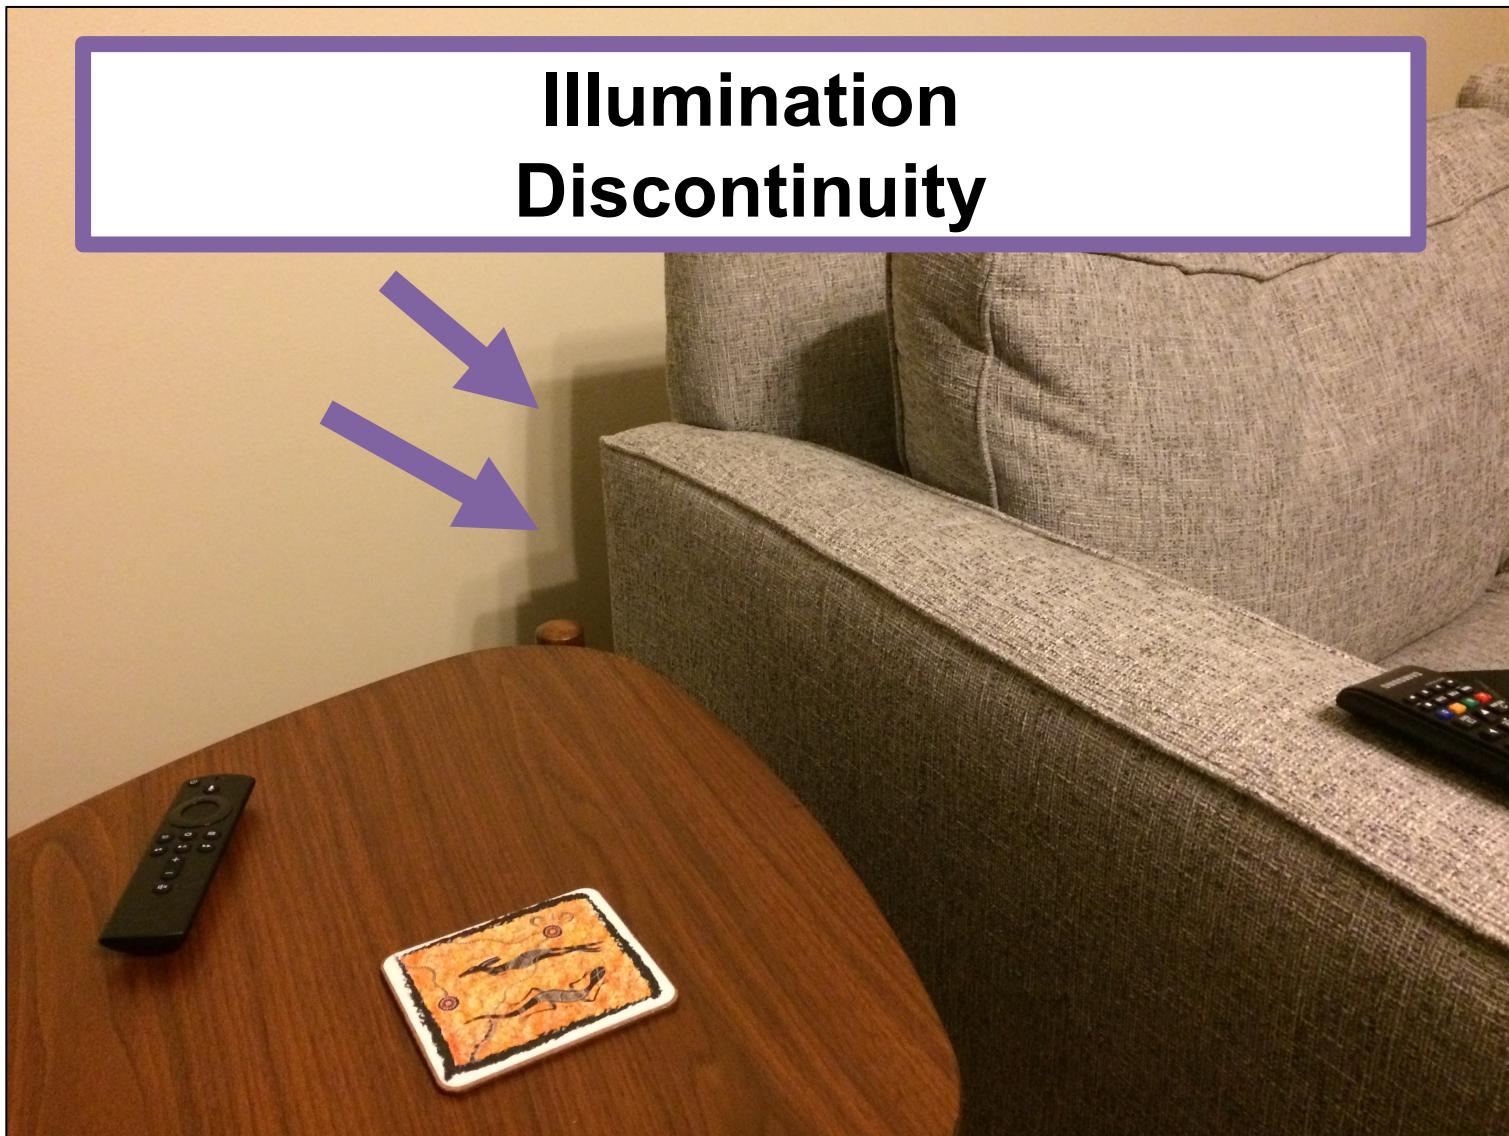
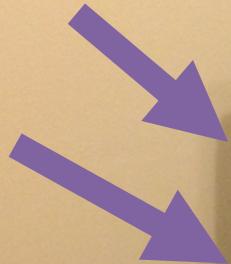
Where do Edges Come From?

**Surface Color / Reflectance
Properties Discontinuity**



Where do Edges Come From?

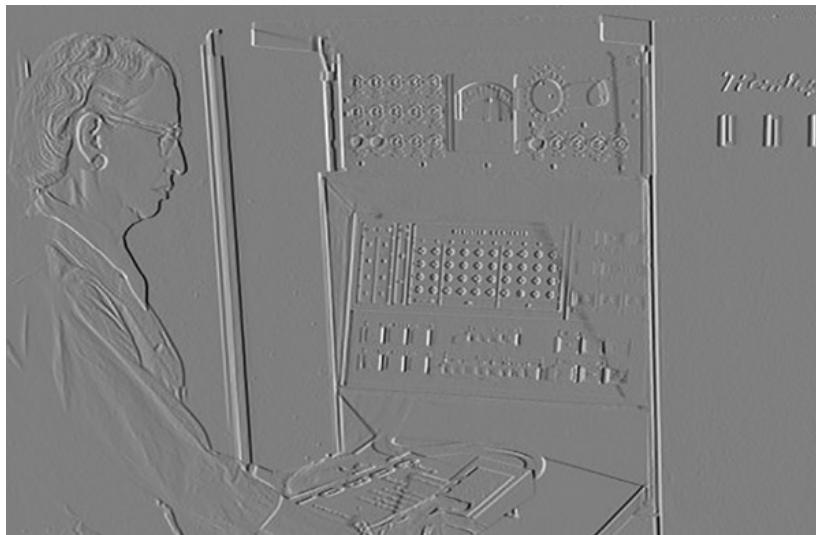
**Illumination
Discontinuity**



Last Time

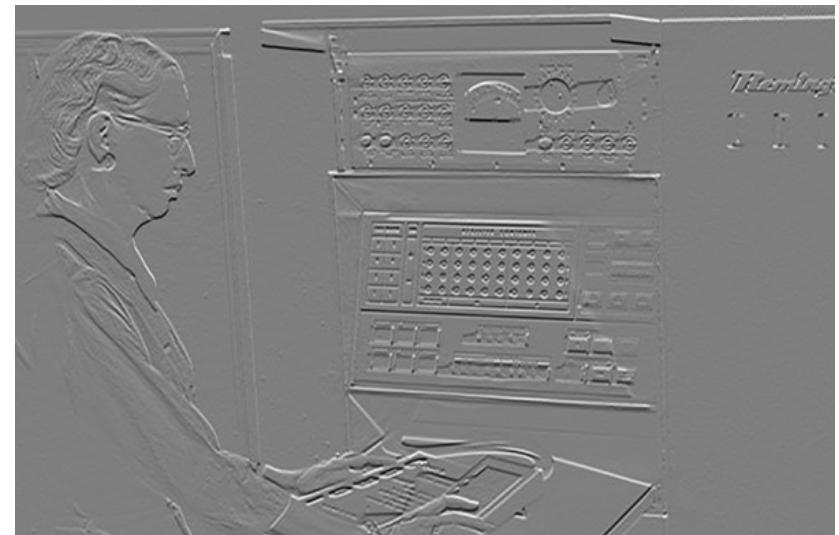
-1	0	1
----	---	---

|x



-1	0	1
----	---	---

|y



Derivatives

Remember derivatives?

Derivative: rate at which a function $f(x)$ changes at a point as well as the direction that increases the function

Gradient: all of the partial derivatives (derivatives in only one direction) stacked together.

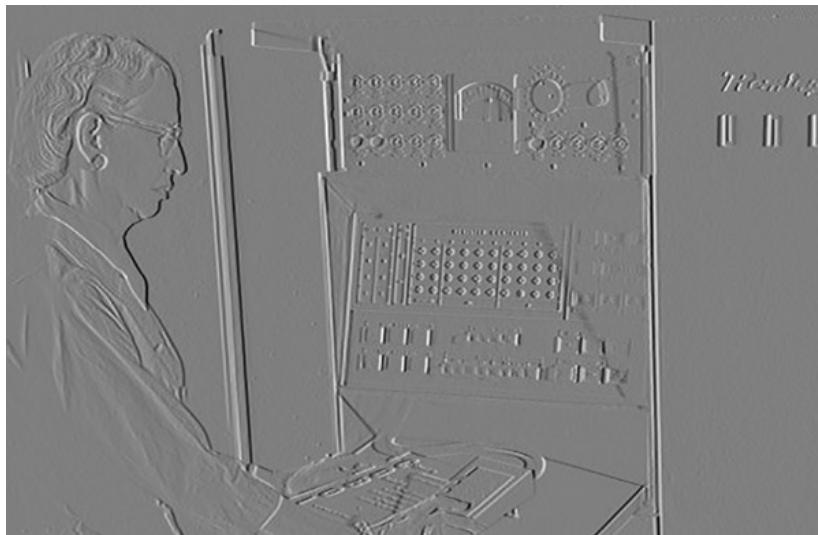
What Should I Know?

- Gradients are simply partial derivatives per-dimension: if x in $f(x)$ has n dimensions, $\nabla_f(x)$ has n dimensions
- Gradients point in direction of ascent and tell the rate of ascent
- If a is minimum of $f(x) \rightarrow \nabla_f(a) = 0$
- Reverse is not true, especially in high-dimensional spaces

Last Time

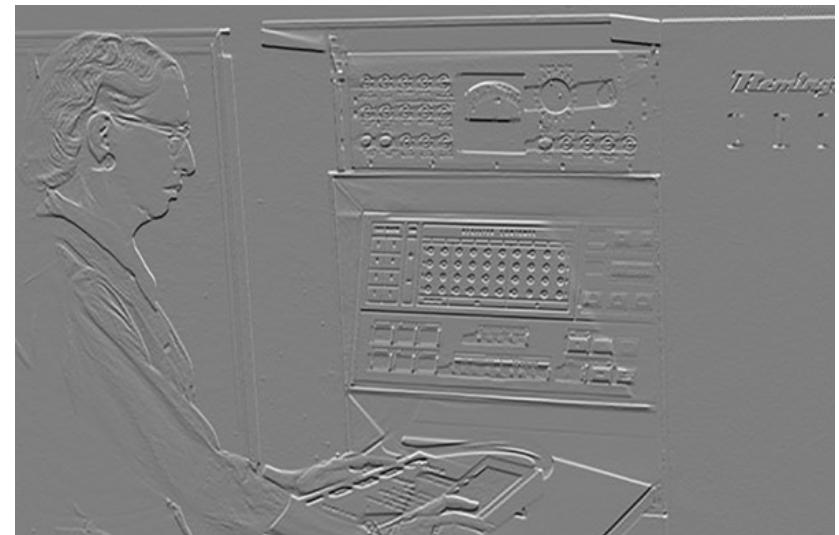
-1	0	1
----	---	---

|x



-1	0	1
----	---	---

|y



Why Does This Work?

Image is function $f(x,y)$

Remember:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\epsilon \rightarrow 0} \frac{f(x + \epsilon, y) - f(x, y)}{\epsilon}$$

Approximate:

-1	1
----	---

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x - 1, y)}{2}$$

Another one:

-1	0	1
----	---	---

The reason why use the 3*3 metric instead of 1*3 vector is that we can rely on more pixels and thus we can more reliably estimate the gradient

Other Differentiation Operations

	Horizontal	Vertical
Prewitt	$\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$	$\begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$
Sobel	$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$	$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Why might people use these compared to [-1,0,1]?

The function $f(x,y)$, input the x and y coordinates and output the intensity of the image at that given x and y position.

Images as Functions

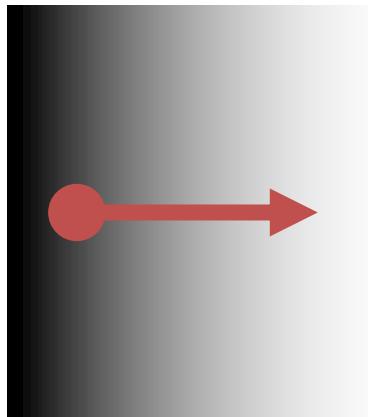
Key idea: can treat image a function of x,y .

$$\nabla I(x, y) = \begin{bmatrix} \frac{\partial I}{\partial x}(x, y) \\ \frac{\partial I}{\partial y}(x, y) \end{bmatrix}$$

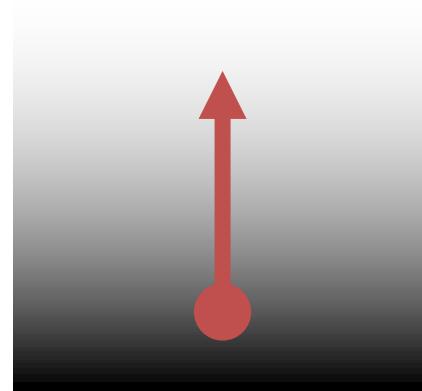
How much the intensity of the image changes as you go horizontally at (x,y)
(Let's call it “ I_x ”)

Image Gradient Direction

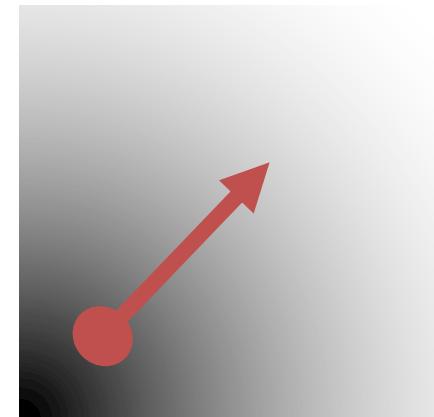
Some gradients



$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$



$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$

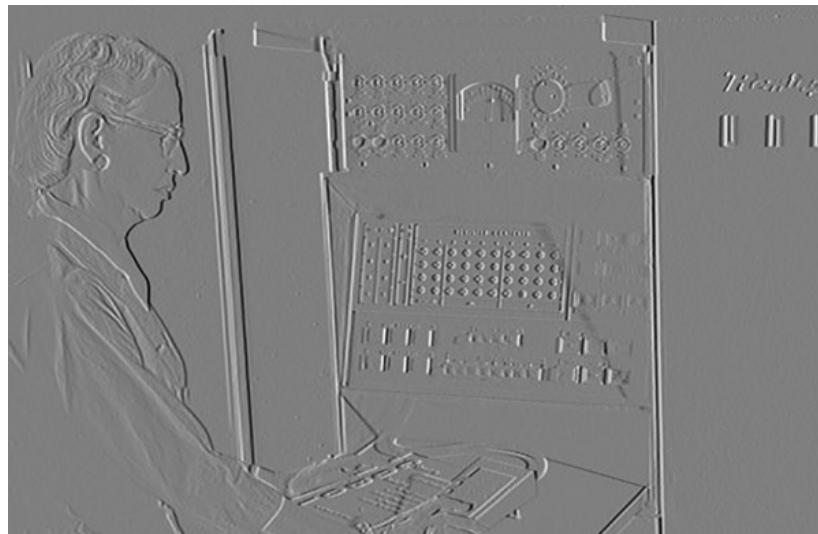


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

Image Gradient

Gradient: direction of maximum change.
What's the relationship to edge direction?

I_x



I_y

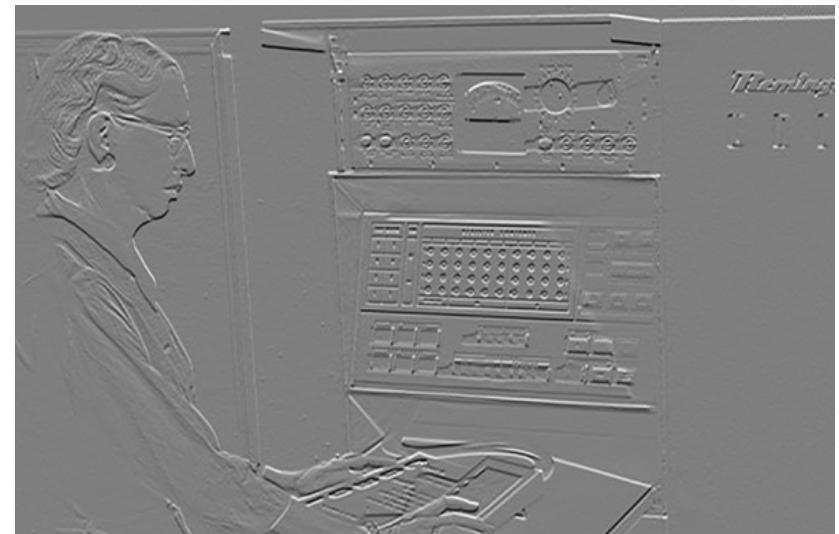


Image Gradient

$(Ix^2 + Iy^2)^{1/2}$: magnitude



Image Gradient

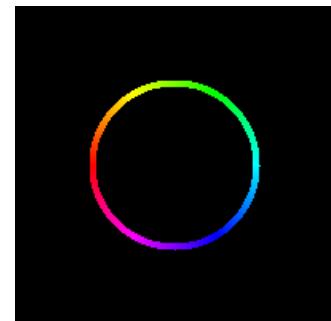
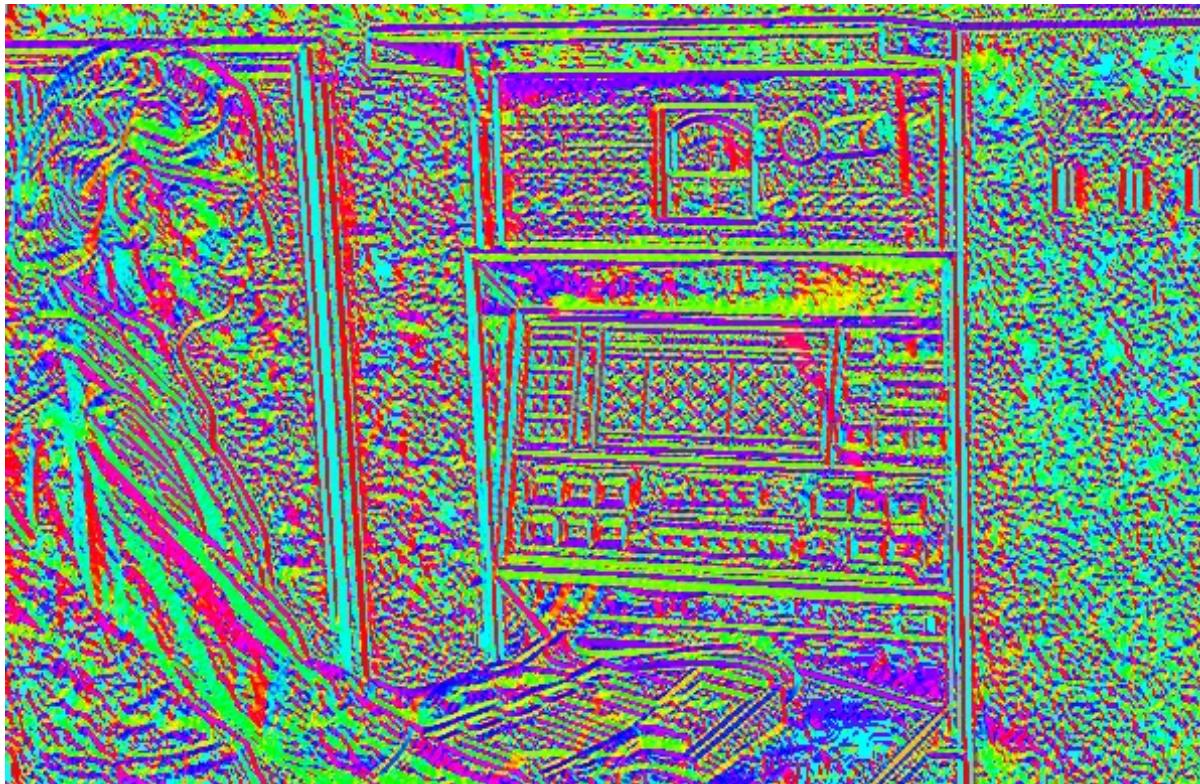
$\text{atan2}(I_y, I_x)$: orientation



I'm making the lightness equal to gradient magnitude

Image Gradient

$\text{atan2}(I_y, I_x)$: orientation

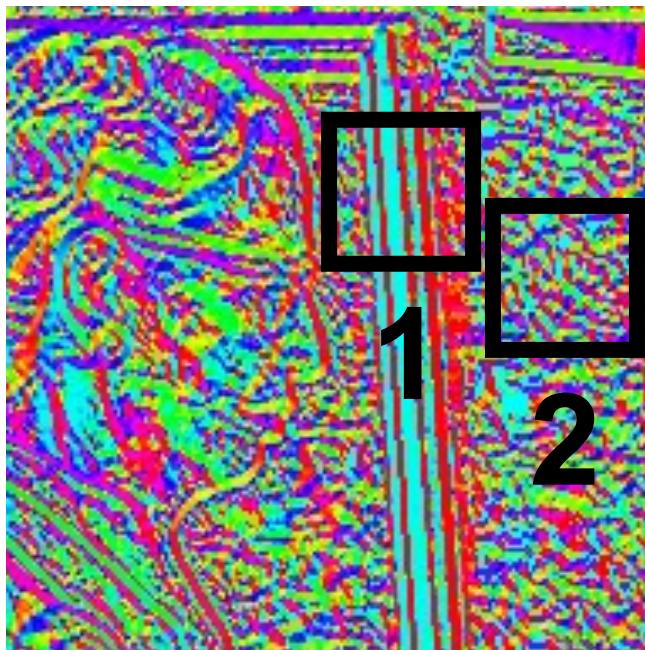


Now I'm showing *all* the gradients

Image Gradient

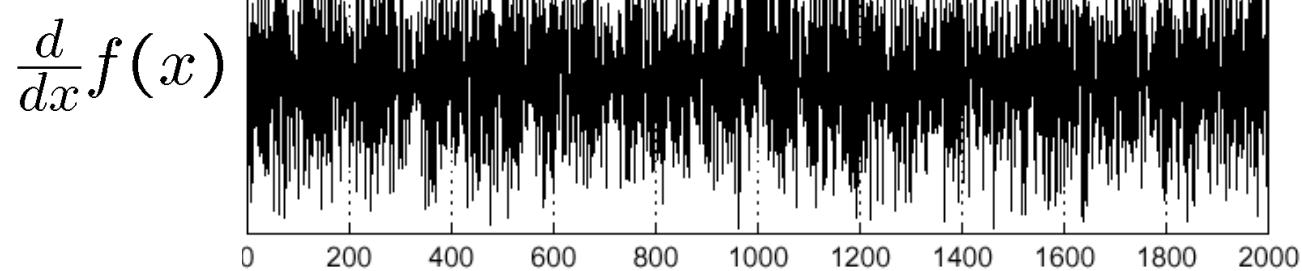
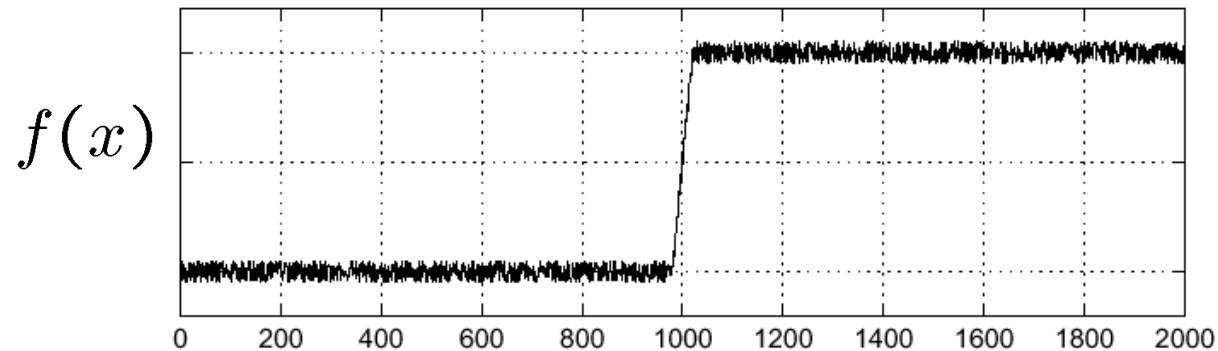
$\text{atan2}(I_y, I_x)$: orientation

Why is there structure at 1 and not at 2?



Noise

Consider a row of $f(x,y)$ (i.e., fix y)



Noise

Conv. image + per-pixel noise with 

$$I_{i,j} = \text{True image} \quad \epsilon_{i,j} \sim N(0, \sigma^2)$$

$$D_{i,j} = (I_{i,j+1} + \epsilon_{i,j+1}) - (I_{i,j-1} + \epsilon_{i,j-1})$$

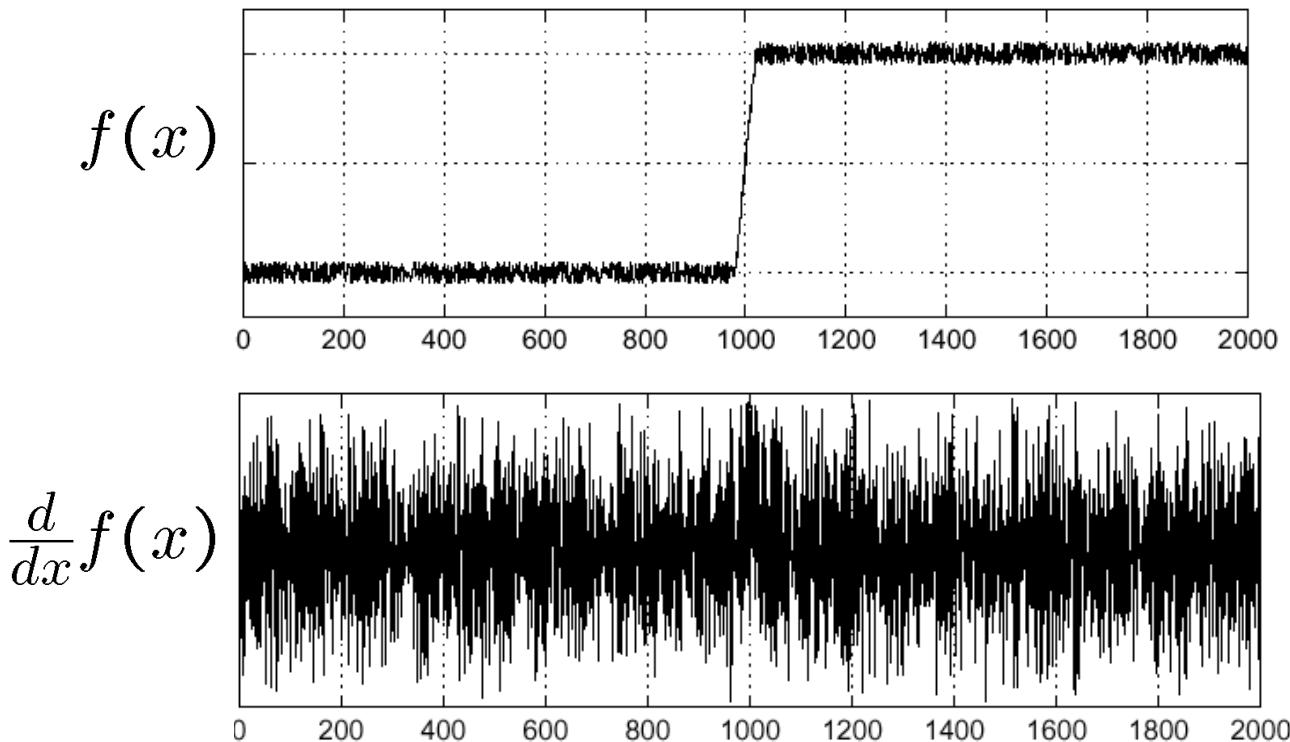
$$D_{i,j} = \underbrace{(I_{i,j+1} - I_{i,j-1})}_{\text{True difference}} + \underbrace{\epsilon_{i,j+1} - \epsilon_{i,j-1}}_{\text{Sum of 2 Gaussians}}$$

True difference Sum of 2 Gaussians

$$\epsilon_{i,j} - \epsilon_{k,l} \sim N(0, 2\sigma^2) \rightarrow \text{Variance doubles!}$$

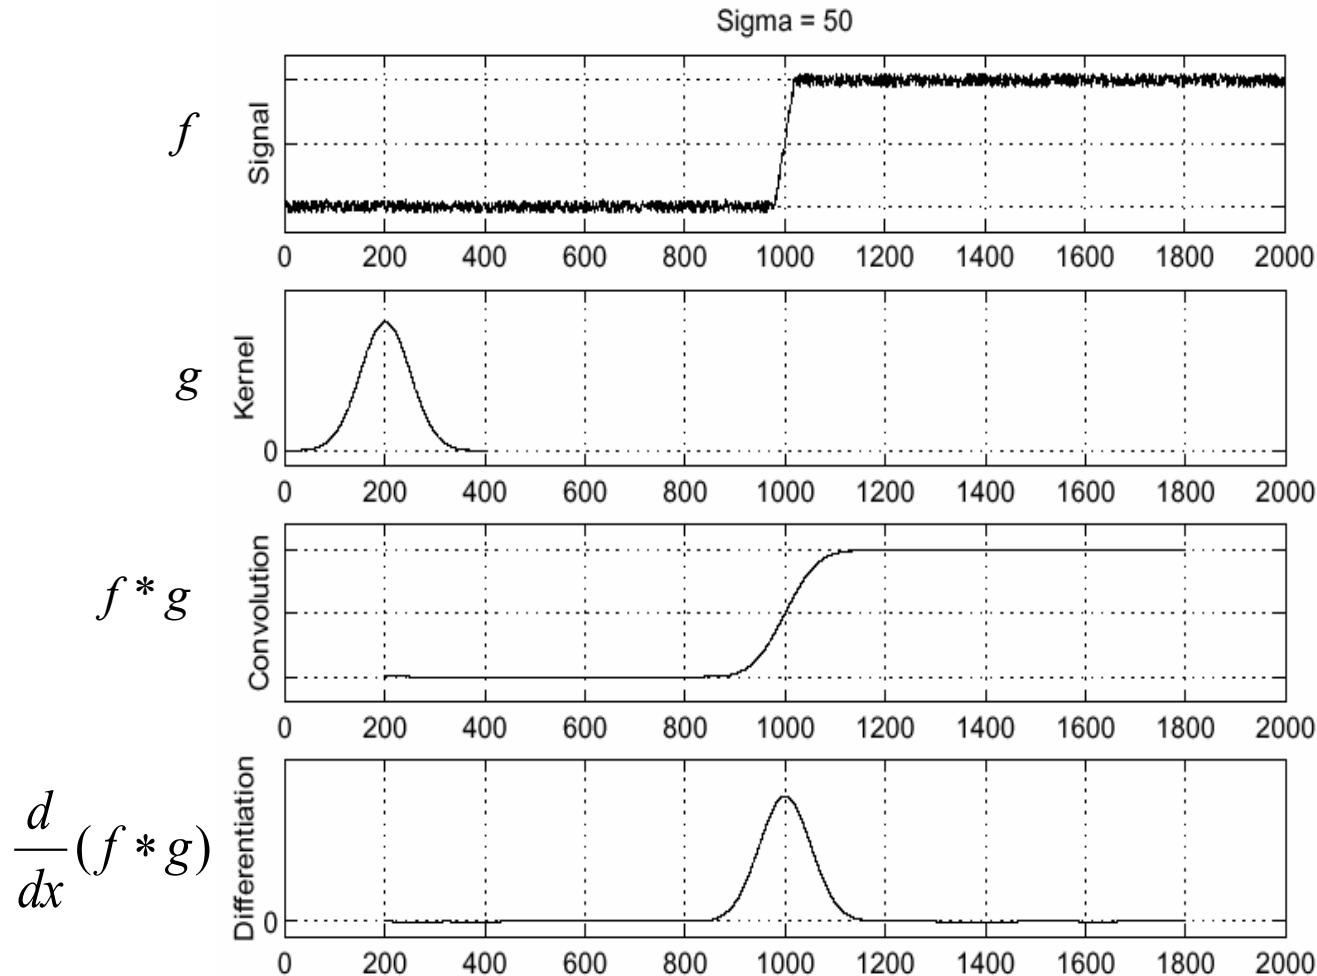
Noise

Consider a row of $f(x,y)$ (i.e., make y constant)



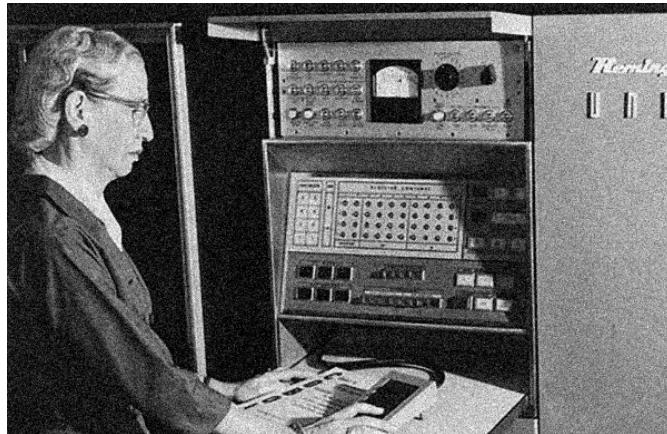
How can we use the last class to fix this?

Handling Noise

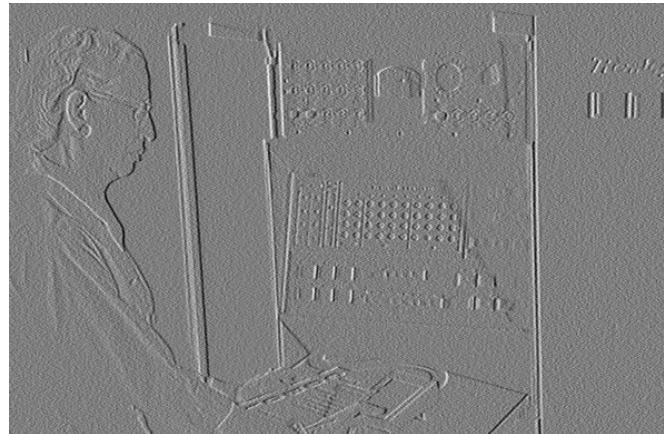


Noise in 2D

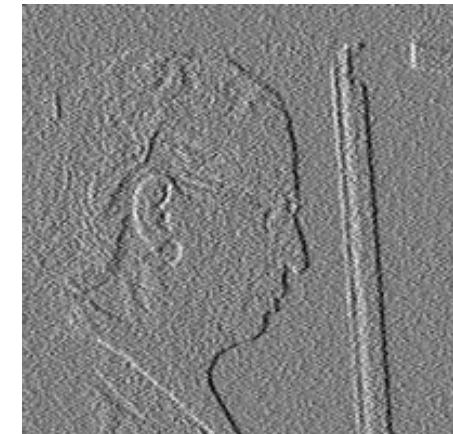
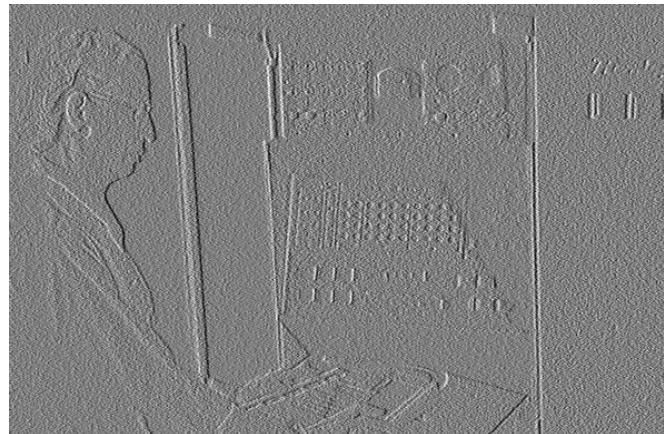
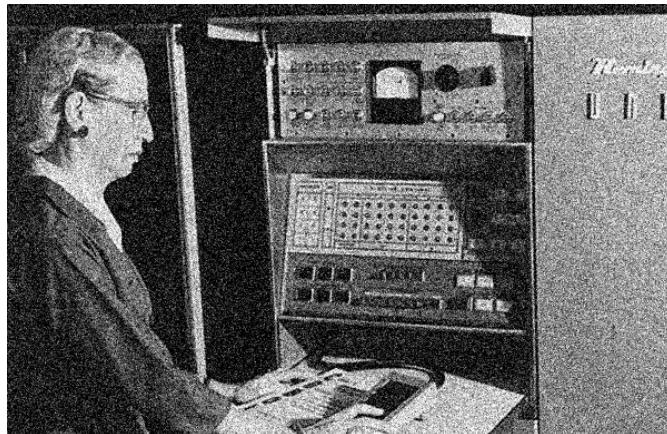
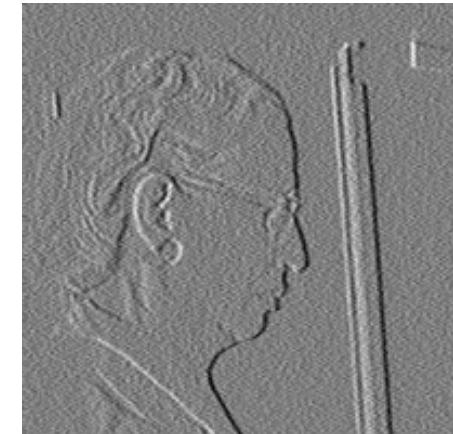
Noisy Input



I_x via $[-1,01]$



Zoom

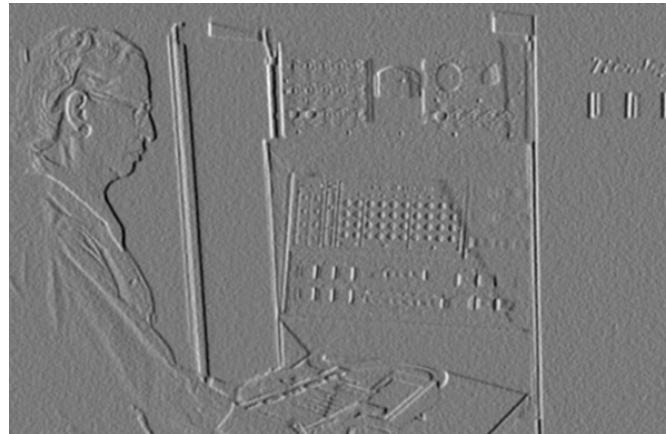


Noise + Smoothing

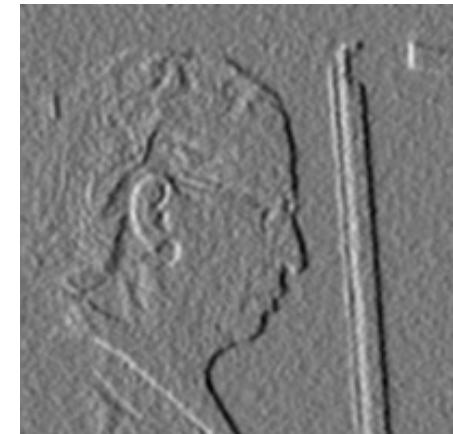
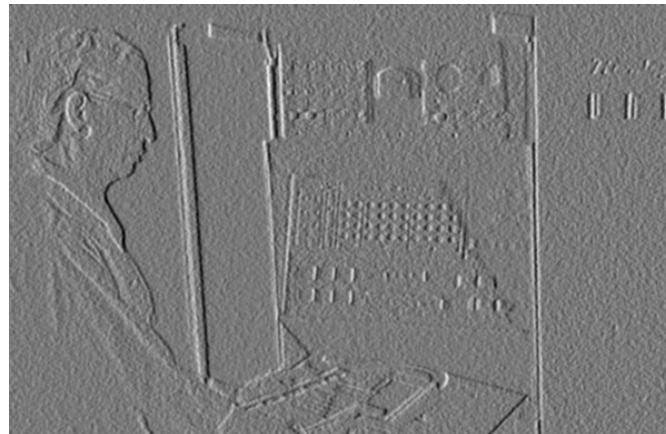
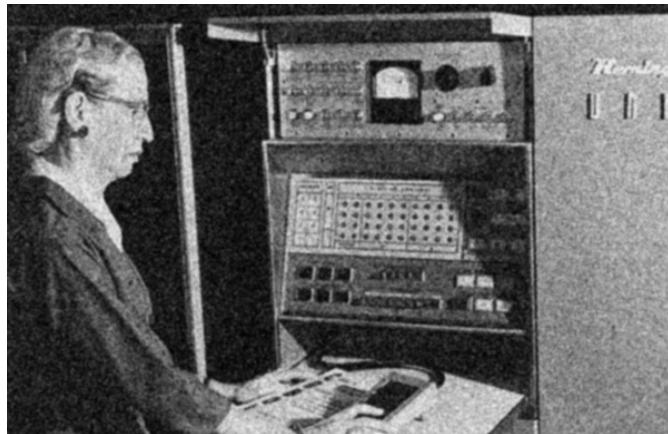
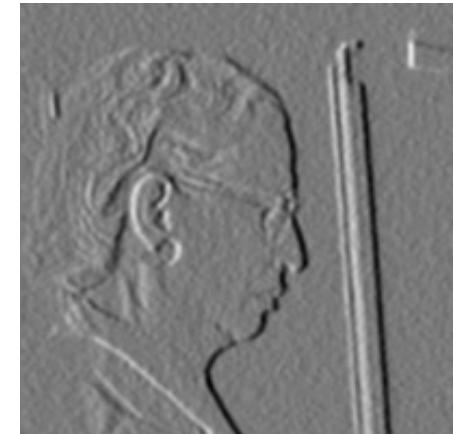
Smoothed Input



Ix via [-1,01]



Zoom

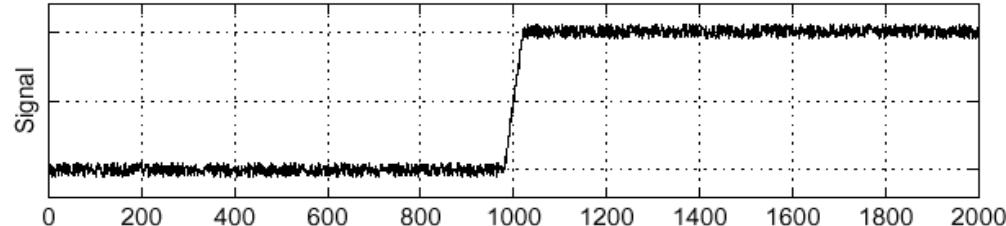


Let's Make It One Pass (1D)

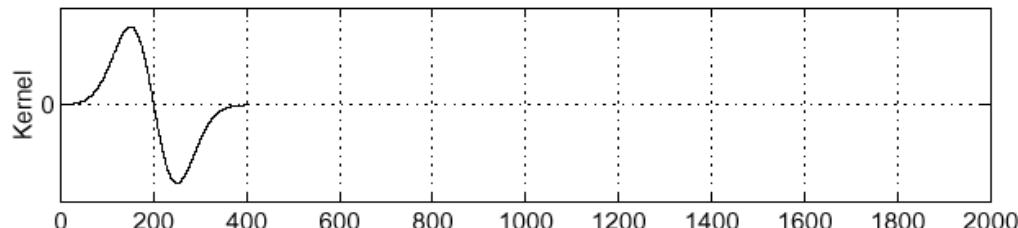
$$\frac{d}{dx} (f * g) = f * \frac{d}{dx} g$$

Sigma = 50

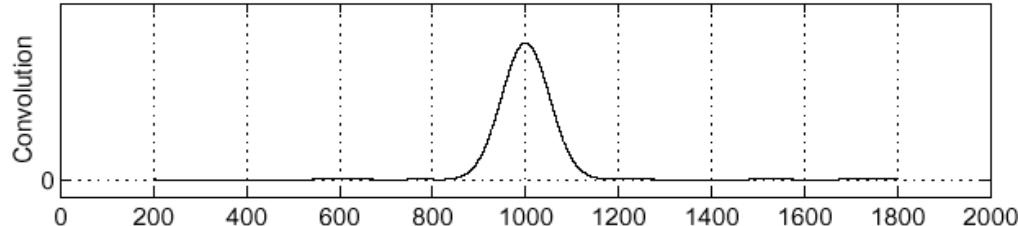
f



$\frac{d}{dx} g$



$f * \frac{d}{dx} g$



Proof of $\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$

Definition of Convolution

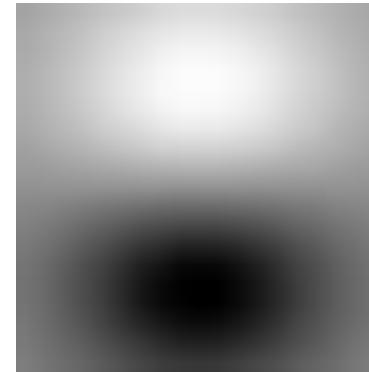
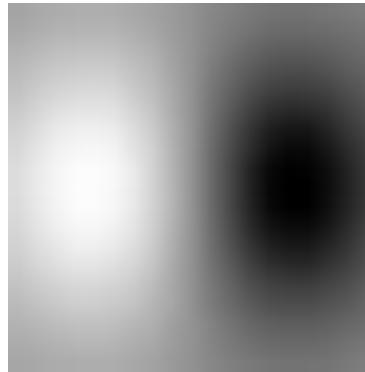
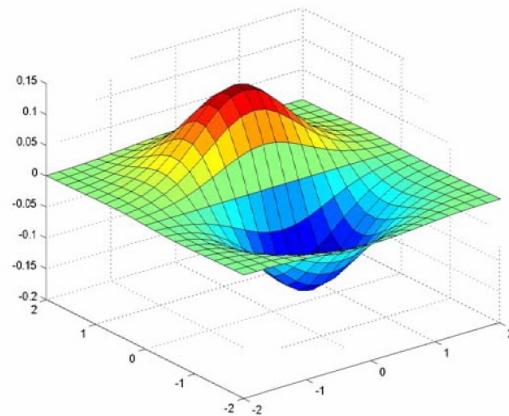
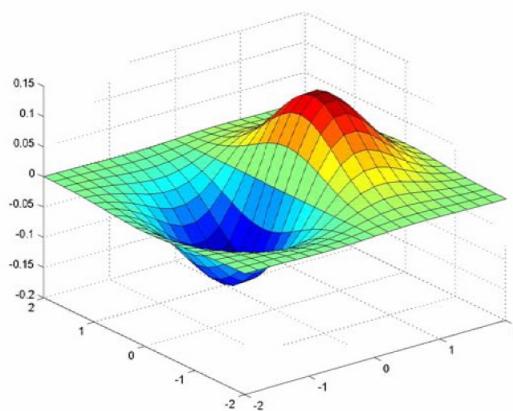
$$(f * g)(x) = \int_{-\infty}^{+\infty} f(t) \cdot g(x - t) dt$$

Using Linearity of Differential Operator

$$\frac{d}{dx}(f * g)(x) = \frac{d}{dx} \int_{-\infty}^{+\infty} f(t) \cdot g(x - t) dt = \int_{-\infty}^{+\infty} f(t) \cdot \frac{d}{dx}g(x - t) dt$$

Let's Make It One Pass (2D)

Gaussian Derivative Filter



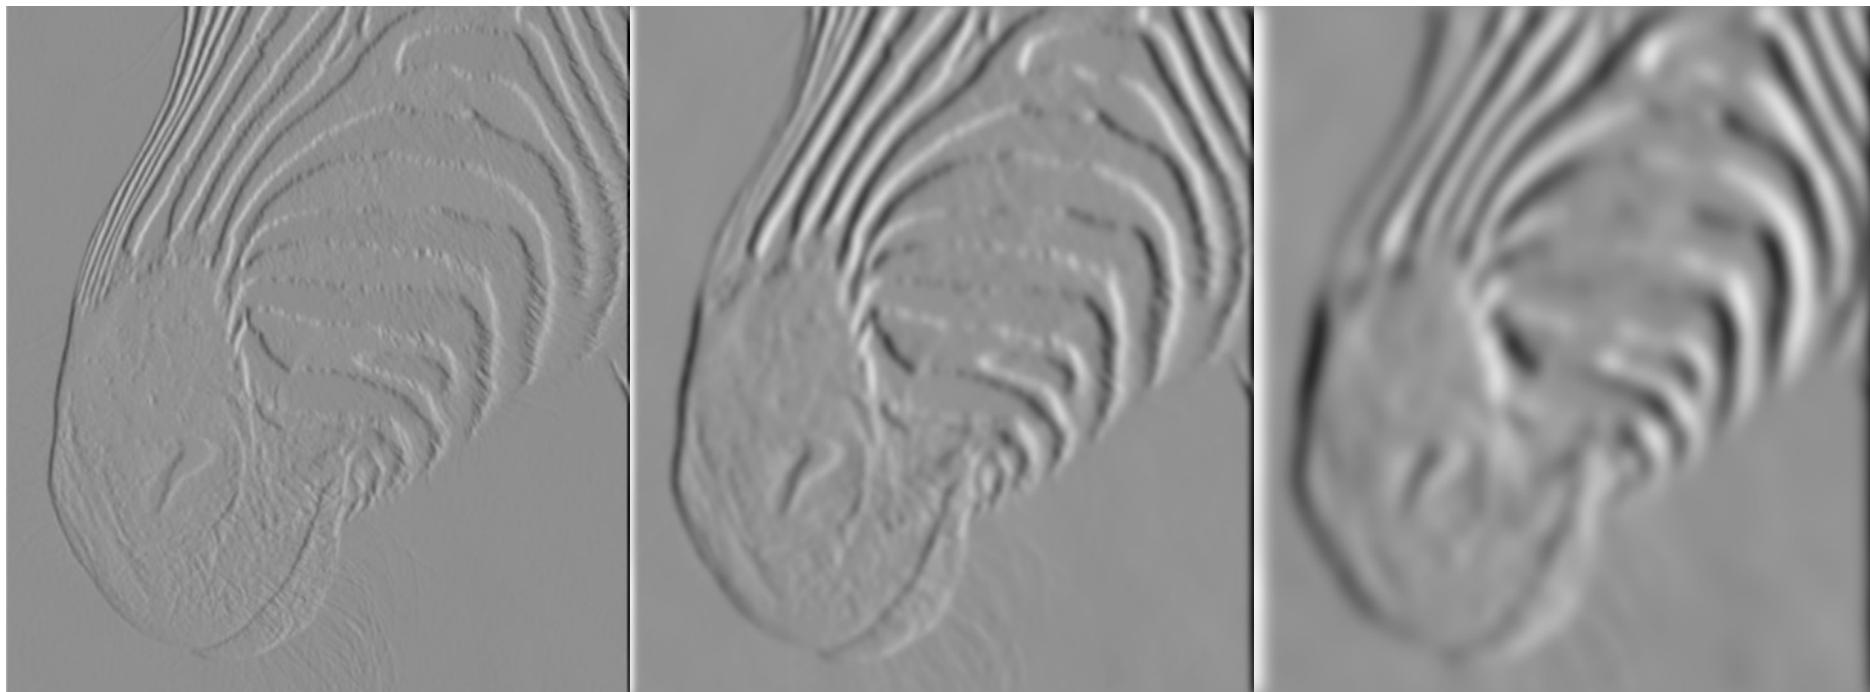
Which one finds the X direction?

Applying the Gaussian Derivative

SD: 1 pixel

3 pixels

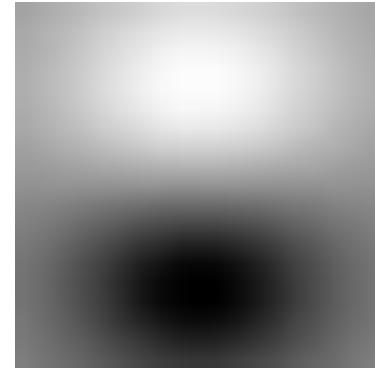
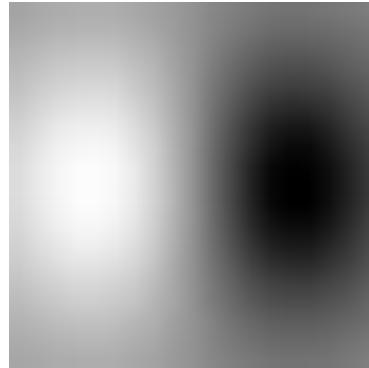
7 pixels



Removes noise, but blurs edge

Compared with the Past

Gaussian
Derivative



Sobel
Filter

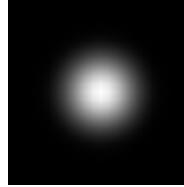
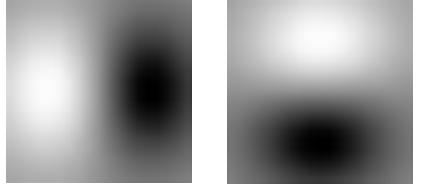
$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Why would anybody use the bottom filter?

1. The Gaussian derivative is expensive to compute.
2. only the liner filter we can do the derivative of the filter first

Filters We've Seen

	Smoothing	Derivative
Example	 A 3x3 Gaussian smoothing filter kernel, centered white pixel with a black background.	 Two 3x3 derivative of Gaussian filter kernels, showing horizontal and vertical gradients.
Goal	Gaussian	Deriv. of gauss
Only +?	Remove noise	Find edges
Sums to	Yes	No
	1	0

Why sum to 1 or 0, intuitively?

Don't want to change the intense of the picture.

Slide Credit: J. Deng

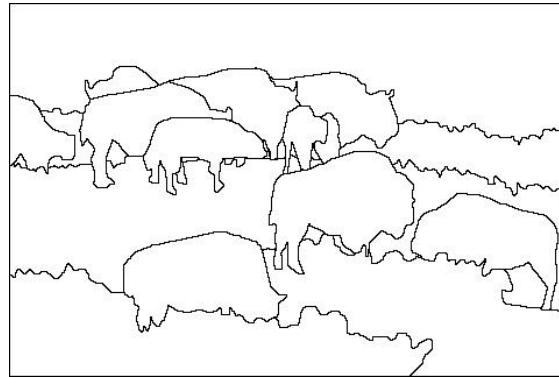
We want to left and right to be symmetric when we do the compare.

Problems

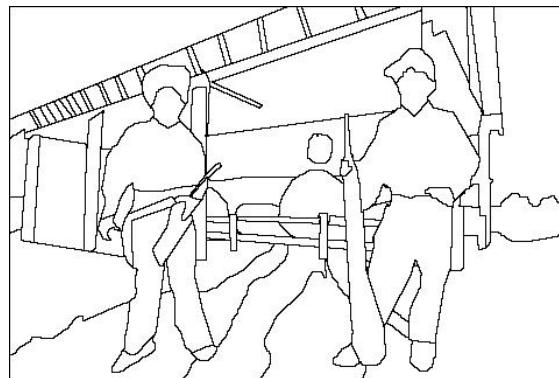
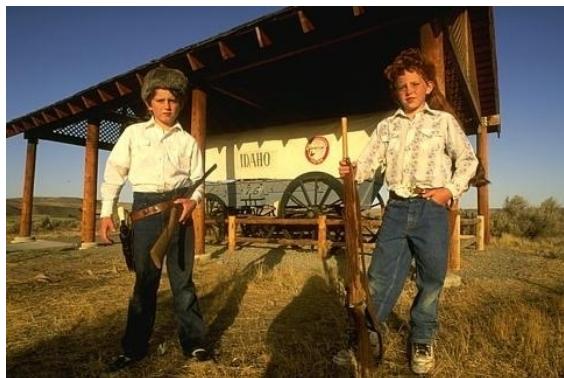
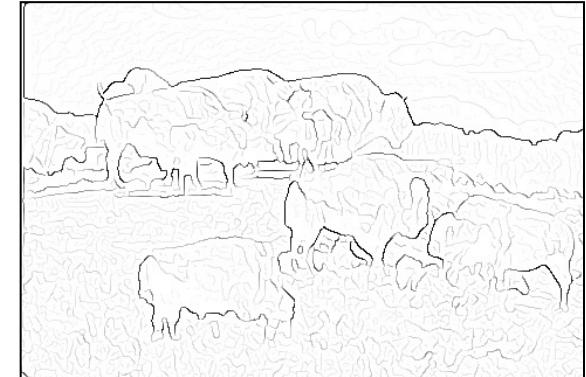
Image



human segmentation



gradient magnitude



Still an unsolved problem

Localizing Reliably

- Suppose you need to meet someone but you can't use your cell phone to coordinate
- Where do you agree to meet?

A: Along the Huron river

B: Along State Street

C: At Liberty and State Street

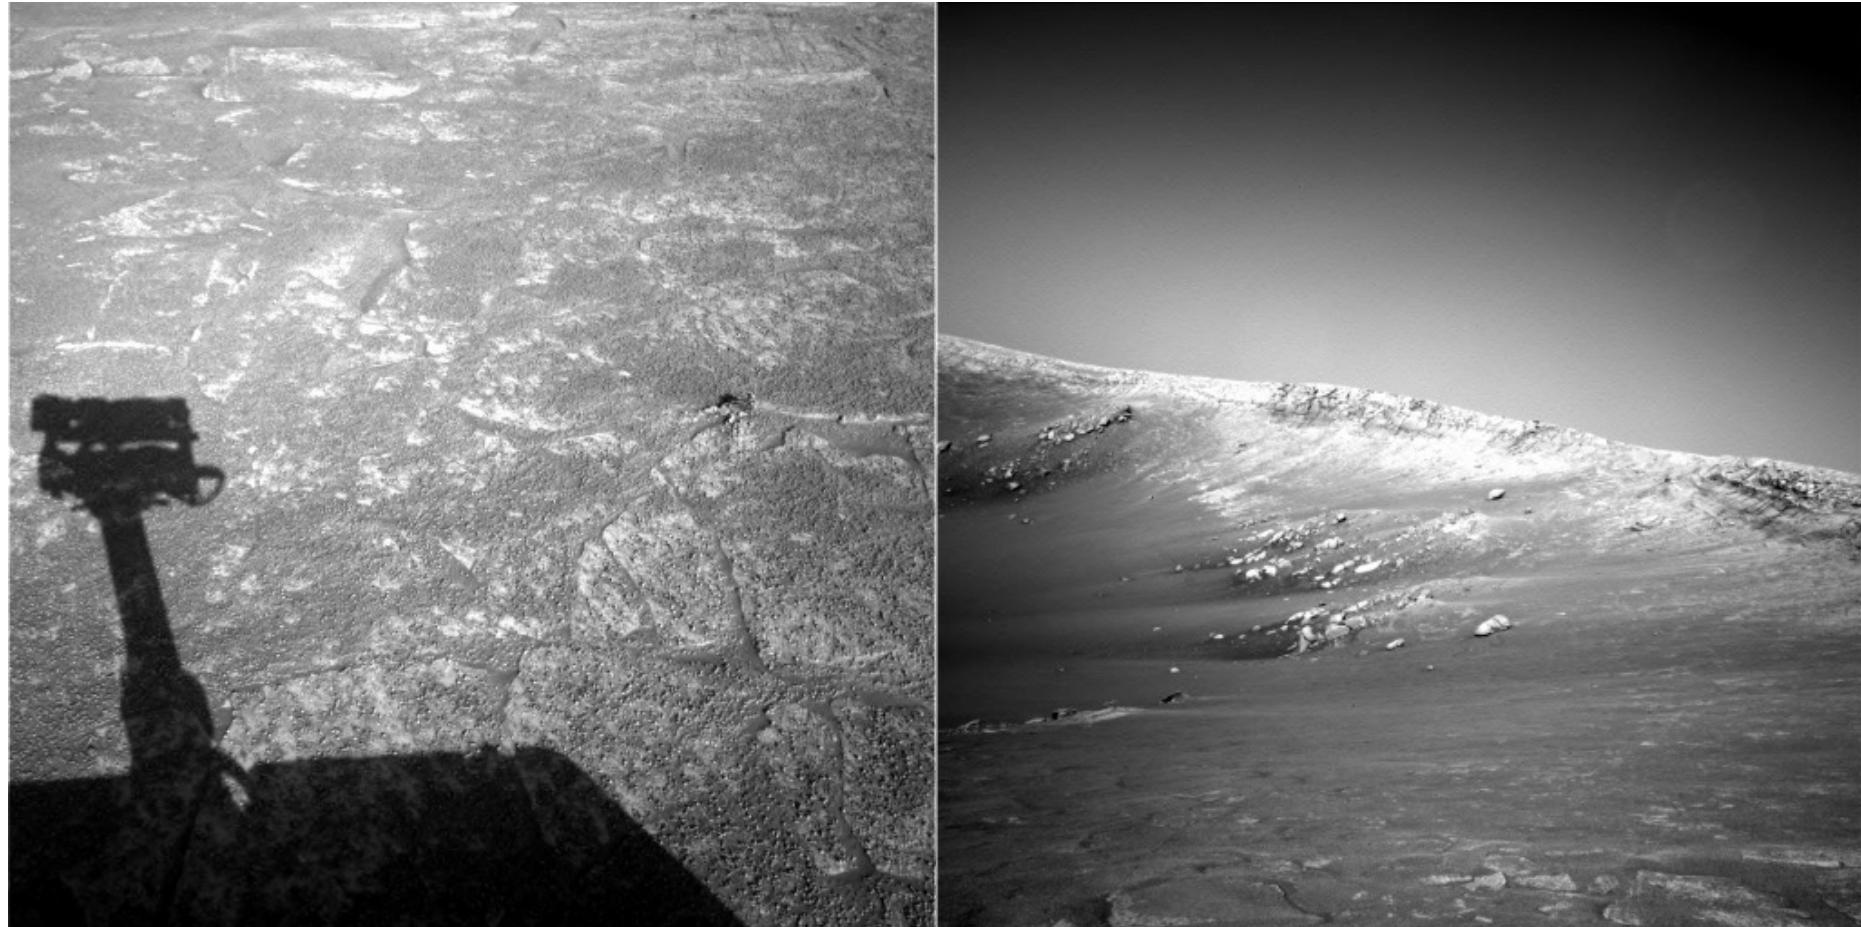
D: On North Campus

Corners over a line!

Desirables

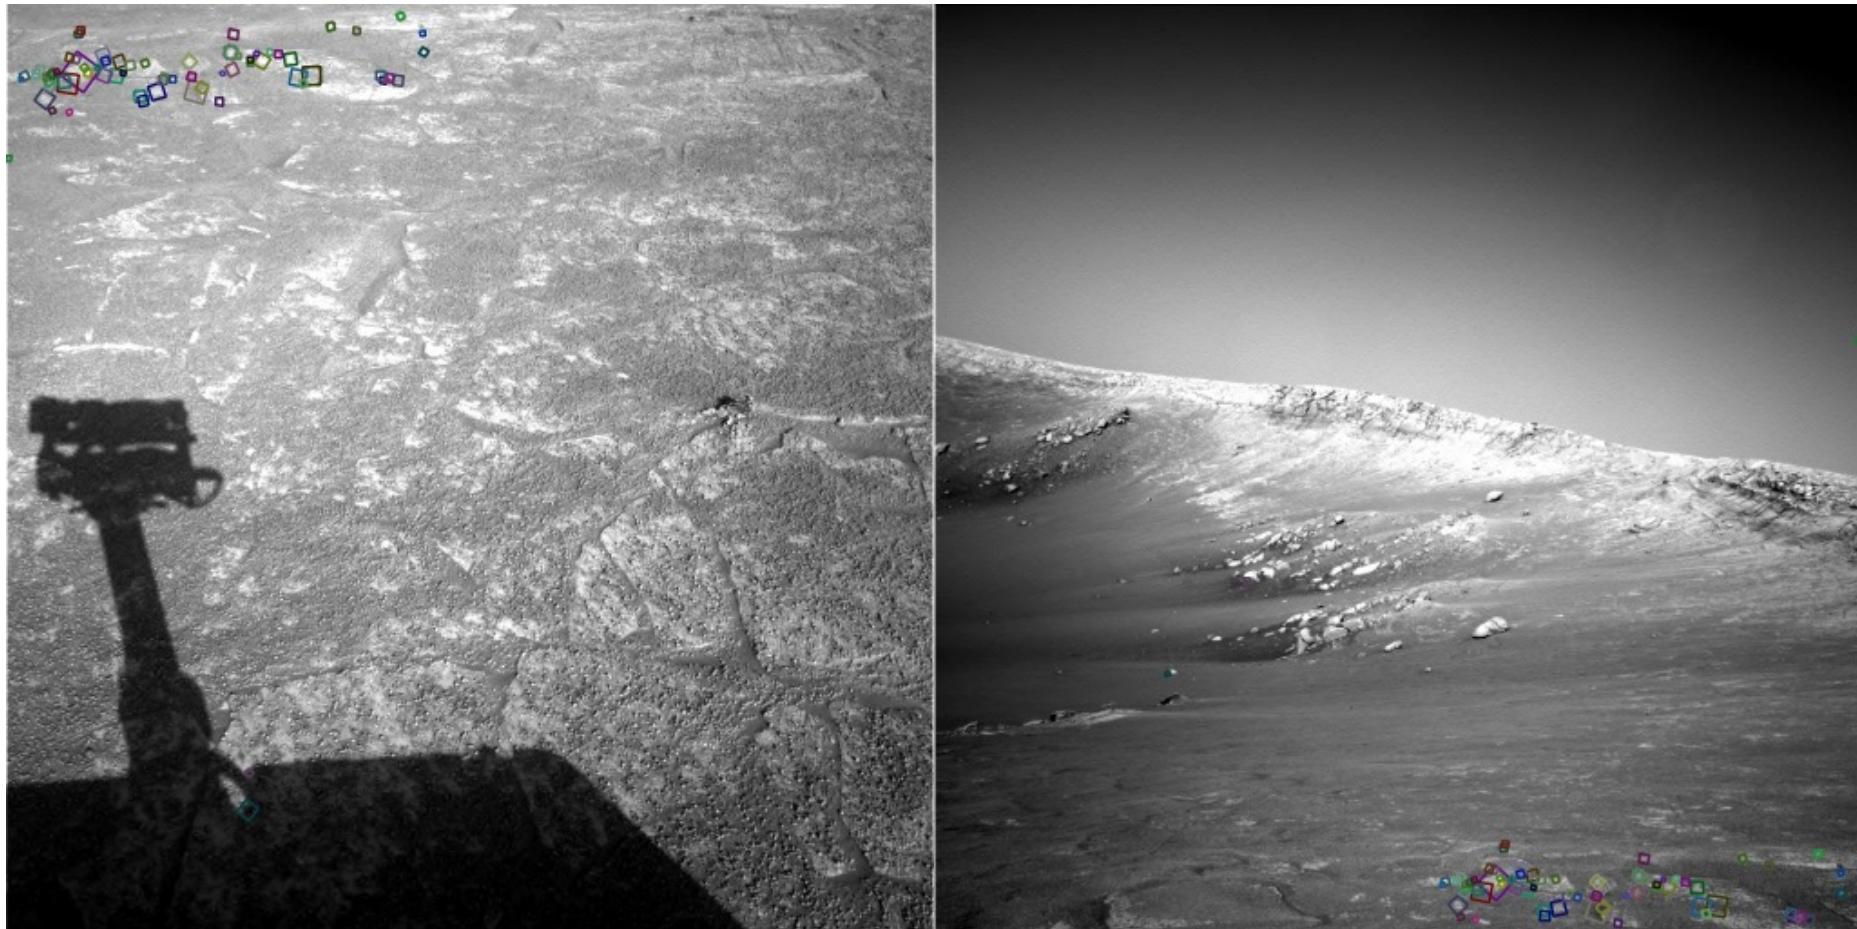
- Repeatable: should find same things even with distortion
- Saliency: each feature should be distinctive
- Compactness: shouldn't just be all the pixels
- Locality: should only depend on local image data

Example



Can you find the correspondences?

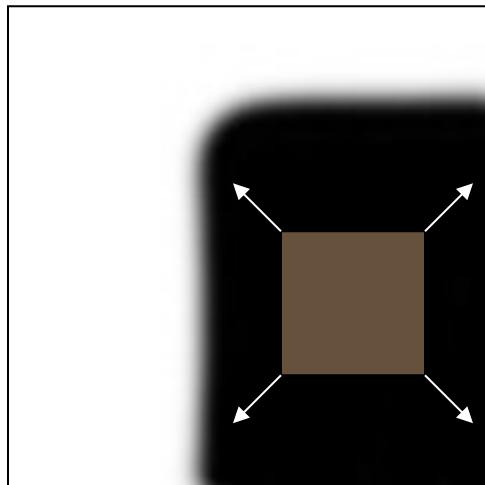
Example Matches



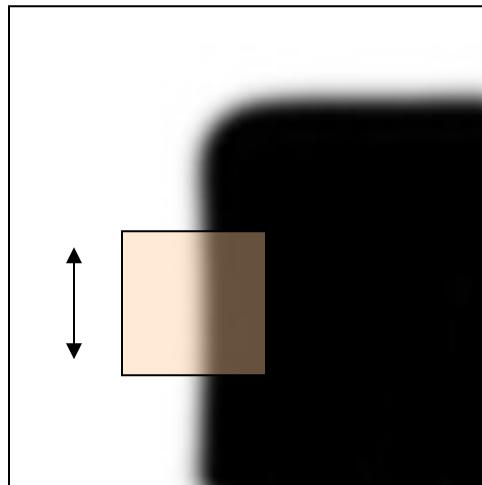
Look for the colored squares

Basic Idea

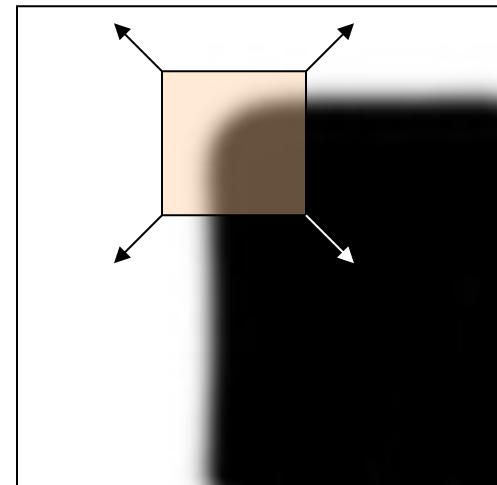
Should see where we are based on small window, or any shift → big intensity change.



“flat” region:
no change in
all directions



“edge”:
no change
along the edge
direction



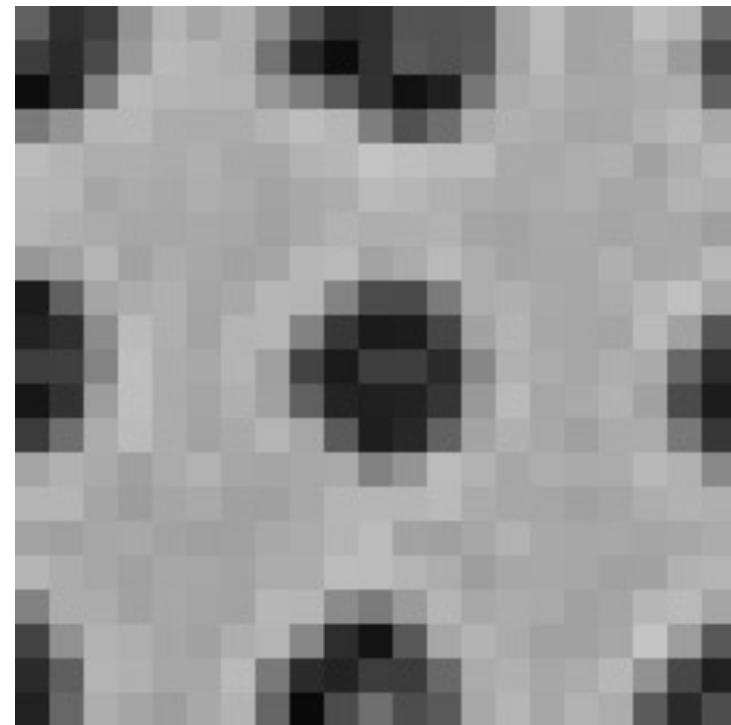
“corner”:
significant
change in all
directions

Formalizing Corner Detection



Formalizing Corner Detection

Zoom-In at x,y

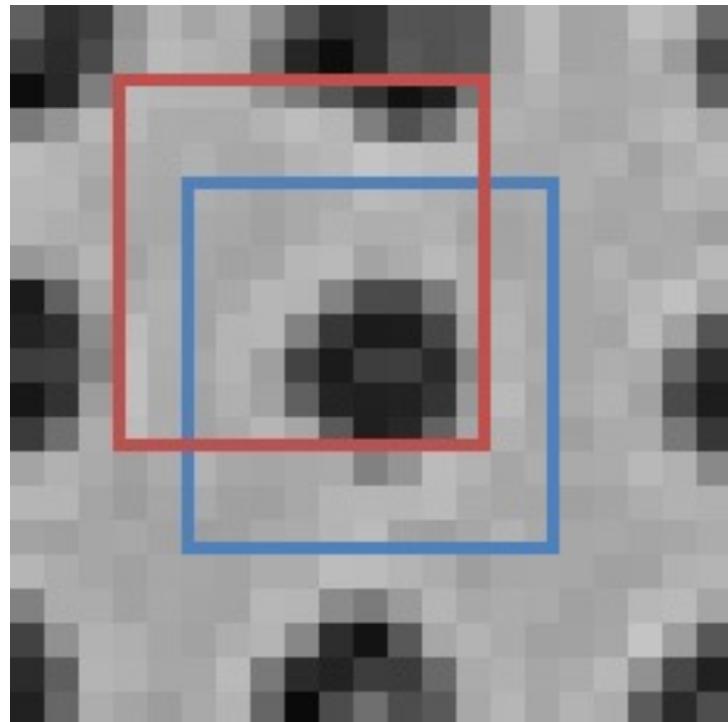


Original Image



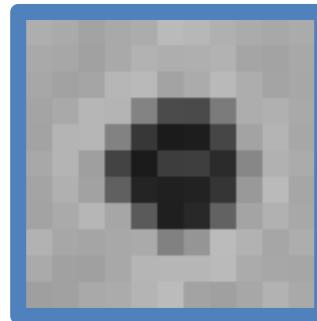
Formalizing Corner Detection

Zoom-In at x, y



Window **without** and **with** Offset

“Window”
At $x+u, y+v$
Here: $u=-2, v=-3$

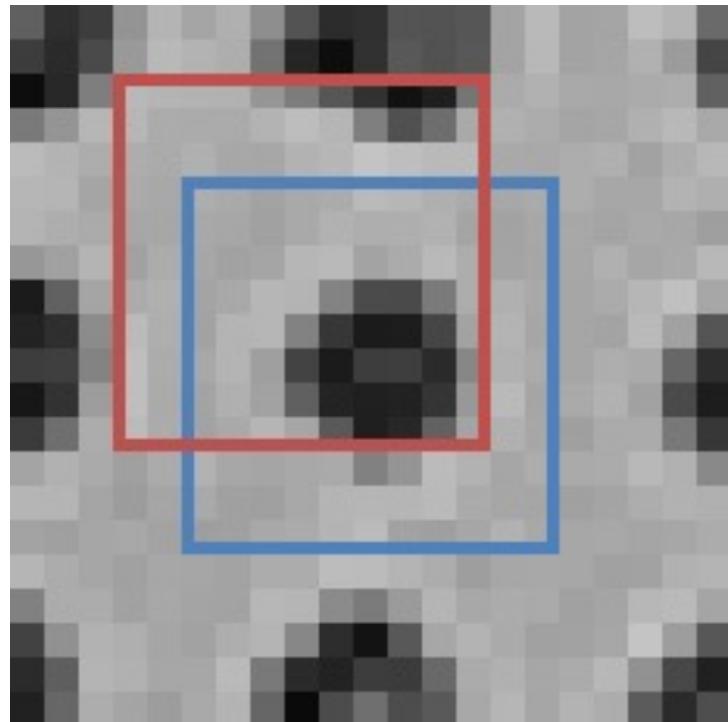


“Window”
At x, y

How might we measure similarity?

Formalizing Corner Detection

Zoom-In at x,y



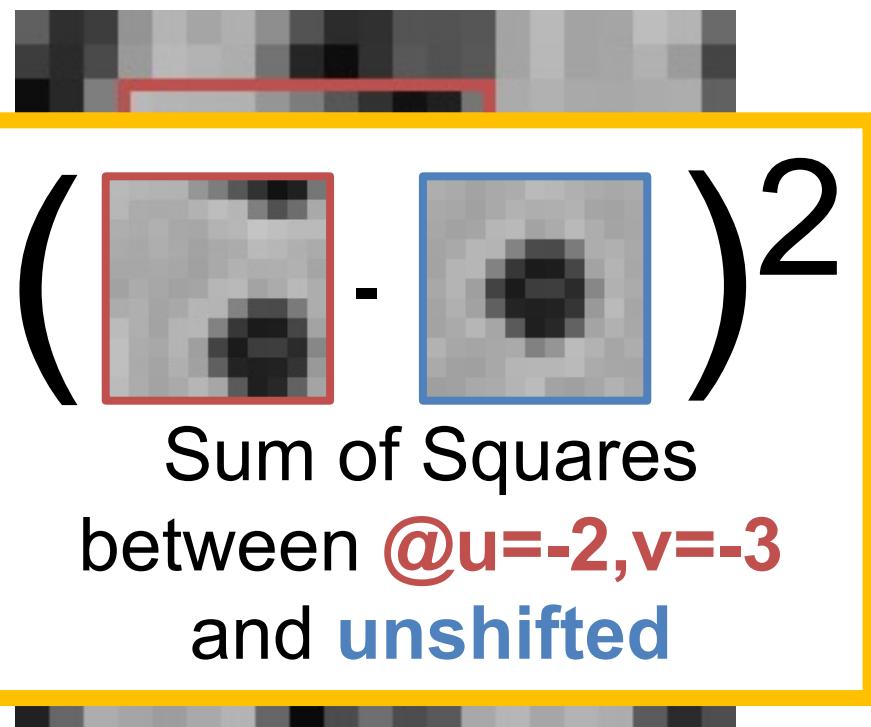
Error (Sum Sq) for u,v offset

$$E(u, v) = \sum_{(x,y) \in W} (I[x + u, y + v] - I[x, y])^2$$

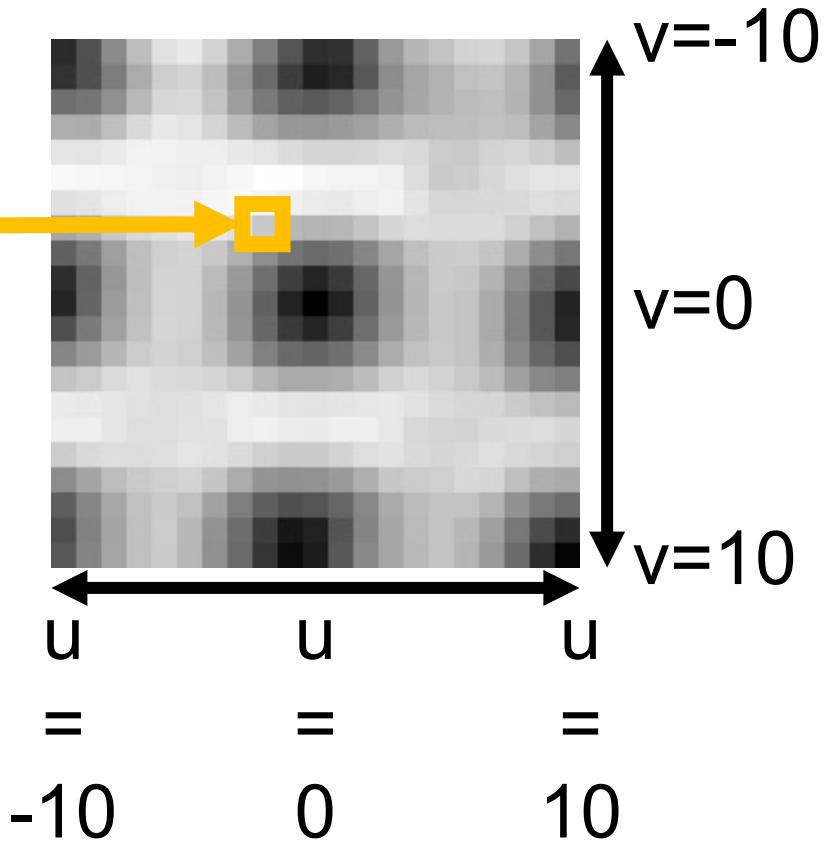
$$\left(\text{[Red Box]} - \text{[Blue Box]} \right)^2$$

Formalizing Corner Detection

Zoom-In at x,y



Error (Sum Sq) for u,v offset



Formalizing Corner Detection

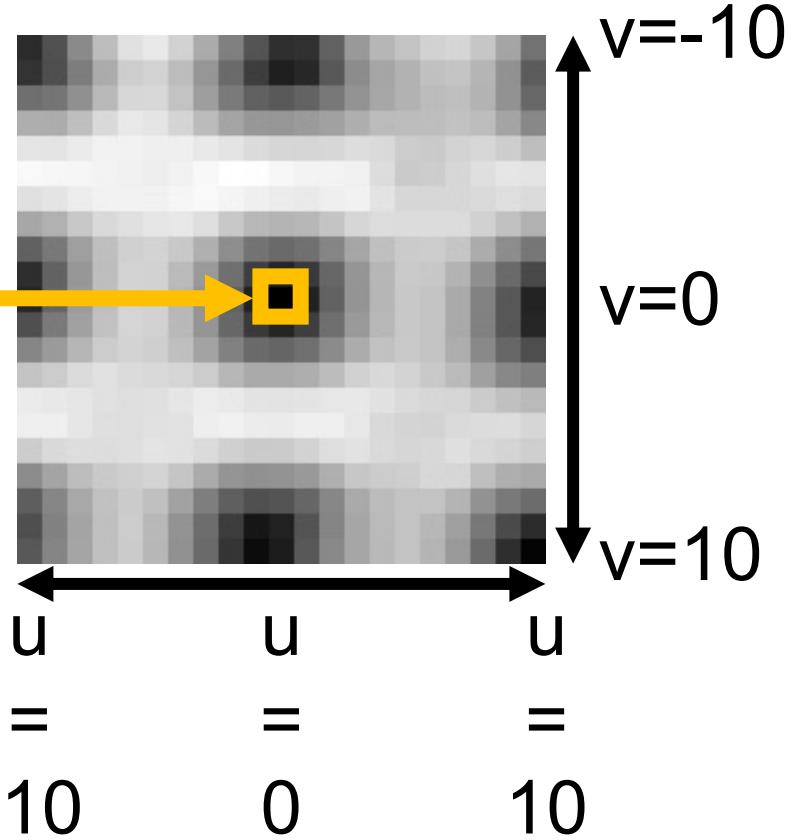
Zoom-In at x,y



Error at $u=0, v=0$ is
always 0. Why?

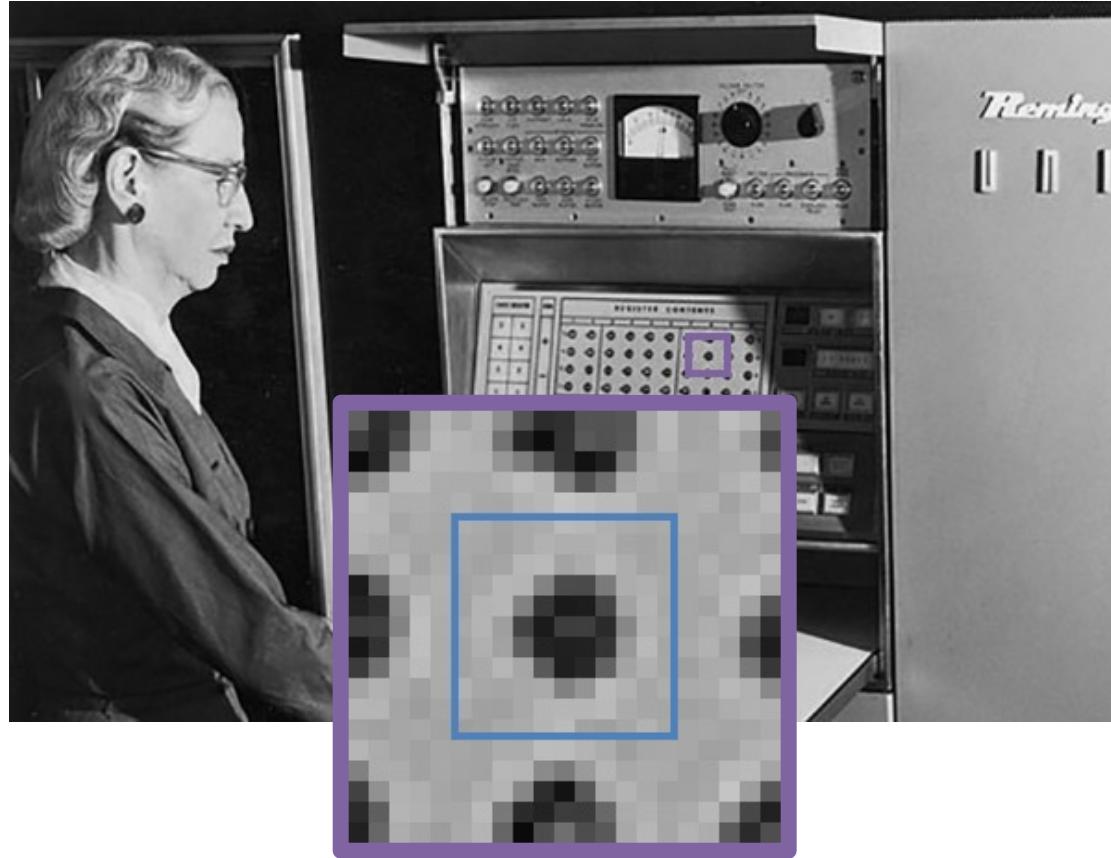


Error (Sum Sq) for u,v offset

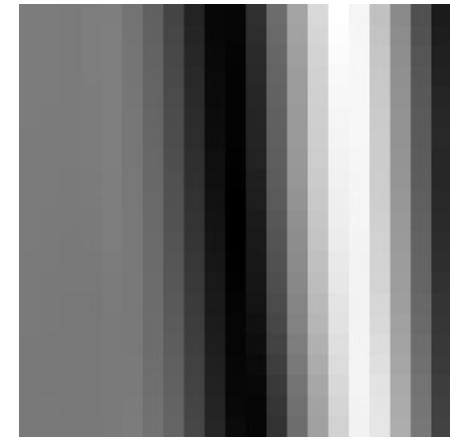


Match The Location and Plot

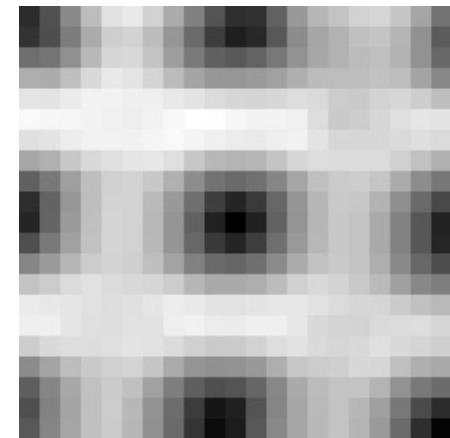
Original Image and Zoom-In



Error Options



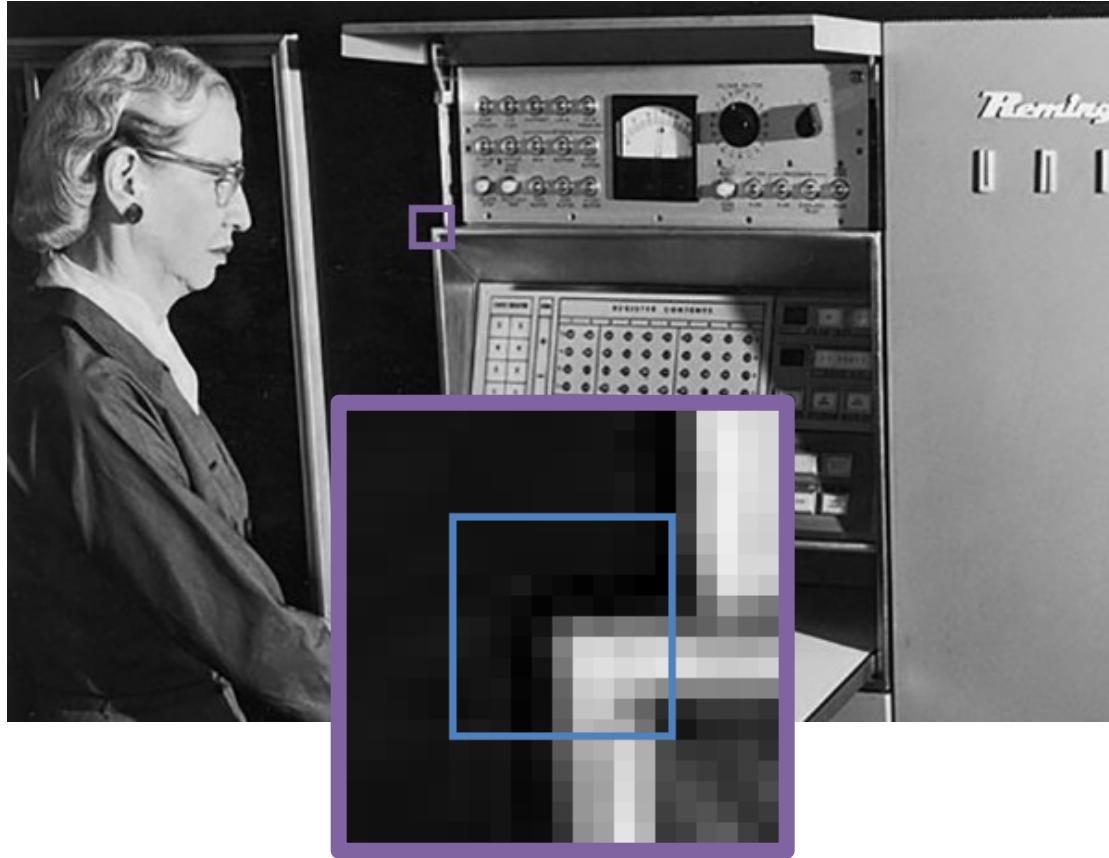
A



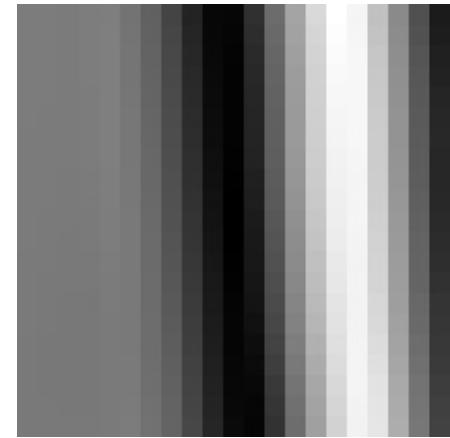
B

Match The Location and Plot

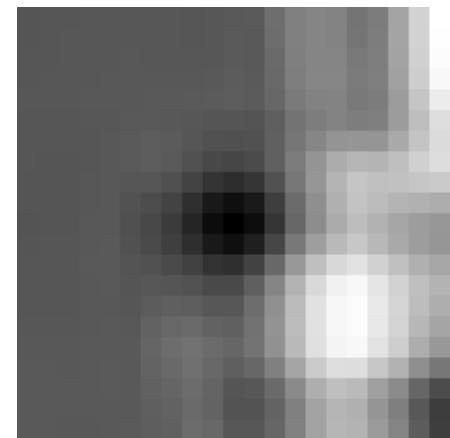
Original Image and Zoom-In



Error Options



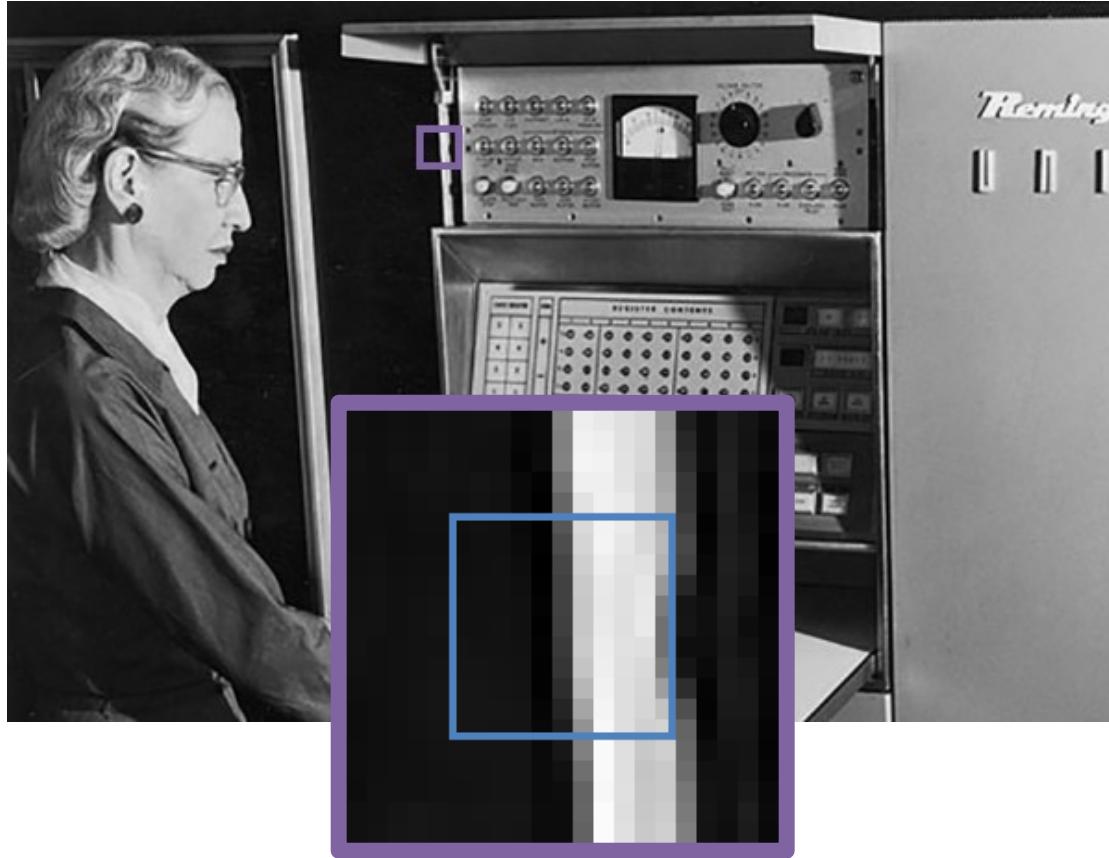
A



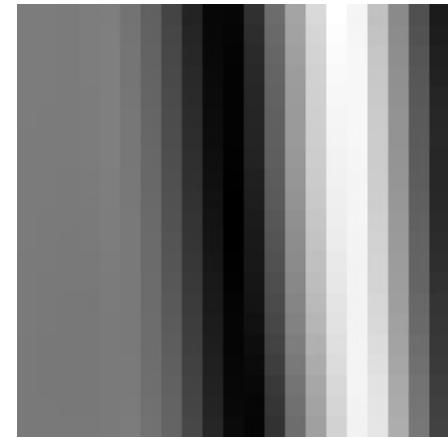
B

Match The Location and Plot

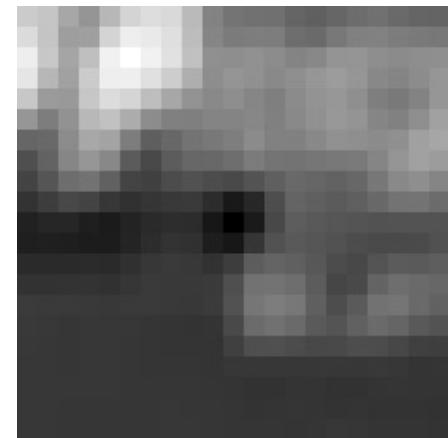
Original Image and Zoom-In



Error Options



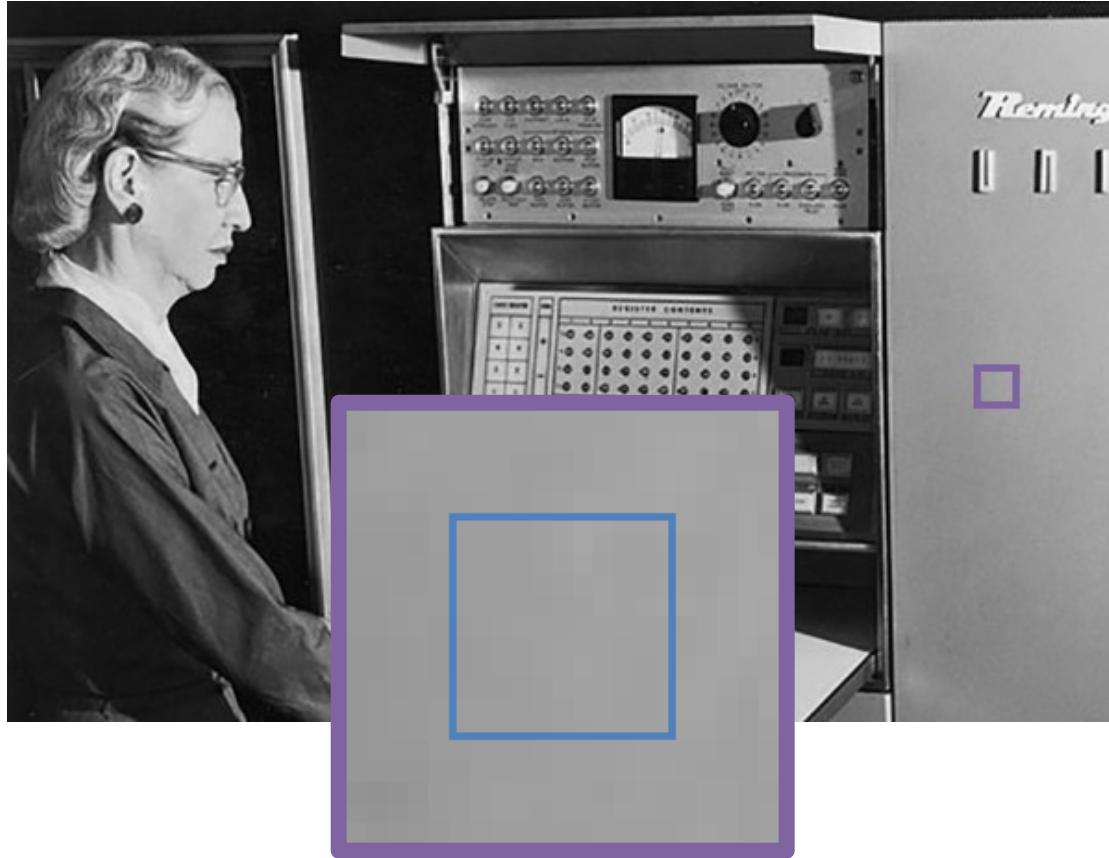
A



B

Match The Location and Plot

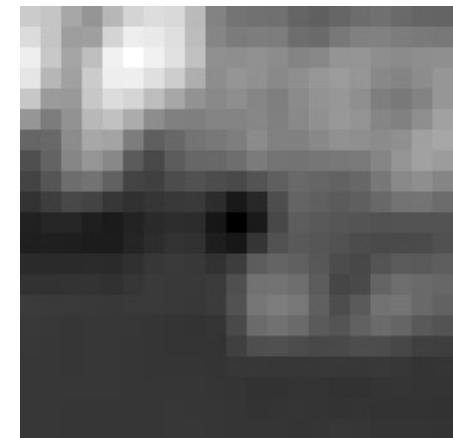
Original Image and Zoom-In



Error Options



A



B

Ok But Back To Math

$$E(u, v) = \sum_{(x,y) \in W} (I[x + u, y + v] - I[x, y])^2$$

$$\left(\begin{array}{c} \text{[Pixelated image with red border]} \\ - \\ \text{[Pixelated image with blue border]} \end{array} \right)^2$$

Shifting windows around is expensive!
We'll find a trick to approximate this.

Note: only need to get the gist

Aside: Taylor Series for Images

Recall Taylor Series – way of *linearizing* a function:

$$f(x + d) \approx f(x) + \frac{\partial f(x)}{\partial x} d$$

$$f(x + u, y + v) \approx f(x, y) + \frac{\partial f(x, y)}{\partial x} u + \frac{\partial f(x, y)}{\partial y} v$$

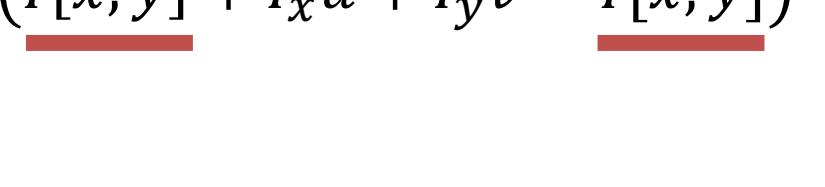
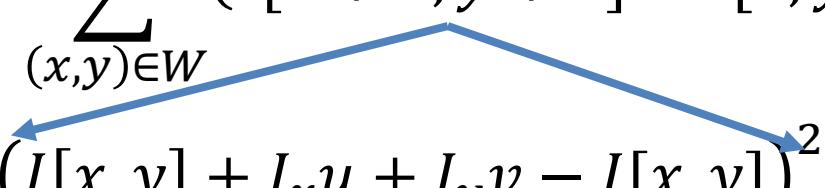
Do the same with images, treating them as
function of x, y

$$I(x + u, y + v) \approx I(x, y) + I_x u + I_y v$$

For brevity: $I_x = I_x(x, y)$ at point (x, y) , $I_y = I_y(x, y)$ at
point (x, y)

Formalizing Corner Detection

Taylor series
expansion for I
at every single
point in window

$$E(u, v) = \sum_{(x,y) \in W} (I[x + u, y + v] - I[x, y])^2$$
$$\approx \sum_{(x,y) \in W} (I[x, y] + I_x u + I_y v - I[x, y])^2$$


Cancel

$$= \sum_{(x,y) \in W} (I_x u + I_y v)^2$$

Expand

$$= \sum_{(x,y) \in W} I_x^2 u^2 + 2I_x I_y u v + I_y^2 v^2$$

For brevity: $I_x = I_x$ at point (x,y) , $I_y = I_y$ at point (x,y)

Formalizing Corner Detection

By linearizing image, we can approximate $E(u,v)$ with quadratic function of u and v

$$\begin{aligned} E(u, v) &\approx \sum_{(x,y) \in W} (I_x^2 u^2 + 2I_x I_y u v + I_y^2 v^2) \\ &= [u, v] \mathbf{M} [u, v]^T \end{aligned}$$

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix}$$

\mathbf{M} is called the second moment matrix
Simple algebra: Try at home!

Intuitively what is M?

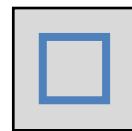
Pretend gradients are *either* vertical or horizontal

Obviously Wrong!

$$M = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} \approx \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$$

a,b both small:

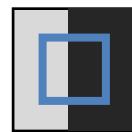
flat



$$\begin{bmatrix} 0.1 & 0 \\ 0 & 0.1 \end{bmatrix}$$

One big,
other small:

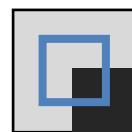
edge



$$\begin{bmatrix} 50 & 0 \\ 0 & 0.1 \end{bmatrix} \text{ or } \begin{bmatrix} 0.1 & 0 \\ 0 & 50 \end{bmatrix}$$

a,b both big:

corner



$$\begin{bmatrix} 50 & 0 \\ 0 & 50 \end{bmatrix}$$

Intuitively what is M?

General case: so $I_x I_y \neq 0$

$$M = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} \approx? \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix}$$

a,b both small: flat

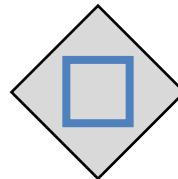
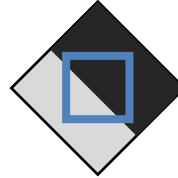
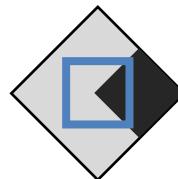


Image might be rotated by rotation $\theta!$

One big,
other small:
edge



a,b both big:
corner

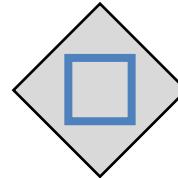


Intuitively what is M?

General case: so $I_x I_y \neq 0$

$$M = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} = V^{-1} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} V$$

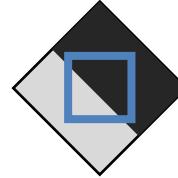
a,b both small: flat



If image rotated by rotation θ / matrix V

One big,
other small:

edge

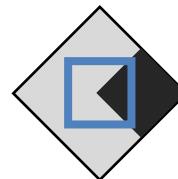


M will look like

$$V^{-1} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} V$$

a,b both big:

corner



So What Now?

Can calculate \mathbf{M} at pixel, by summing nearby gradients, but need to find a and b .

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} = \mathbf{V}^{-1} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} \mathbf{V}$$

Given \mathbf{M} , can decompose it into eigenvectors \mathbf{V} and eigenvalues λ_1, λ_2 with $\mathbf{M} = \mathbf{V}^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} \mathbf{V}$.

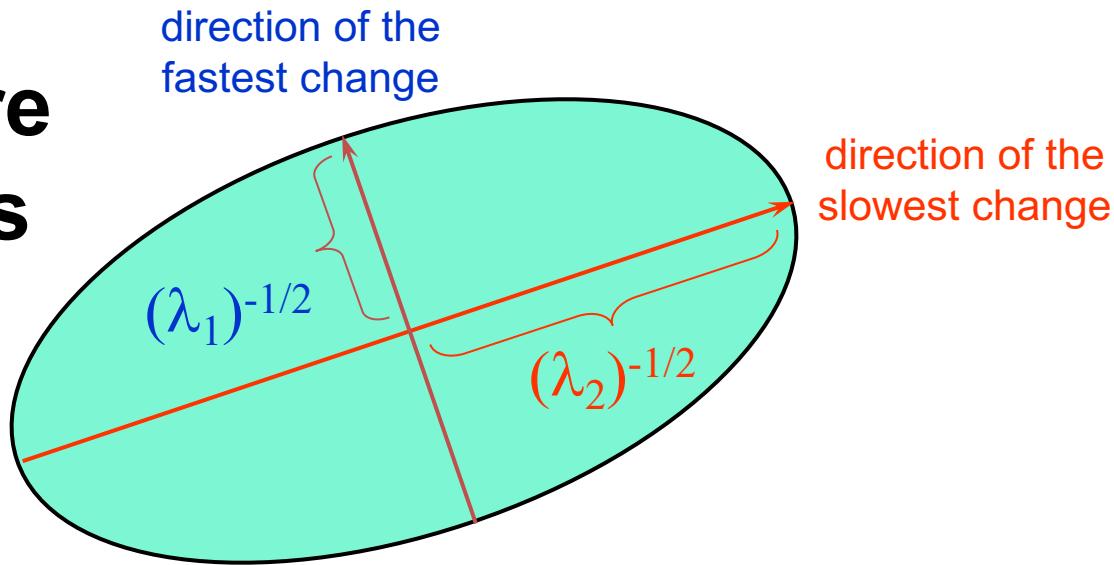
**Eigenvalue
Decomposition**

Review: Quadratic Forms

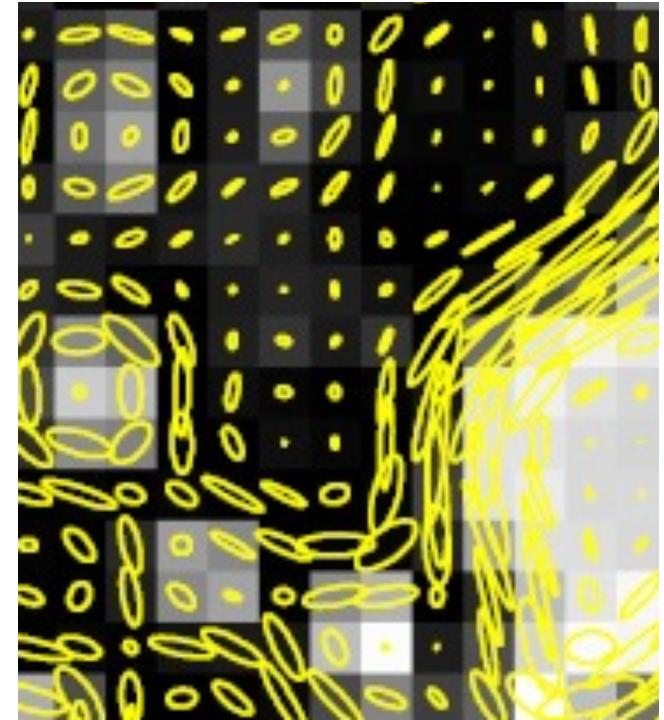
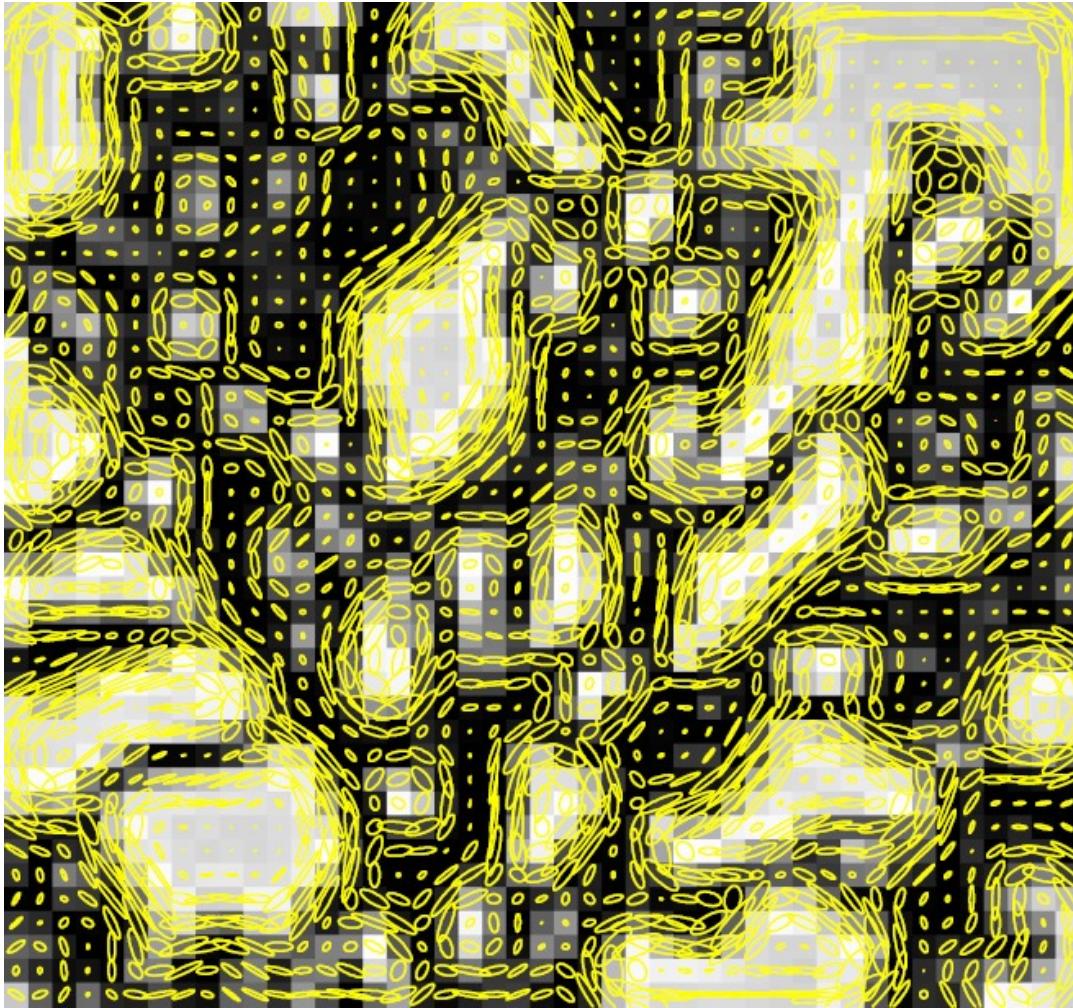
We can look at the shape of this ellipse by decomposing a symmetric matrix M into a rotation + scaling

$$M = V^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} V$$

λ_1 and λ_2 are eigenvalues



Visualizing M



Technical note: M is often best *visualized* by first taking inverse, so long edge of ellipse goes along edge

So What Now?

Can calculate M at pixel, by summing nearby gradients, but need to find a and b .

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} = \mathbf{V}^{-1} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} \mathbf{V}$$

Given \mathbf{M} , can decompose it into eigenvectors \mathbf{V} and eigenvalues λ_1, λ_2 with $\mathbf{M} = \mathbf{V}^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} \mathbf{V}$.

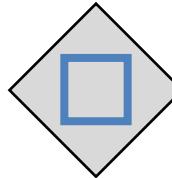
Rotation direction factored out by \mathbf{V}
 λ_1, λ_2 only describe magnitude!

Intuitively what is M?

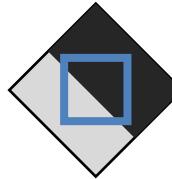
Can calculate M at pixel, by summing nearby gradients, but need to find λ_1 and λ_2 .

$$M = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} = V^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} V$$

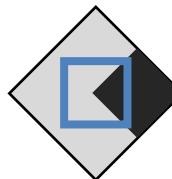
λ_1, λ_2 both small: flat



One big,
other small:
edge



λ_1, λ_2 both big:
corner



So What Now?

Can calculate \mathbf{M} at pixel, by summing nearby gradients, but need to find λ_1 and λ_2 .

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} = \mathbf{V}^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} \mathbf{V}$$

Given \mathbf{M} , can decompose it into eigenvectors \mathbf{V} and eigenvalues λ_1, λ_2 with $\mathbf{M} = \mathbf{V}^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} \mathbf{V}$.

Slow. Why?

So What Now?

Can calculate M at pixel, by summing nearby gradients, but need to find a and b .

$$M = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix} = V^{-1} \begin{bmatrix} a & 0 \\ 0 & b \end{bmatrix} V$$

Instead: compute quantity R from M

$$R = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

Easy fast formula
for 2x2

Fast – sum the diagonal

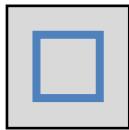
Empirical value,
usually 0.04-0.06

So What Now?

R tells us whether we're at a corner, edge, or flat

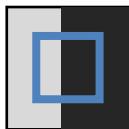
$$R = \det(\mathbf{M}) - \alpha \operatorname{trace}(\mathbf{M})^2 = \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2$$

flat



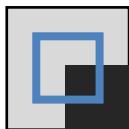
$$\lambda_1, \lambda_2 \approx 0$$

edge

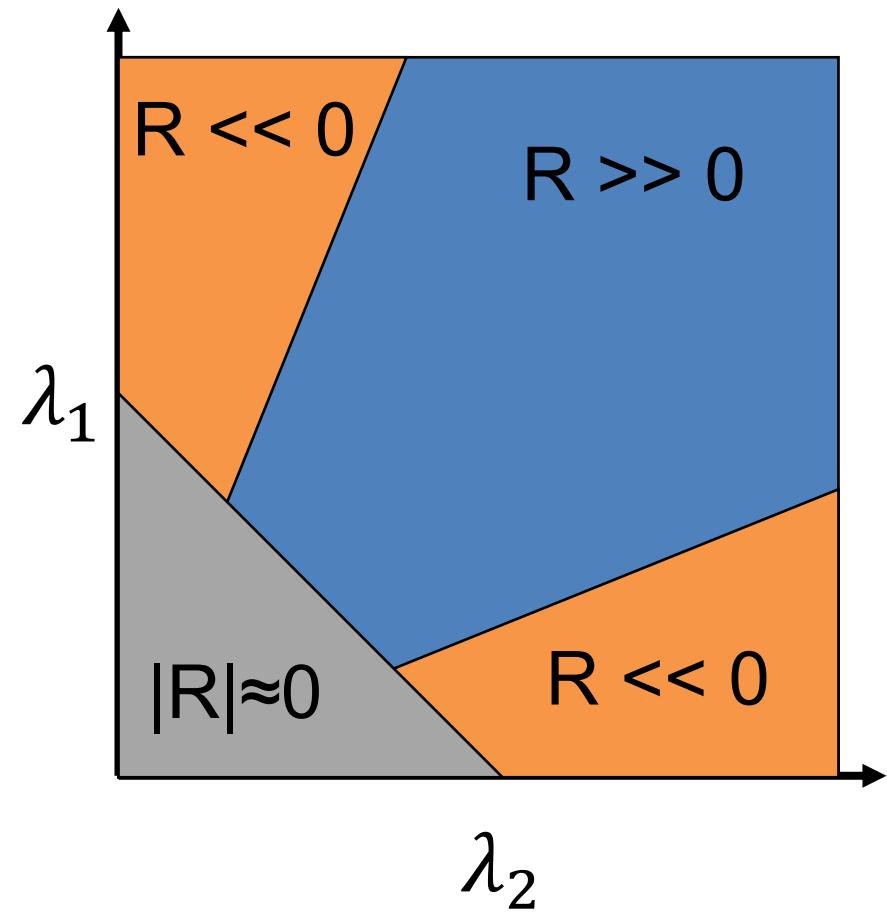


$$\begin{aligned}\lambda_1 &\gg \lambda_2 \gg 0 \\ \lambda_2 &\gg \lambda_1 \gg 0\end{aligned}$$

corner



$$\lambda_1 \approx \lambda_2 \gg 0$$



What Do I Need To Know?

- Need to be able to take derivatives of image
- Need to be able to compute the entries of \mathbf{M} at every pixel.
- Should know that some properties of \mathbf{M} indicate whether a pixel is a corner or not.

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix}$$

In Practice

1. Compute partial derivatives I_x , I_y per pixel
2. Compute \mathbf{M} at each pixel, using Gaussian weighting w

$$\mathbf{M} = \begin{bmatrix} \sum_{x,y \in W} w(x,y) I_x^2 & \sum_{x,y \in W} w(x,y) I_x I_y \\ \sum_{x,y \in W} w(x,y) I_x I_y & \sum_{x,y \in W} w(x,y) I_y^2 \end{bmatrix}$$

C.Harris and M.Stephens. [A Combined Corner and Edge Detector.](#)
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

In Practice

1. Compute partial derivatives I_x, I_y per pixel
2. Compute \mathbf{M} at each pixel, using Gaussian weighting w
3. Compute response function R

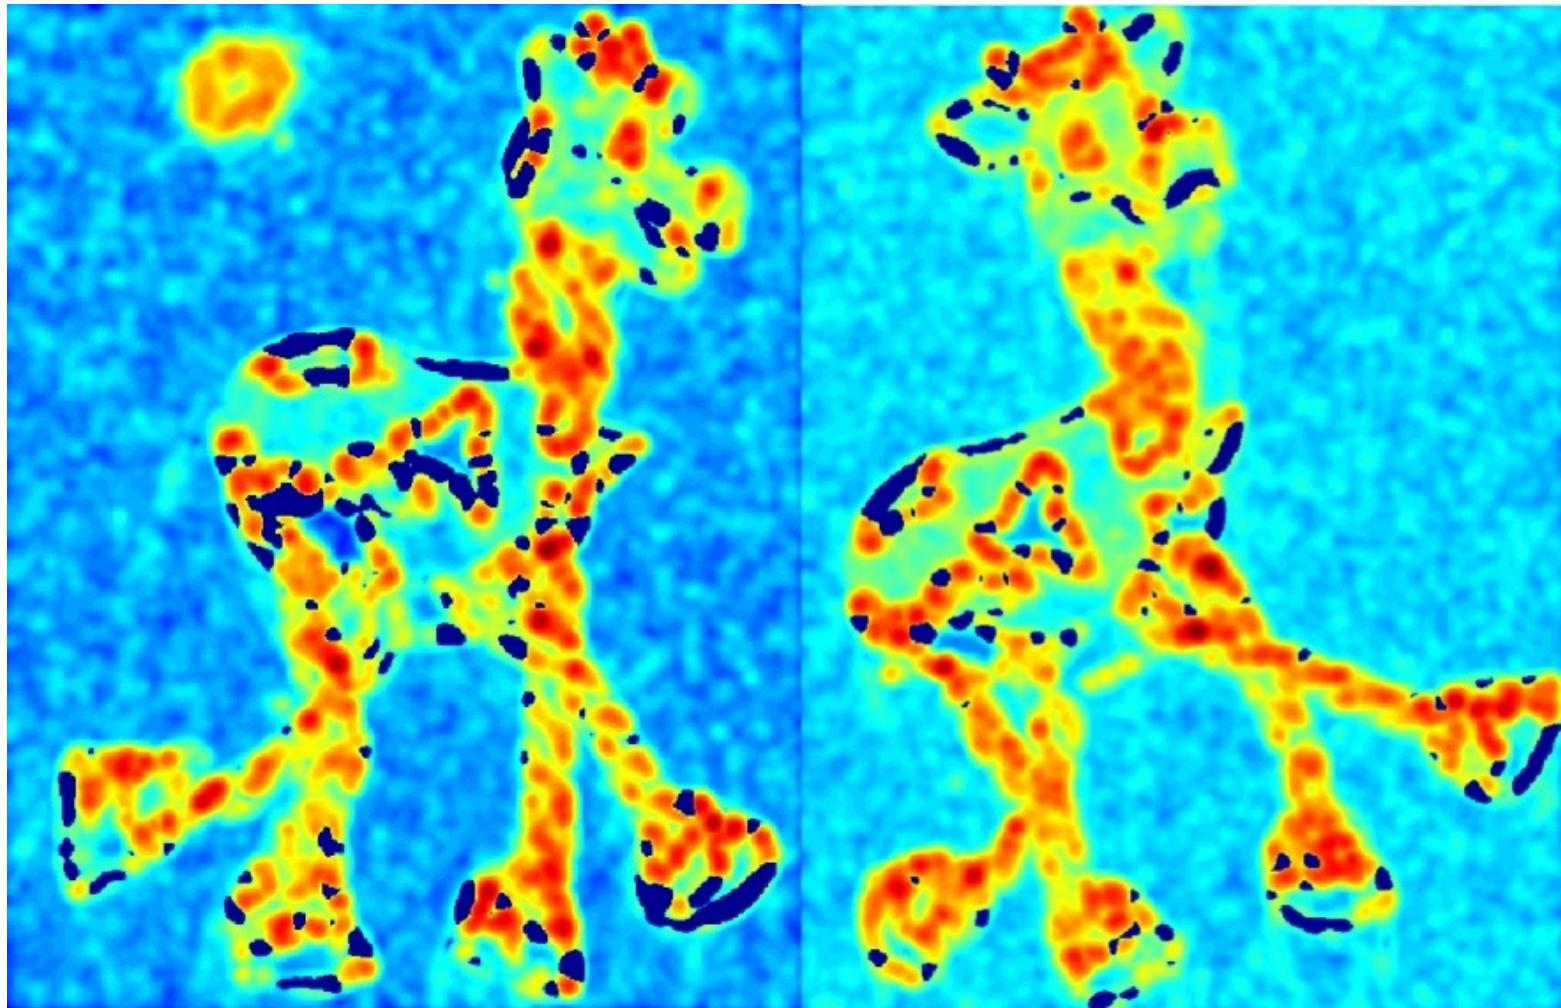
$$\begin{aligned} R &= \det(\mathbf{M}) - \alpha \operatorname{trace}(\mathbf{M})^2 \\ &= \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2 \end{aligned}$$

C.Harris and M.Stephens. ["A Combined Corner and Edge Detector."](#)
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Computing R



Computing R



In Practice

1. Compute partial derivatives I_x , I_y per pixel
2. Compute \mathbf{M} at each pixel, using Gaussian weighting w
3. Compute response function R
4. Threshold R

C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)"
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Thresholded R

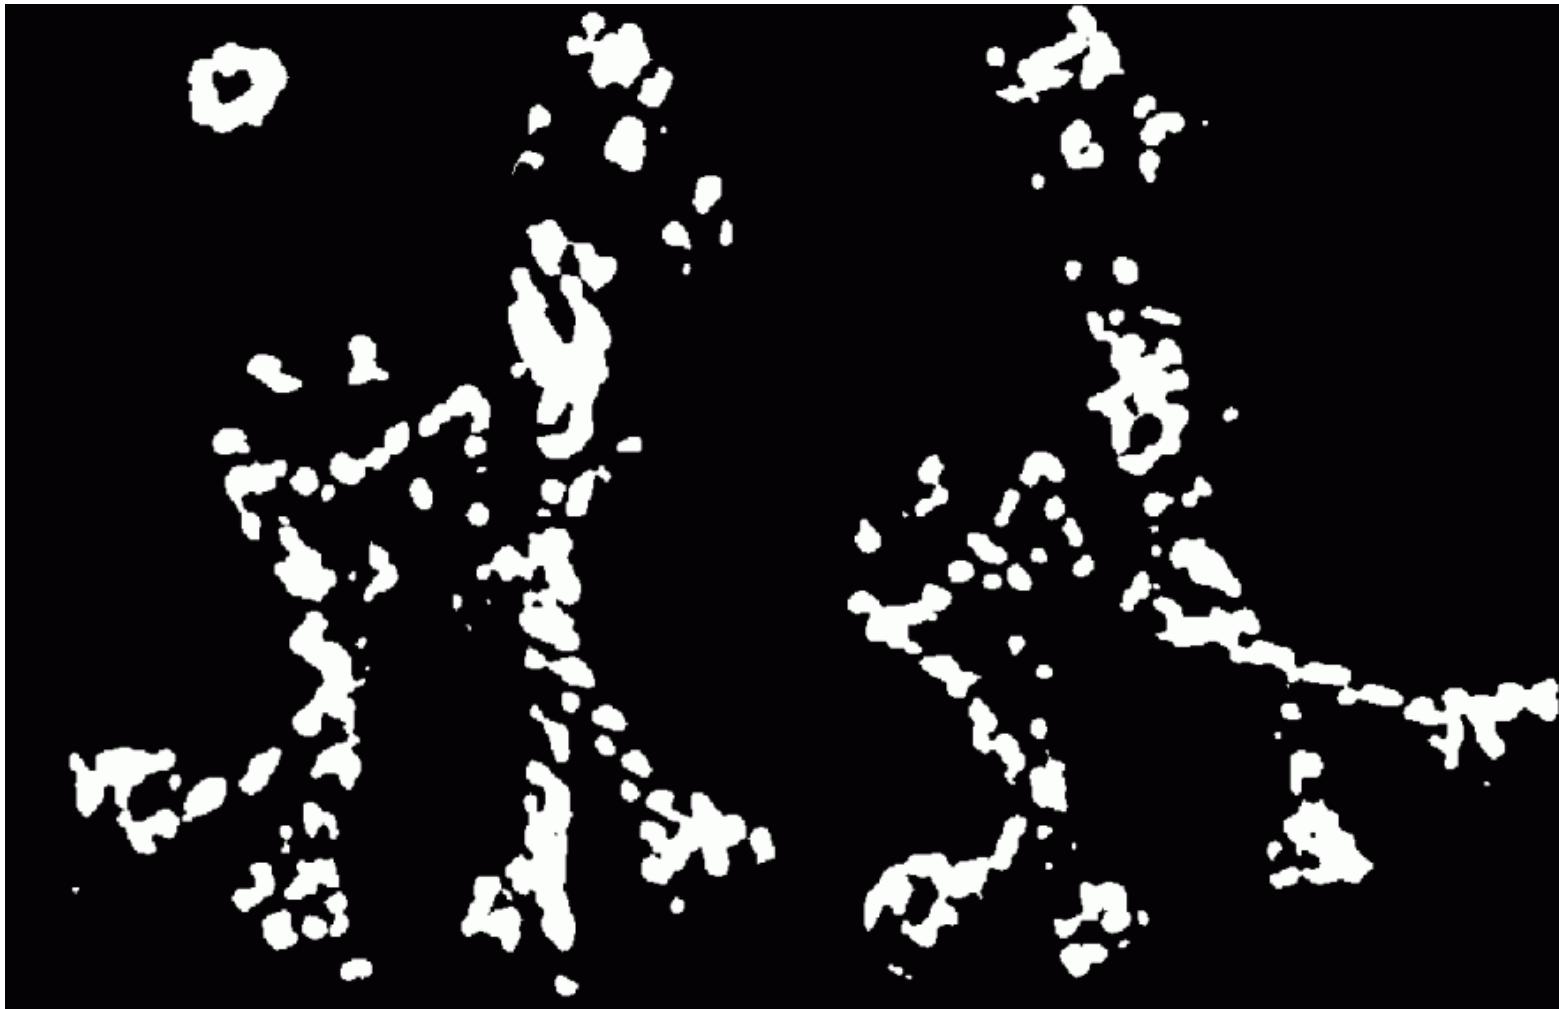


In Practice

1. Compute partial derivatives I_x , I_y per pixel
2. Compute \mathbf{M} at each pixel, using Gaussian weighting w
3. Compute response function R
4. Threshold R
5. Take only local maxima (called non-maxima suppression)

C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)"
Proceedings of the 4th Alvey Vision Conference: pages 147—151, 1988.

Thresholded, NMS R



Final Results



Desirable Properties

If our detectors are repeatable, they should be:

- **Invariant** to some things: image is transformed and corners remain the same
- **Covariant/equivariant** with some things: image is transformed and corners transform with it.

Recall Motivating Problem

Images may be different in lighting and geometry



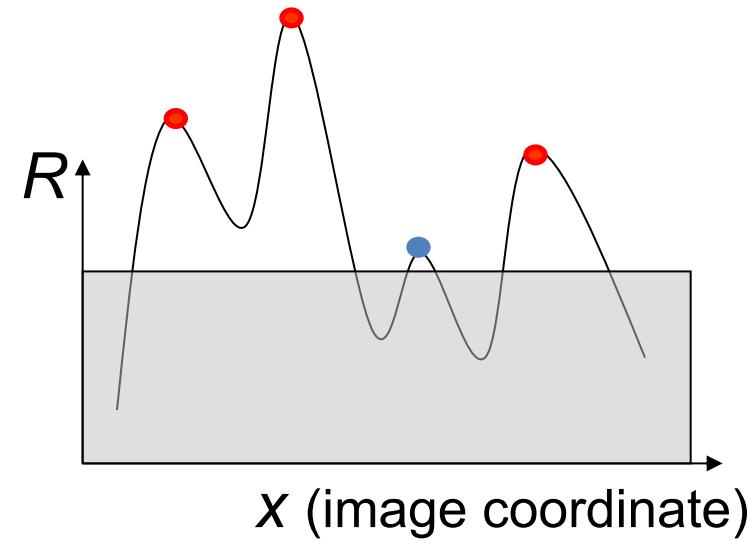
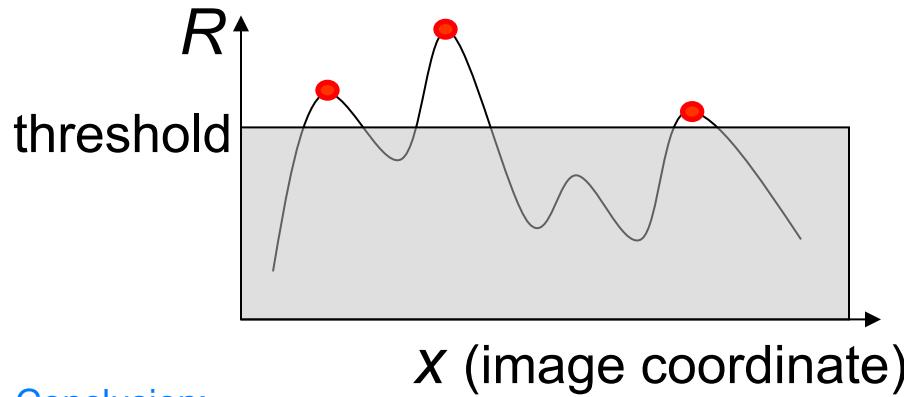
These M matrix is resilient to intensity changes, mostly because these peaks are preserved, but the constant a might pull out some of the peaks above the threshold, and you will have the new corner detected.

Affine Intensity Change

$$I_{new} = aI_{old} + b$$

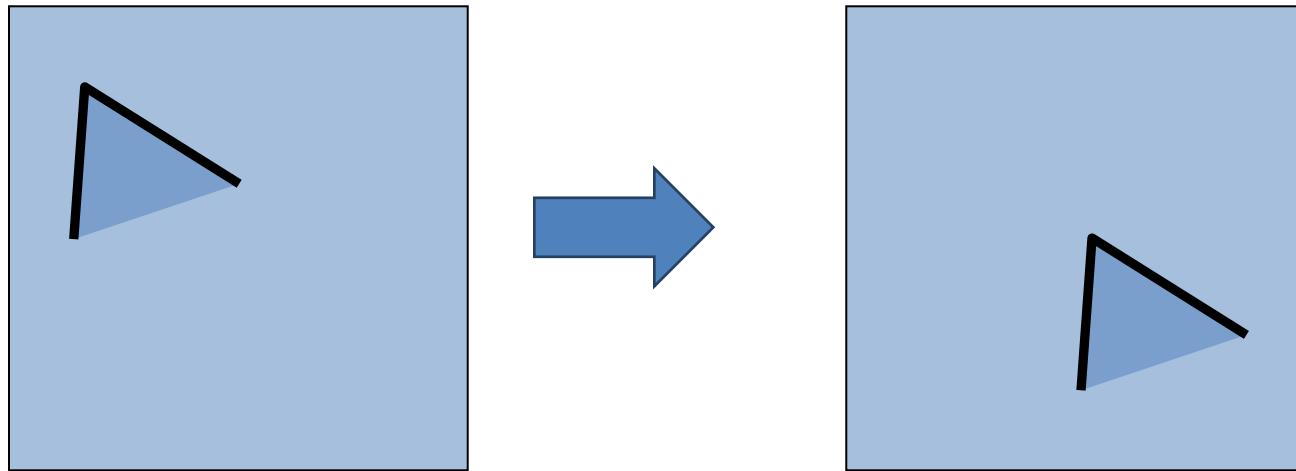
M only depends on derivatives, so b is irrelevant

But a scales derivatives and there's a threshold



Partially invariant to affine intensity changes

Image Translation

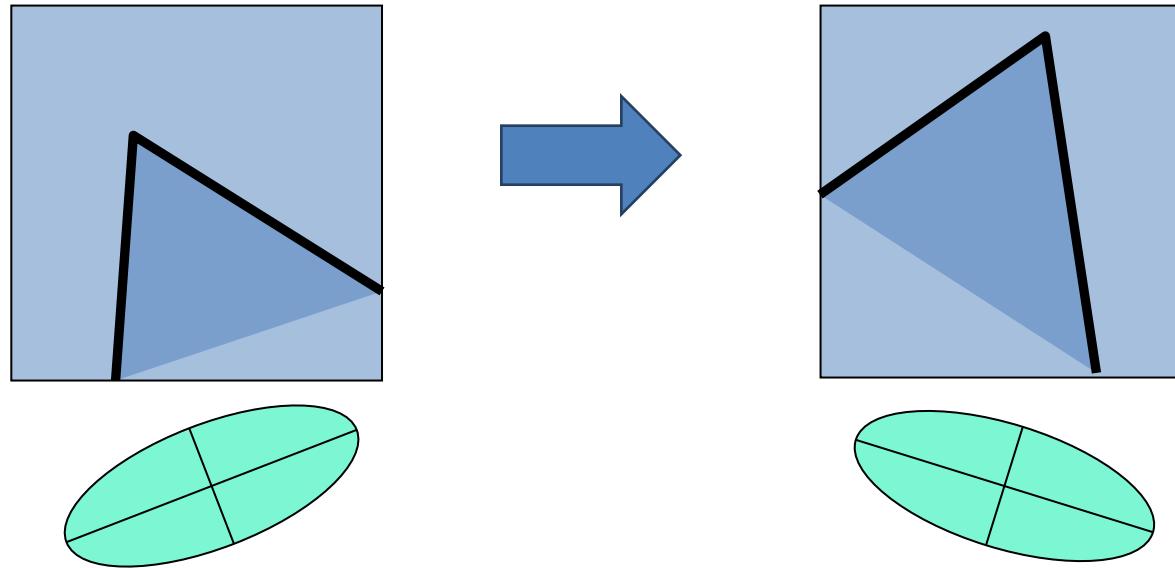


All done with convolution. Convolution is
translation invariant.

Equivariant with translation

It will change the eigenvector of M but will not change the eigenvalue of the M. So if you just pick the pixel with the high eight value then you still pick the same corner.

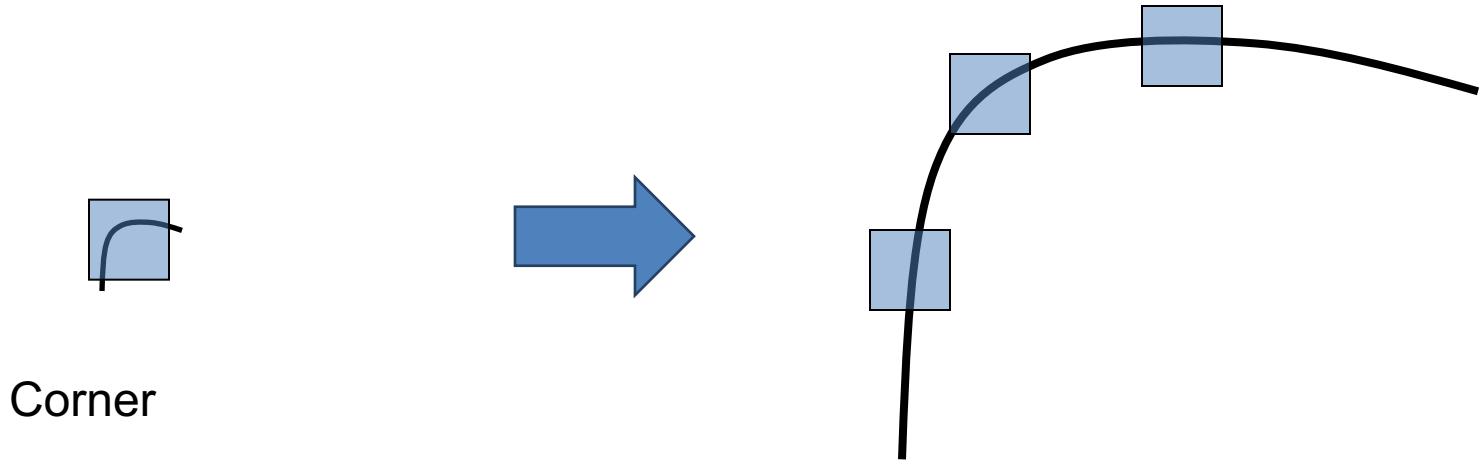
Image Rotation



Rotations just cause the corner rotation to change.
Eigenvalues remain the same.

Equivariant with rotation

Image Scaling



One pixel can become many pixels and vice-versa.

Not equivariant with scaling
Fix: Next Class

For the Curious

Review: Quadratic Forms

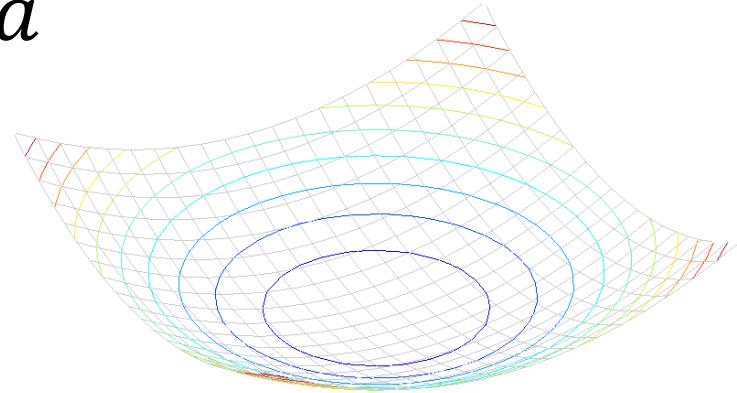
Suppose have symmetric matrix \mathbf{M} , scalar a , vector $[u, v]$:

$$E([u, v]) = [u, v] \mathbf{M} [u, v]^T$$

Then the isocontour / slice-through of E , i.e.

$$E([u, v]) = a$$

is an ellipse.

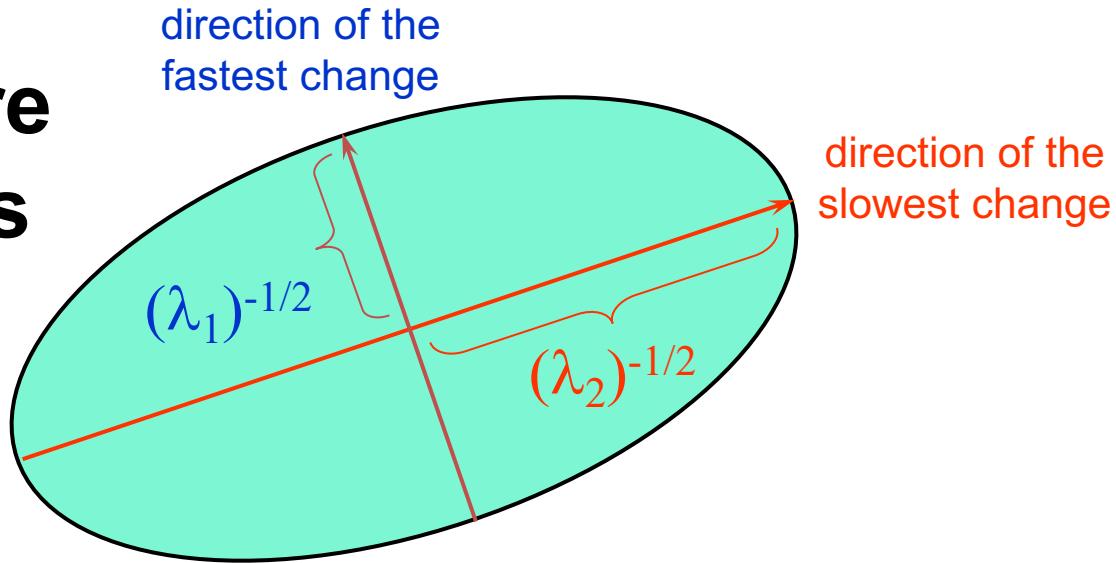


Review: Quadratic Forms

We can look at the shape of this ellipse by decomposing M into a rotation + scaling

$$M = V^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} V$$

λ_1 and λ_2 are eigenvalues



Interpreting The Matrix M

The second moment matrix tells us how quickly the image changes and in which directions.

Can compute at each pixel

$$M = \overbrace{\begin{bmatrix} \sum_{x,y \in W} I_x^2 & \sum_{x,y \in W} I_x I_y \\ \sum_{x,y \in W} I_x I_y & \sum_{x,y \in W} I_y^2 \end{bmatrix}}^{\text{Directions}} = V^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} V$$

Amounts

Directions

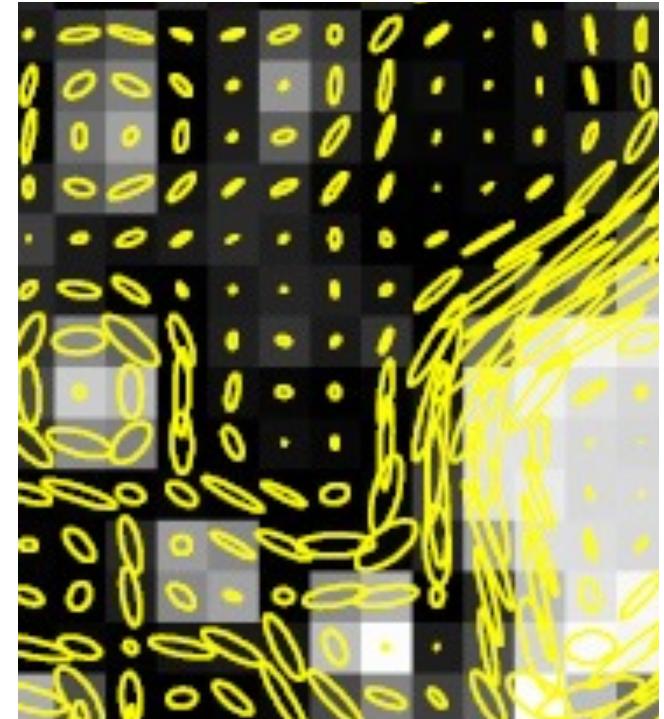
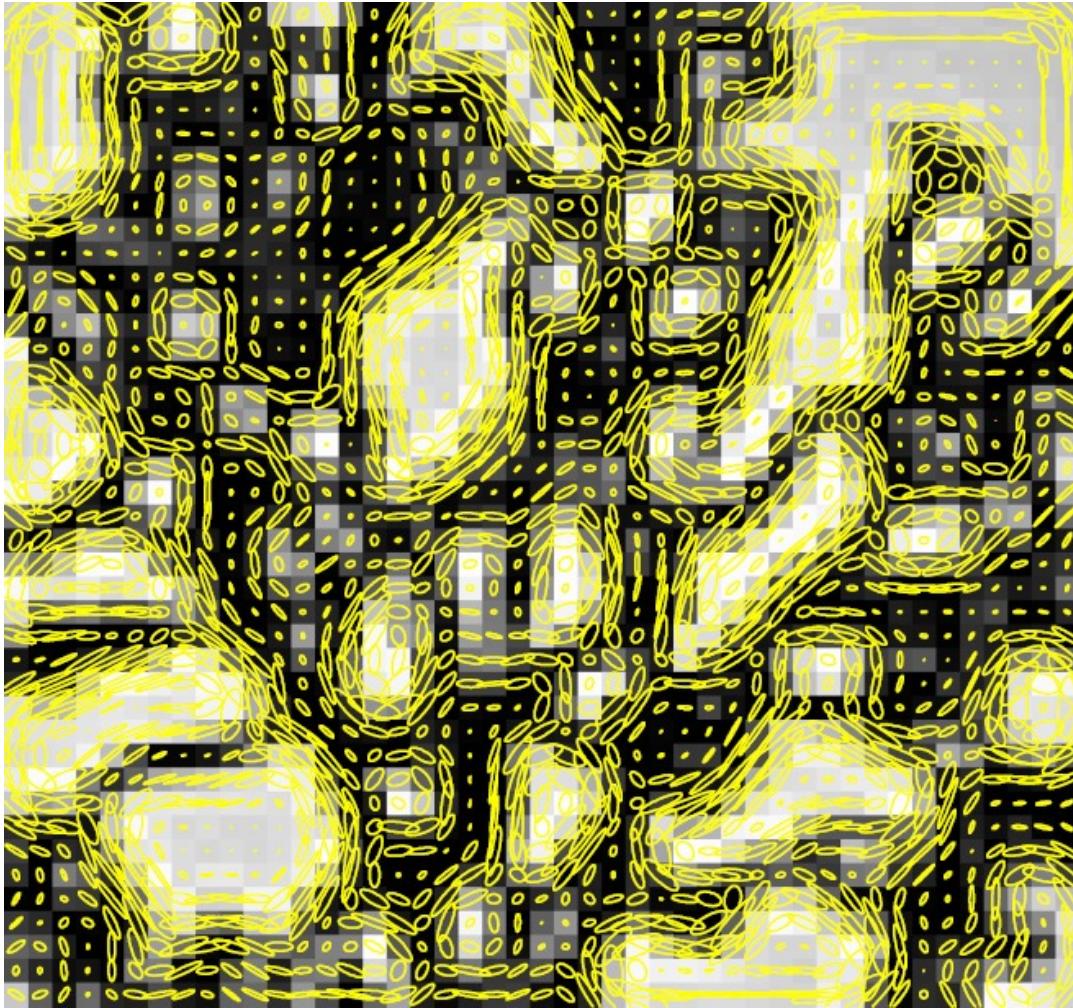
Amounts

Visualizing M



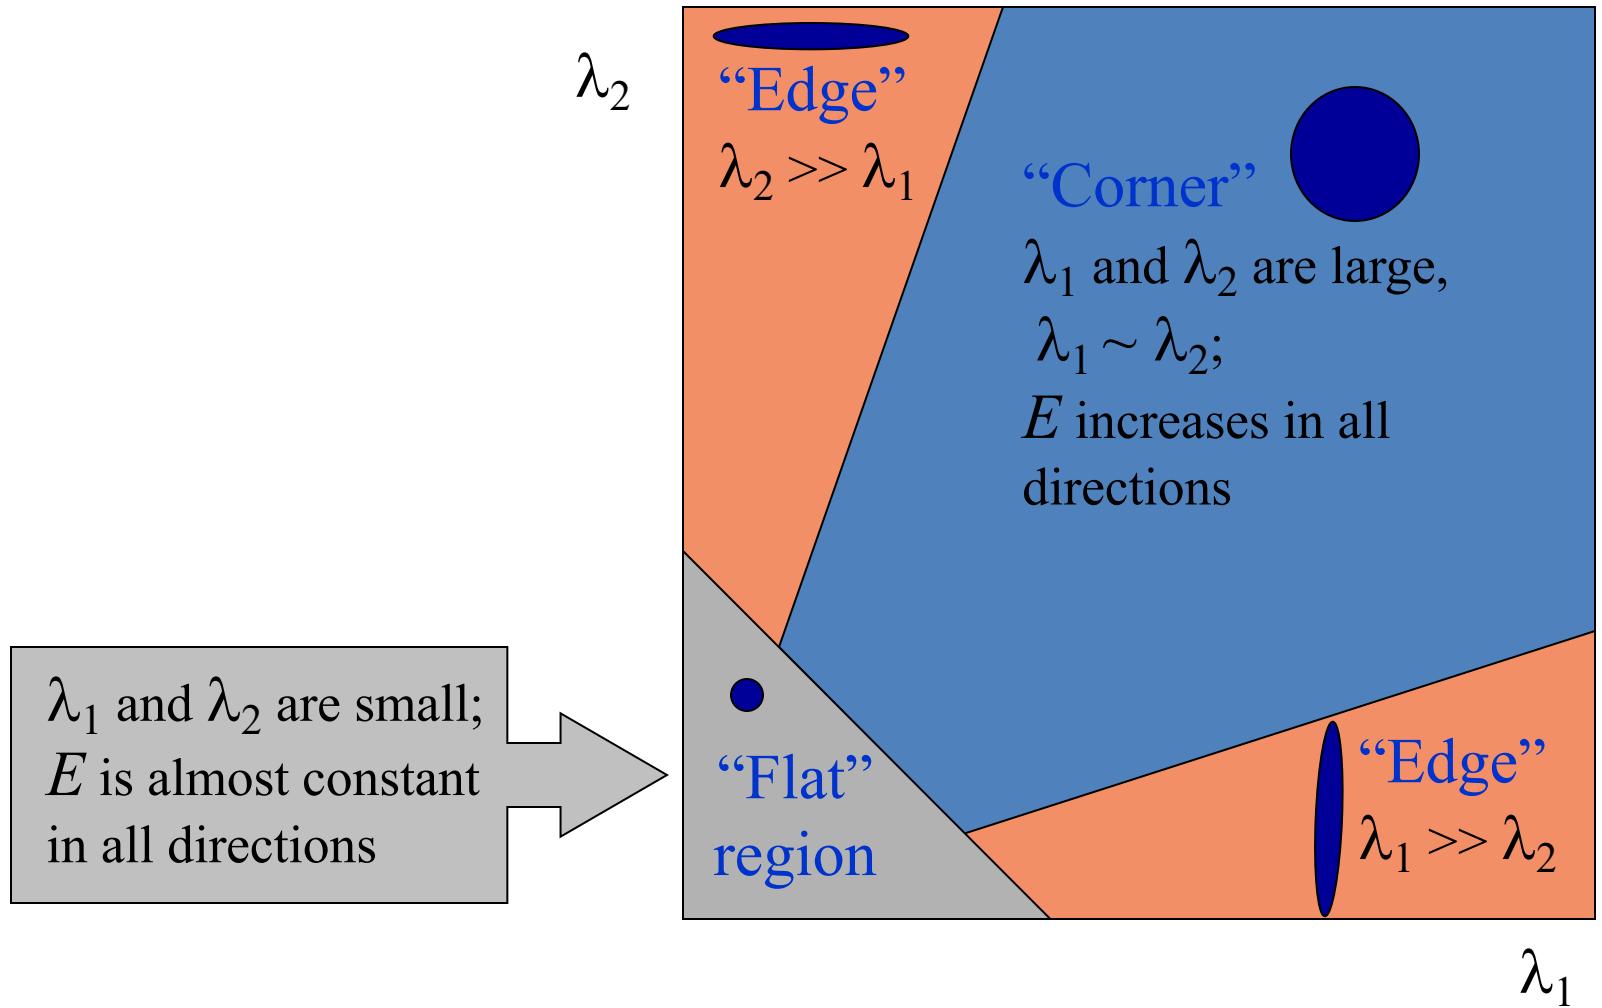
Slide credit: S. Lazebnik

Visualizing M



Technical note: M is often best *visualized* by first taking inverse, so long edge of ellipse goes along edge

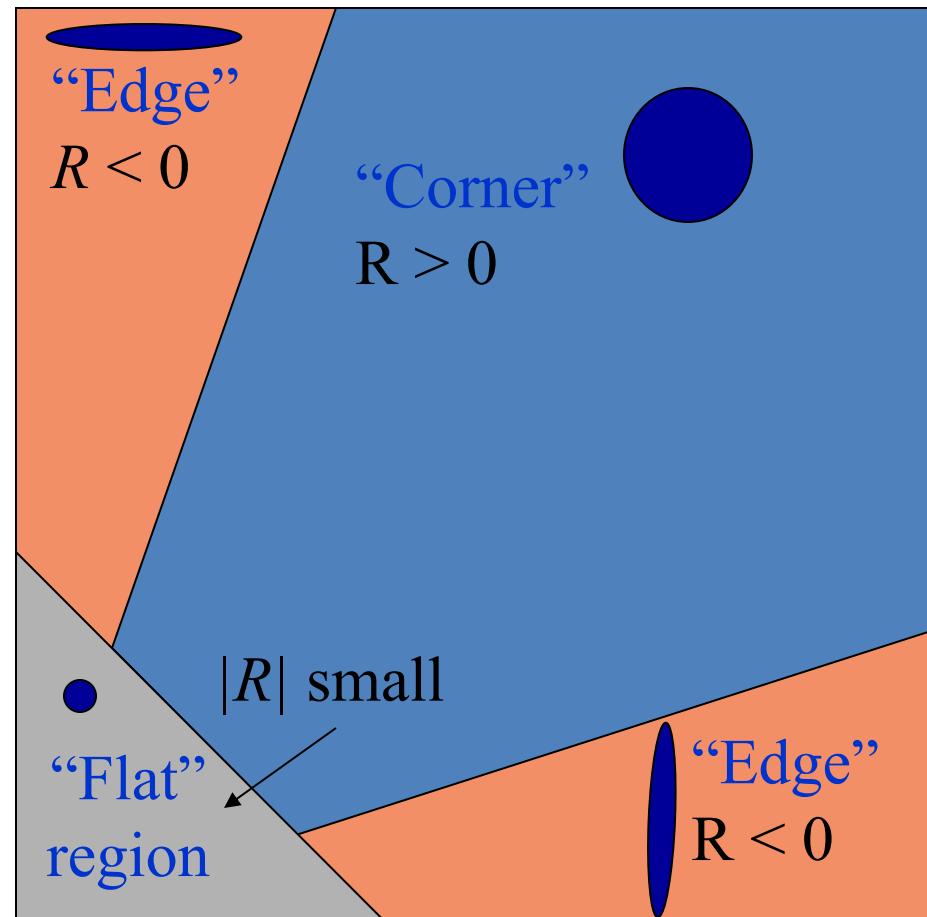
Interpreting Eigenvalues of M



Putting Together The Eigenvalues

$$\begin{aligned} R &= \det(\mathbf{M}) - \alpha \operatorname{trace}(\mathbf{M})^2 \\ &= \lambda_1 \lambda_2 - \alpha (\lambda_1 + \lambda_2)^2 \end{aligned}$$

α : constant (0.04 to 0.06)



Corners

9300 Harris Corners Pkwy, Charlotte, NC



Derivatives Review

Given quadratic function $f(x)$

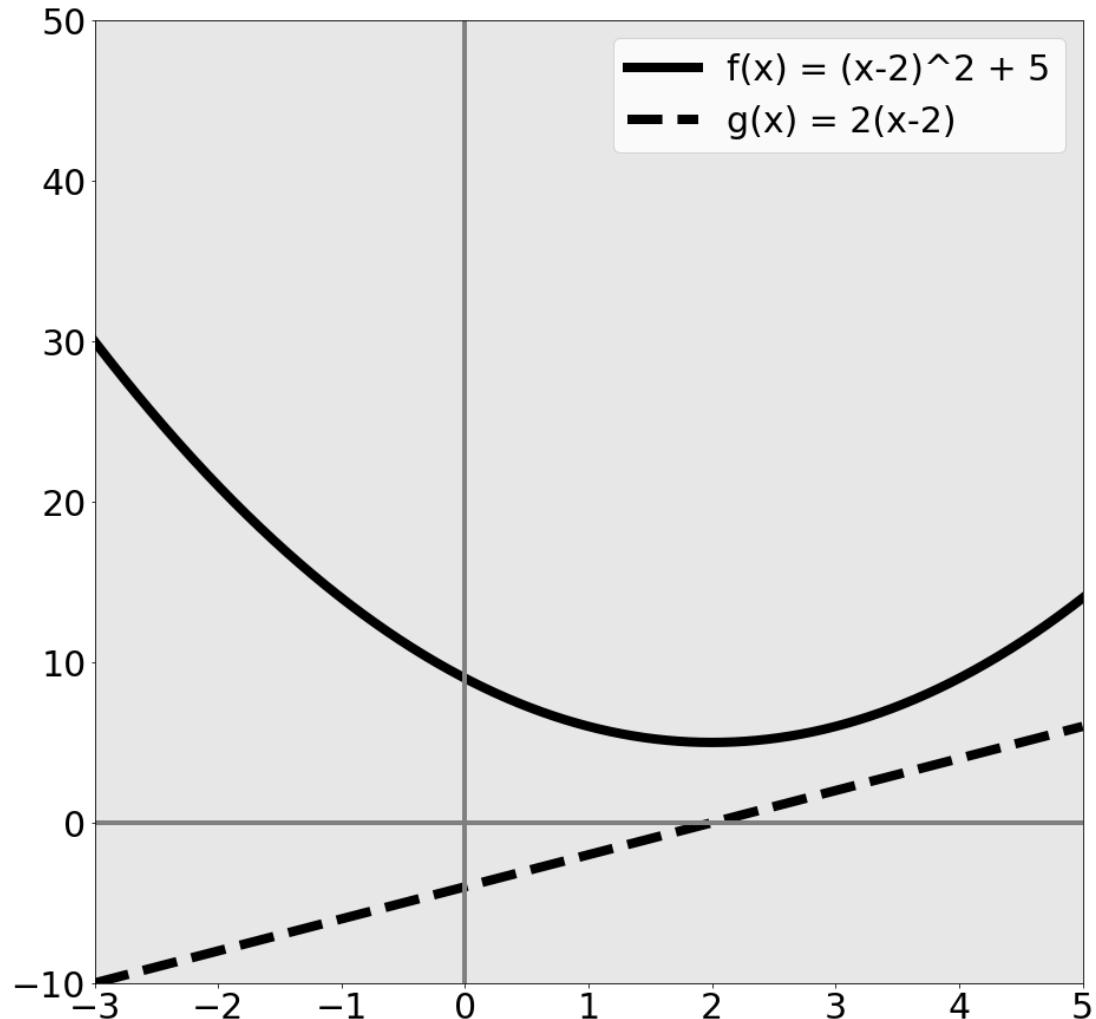
$$f(x) = (x - 2)^2 + 5$$

$f(x)$ is function

$$g(x) = f'(x)$$

aka

$$g(x) = \frac{d}{dx} f(x)$$



Given quadratic function $f(x)$

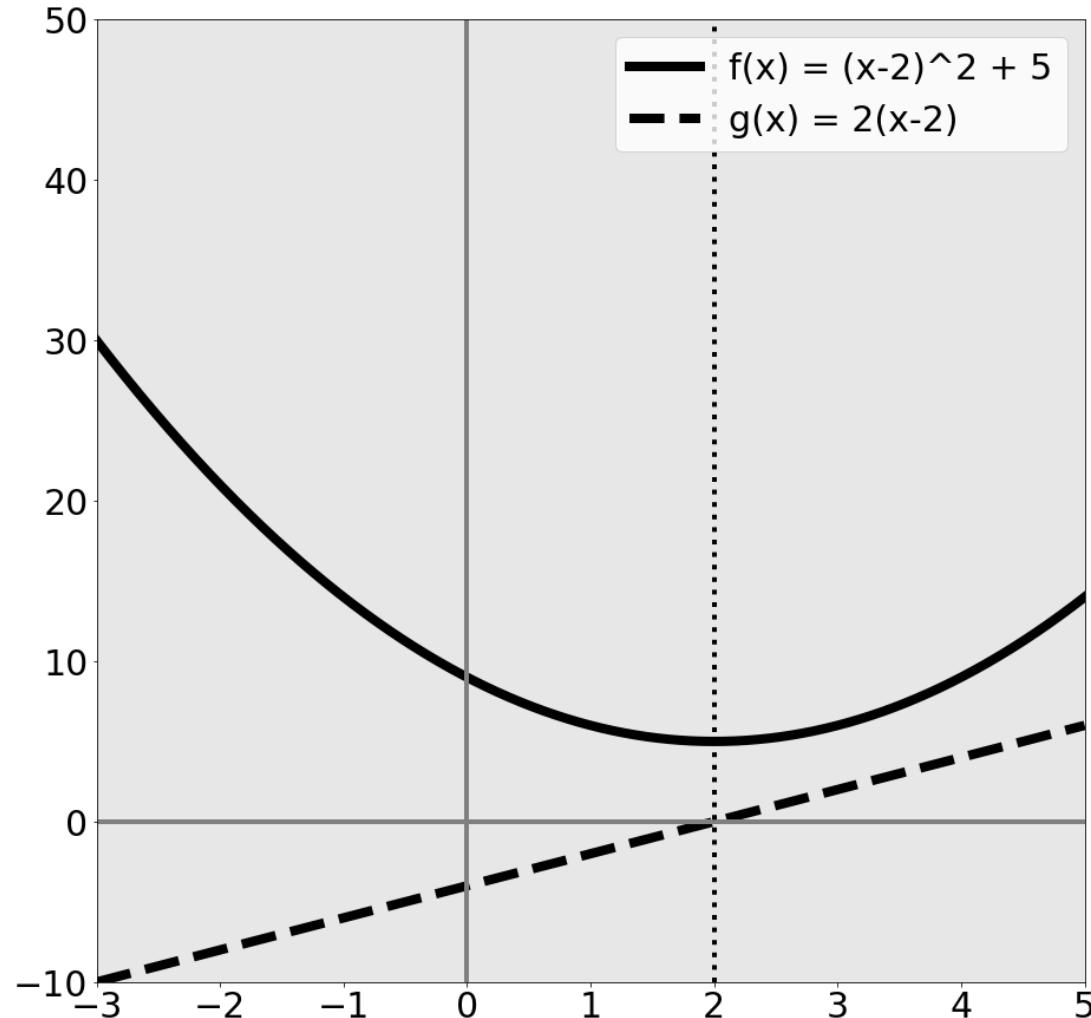
$$f(x) = (x - 2)^2 + 5$$

What's special about $x=2$?

$f(x)$ minim. at 2
 $g(x) = 0$ at 2

$a = \text{minimum of } f \rightarrow$
 $g(a) = 0$

Reverse is **not true**



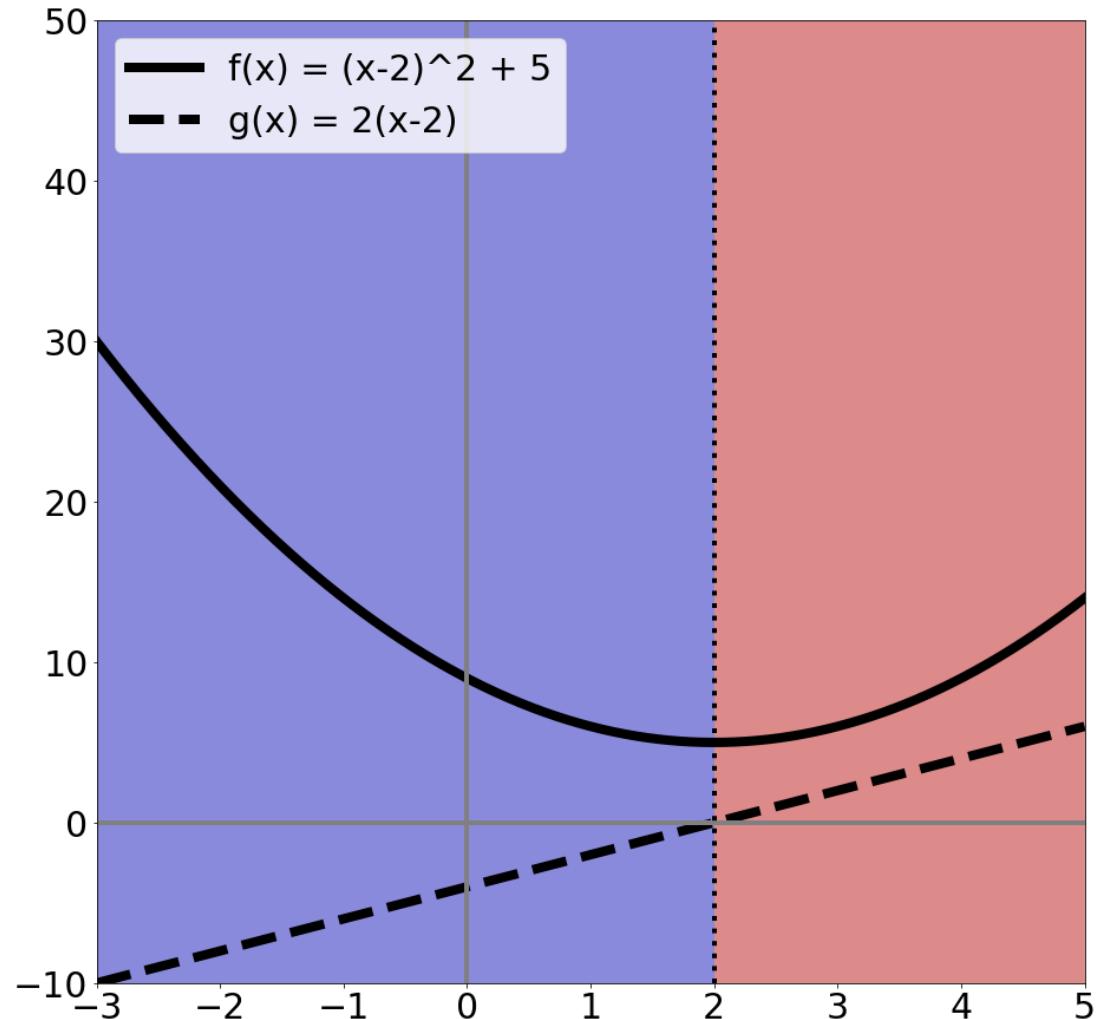
Rates of change

$$f(x) = (x - 2)^2 + 5$$

Suppose I want to increase $f(x)$ by changing x :

Blue area: move left
Red area: move right

Derivative tells you direction of ascent and rate



What Calculus Should I Know

- Really need intuition
- Need chain rule
- Rest you should look up / use a computer algebra system / use a cookbook
- Partial derivatives (and that's it from multivariable calculus)

Partial Derivatives

- Pretend other variables are constant, take a derivative. That's it.
- Make our function a function of two variables

$$f(x) = (x - 2)^2 + 5$$

$$\frac{\partial}{\partial x} f(x) = 2(x - 2) * 1 = 2(x - 2)$$

$$f_2(x, y) = (x - 2)^2 + 5 + (y + 1)^2$$

$$\frac{\partial}{\partial x} f_2(x) = 2(x - 2)$$

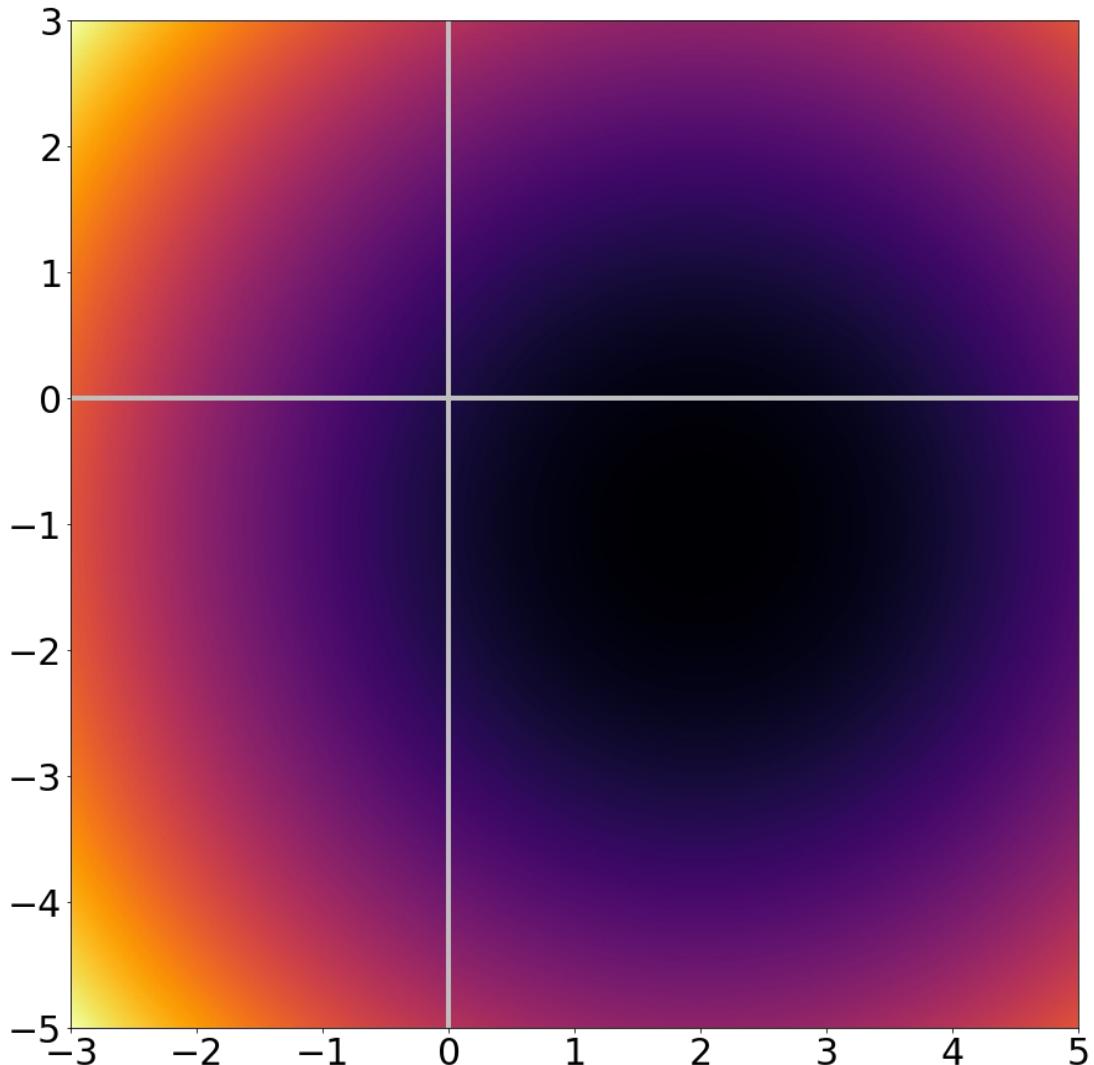
Pretend it's
constant →
derivative = 0

Zooming Out

$$f_2(x, y) = (x - 2)^2 + 5 + (y + 1)^2$$

Dark = $f(x,y)$ low

Bright = $f(x,y)$ high



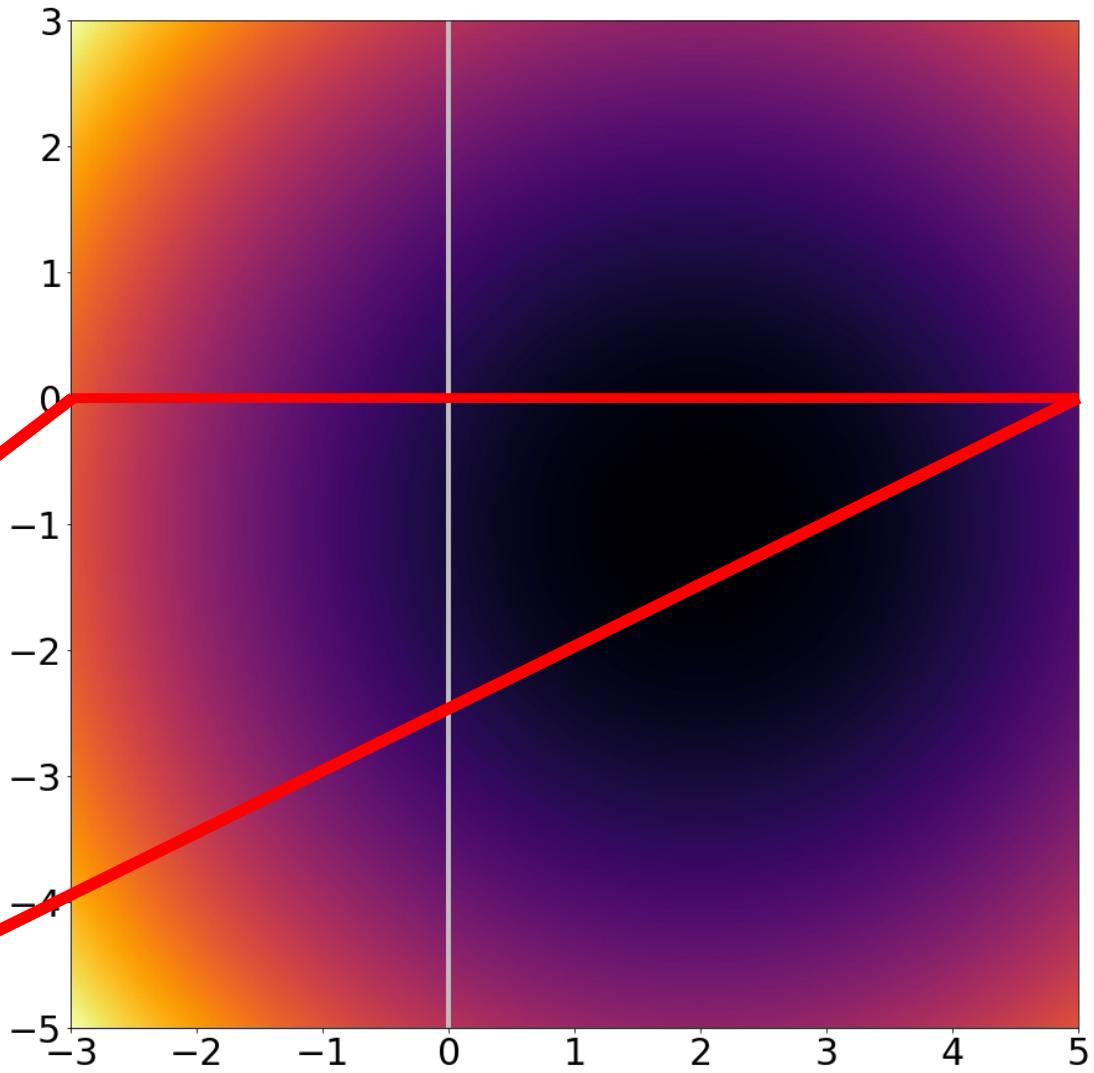
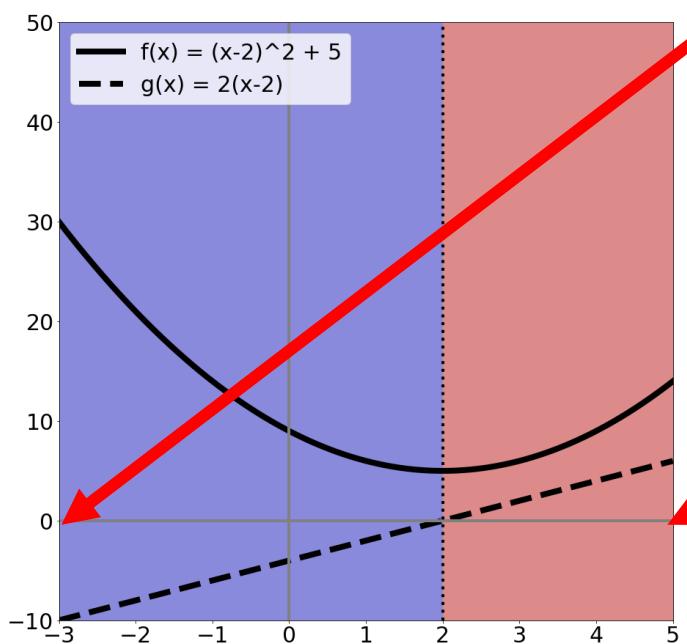
Taking a slice of

$$f_2(x, y) = (x - 2)^2 + 5 + (y + 1)^2$$

Slice of $y=0$ is the function from before:

$$f(x) = (x - 2)^2 + 5$$

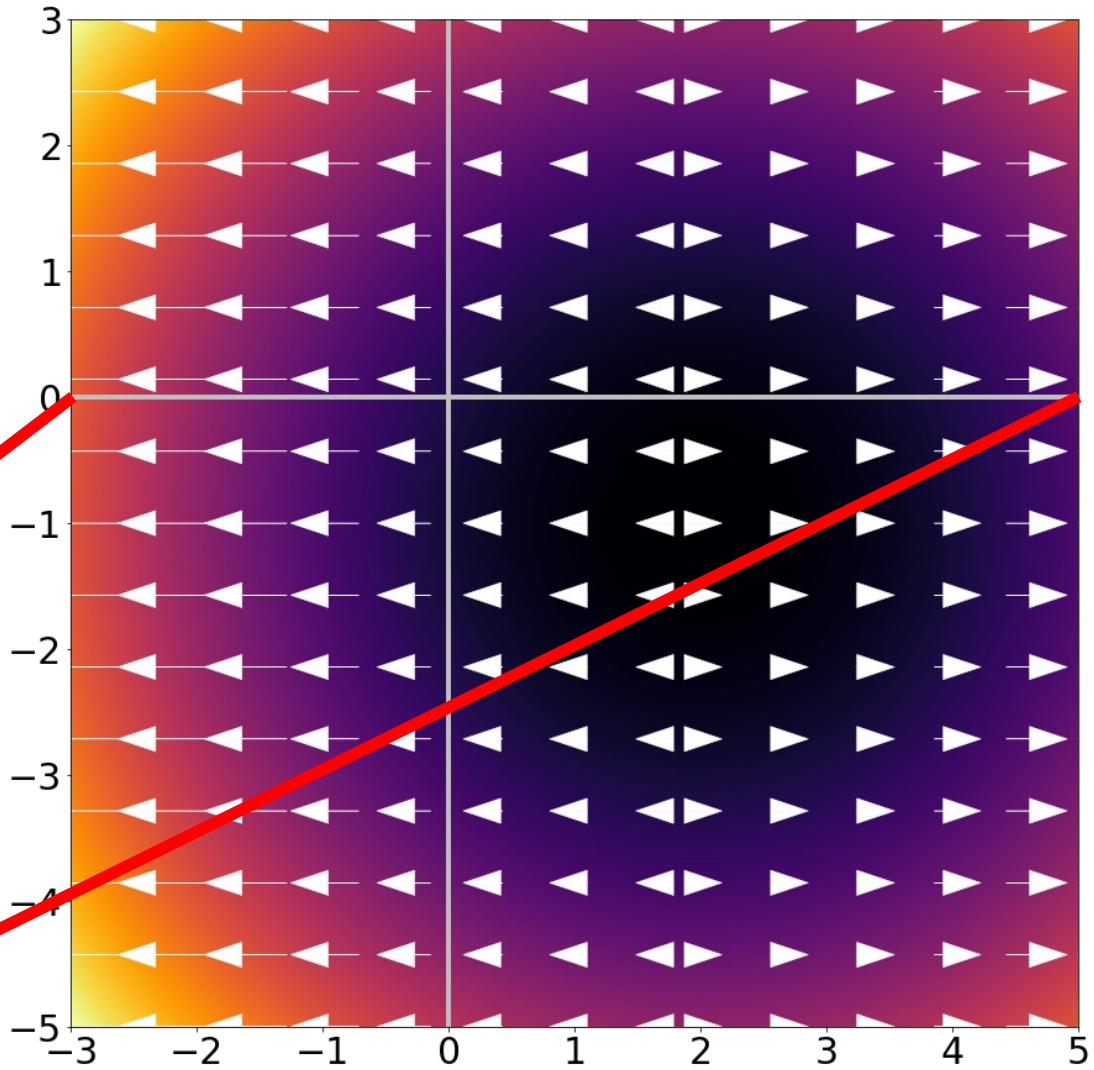
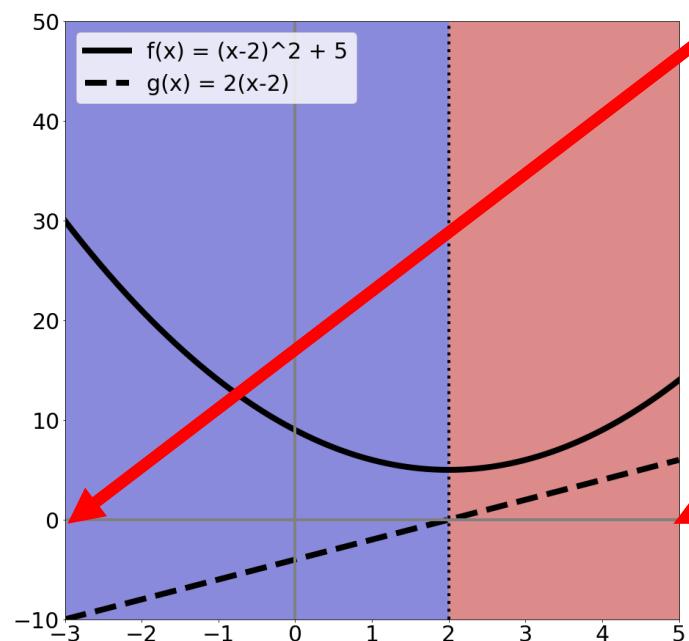
$$f'(x) = 2(x - 2)$$



Taking a slice of

$$f_2(x, y) = (x - 2)^2 + 5 + (y + 1)^2$$

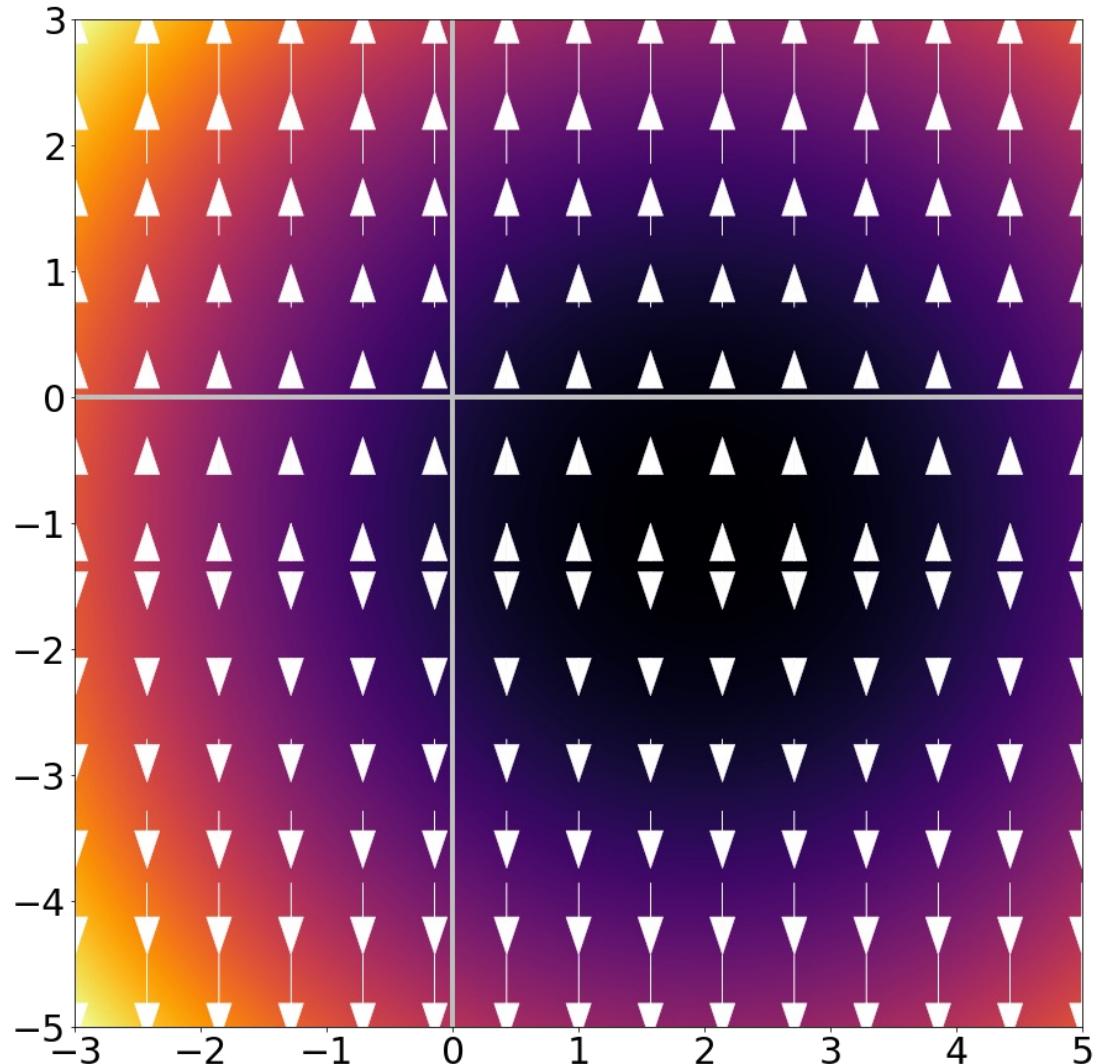
$\frac{\partial}{\partial x} f_2(x, y)$ is rate of change & direction in x dimension



Zooming Out

$$f_2(x, y) = (x - 2)^2 + 5 + (y + 1)^2$$

$\frac{\partial}{\partial y} f_2(x, y)$ is
 $2(y + 1)$
and is the rate of
change & direction in
y dimension



Zooming Out

$$f_2(x, y) = (x - 2)^2 + 5 + (y + 1)^2$$

Gradient/Jacobian:

Making a vector of

$$\nabla_f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

gives rate and
direction of change.

Arrows point OUT of
minimum / basin.

