

# Image Filtering

EECS 442 – Jeong Joon Park  
Winter 2024, University of Michigan

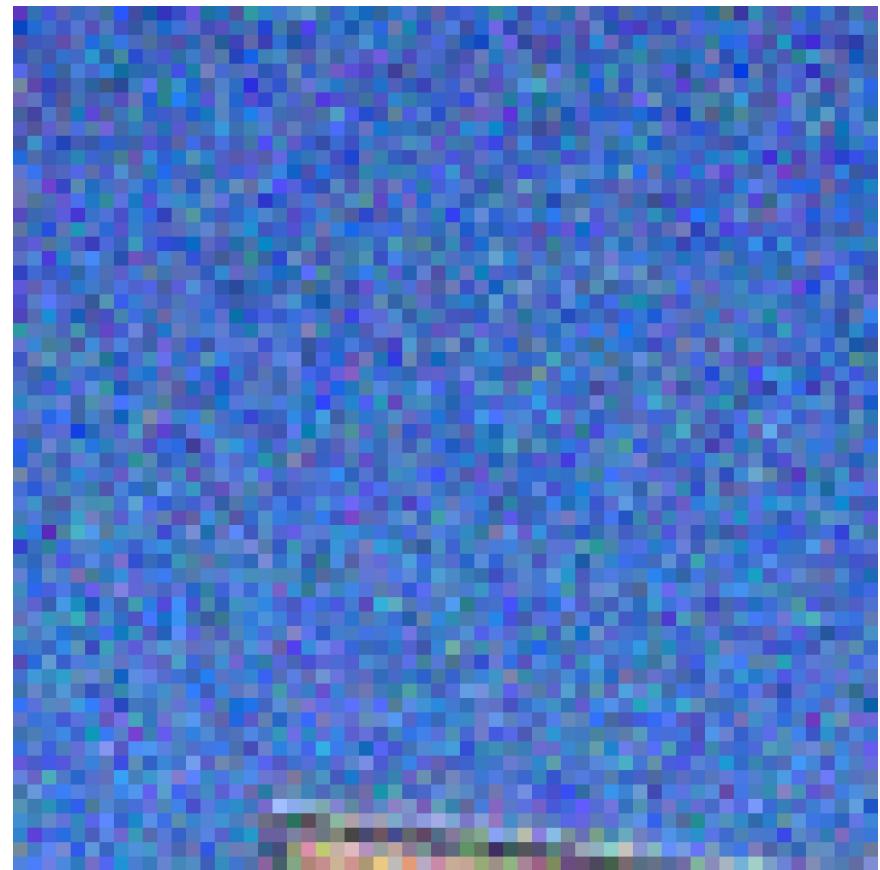
# Administrivia

- HW1 due Wednesday at 11:59pm (although you do have late tokens...)
  - TAs work 9 AM-5 PM. So don't expect answers late
- Next few classes and homeworks: building momentum towards stuff
- Good news: probably won't be as tricky implementation wise now that you've seen the bugs
- Bad news: will be trickier conceptual-wise

# Administrivia

- Piazza and Office Hours
  1. Don't post code publicly (Or pictures)
  2. Ask *specific questions* about small code snippets (1 or 2 lines)
    - We cannot debug your code. 7TAs for 280 students
    - Prepare specific questions for OH
  3. No "does this look right?" questions
  4. Show respect. No "Hey, can you help on this, lol?"
  5. Carefully Read the Announcement

# Let's Take An Image



# Let's Fix Things

- We have noise in our image
- Let's replace each pixel with a *weighted* average of its neighborhood
- Weights are *filter kernel*

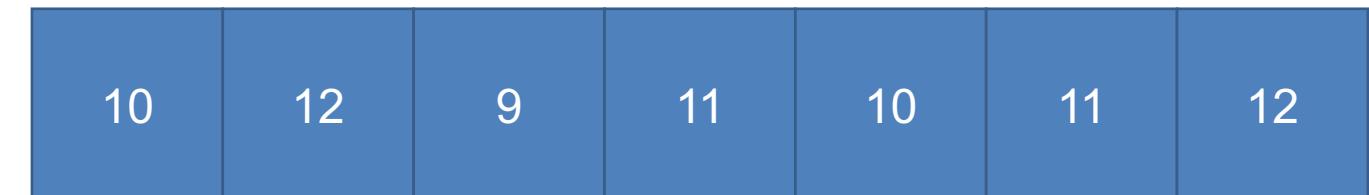
	Out	

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

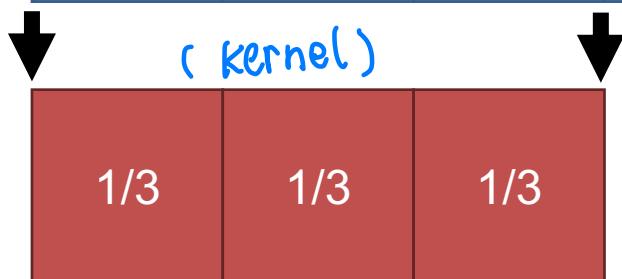
`np.dot(input, filter)`

# 1D Case

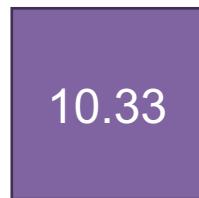
Signal



Filter



Output

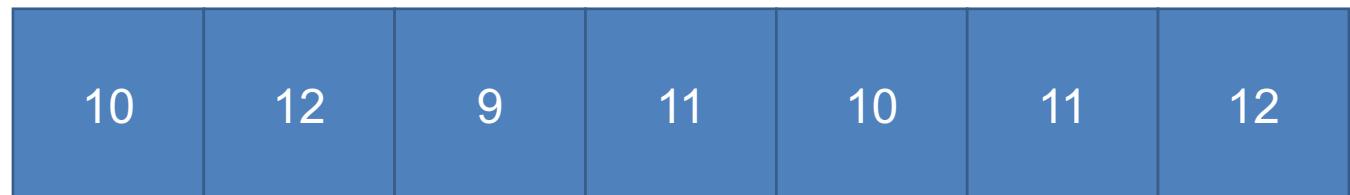


What's the average  
of 9, 10, 12?

- (a) 9
- (b) 11.5
- (c) 10.33
- (d) 11.66

# 1D Case

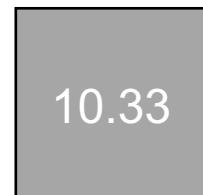
Signal



Filter



Output

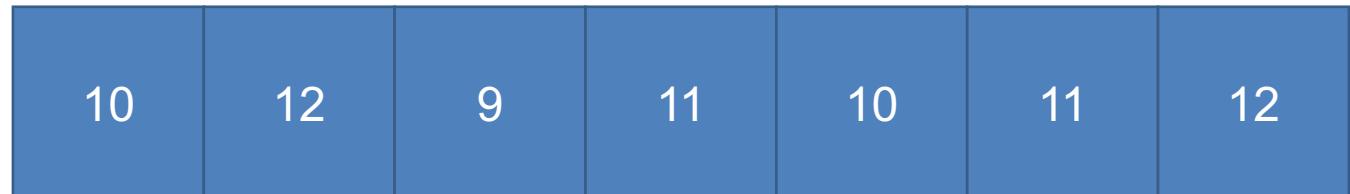


Done! Next?

# 1D Case

(1) 10.66 (2) 9.33  
(3) 14.2 (4) 11.33

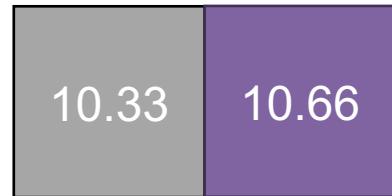
Signal



Filter



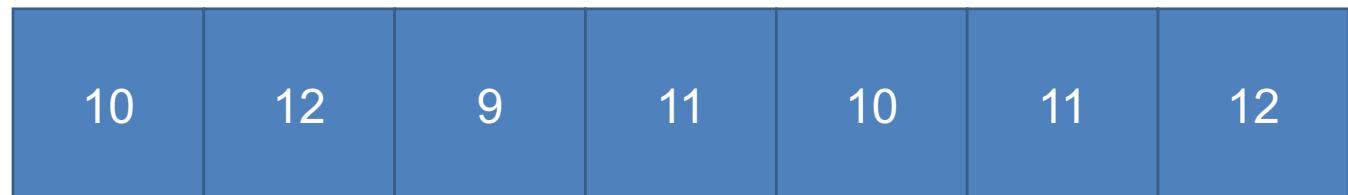
Output



# 1D Case

(1) 10.33 (2) 11.33  
(3) 10 (4) 9.1

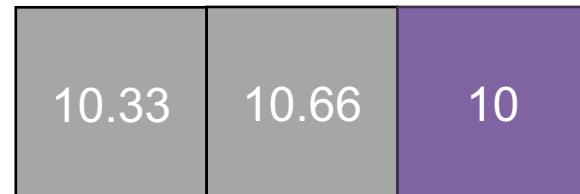
Signal



Filter

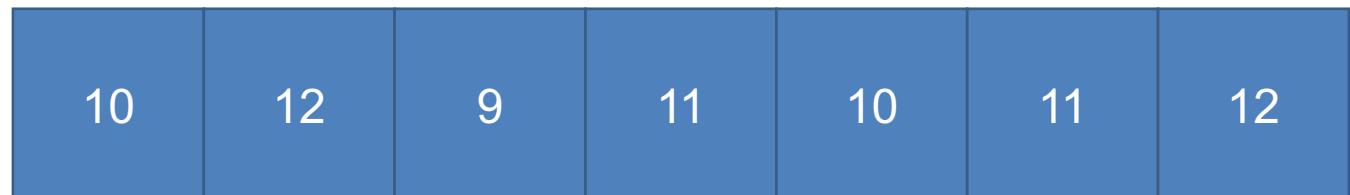


Output



# 1D Case

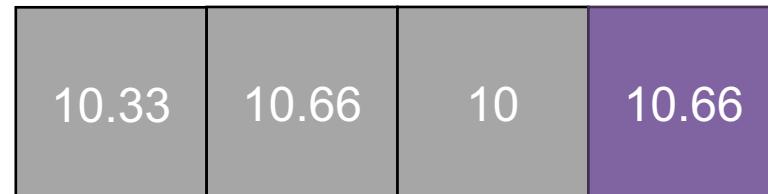
Signal



Filter

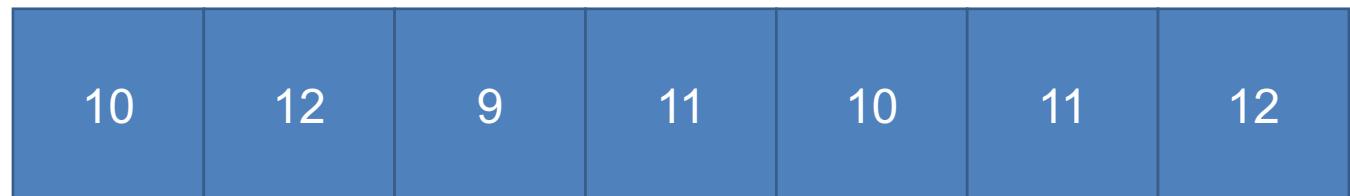


Output



# 1D Case

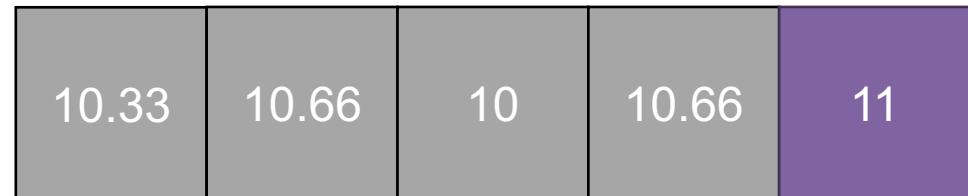
Signal



Filter

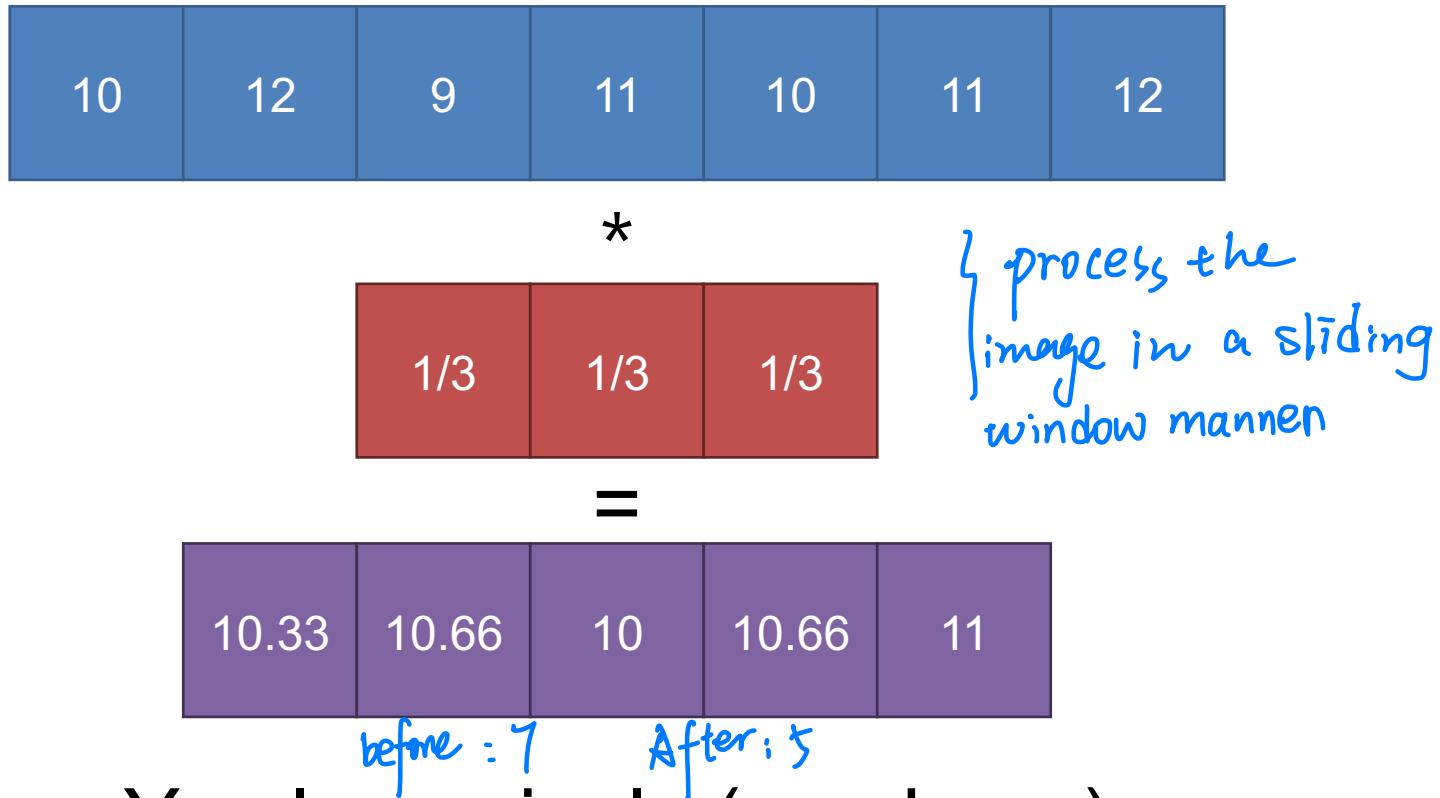


Output



This concept is called sliding window.

## 1D Case



You lose pixels (maybe...)

Filter “sees” only a few pixels at a time

Sliding window    Here, only see 3 at each time .

# Applying a Linear Filter

Input

I11	I12	I13	I14	I15	I16
I21	I22	I23	I24	I25	I26
I31	I32	I33	I34	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Filter

F11	F12	F13
F21	F22	F23
F31	F32	F33

Output

O11	O12	O13	O14
O21	O22	O23	O24
O31	O32	O33	O34

# Applying a Linear Filter

Input & Filter

F11	F12	F13	I14	I15	I16
F21	F22	F23	I24	I25	I26
F31	F32	F33	I34	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Output

O11

$$O_{11} = I_{11} * F_{11} + I_{12} * F_{12} + \dots + I_{33} * F_{33}$$

# Applying a Linear Filter

Input & Filter

I11	F11	F12	F13	I15	I16
I21	F21	F22	F23	I25	I26
I31	F31	F32	F33	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Output

O11	O12
-----	-----

$$O_{12} = I_{12} * F_{11} + I_{13} * F_{12} + \dots + I_{34} * F_{33}$$

# Applying a Linear Filter

Input

I11	I12	I13	I14	I15	I16
I21	I22	I23	I24	I25	I26
I31	I32	I33	I34	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Filter

F11	F12	F13
F21	F22	F23
F31	F32	F33

Output

How many times can we apply a  
3x3 filter to a 5x6 image?

# Applying a Linear Filter

Input

I11	I12	I13	I14	I15	I16
I21	I22	I23	I24	I25	I26
I31	I32	I33	I34	I35	I36
I41	I42	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56

Filter

F11	F12	F13
F21	F22	F23
F31	F32	F33

Output

O11	O12	O13	O14
O21	O22	O23	O24
O31	O32	O33	O34

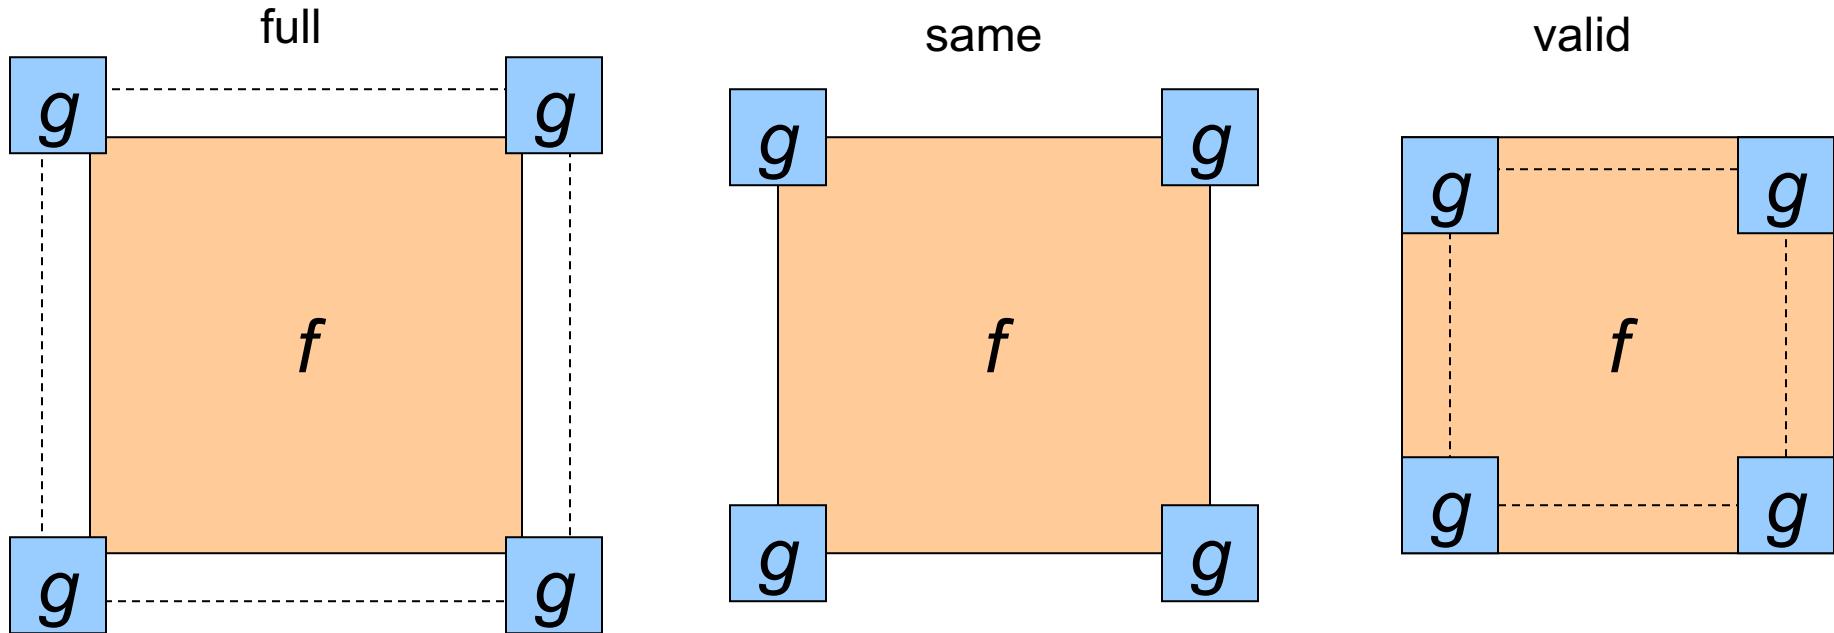
$$O_{ij} = I_{ij} * F_{11} + I_{i(j+1)} * F_{12} + \dots + I_{(i+2)(j+2)} * F_{33}$$

# Painful Details – Edge Cases

Filtering doesn't keep the whole image.

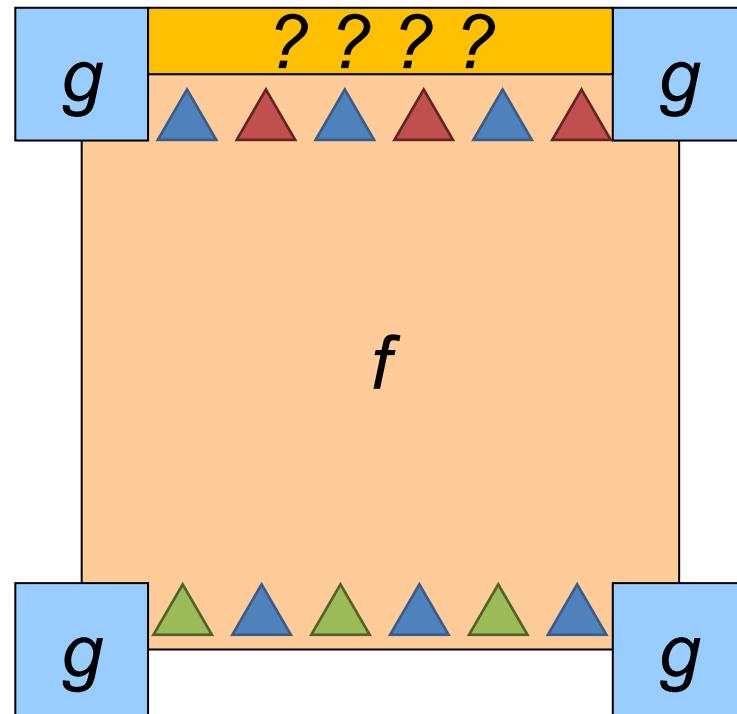
Suppose **f** is the image and **g** the filter.

**Full** – any part of g touches f. **Same** – same size as f;  
**Valid** – only when filter doesn't fall off edge.

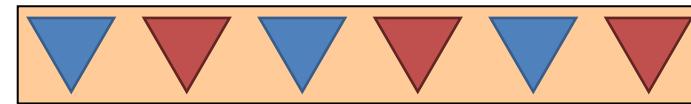


# Painful Details – Edge Cases

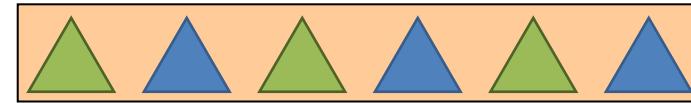
What to about the “?” region?



Symm: fold sides over



Circular/Wrap: wrap around



pad/fill: add value, often 0

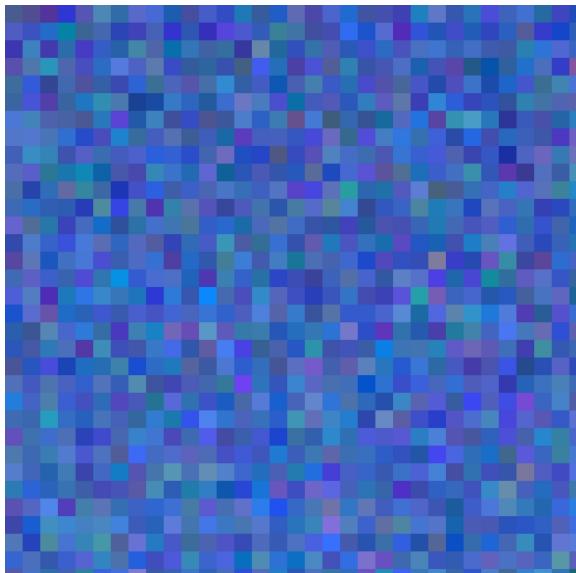


# Painful Details – Does it Matter?

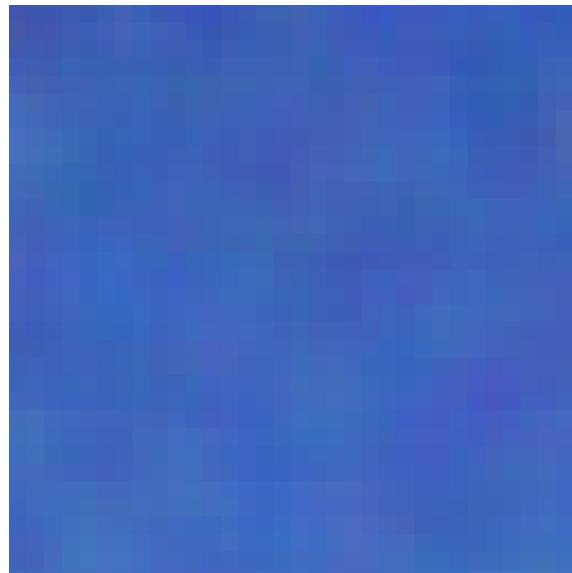
(I've applied the filter per-color channel)

**Which padding did I use and why?**

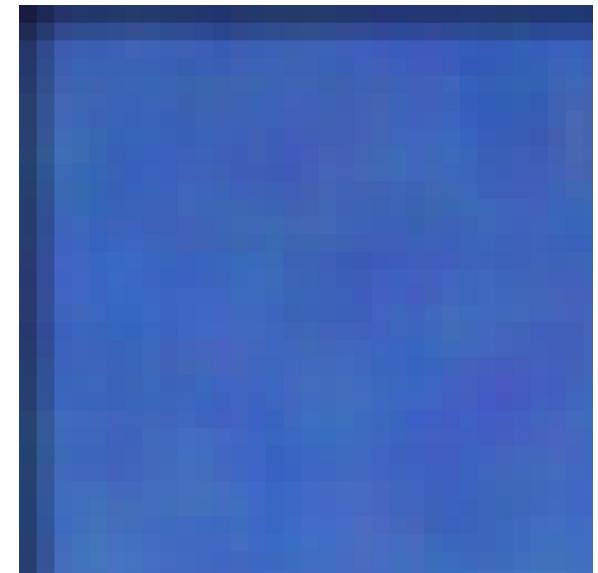
Input  
Image



Filtered  
???



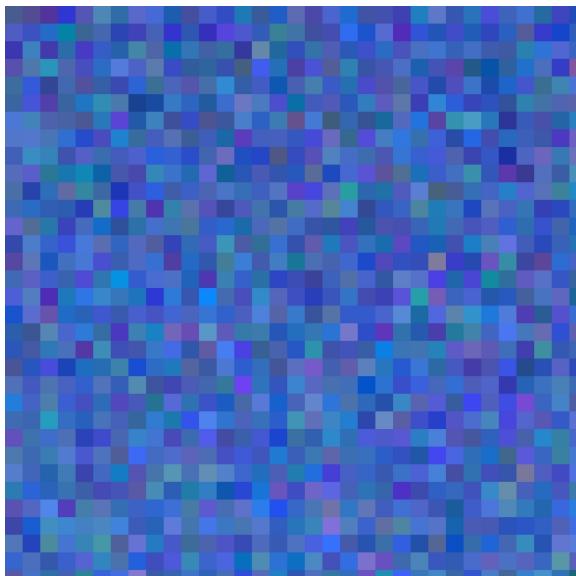
Filtered  
???



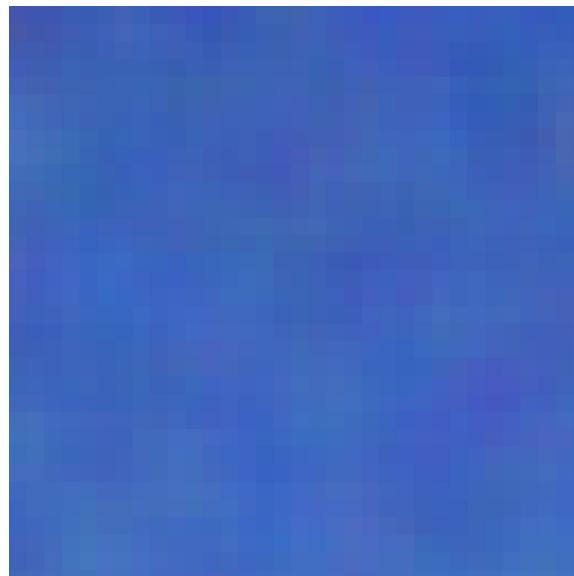
# Painful Details – Does it Matter?

(I've applied the filter per-color channel)

Input  
Image



Filtered  
Symm Pad



Filtered  
Zero Pad



# Practice with Linear Filters



Original

0	0	0
0	1	0
0	0	0

?

# Practice with Linear Filters



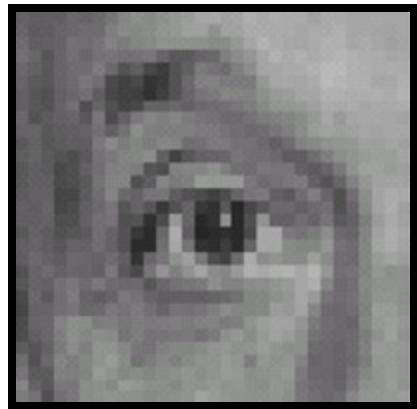
Original

0	0	0
0	1	0
0	0	0



The Same!

# Practice with Linear Filters



Original

0	0	0
0	0	1
0	0	0

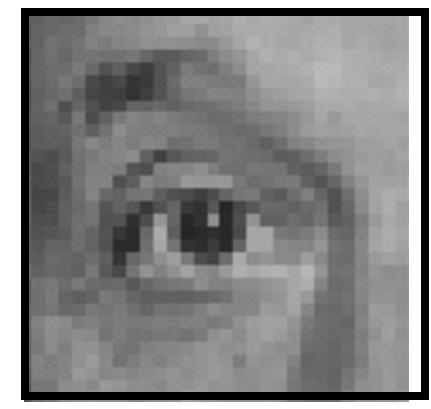
?

# Practice with Linear Filters



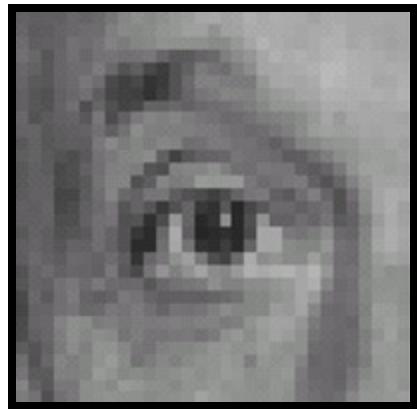
Original

0	0	0
0	0	1
0	0	0



Shifted  
LEFT  
1 pixel

# Practice with Linear Filters



Original

0	1	0
0	0	0
0	0	0

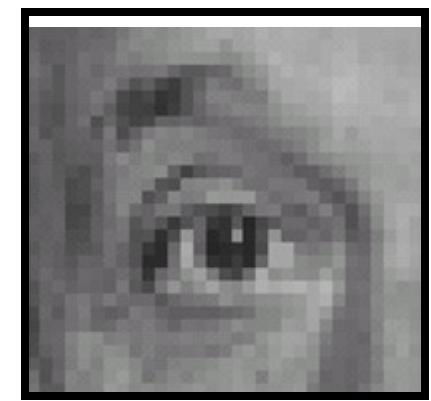
?

# Practice with Linear Filters



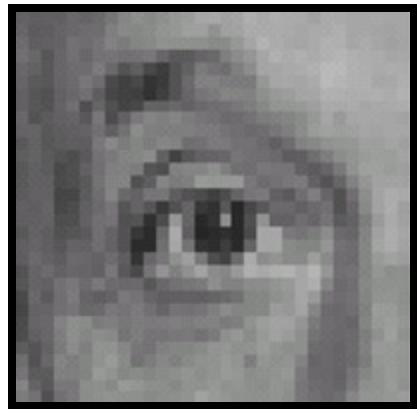
Original

0	1	0
0	0	0
0	0	0



Shifted  
**DOWN**  
1 pixel

# Practice with Linear Filters



Original

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

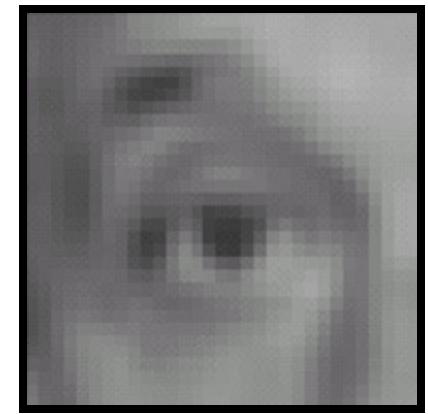
?

# Practice with Linear Filters



Original

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9



Blur  
(Box Filter)

# Practice with Linear Filters



0	0	0
0	2	0
0	0	0

-

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

?

# Practice with Linear Filters

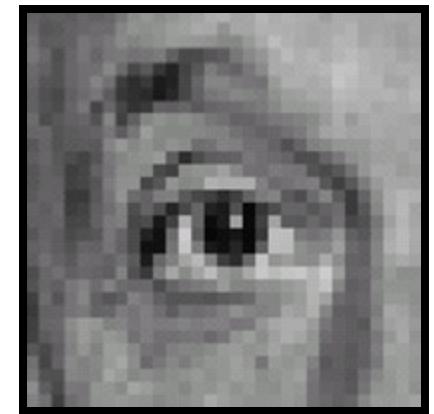


Original

0	0	0
0	2	0
0	0	0

-

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

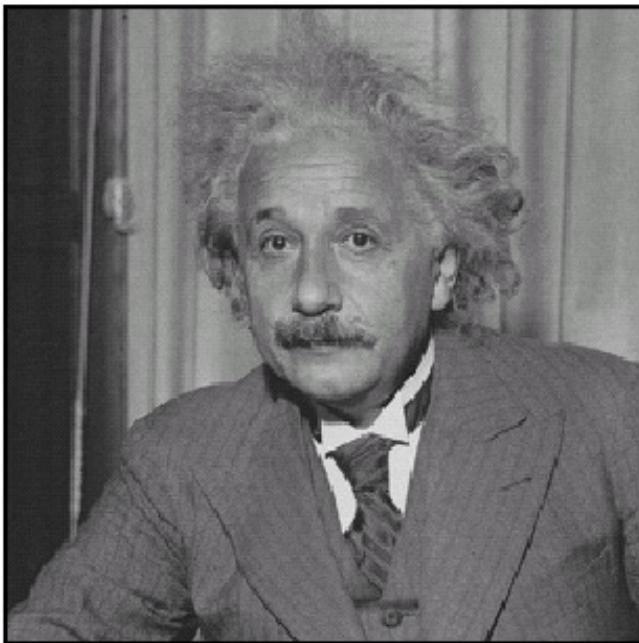


Sharpened

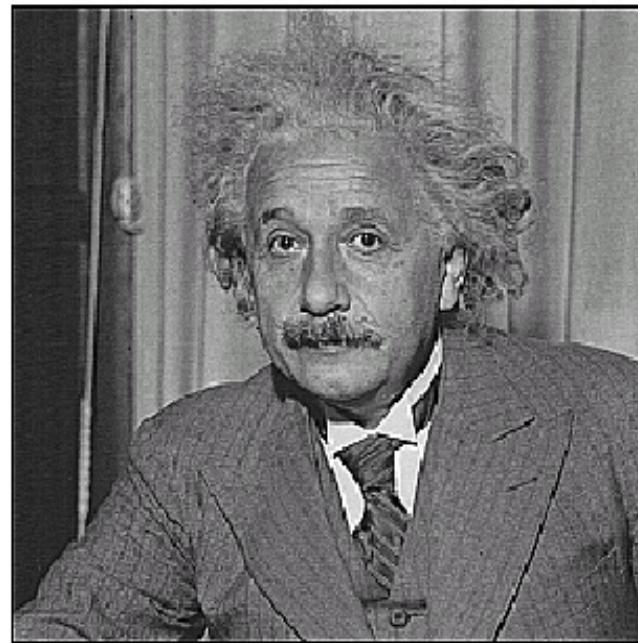
(Accentuates  
difference from  
local average)

*emphasizing the high-frequency  
details of this image  
↓  
gives me a more sharpened image*

# Sharpening



**before**



**after**

# Properties – Linear

Input

I11	I12	I13	I14
I21	I22	I23	I24
I31	I32	I33	I34
I41	I42	I43	I44

Filter

F11	F12	F13
F21	F22	F23
F31	F32	F33

Output

O11	O12
O21	O22

We want to do by simple matrix multiplication.

Flatten Input and multiply some matrix W  
Can exactly mimic the sliding window filtering  
Linear filtering is linear!

# Properties – Linear

Input

I11	I12	I13	I14
I21	I22	I23	I24
I31	I32	I33	I34
I41	I42	I43	I44

Filter

F11	F12	F13
F21	F22	F23
F31	F32	F33

Output

O11	O12
O21	O22

$$\begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} & w_{1,5} & w_{1,6} & w_{1,7} & w_{1,8} & w_{1,9} & w_{1,10} & w_{1,11} & w_{1,12} & w_{1,13} & w_{1,14} & w_{1,15} & w_{1,16} \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} & w_{2,5} & w_{2,6} & w_{2,7} & w_{2,8} & w_{2,9} & w_{2,10} & w_{2,11} & w_{2,12} & w_{2,13} & w_{2,14} & w_{2,15} & w_{2,16} \\ w_{3,1} & w_{3,2} & w_{3,3} & w_{3,4} & w_{3,5} & w_{3,6} & w_{3,7} & w_{3,8} & w_{3,9} & w_{3,10} & w_{3,11} & w_{3,12} & w_{3,13} & w_{3,14} & w_{3,15} & w_{3,16} \\ w_{4,1} & w_{4,2} & w_{4,3} & w_{4,4} & w_{4,5} & w_{4,6} & w_{4,7} & w_{4,8} & w_{4,9} & w_{4,10} & w_{4,11} & w_{4,12} & w_{4,13} & w_{4,14} & w_{4,15} & w_{4,16} \end{bmatrix}$$

W

# Properties – Linear

Assume: I image f1, f2 filters

**Linear:**  $\text{apply}(I, f1 + f2) = \text{apply}(I, f1) + \text{apply}(I, f2)$

I is a white box on black, and f1, f2 are rectangles

$$A(\begin{array}{|c|}\hline \bullet \\ \hline\end{array}, \begin{array}{|c|}\hline \square \\ \hline\end{array} + \begin{array}{|c|}\hline \square \\ \hline\end{array}) = A(\begin{array}{|c|}\hline \bullet \\ \hline\end{array}, \begin{array}{|c|}\hline \square \\ \hline\end{array}) = \begin{array}{|c|}\hline \bullet \\ \hline\end{array}$$

$$A(\begin{array}{|c|}\hline \bullet \\ \hline\end{array}, \begin{array}{|c|}\hline \square \\ \hline\end{array}) + A(\begin{array}{|c|}\hline \bullet \\ \hline\end{array}, \begin{array}{|c|}\hline \square \\ \hline\end{array}) = \begin{array}{|c|}\hline \square \\ \hline\end{array} + \begin{array}{|c|}\hline \square \\ \hline\end{array} = \begin{array}{|c|}\hline \bullet \\ \hline\end{array}$$

Note: I am showing filters un-normalized and blown up. They're a smaller box filter (i.e., each entry is  $1/(\text{size}^2)$ )

# Properties – Shift-Invariant

Assume: I image, f filter

**Shift-invariant:**  $\text{shift}(\text{apply}(I, f)) = \text{apply}(\text{shift}(I, f))$

Intuitively: only depends on filter neighborhood

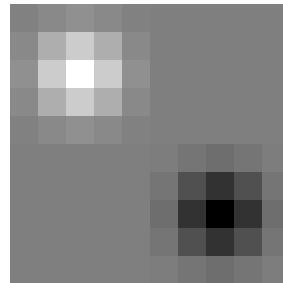
$$A( \begin{array}{|c|} \hline \square \\ \hline \end{array}, \boxed{\phantom{0}} ) = \begin{array}{|c|c|c|} \hline & & \\ \hline & \square & \\ \hline & & \\ \hline \end{array}$$

$$A( \begin{array}{|c|} \hline \\ \hline \square \\ \hline \end{array}, \boxed{\phantom{0}} ) = \begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & \square & \\ \hline \end{array}$$

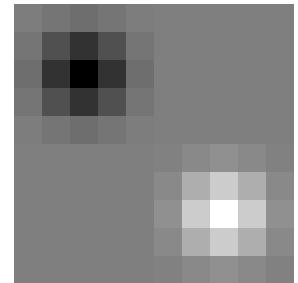
# Painful Details – Signal Processing

What I just showed is often called “convolution”.  
*Actually* cross-correlation. Source of **terrible** confusion.

Cross-Correlation  
(Original Orientation)



Convolution  
(Flip filter in x,y first)



# Convolution

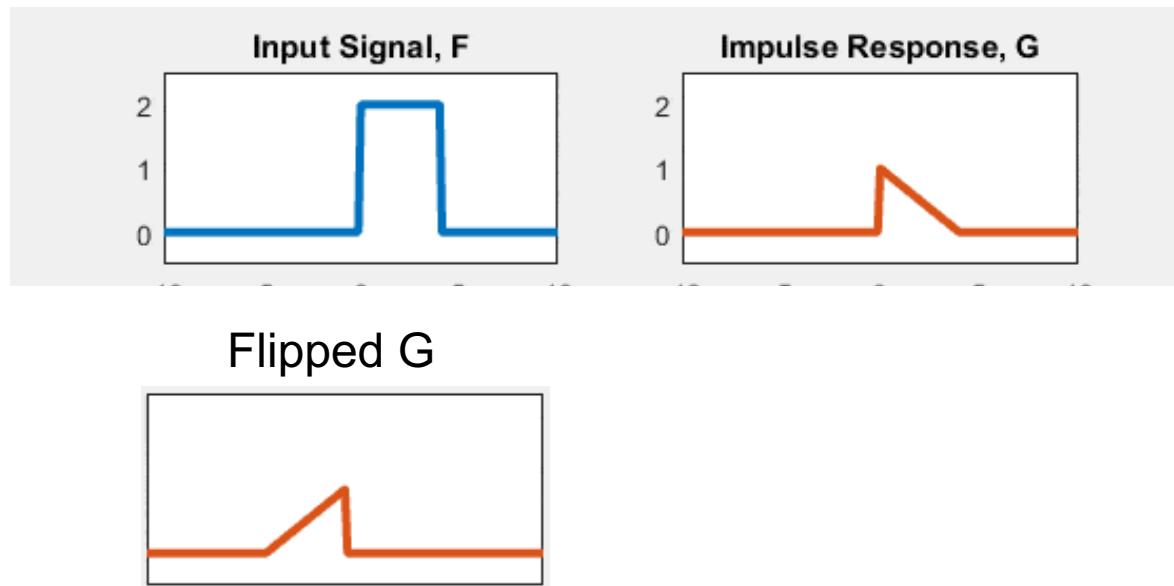
To be more clear:

```
def imageFilter2D(image, filter):
    for y in range(...): for x in range(...)
        # you'll fill this in
    return filtered
```

```
def imageConvolve2D(image, filter):
    # flip the filter left/right and up/down
    filterUse = np.fliplr(np.flipud(filter))
    return imageFilter2D(image, filterUse)
```

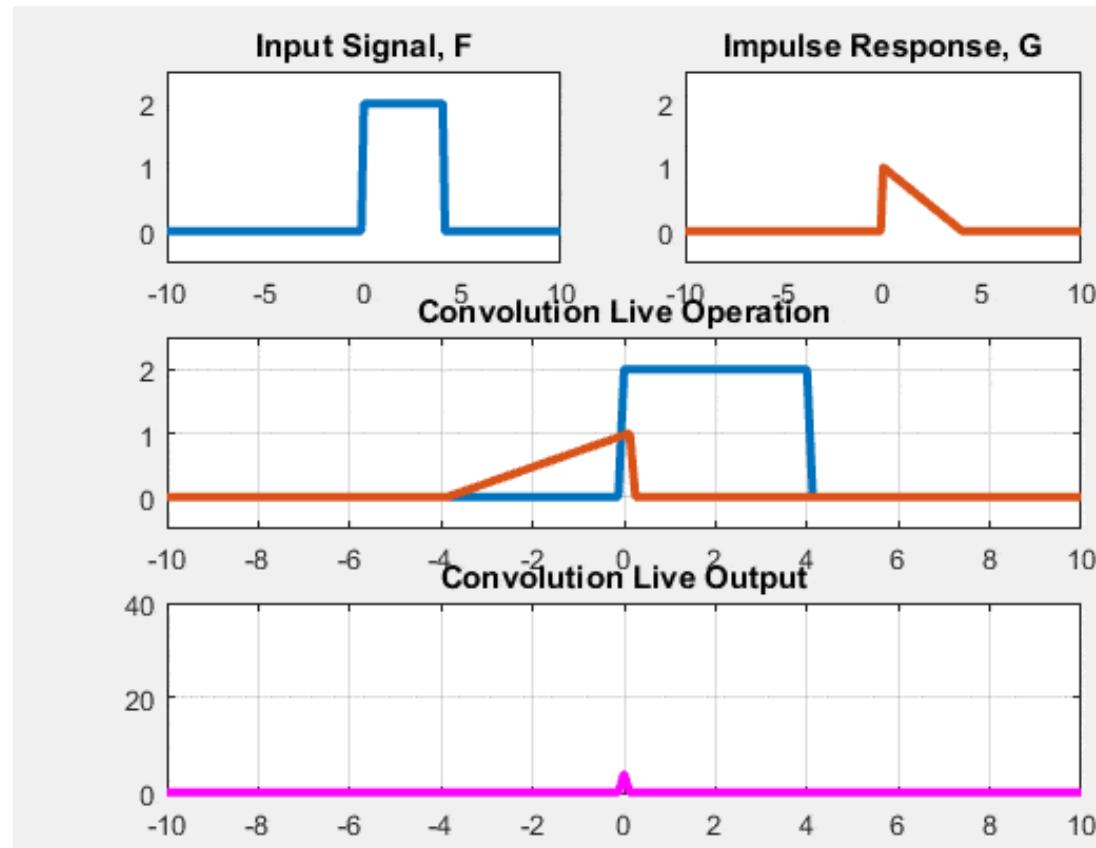
# Convolutions

- $(f * g)[t] = \sum_{\tau=-\infty}^{\infty} f[\tau]g[t - \tau]$



# Convolutions

- $(f * g)[t] = \sum_{\tau=-\infty}^{\infty} f[\tau]g[t - \tau]$



# Convolutions is Commutative

- $(f * g)[t] = \sum_{\tau=-\infty}^{\infty} f[\tau]g[t - \tau]$
- Replace  $\tau$  with  $t - \tau$ :
- $(f * g)[t] = \sum_{\tau=-\infty}^{\infty} f[t - \tau]g[\tau] = (g * f)[t]$

# Cross-Correlation is Not

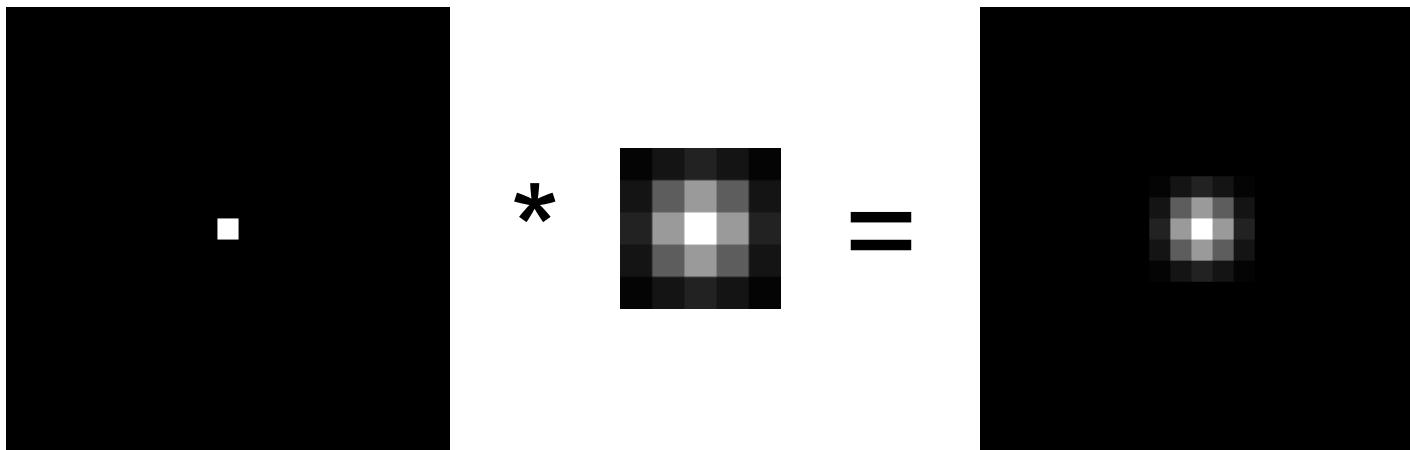
- $(f \star g)[t] = \sum_{\tau=-\infty}^{\infty} f[\tau]g[t + \tau]$
- Replace  $\tau$  with  $t + \tau$ :
- $(f \star g)[t] = \sum_{\tau=-\infty}^{\infty} f[t + \tau]g[2t + \tau]$   
 $\neq (g \star f)[t]$
- Not Commutative!

# Convolution vs Cross-Correlation

- What we actually do is Cross-Correlation
- But they are very similar
- And Computer Vision people just use the term “Convolution.”

# Properties of Convolution

- Convolution ( $*$ ) is a shift-invariant linear operator
- Commutative:  $f * g = g * f$
- Associative:  $(f * g) * h = f * (g * h)$
- Distributes over  $+$ :  $f * (g + h) = f * g + f * h$
- Scalars factor out:  $k f * g = f * kg = k (f * g)$
- Identity (a single one with all zeros):



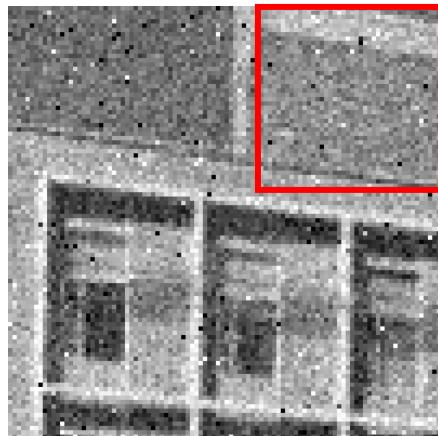
# Questions?

- Nearly everything onwards is a convolution.
- This is important to get right.

# Smoothing With A Box

Intuition: if filter touches it, it gets a contribution.

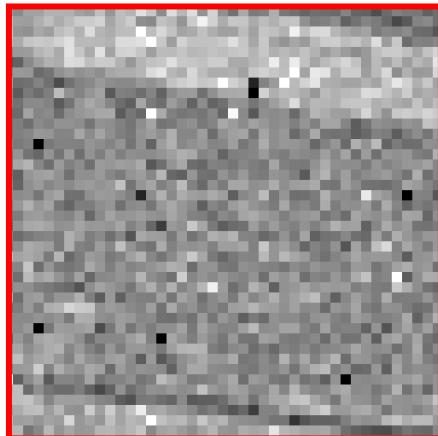
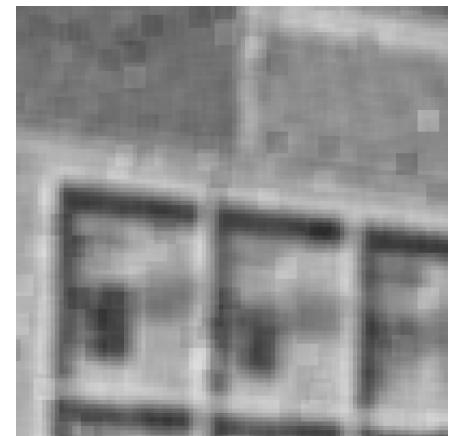
Input



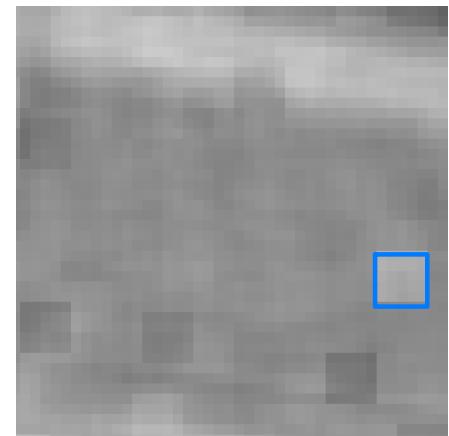
Filter

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

Output



we don't like this  
boxy thing  
We want to remove  
them



# Solution – Weighted Combination

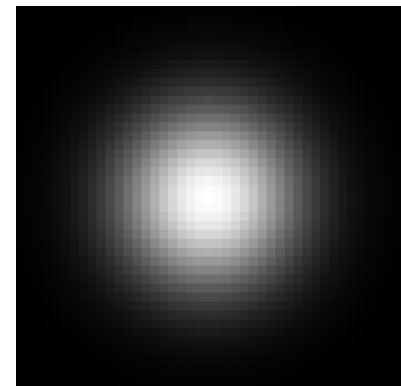
Intuition: weight according to closeness to center. Define 0,0 to be center pixel, then:

$$Filter_{ij} \propto 1$$

What's this?



$$Filter_{ij} \propto \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$



# Recognize the Filter?

It's a Gaussian!

$$Filter_{ij} \propto \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

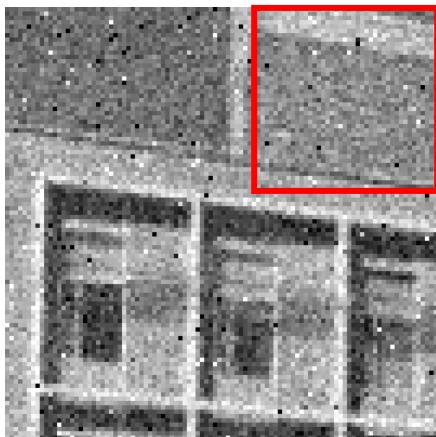
$$\begin{bmatrix} 0.003 & 0.013 & 0.022 & 0.013 & 0.003 \\ 0.013 & 0.060 & 0.098 & 0.060 & 0.013 \\ 0.022 & 0.098 & 0.162 & 0.098 & 0.022 \\ 0.013 & 0.060 & 0.098 & 0.060 & 0.013 \\ 0.003 & 0.013 & 0.022 & 0.013 & 0.003 \end{bmatrix}$$



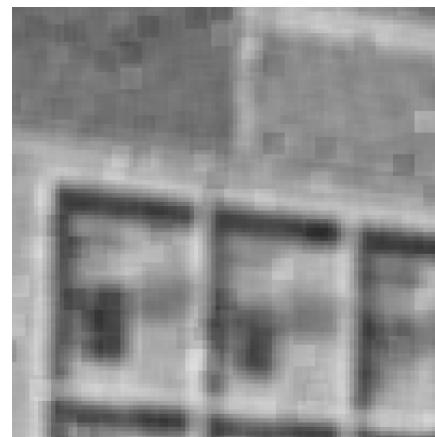
# Smoothing With A Box & Gauss

Still have some speckles, but it's not a big box

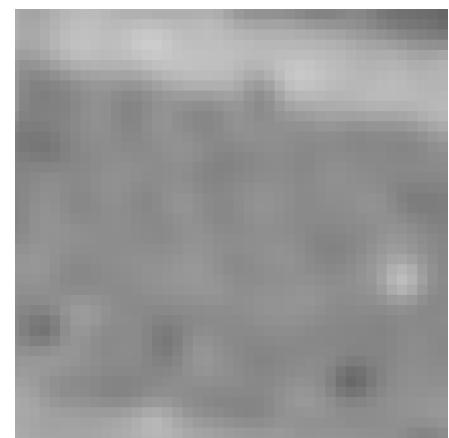
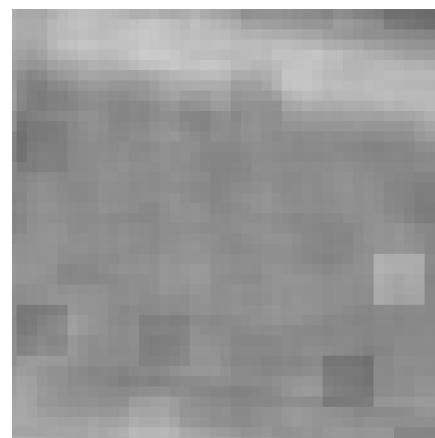
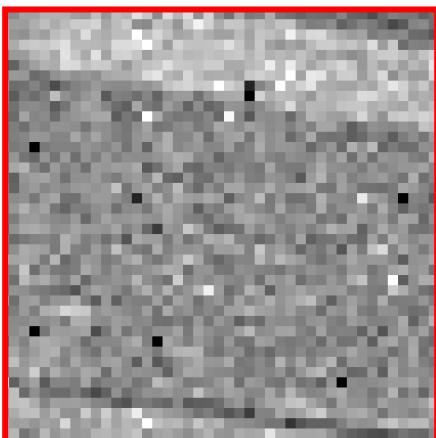
Input



Box Filter



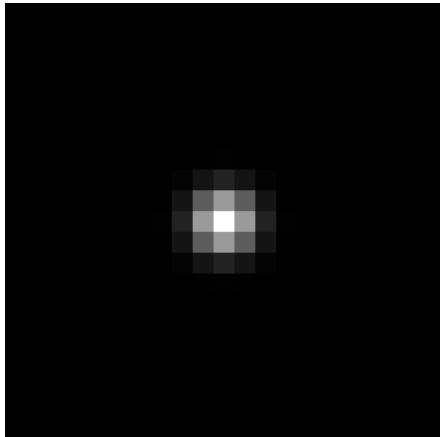
Gauss. Filter



# Gaussian Filters

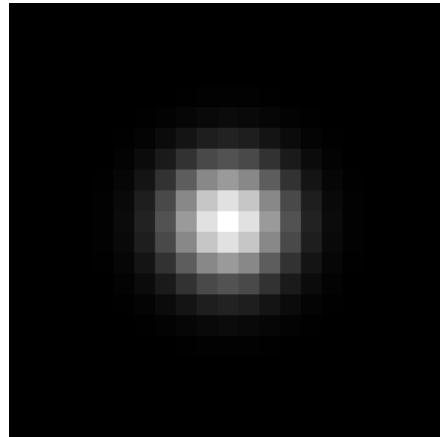
$\sigma = 1$

filter = 21x21



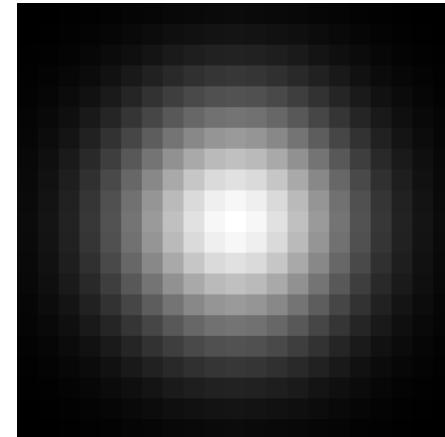
$\sigma = 2$

filter = 21x21



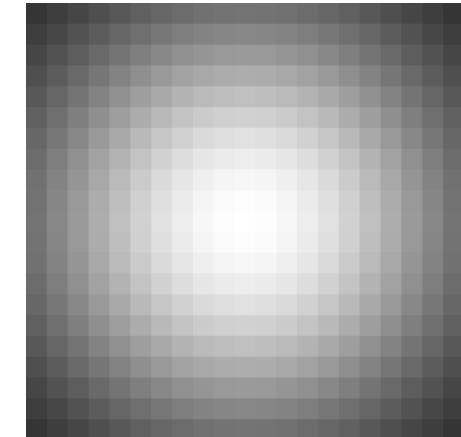
$\sigma = 4$

filter = 21x21



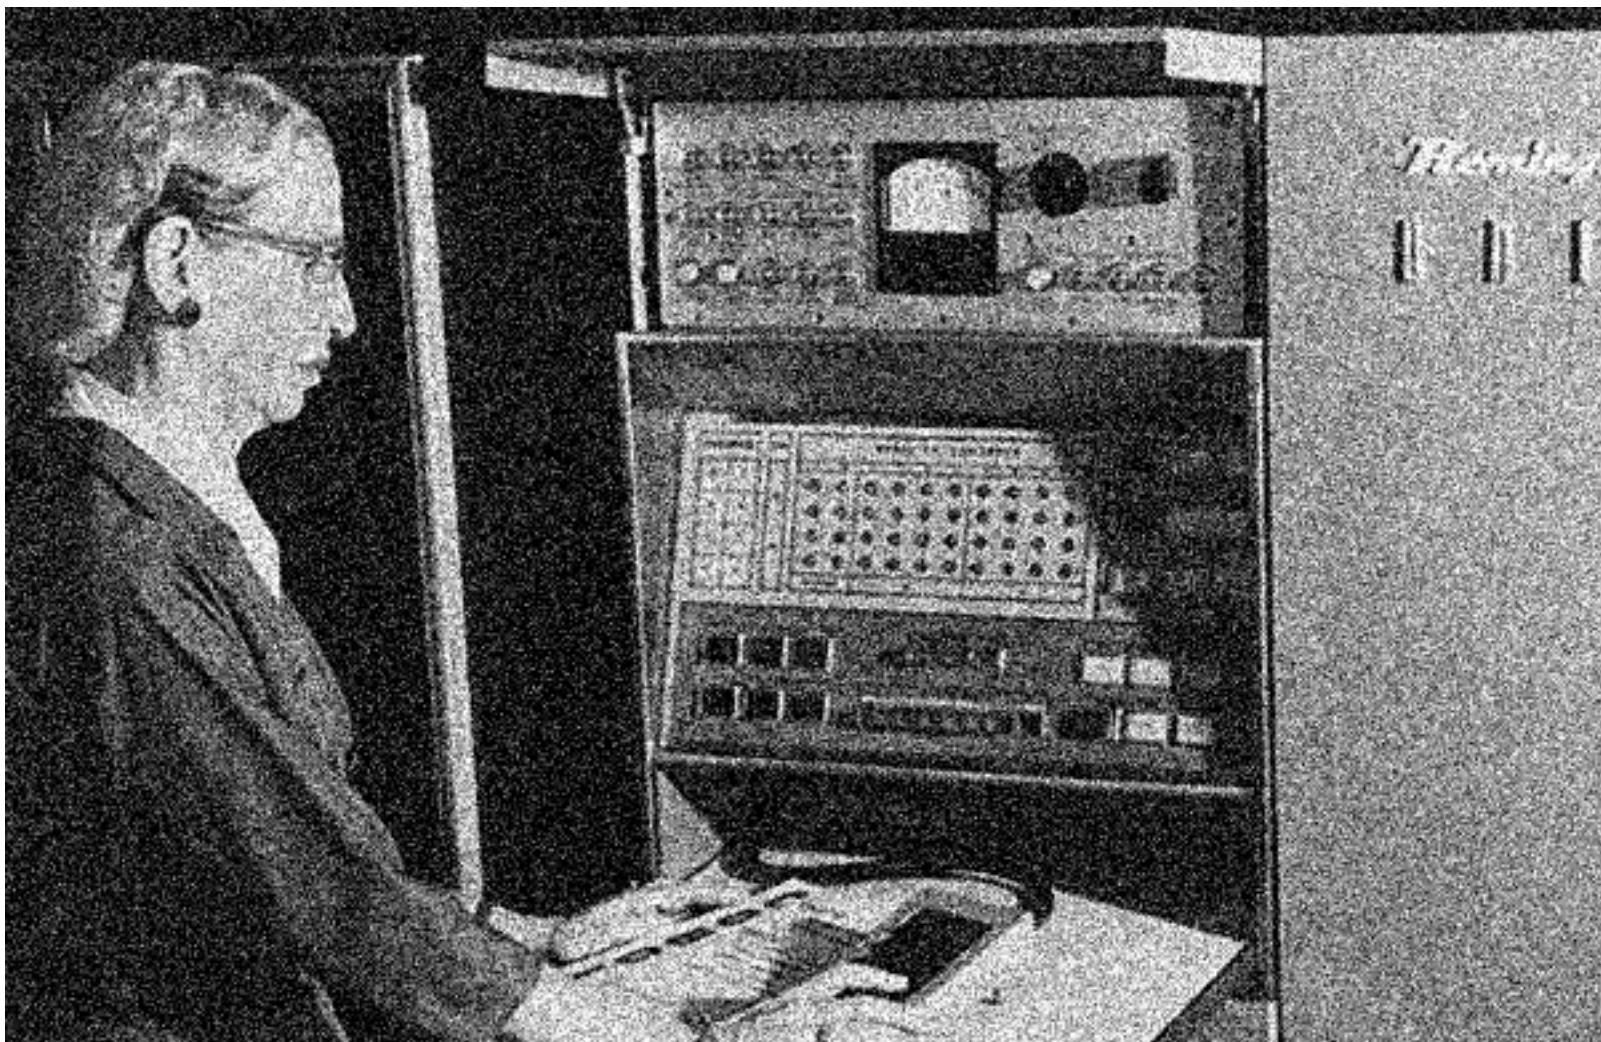
$\sigma = 8$

filter = 21x21



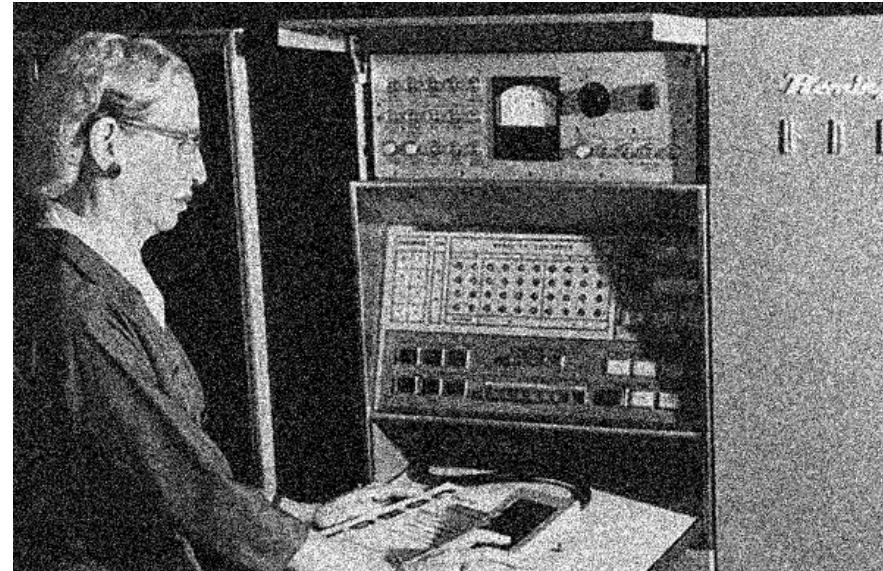
Note: filter visualizations are independently normalized throughout the slides so you can see them better

# Applying Gaussian Filters



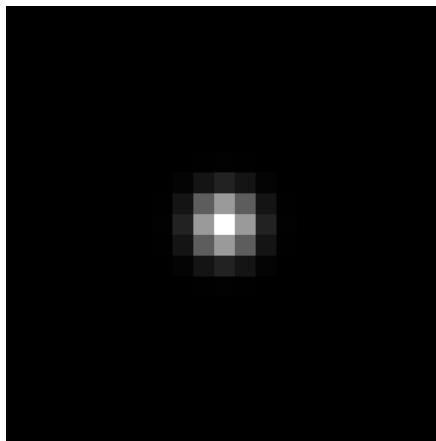
# Applying Gaussian Filters

Input Image  
(no filter)



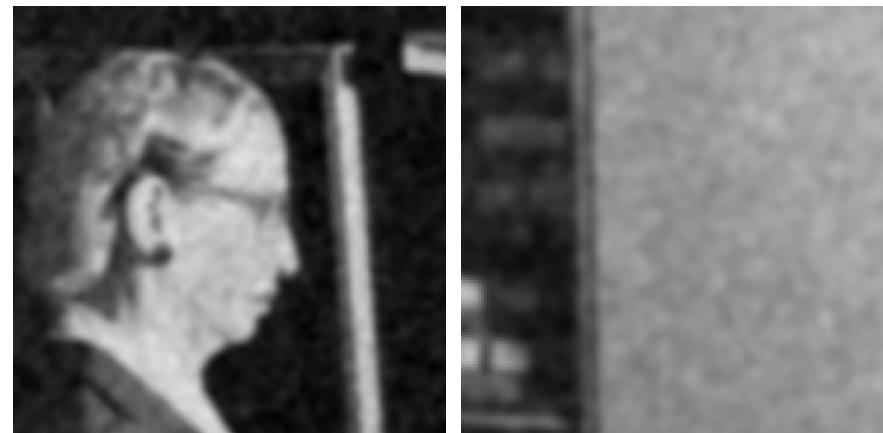
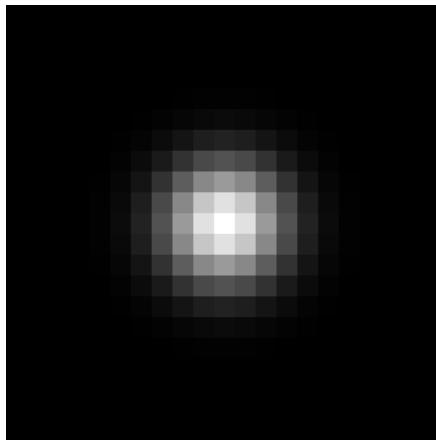
# Applying Gaussian Filters

$$\sigma = 1$$



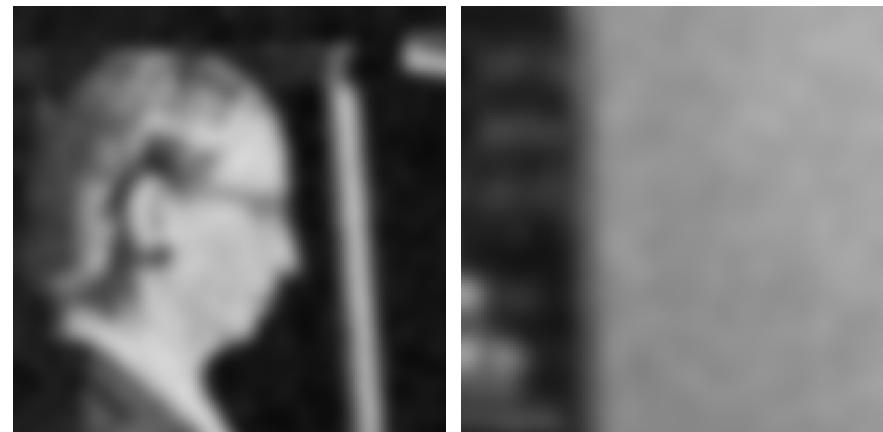
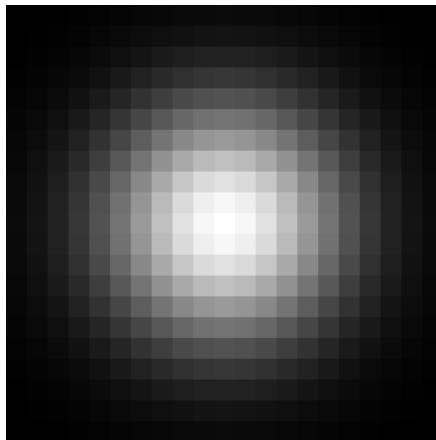
# Applying Gaussian Filters

$\sigma = 2$



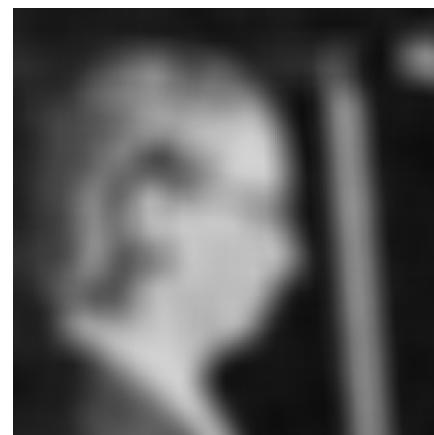
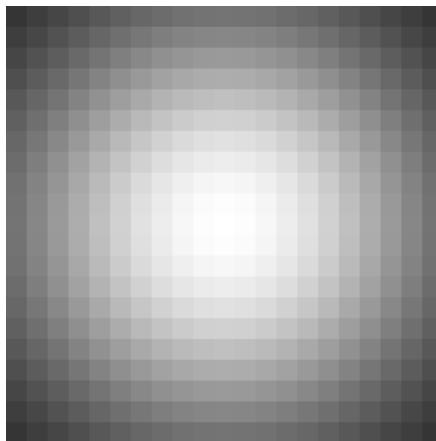
# Applying Gaussian Filters

$$\sigma = 4$$



# Applying Gaussian Filters

$\sigma = 8$



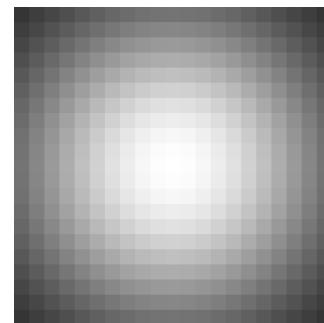
# Picking a Filter Size

Too small filter → bad approximation

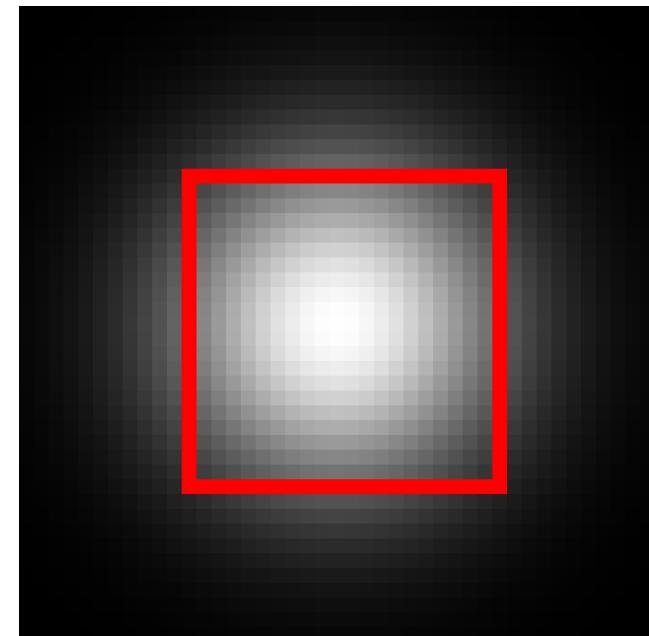
Want size  $\approx 6\sigma$  (99.7% of energy)

Left far too small; right slightly too small!

$\sigma = 8$ , size = 21



$\sigma = 8$ , size = 43



# Runtime Complexity

Image size =  $N \times N = 6 \times 6$

Filter size =  $M \times M = 3 \times 3$

I11	I12	I13	I14	I15	I16
I21	F11	F12	F13	I25	I26
I31	F21	F22	F23	I35	I36
I41	F31	F32	F33	I45	I46
I51	I52	I53	I54	I55	I56
I61	I62	I63	I64	I65	I66

```
for ImageY in range(N):  
    for ImageX in range(N):  
        for FilterY in range(M):  
            for FilterX in range(M):  
                ...
```

Time:  $O(N^2M^2)$

# Separability

$\text{Conv}(\text{vector}, \text{transposed vector}) \rightarrow \text{outer product}$

$$\begin{array}{c|c|c|c} Fy1 & & & \\ \hline Fy2 & * & Fx1 & Fx2 & Fx3 \\ \hline Fy3 & & & & \end{array} = \begin{array}{c|c|c|c} Fx1 * Fy1 & Fx2 * Fy1 & Fx3 * Fy1 \\ \hline Fx1 * Fy2 & Fx2 * Fy2 & Fx3 * Fy2 \\ \hline Fx1 * Fy3 & Fx2 * Fy3 & Fx3 * Fy3 \end{array}$$

# Separability

$$Filter_{ij} \propto \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$

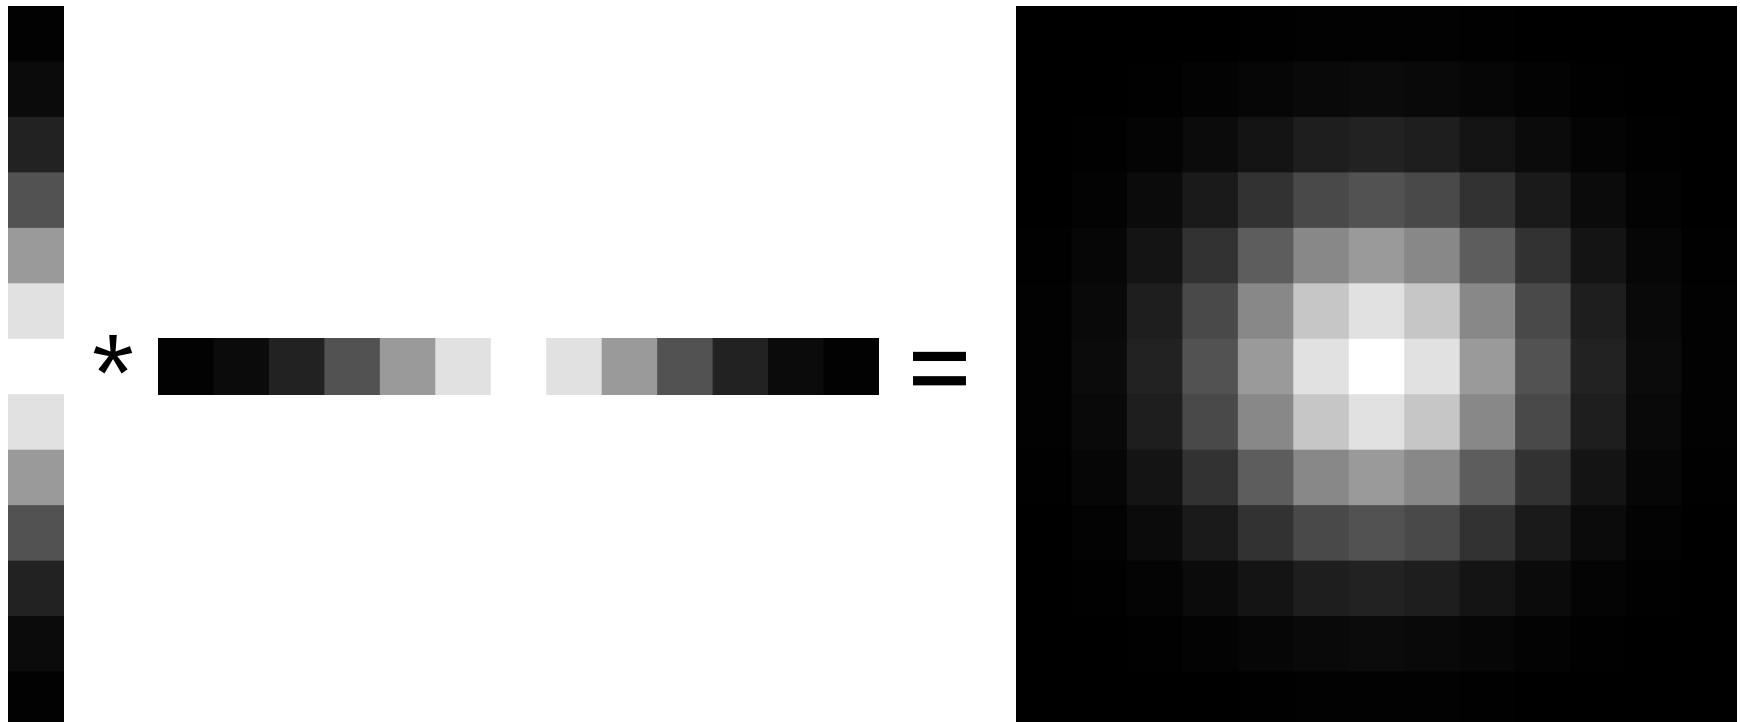


$$Filter_{ij} \propto \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{x^2}{2\sigma^2}\right) \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{y^2}{2\sigma^2}\right)$$

# Separability

$1D \text{ Gaussian} * 1D \text{ Gaussian} = 2D \text{ Gaussian}$

$\text{Image} * 2D \text{ Gauss} = \text{Image} * (1D \text{ Gauss} * 1D \text{ Gauss})$   
 $= (\text{Image} * 1D \text{ Gauss}) * 1D \text{ Gauss}$



# Runtime Complexity

Image size =  $N \times N = 6 \times 6$

Filter size =  $M \times 1 = 3 \times 1$

I11	I12	I13	I14	I15	I16
I21	F1	I23	I24	I25	I26
I31	F2	I33	I34	I35	I36
I41	F3	I43	I44	I45	I46
I51	I52	I53	I54	I55	I56
I61	I62	I63	I64	I65	I66

What are my compute savings for a  $13 \times 13$  filter?

```
for ImageY in range(N):  
    for ImageX in range(N):  
        for FilterY in range(M):  
            ...  
        for ImageY in range(N):  
            for ImageX in range(N):  
                for FilterX in range(M):  
                    ...
```

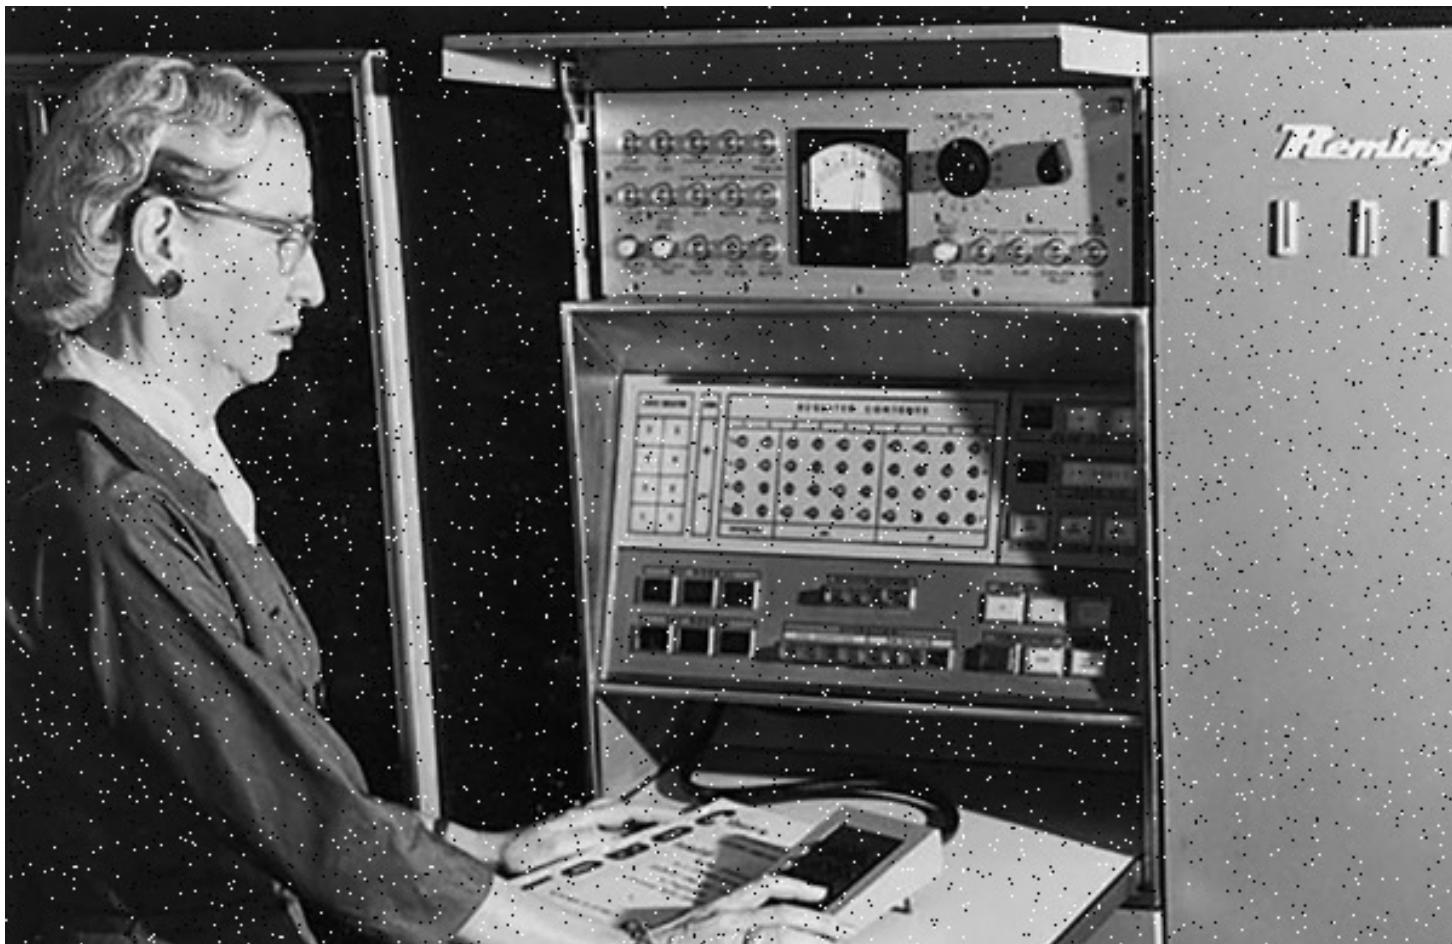
Time:  $O(N^2M)$

# Why Gaussian?

Gaussian filtering removes parts of the signal above a certain frequency. Often noise is high frequency and signal is low frequency.

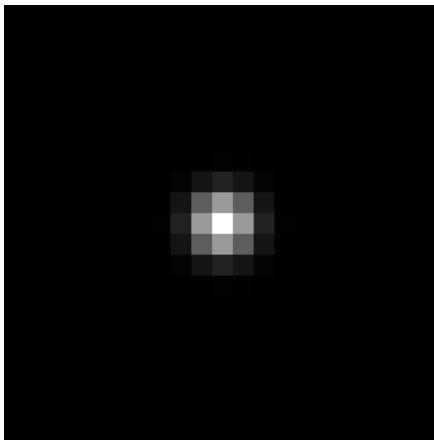


# Where Gaussian Fails



# Applying Gaussian Filters

$$\sigma = 1$$



# Why Does This Fail?

Means can be arbitrarily distorted by outliers

Signal

10	12	9	8	1000	11	10	12
----	----	---	---	------	----	----	----

Filter

0.1	0.8	0.1
-----	-----	-----

Output

11.5	9.2	107.3	801.9	109.8	10.3
------	-----	-------	-------	-------	------

What else is an “average” other than a mean?

# Non-linear Filters (2D)

40	81	13	22
125	830	76	80
144	92	108	95
132	102	106	87

[040, 081, 013, 125, 830, 076, 144, 092, 108]

↓ Sort ↓

[013, 040, 076, 081, 092, 108, 125, 144, 830]

↓  
92

[830, 076, 080, 092, 108, 095, 102, 106, 087]

↓ Sort ↓

[076, 080, 087, 092, 095, 102, 106, 108, 830]

↓  
95

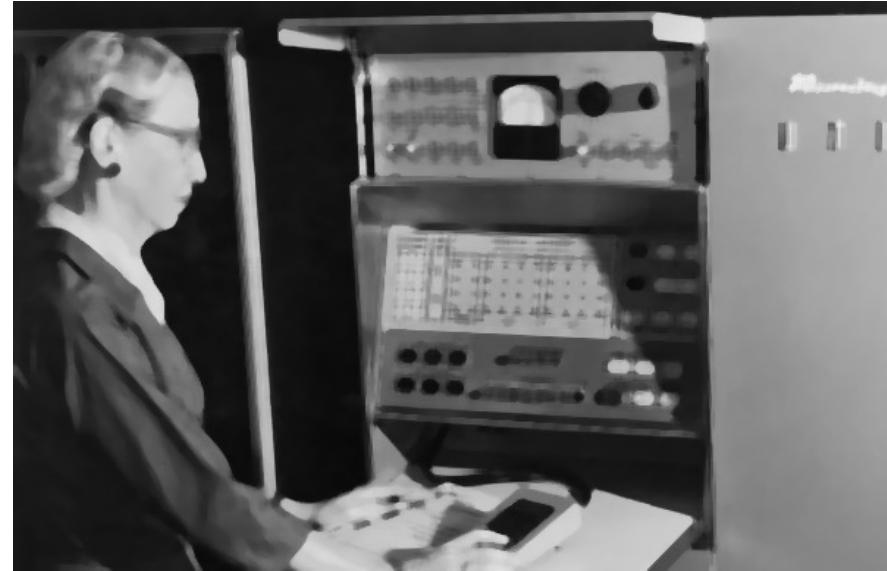
# Applying Median Filter

Median  
Filter  
(size=3)



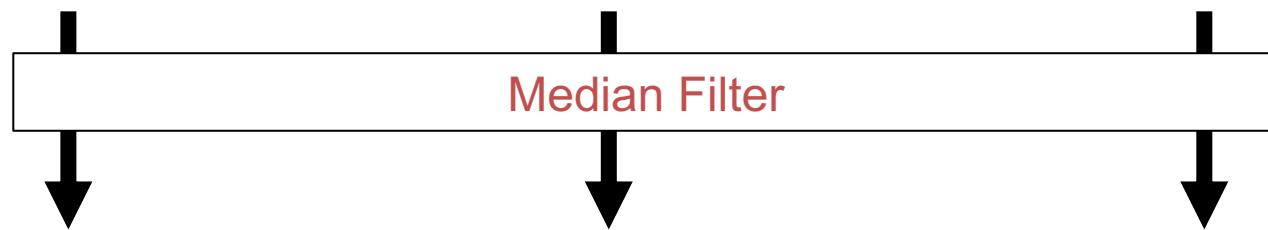
# Applying Median Filter

Median  
Filter  
(size = 7)



# Is Median Filtering Linear?

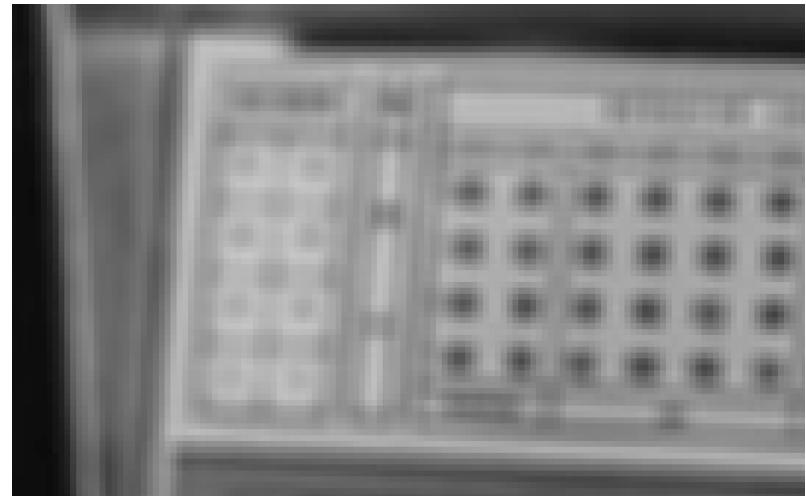
$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 2 \\ 2 & 2 & 2 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 2 & 2 \\ 2 & 2 & 2 \end{bmatrix}$$



$$1 + 0 = 2$$

# Bilateral Filtering

- Gaussian Filtering
- Pro: Removes noise
- Con: Removes details (blurry output)
- Alternative:
- Bilateral Filtering!



# Bilateral Filtering

- Gaussian Filtering

- $I^G(x) = \frac{1}{W} \sum_{x_i \in \Omega} I(x_i) g_\sigma(||x_i - x||)$

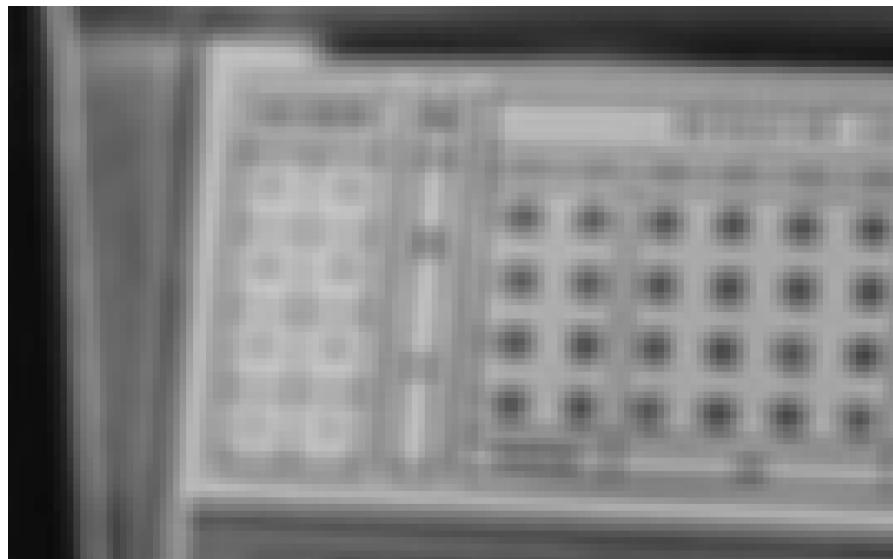
- Bilateral Filtering:

- Weight the kernel based on content similarity!
- Preserves Edges!

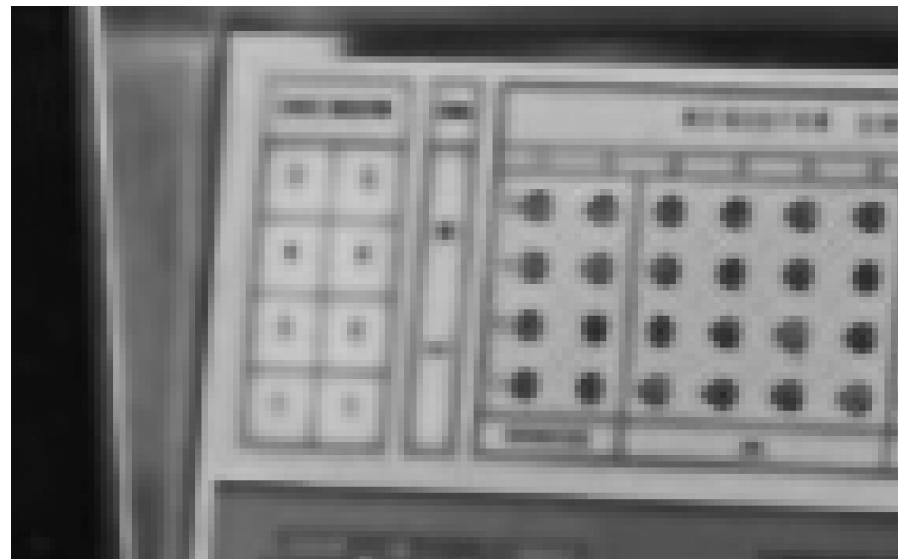
- $I^B(x) = \frac{1}{W} \sum_{x_i \in \Omega} I(x_i) g_\sigma(||x_i - x||) f_\eta(||I(x_i) - I(x)||)$

# Bilateral Filtering

- Bilateral Filtering Preserves Edges!
- Better edges and texture than Gaussian or Median
- Computationally Heavy



Gaussian Filtering



Bilateral Filtering

# Some Examples of Filtering Effects

# Filtering – Sharpening

## Image                                  Smoothed

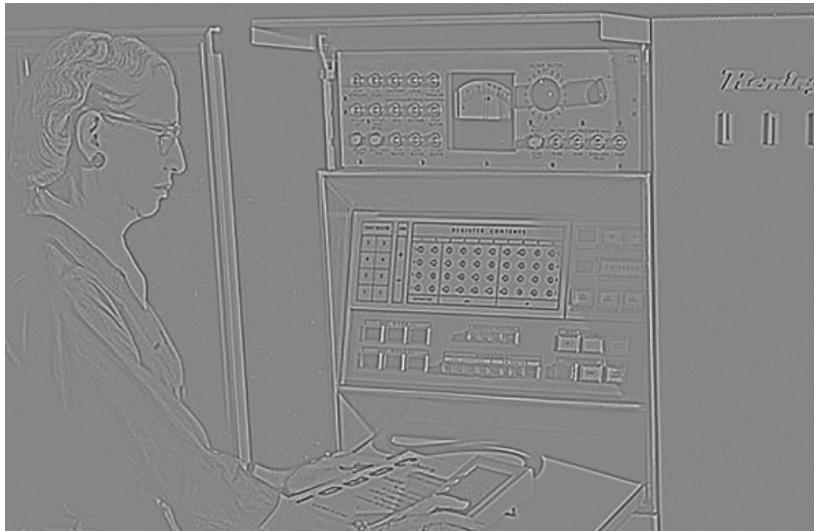


=



## Details

=

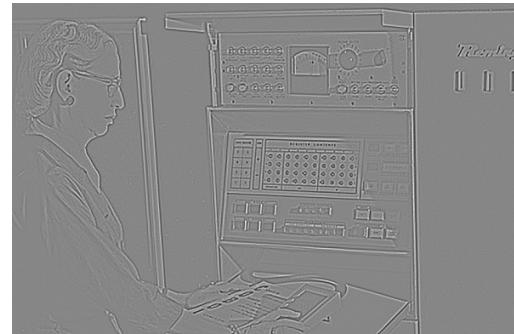


# Filtering – Sharpening

## Image Details



$+ \alpha$



“Sharpened”  $\alpha=1$

=

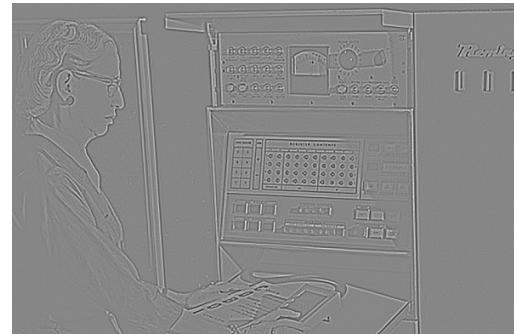


# Filtering – Sharpening

## Image Details



$+ \alpha$



“Sharpened”  $\alpha=0$

=

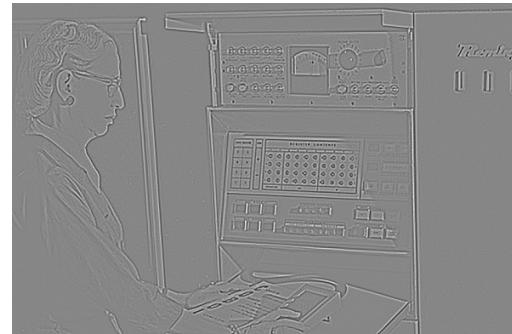


# Filtering – Sharpening

## Image Details



$+ \alpha$



“Sharpened”  $\alpha=2$

=

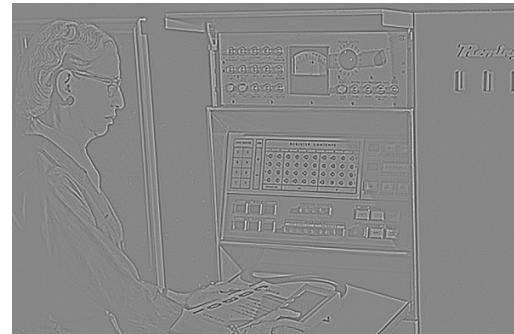


# Filtering – Sharpening

## Image Details



$+ \alpha$



“Sharpened”  $\alpha=0$

=

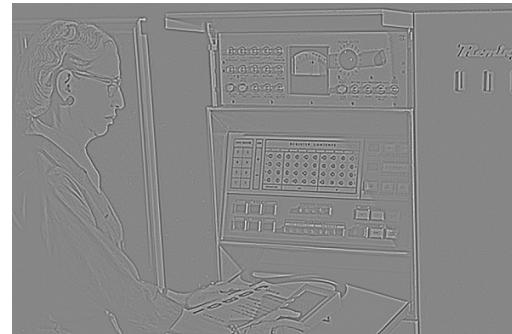


# Filtering – Extreme Sharpening

## Image Details



$+ \alpha$



“Sharpened”  $\alpha=10$

=



# Filtering

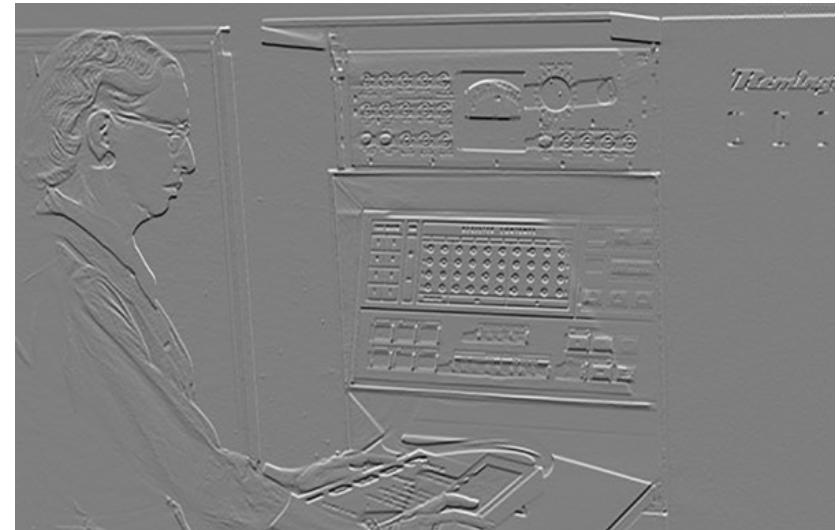
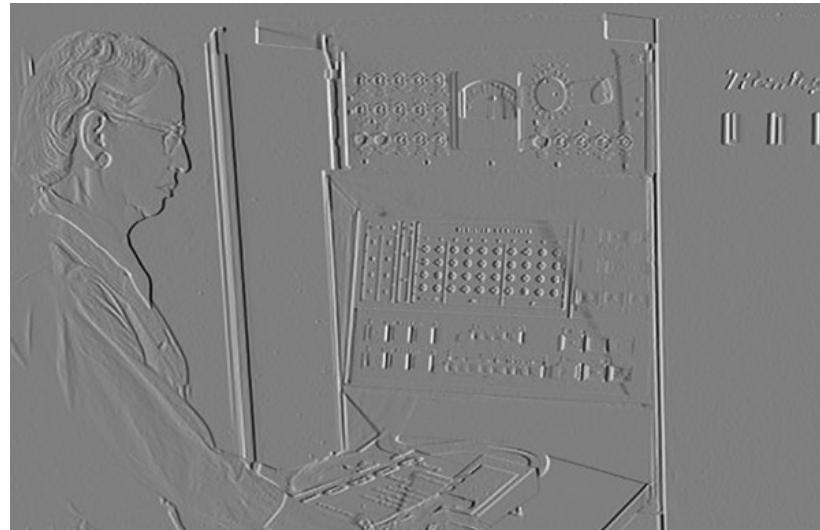
**What's this Filter?**

-1	0	1
----	---	---

**Dx**

-1	0	1
----	---	---

**Dy**



# Filtering – Derivatives

$$(Dx^2 + Dy^2)^{1/2}$$



# Filtering – Bonus

- If you're curious, you can use filters to accomplish a surprisingly large number of things.

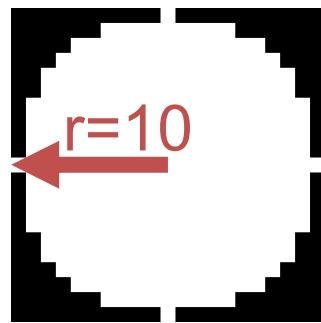
# Filtering – Counting

How many “on” pixels have  
10+ neighbors within 10 pixels?

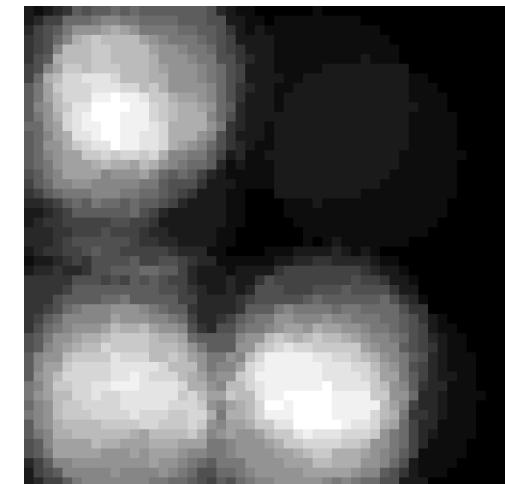
Pixels



Disk



???



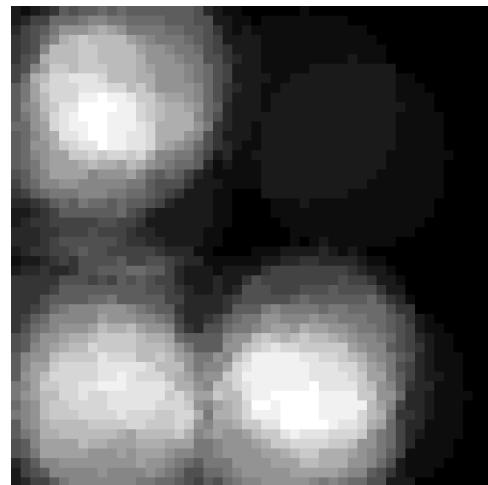
# Filtering – Counting

How many “on” pixels have  
10+ neighbors within 10 pixels?

Pixels



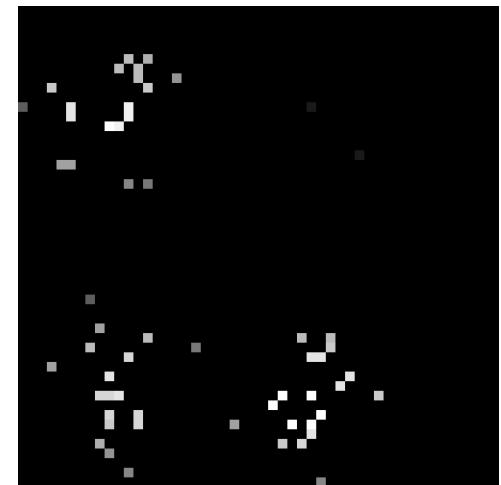
Density



X

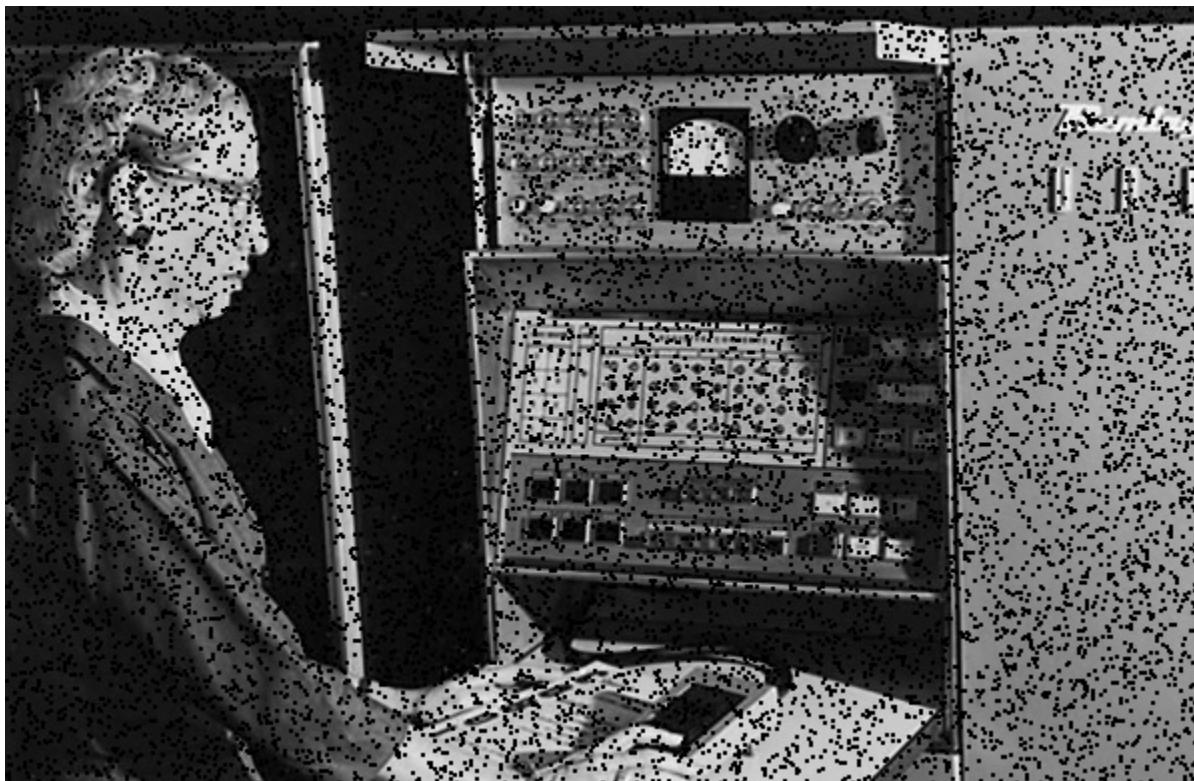
=

Answer



# Filtering – Missing Data

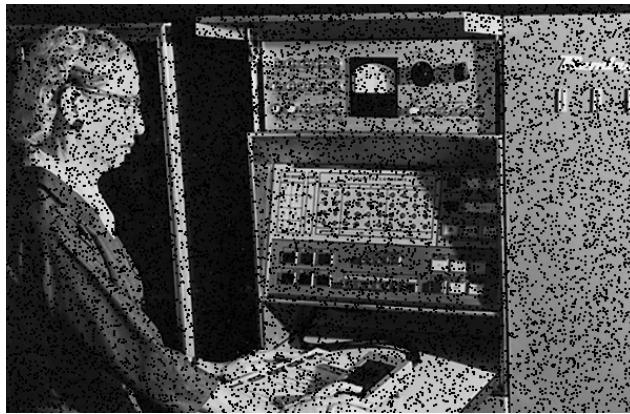
Oh no! Missing data!  
(and we know where)



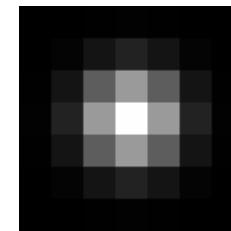
Common with many non-normal cameras (e.g., depth cameras)

# Filtering – Missing Data

Image

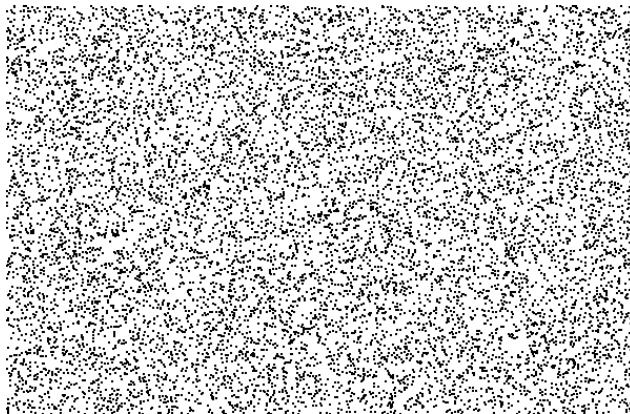


\*

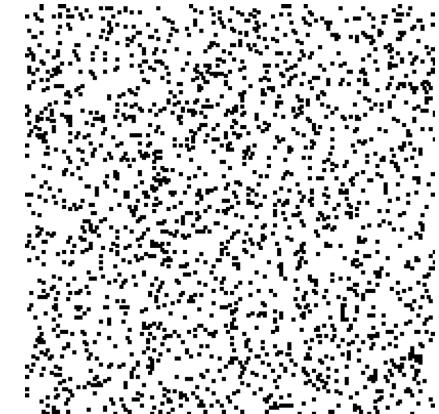
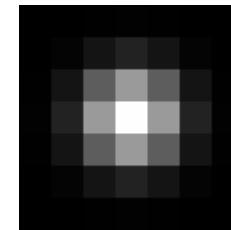


Per-element /

Binary  
Mask

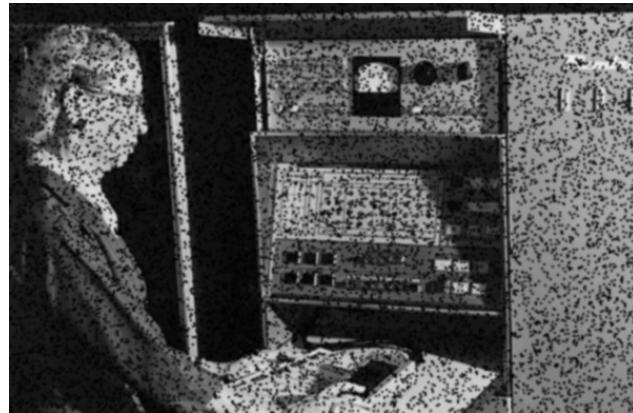


\*



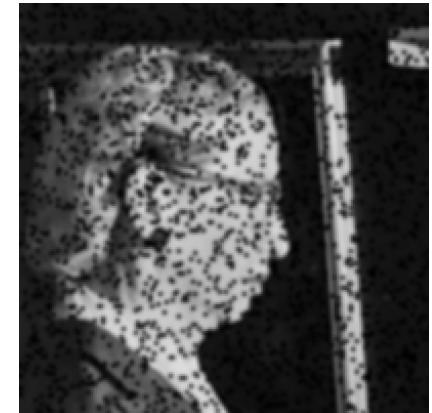
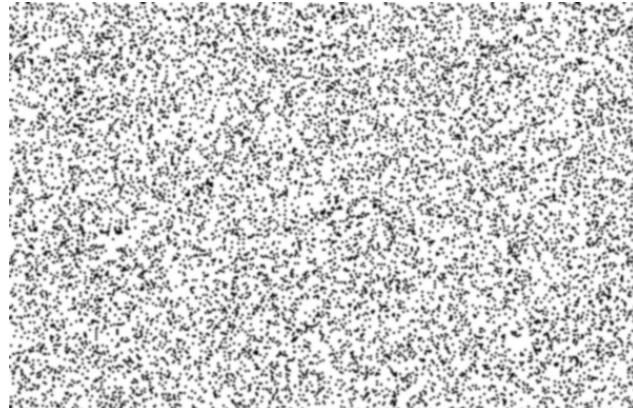
# Filtering – Missing Data

Image



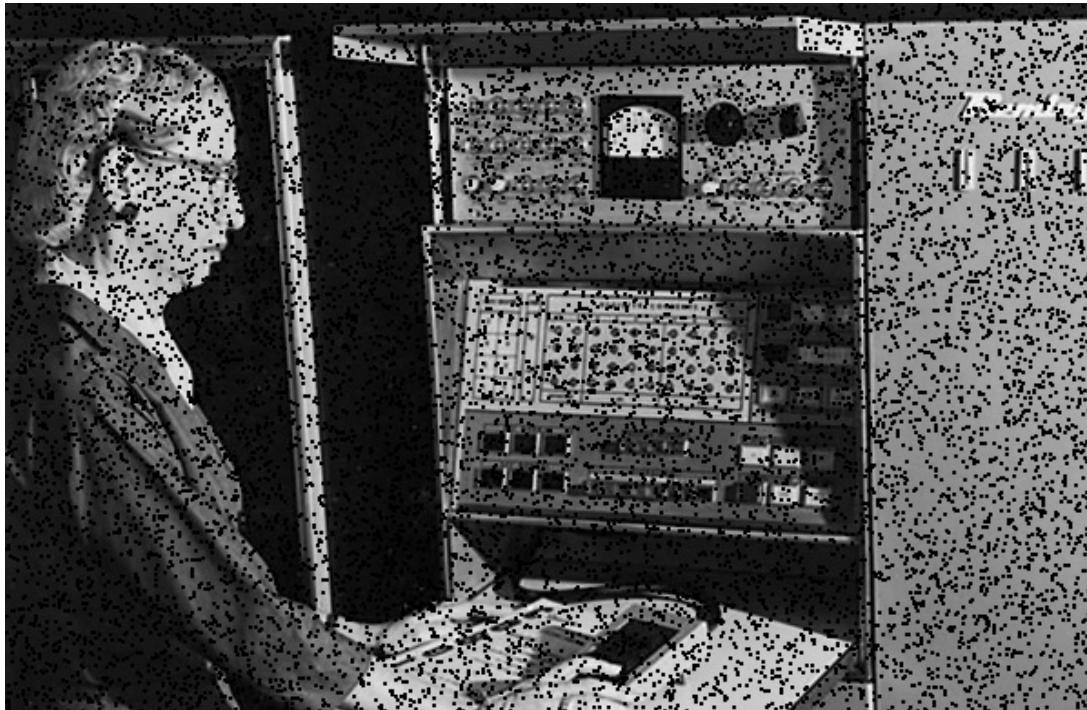
Per-element /

Binary  
Mask



# Filtering – Missing Data

Before



# Filtering – Missing Data

After



# Filtering – Missing Data

After (without missing data)

