

Python Code

Your Name

December 20, 2024

1 Task 1: Implementing Computational Graphs

Q1 :Implement the forward and backward passes for the computational graph f3 below.

A1 :

```
1  def f3(s1, s2, y):
2      """
3      Computes the forward and backward pass through the computational graph f3
4      from the homework PDF.
5
6      A few clarifications about the graph:
7      - The input y is an integer with y == 1 or y == 2; you do not need to
8        compute a gradient for this input.
9      - The division nodes compute p1 = e1 / d and p2 = e2 / d
10     - The choose(p1, p2, y) node returns p1 if y is 1, or p2 if y is 2.
11
12     Inputs:
13     - s1, s2: Python floats
14     - y: Python integer, either equal to 1 or 2
15
16     Returns a tuple of:
17     - L: Python scalar giving the output of the graph
18     - grads: A tuple (grad_s1, grad_s2) giving the derivative of the output L
19       with respect to the inputs s1 and s2.
20     """
21     assert y == 1 or y == 2
22     # Forward pass: Compute loss
23     L = None
24     #####
25     # TODO: Implement the forward pass for the computational graph f3 shown #
26     # in the homework description. Store the loss in the variable L.      #
27     #####
28     e1 = math.exp(s1)
29     e2 = math.exp(s2)
30     e11 = e1
31     e12 = e1
32     e21 = e2
33     e22 = e2
34     d = e11 + e21
35     d1 = d
36     d2 = d
37     p1 = e12 / d1
38     p2 = e22 / d2
39     p_plus = p1*(2-y) - p2*(1-y)
40     L = -1 * math.log(p_plus)
41
42     #####
43     #                                     END OF YOUR CODE
44     #####
45
46     # Backward pass: Compute gradients
47     grad_s1, grad_s2 = None, None
48     #####
49     # TODO: Implement the backward pass for the computational graph f3 shown #
50     # in the homework description. Store the gradients for each input      #
51     # variable in the corresponding grad variagbles defined above. You do not #
52     # need to compute a gradient for the input y since it is an integer.     #
53     #
54     # HINT: You may need an if statement to backprop through the choose node #
```

```

55 #####
56 grad_L = 1
57 grad_p_plus = -1*(grad_L / p_plus)
58 grad_p1 = grad_p_plus * (2-y)
59 grad_p2 = grad_p_plus * (y-1)
60 grad_e12 = grad_p1 / d1
61 grad_e22 = grad_p2 / d2
62 grad_d1 = grad_p1 * e12 * -1 * (1 / d1**2)
63 grad_d2 = grad_p2 * e22 * -1 * (1 / d2**2)
64 grad_d = grad_d1 + grad_d2
65 grad_e11 = grad_d
66 grad_e21 = grad_d
67 grad_e1 = grad_e11 + grad_e12
68 grad_e2 = grad_e21 + grad_e22
69 grad_s1 = math.exp(s1) * grad_e1
70 grad_s2 = math.exp(s2) * grad_e2
71 grad_y = grad_p_plus * (-p1 + p2)
72 #####
73 #                               END OF YOUR CODE                               #
74 #####
75
76 grads = (grad_s1, grad_s2)
77 return L, grads

```

2 Task 2: Modular Backprop API

Q1 :Fully-connected layer: `fc_forward` and `fc_backward`.

A1 :

```

1  def fc_forward(x, w, b):
2  """
3  Computes the forward pass for a fully-connected layer.
4
5  The input x has shape (N, Din) and contains a minibatch of N
6  examples, where each example x[i] has shape (Din,).
7
8  Inputs:
9  - x: A numpy array of shape (N, Din) giving input data
10 - w: A numpy array of shape (Din, Dout) giving weights
11 - b: A numpy array of shape (Dout,) giving biases
12
13 Returns a tuple of:
14 - out: output, of shape (N, Dout)
15 - cache: (x, w, b)
16 """
17 out = None
18 #####
19 # TODO: Implement the forward pass. Store the result in out. #
20 #####
21 out = x @ w + b
22 #####
23 #                               END OF YOUR CODE                               #
24 #####
25 cache = (x, w, b)
26 return out, cache
27

```

```

1  def fc_backward(grad_out, cache):
2  """
3  Computes the backward pass for a fully-connected layer.
4
5  Inputs:
6  - grad_out: Numpy array of shape (N, Dout) giving upstream gradients
7  - cache: Tuple of:
8    - x: A numpy array of shape (N, Din) giving input data
9    - w: A numpy array of shape (Din, Dout) giving weights
10   - b: A numpy array of shape (Dout,) giving biases
11
12 Returns a tuple of downstream gradients:
13 - grad_x: A numpy array of shape (N, Din) of gradient with respect to x
14 - grad_w: A numpy array of shape (Din, Dout) of gradient with respect to w

```

```

15 - grad_b: A numpy array of shape (Dout,) of gradient with respect to b
16 """
17 x, w, b = cache
18 grad_x, grad_w, grad_b = None, None, None
19 #####
20 # TODO: Implement the backward pass for the fully-connected layer #
21 #####
22 grad_x = grad_out @ w.T
23 grad_w = x.T @ grad_out
24 grad_b = np.sum(grad_out,axis=0)
25
26 #####
27 #                               END OF YOUR CODE                               #
28 #####
29 return grad_x, grad_w, grad_b
30

```

Q2 :relu_forward and relu_backward

A2 :

```

1 def relu_forward(x):
2     """
3     Computes the forward pass for the Rectified Linear Unit (ReLU) nonlinearity
4
5     Input:
6     - x: A numpy array of inputs, of any shape
7
8     Returns a tuple of:
9     - out: A numpy array of outputs, of the same shape as x
10    - cache: x
11    """
12    out = None
13    #####
14    # TODO: Implement the ReLU forward pass. #
15    #####
16    out = np.maximum(0,x)
17
18    #####
19    #                               END OF YOUR CODE                               #
20    #####
21    cache = x
22    return out, cache
23

```

```

1 def relu_backward(grad_out, cache):
2     """
3     Computes the backward pass for a Rectified Linear Unit (ReLU) nonlinearity
4
5     Input:
6     - grad_out: Upstream derivatives, of any shape
7     - cache: Input x, of same shape as dout
8
9     Returns:
10    - grad_x: Gradient with respect to x
11    """
12    grad_x, x = None, cache
13    #####
14    # TODO: Implement the ReLU backward pass. #
15    #####
16    grad_x = grad_out * (x > 0)
17    #####
18    #                               END OF YOUR CODE                               #
19    #####
20    return grad_x
21

```

Q3 :Softmax Loss Function:softmax_loss

A3 :

```

1     def softmax_loss(x, y):
2         """
3         Computes the loss and gradient for softmax (cross-entropy) loss function.
4
5         Inputs:
6         - x: Numpy array of shape (N, C) giving predicted class scores, where
7           x[i, c] gives the predicted score for class c on input sample i
8         - y: Numpy array of shape (N,) giving ground-truth labels, where
9           y[i] = c means that input sample i has ground truth label c, where
10            0 <= c < C.
11
12         Returns a tuple of:
13         - loss: Scalar giving the loss
14         - grad_x: Numpy array of shape (N, C) giving the gradient of the loss with
15           with respect to x
16         """
17         loss, grad_x = None, None
18         #####
19         # TODO: Implement softmax loss
20         #####
21         M = np.max(x,axis = 1,keepdims=True)
22         x_minus_M = x - M
23         exp_x = np.exp(x_minus_M)
24         exp_x_each_row_sum = np.sum(exp_x,axis=1, keepdims=True)
25         p_x = exp_x / exp_x_each_row_sum
26
27         n = y.shape[0]
28
29         loss = 0
30         for i in range(n):
31             true_label = y[i]
32             p_true_label = p_x[i][true_label]
33             log_p_true_label = np.log(p_true_label)
34             loss += log_p_true_label
35         loss = loss * -1 / n
36
37         grad_x = p_x.copy() # Make a copy of p_x
38         grad_x[np.arange(n), y] -= 1
39         grad_x /= n
40         #####
41         #                                     END OF YOUR CODE
42         #####
43         return loss, grad_x
44
45

```

Q4 :L2 Regularization: l2_regularization which implements the L2 regularization loss

A4 :

```

1     def l2_regularization(w, reg):
2         """
3         Computes loss and gradient for L2 regularization of a weight matrix:
4
5         loss = (reg / 2) * sum_i w_i^2
6
7         Where the sum ranges over all elements of w.
8
9         Inputs:
10        - w: Numpy array of any shape
11        - reg: float giving the regularization strength
12
13        Returns:
14        """
15        loss, grad_w = None, None
16        #####
17        # TODO: Implement L2 regularization.
18        #####
19        sum_Wi_squire = np.sum(w**2)
20        loss = (reg / 2) * sum_Wi_squire
21        grad_w = reg * w
22        #####
23        #                                     END OF YOUR CODE
24        #####

```

```

25     return loss, grad_w
26
27

```

3 Task 3: Implement a Two-Layer Network

Q1 :Complete the implementation of the TwoLayerNet class. Your implementations for the forward and backward methods should use the modular forward and backward functions that you implemented in the previous task.

A1 :

```

1     class TwoLayerNet(Classifier):
2         """
3         A neural network with two layers, using a ReLU nonlinearity on its one
4         hidden layer. That is, the architecture should be:
5
6         input -> FC layer -> ReLU layer -> FC layer -> scores
7         """
8         def __init__(self, input_dim=3072, num_classes=10, hidden_dim=512,
9                       weight_scale=1e-3):
10             """
11             Initialize a new two layer network.
12
13             Inputs:
14             - input_dim: The number of dimensions in the input.
15             - num_classes: The number of classes over which to classify
16             - hidden_dim: The size of the hidden layer
17             - weight_scale: The weight matrices of the model will be initialized
18                           from a Gaussian distribution with standard deviation equal to
19                           weight_scale. The bias vectors of the model will always be
20                           initialized to zero.
21             """
22             #####
23             # TODO: Initialize the weights and biases of a two-layer network.      #
24             #####
25             self.W1 = weight_scale * np.random.randn(input_dim, hidden_dim)
26             self.b1 = np.zeros(hidden_dim)
27             self.W2 = weight_scale * np.random.randn(hidden_dim, num_classes)
28             self.b2 = np.zeros(num_classes)
29             #####
30             #                               END OF YOUR CODE                               #
31             #####
32
33         def parameters(self):
34             params = None
35             #####
36             # TODO: Build a dict of all learnable parameters of this model.      #
37             #####
38             params = {
39                 'W1': self.W1,
40                 'b1': self.b1,
41                 'W2': self.W2,
42                 'b2': self.b2,
43             }
44
45             #####
46             #                               END OF YOUR CODE                               #
47             #####
48             return params
49
50         def forward(self, X):
51             scores, cache = None, None
52             #####
53             # TODO: Implement the forward pass to compute classification scores    #
54             # for the input data X. Store into cache any data that will be needed  #
55             # during the backward pass.                                           #
56             #####
57             W1 = self.W1
58             W2 = self.W2
59             b1 = self.b1
60             b2 = self.b2

```

```

62     #first fc layer
63     fc1_out, fc1_cache = fc_forward(X,W1,b1)
64     #apply the Relu layer
65     relu_out, relu_cache = relu_forward(fc1_out)
66     #second fc layer
67     fc2_out, fc2_cache = fc_forward(relu_out,W2, b2)
68     cache = (fc1_cache, relu_cache, fc2_cache)
69     scores = fc2_out
70
71
72     #####
73     #                               END OF YOUR CODE                               #
74     #####
75     return scores, cache
76
77 def backward(self, grad_scores, cache):
78     grads = None
79     #####
80     # TODO0: Implement the backward pass to compute gradients for all      #
81     # learnable parameters of the model, storing them in the grads dict    #
82     # above. The grads dict should give gradients for all parameters in    #
83     # the dict returned by model.parameters().                               #
84     #####
85     W1 = self.W1
86     W2 = self.W2
87     b1 = self.b1
88     b2 = self.b2
89
90     fc1_cache, relu_cache, fc2_cache = cache
91     grad_relu, grads_W2, grads_b2 = fc_backward(grad_scores, fc2_cache)
92
93     grad_fc1 = relu_backward(grad_relu,relu_cache)
94
95     trash, grads_W1, grads_b1 = fc_backward(grad_fc1,fc1_cache)
96     grads = {
97         'W1' : grads_W1,
98         'b1' : grads_b1,
99         'W2' : grads_W2,
100        'b2' : grads_b2,
101    }
102
103
104     #####
105     #                               END OF YOUR CODE                               #
106     #####
107     return grads
108

```

4 Task4: Training Two-Layer Networks

Q1 : Implement the training_step function in the file neuralnettrainpy

A1 :

```

1     def training_step(model, X_batch, y_batch, reg):
2         """
3         Compute the loss and gradients for a single training iteration of a model
4         given a minibatch of data. The loss should be a sum of a cross-entropy loss
5         between the model predictions and the ground-truth image labels, and
6         an L2 regularization term on all weight matrices in the fully-connected
7         layers of the model. You should not regularize the bias vectors.
8
9         Inputs:
10        - model: A Classifier instance
11        - X_batch: A numpy array of shape (N, D) giving a minibatch of images
12        - y_batch: A numpy array of shape (N,) where 0 <= y_batch[i] < C is the
13          ground-truth label for the image X_batch[i]
14        - reg: A float giving the strength of L2 regularization to use.
15
16        Returns a tuple of:
17        - loss: A float giving the loss (data loss + regularization loss) for the
18          model on this minibatch of data

```

```

19 - grads: A dictionary giving gradients of the loss with respect to the
20     parameters of the model. In particular grads[k] should be the gradient
21     of the loss with respect to model.parameters()[k].
22 """
23 loss, grads = None, None
24 #####
25 # TODO: Compute the loss and gradient for one training iteration.      #
26 #####
27 prediction, cache = model.forward(X_batch)
28
29 #compute the data loss
30 data_loss, grad_x = softmax_loss(prediction, y_batch)
31
32 #compute the regularization loss
33 reg_loss = 0
34 for param_name, param in model.parameters().items():
35     if param_name.startswith('W'):
36         reg_loss += l2_regularization(param, reg)[0]
37 #calculate the total loss
38 loss = data_loss + reg_loss
39
40 grads = model.backward(grad_x, cache)
41
42 for param_name in grads:
43     if param_name.startswith('W'):
44         grads[param_name] += reg * model.parameters()[param_name]
45
46 #####
47 #                               END OF YOUR CODE                       #
48 #####
49 return loss, grads
50

```


Q2 :include the loss / accuracy plot for your best model, describe the hyperparameter settings you used, and give the final test-set performance of your model.

A2 :
the parameter I use here is:

```

1  # How much data to use for training
2  num_train = 20000
3
4  # Model architecture hyperparameters.
5  hidden_dim = 64
6
7  # Optimization hyperparameters.
8  batch_size = 64
9  num_epochs = 100
10 learning_rate = 0.005
11 reg = 0.01
12

```




Q3 : include the loss / accuracy plot for your overfit model and describe the hyperparameter settings you used.

A3 :
the parameter I use here is:

```

1      # How much data to use for training
2      num_train = 3000
3
4      # Model architecture hyperparameters.
5      hidden_dim = 64
6
7      # Optimization hyperparameters.
8      batch_size = 32
9      num_epochs = 500
10     learning_rate = 0.01
11     reg = 0.01
12

```



5 Task 5: Train Your Own Classification Model

Q1 :Submit the notebook (with outputs) that trains with your best combination of model architecture, optimizer and training parameters, and evaluates on the test set to report an accuracy at the end.

A1 :

EECS442_HW4_task5_report

March 19, 2024

1 EECS 442 Homework 4: Fashion-MNIST Classification

In this part, you will implement and train Convolutional Neural Networks (ConvNets) in PyTorch to classify images. Unlike HW4 Secion 1, backpropagation is automatically inferred by PyTorch, so you only need to write code for the forward pass.

Before we start, please put your name and UMID in following format Firstname
LASTNAME, #00000000 // e.g.) David FOUHEY, #12345678

Your Answer:

Hello EECS442 #12345678

1.1 Setup

```
[ ]: # Run the command in the terminal if it failed on local Jupyter Notebook,  
      ↪ remove "!" before each line  
      !pip install torchsummary
```

```
[ ]: import numpy as np  
      import matplotlib.pyplot as plt  
      from tqdm import tqdm # Displays a progress bar  
  
      import torch  
      from torch import nn  
      from torch import optim  
      import torch.nn.functional as F  
      from torchsummary import summary  
      from torchvision import datasets, transforms  
      from torch.utils.data import Dataset, Subset, DataLoader, random_split
```

```
[ ]: if torch.cuda.is_available():  
      print("Using the GPU. You are good to go!")  
      device = 'cuda'  
else:  
      print("Using the CPU. Overall speed may be slowed down")  
      device = 'cpu'
```

Using the CPU. Overall speed may be slowed down

1.2 Loading Dataset

The dataset we use is Fashion-MNIST dataset, which is available at <https://github.com/zalandoresearch/fashion-mnist> and in torchvision.datasets. Fashion-MNIST has 10 classes, 60000 training+validation images (we have splitted it to have 50000 training images and 10000 validation images, but you can change the numbers), and 10000 test images.

```
[ ]: # Load the dataset and train, val, test splits
print("Loading datasets...")
# Transform from [0,255] uint8 to [0,1] float,
# then normalize to zero mean and unit variance
FASHION_transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize([0.2859], [0.3530])
])
FASHION_trainval = datasets.FashionMNIST('.', download=True, train=True,
                                         transform=FASHION_transform)
FASHION_train = Subset(FASHION_trainval, range(50000))
FASHION_val = Subset(FASHION_trainval, range(50000, 60000))
FASHION_test = datasets.FashionMNIST('.', download=True, train=False,
                                       transform=FASHION_transform)

print("Done!")
```

Loading datasets...

Done!

Now, we will create the dataloader for train, val and test dataset. You are free to experiment with different batch sizes.

```
[ ]: # Create dataloaders
#####
# TODO: Experiment with different batch sizes
#####
batch_size = 50
#####
#                               END OF YOUR CODE
#####
trainloader = DataLoader(FASHION_train, batch_size=batch_size, shuffle=True)
valloader = DataLoader(FASHION_val, batch_size=batch_size, shuffle=True)
testloader = DataLoader(FASHION_test, batch_size=batch_size, shuffle=True)
```

1.3 Model

Initialize your model and experiment with with different optimizers, parameters (such as learning rate) and number of epochs.

```
[ ]: class Network(nn.Module):
      def __init__(self):
          super().__init__()
```

```

    ↳ #####
        # TODO: Design your own network, define layers here. ┐
    ↳     #
        # Here We provide a sample of two-layer fc network from HW4 Part3. ┐
    ↳     #
        # Your solution, however, should contain convolutional layers. ┐
    ↳     #
        # Refer to PyTorch documentations of torch.nn to pick your layers. ┐
    ↳     #
        # (https://pytorch.org/docs/stable/nn.html) ┐
    ↳     #
        # Some common choices: Linear, Conv2d, ReLU, MaxPool2d, AvgPool2d, ┐
    ↳ Dropout #
        # If you have many layers, use nn.Sequential() to simplify your code ┐
    ↳     #
    ┐
    ↳ #####
        self.conv1 = nn.
    ↳ Conv2d(in_channels=1,out_channels=16,kernel_size=3,stride=1,padding=1)
        self.conv2 = nn.Conv2d(in_channels=16, out_channels=32,kernel_size=3, ┐
    ↳ stride=1, padding=1)

        self.relu = nn.ReLU()
        self.pool = nn.MaxPool2d(kernel_size=2,stride=2)

        self.fc1 = nn.Linear(32*7*7, 128)
        self.fc2 = nn.Linear(128, 10)

    ┐
    ↳ #####
        # END OF YOUR CODE ┐
    ↳     #
    ┐
    ↳ #####

    def forward(self, x):
    ┐
    ↳ #####
        # TODO: Design your own network, implement forward pass here ┐
    ↳     #
    ┐
    ↳ #####
        x = self.pool(self.relu(self.conv1(x)))
        x = self.pool(self.relu(self.conv2(x)))
        x = x.view(-1,32*7*7)

```

```

        x = self.relu(self.fc1(x))
        x = self.fc2(x)
        return x

    ↵
↪ #####
    #                                END OF YOUR CODE                                ↵
↪    #
    ↵
↪ #####

model = Network().to(device)
criterion = nn.CrossEntropyLoss() # Specify the loss layer
print('Your network:')
print(summary(model, (1,28,28), device=device)) # visualize your model

#####
# TODO: Modify the lines below to experiment with different optimizers,      #
# parameters (such as learning rate) and number of epochs.                  #
#####
# Set up optimization hyperparameters
learning_rate, weight_decay, num_epoch = 0.001, 0.001, 10
optimizer = optim.Adam(model.parameters(), lr = learning_rate, ↵
    ↪weight_decay=weight_decay)
#####
#                                END OF YOUR CODE                                #
#####

```

Your network:

```

-----

```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 16, 28, 28]	160
ReLU-2	[-1, 16, 28, 28]	0
MaxPool2d-3	[-1, 16, 14, 14]	0
Conv2d-4	[-1, 32, 14, 14]	4,640
ReLU-5	[-1, 32, 14, 14]	0
MaxPool2d-6	[-1, 32, 7, 7]	0
Linear-7	[-1, 128]	200,832
ReLU-8	[-1, 128]	0
Linear-9	[-1, 10]	1,290

```

=====
Total params: 206,922
Trainable params: 206,922
Non-trainable params: 0
-----

Input size (MB): 0.00
Forward/backward pass size (MB): 0.33

```

Params size (MB): 0.79
Estimated Total Size (MB): 1.12

None

Run the cell below to start your training, we expect you to achieve over **85%** on the test set. A valid solution that meet the requirement take no more than **10 minutes** on normal PC Intel core CPU setting. If your solution takes too long to train, try to simplify your model or reduce the number of epochs.

```
[ ]: %%time
def train(model, trainloader, valloader, num_epoch=10): # Train the model
    print("Start training...")
    trn_loss_hist = []
    trn_acc_hist = []
    val_acc_hist = []
    model.train() # Set the model to training mode
    for i in range(num_epoch):
        running_loss = []
        print('-----Epoch = %d-----' % (i+1))
        for batch, label in tqdm(trainloader):
            batch = batch.to(device)
            label = label.to(device)
            optimizer.zero_grad() # Clear gradients from the previous iteration
            # This will call Network.forward() that you implement
            pred = model(batch)
            loss = criterion(pred, label) # Calculate the loss
            running_loss.append(loss.item())
            loss.backward() # Backprop gradients to all tensors in the network
            optimizer.step() # Update trainable weights
        print("\n Epoch {} loss:{}".format(i+1, np.mean(running_loss)))

        # Keep track of training loss, accuracy, and validation loss
        trn_loss_hist.append(np.mean(running_loss))
        trn_acc_hist.append(evaluate(model, trainloader))
        print("\n Evaluate on validation set...")
        val_acc_hist.append(evaluate(model, valloader))
    print("Done!")
    return trn_loss_hist, trn_acc_hist, val_acc_hist

def evaluate(model, loader): # Evaluate accuracy on validation / test set
    model.eval() # Set the model to evaluation mode
    correct = 0
    with torch.no_grad(): # Do not calculate gradient to speed up computation
        for batch, label in tqdm(loader):
            batch = batch.to(device)
            label = label.to(device)
```

```

        pred = model(batch)
        correct += (torch.argmax(pred, dim=1) == label).sum().item()
    acc = correct/len(loader.dataset)
    print("\n Evaluation accuracy: {}".format(acc))
    return acc

trn_loss_hist, trn_acc_hist, val_acc_hist = train(model, trainloader,
                                                  valloader, num_epoch)

#####
# TODO: Note down the evaluation accuracy on test set                                     #
#####
print("\n Evaluate on test set")
evaluate(model, testloader)

```

Start training...

-----Epoch = 1-----

100%| | 1000/1000 [00:11<00:00, 88.20it/s]

Epoch 1 loss:0.45427982886135576

100%| | 1000/1000 [00:08<00:00, 116.68it/s]

Evaluation accuracy: 0.88322

Evaluate on validation set...

100%| | 200/200 [00:01<00:00, 112.52it/s]

Evaluation accuracy: 0.8756

-----Epoch = 2-----

100%| | 1000/1000 [00:14<00:00, 68.83it/s]

Epoch 2 loss:0.3132331562563777

100%| | 1000/1000 [00:13<00:00, 75.81it/s]

Evaluation accuracy: 0.90188

Evaluate on validation set...

100%| | 200/200 [00:02<00:00, 76.01it/s]

Evaluation accuracy: 0.8898

```

-----Epoch = 3-----
100%|      | 1000/1000 [00:20<00:00, 47.72it/s]

Epoch 3 loss:0.27816830180585383
100%|      | 1000/1000 [00:20<00:00, 48.48it/s]

Evaluation accuracy: 0.90692

Evaluate on validation set...
100%|      | 200/200 [00:04<00:00, 48.26it/s]

Evaluation accuracy: 0.8909
-----Epoch = 4-----
100%|      | 1000/1000 [00:32<00:00, 30.58it/s]

Epoch 4 loss:0.25694305121526123
100%|      | 1000/1000 [00:23<00:00, 43.44it/s]

Evaluation accuracy: 0.9169

Evaluate on validation set...
100%|      | 200/200 [00:04<00:00, 42.62it/s]

Evaluation accuracy: 0.8977
-----Epoch = 5-----
100%|      | 1000/1000 [00:33<00:00, 29.63it/s]

Epoch 5 loss:0.24221466190367938
100%|      | 1000/1000 [00:26<00:00, 38.06it/s]

Evaluation accuracy: 0.92516

Evaluate on validation set...
100%|      | 200/200 [00:05<00:00, 35.77it/s]

Evaluation accuracy: 0.906
-----Epoch = 6-----
100%|      | 1000/1000 [00:34<00:00, 28.83it/s]

```

```

Epoch 6 loss:0.22939964060112833
100%|      | 1000/1000 [00:26<00:00, 37.96it/s]

Evaluation accuracy: 0.92846

Evaluate on validation set...
100%|      | 200/200 [00:04<00:00, 42.81it/s]

Evaluation accuracy: 0.906
-----Epoch = 7-----
100%|      | 1000/1000 [00:35<00:00, 28.18it/s]

Epoch 7 loss:0.21812415209040045
100%|      | 1000/1000 [00:26<00:00, 38.35it/s]

Evaluation accuracy: 0.93176

Evaluate on validation set...
100%|      | 200/200 [00:04<00:00, 43.88it/s]

Evaluation accuracy: 0.9114
-----Epoch = 8-----
100%|      | 1000/1000 [00:33<00:00, 30.26it/s]

Epoch 8 loss:0.21002072999626398
100%|      | 1000/1000 [00:24<00:00, 41.14it/s]

Evaluation accuracy: 0.93376

Evaluate on validation set...
100%|      | 200/200 [00:04<00:00, 41.02it/s]

Evaluation accuracy: 0.9128
-----Epoch = 9-----
100%|      | 1000/1000 [00:35<00:00, 28.40it/s]

Epoch 9 loss:0.20510010103322565

```



```
100%|      | 1000/1000 [00:25<00:00, 39.44it/s]
```

```
Evaluation accuracy: 0.92936
```

```
Evaluate on validation set...
```

```
100%|      | 200/200 [00:04<00:00, 41.56it/s]
```

```
Evaluation accuracy: 0.9046
```

```
-----Epoch = 10-----
```

```
100%|      | 1000/1000 [00:36<00:00, 27.54it/s]
```

```
Epoch 10 loss:0.19699719595909118
```

```
100%|      | 1000/1000 [00:22<00:00, 44.43it/s]
```

```
Evaluation accuracy: 0.93926
```

```
Evaluate on validation set...
```

```
100%|      | 200/200 [00:04<00:00, 43.24it/s]
```

```
Evaluation accuracy: 0.9157
```

```
Done!
```

```
Evaluate on test set
```

```
100%|      | 200/200 [00:04<00:00, 43.46it/s]
```

```
Evaluation accuracy: 0.9112
```

```
CPU times: user 2h 1min 28s, sys: 2.4 s, total: 2h 1min 31s
```


```
Wall time: 9min 11s
```

```
[ ]: 0.9112
```

Once your training is complete, run the cell below to visualize the training and validation accuracies across iterations.

```
[ ]: #####
# TODO: Submit the accuracy plot                                     #
#####
# visualize the training / validation accuracies
x = np.arange(num_epoch)
# train/val accuracies for MiniVGG
plt.figure()
```

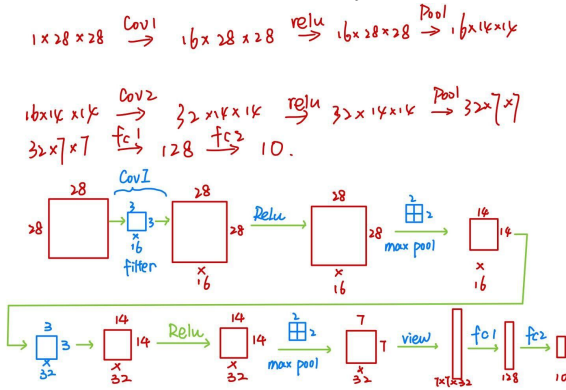
```
plt.plot(x, trn_acc_hist)
plt.plot(x, val_acc_hist)
plt.legend(['Training', 'Validation'])
plt.xticks(x)
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.title('fashion MNIST Classification')
plt.gcf().set_size_inches(10, 5)
plt.savefig('part1.png', dpi=300)
plt.show()
```



Q2 :Report the detailed architecture of your best model. Include information on hyperparameters chosen for training and a plot showing both training and validation accuracy across iterations.

A2 :

The detaied architecture of my best model:




The hyperparameters i chosen to build this cnn:

```

1      #set up the batch_size
2      batch_size = 50
3      # Set up optimization hyperparameters
4      learning_rate, weight_decay, num_epoch = 0.001, 0.001, 10
5      optimizer = optim.Adam(model.parameters(), lr = learning_rate, weight_decay=weight_decay)
6

```

plot showing both training and validation accuracy:



Q3 : Report the accuracy of your best model on the test set.

A3 : My model final test accuracy is 91.12%

```

Evaluate on test set
100%|██████████| 200/200 [00:04<00:00, 43.46it/s]

Evaluation accuracy: 0.9112

```

6 Task 6: Pre-trained NN

Q1 :One image (img1) where the pretrained model gives reasonable predictions, and produces a category label that seems to correctly describe the content of the image

A1 :

Predicted Class: convertible



Q2 : One image (img2) where the pretrained model gives unreasonable predictions, and produces a category label that does not correctly describe the content of the image.

A2 :

Predicted Class: parachute

