

HW5

Yuzhen Chen

December 20, 2024

1 Section1: Pix2Pix

- Task 1.1: Dataloading

Q1 : Implement the Edges2Image class and fill in the TODOs in that cell.

A1 :

```
1  class Edges2Image(Dataset):
2  def __init__(self, root_dir, split='train', transform=None):
3      """
4      Args:
5          root_dir: the directory of the dataset
6          split: "train" or "val"
7          transform: pytorch transformations.
8      """
9
10     self.transform = transform
11
12     self.files = glob.glob(os.path.join(root_dir, split, '*.jpg'))
13
14     def __len__(self):
15         return len(self.files)
16
17     def __getitem__(self, idx):
18         img = Image.open(self.files[idx])
19         img = np.asarray(img)
20         if self.transform:
21             img = self.transform(img)
22         return img
23
24 transform = transforms.Compose([
25     transforms.ToTensor(),
26     transforms.Normalize(mean=(0.5, 0.5, 0.5), std=(0.5, 0.5, 0.5))
27 ])
28
29 #####
30 # TODO: Construct the dataloader
31 # For the train_loader, please use a batch size of 4 and set shuffle True
32 # For the val_loader, please use a batch size of 5 and set shuffle False
33 # Hint: You'll need to create instances of the class above, name them as
34 # tr_dt and te_dt. The dataloaders should be named as train_loader and
35 # test_loader. You also need to include transform in your class
36 #instances
37 #####
38
39 tr_dt = Edges2Image(root_dir="/home/yuzhen/Desktop/EECS442/hw5_updated/hw5-easy-release/mini-edges2shoes",split="train",transform=transform)
40 te_dt = Edges2Image(root_dir="/home/yuzhen/Desktop/EECS442/hw5_updated/hw5-easy-release/mini-edges2shoes",split="val",transform=transform)
41
42 train_loader = DataLoader(dataset=tr_dt,batch_size = 4,shuffle=True)
43 test_loader = DataLoader(dataset=te_dt,batch_size = 5,shuffle=False)
44
45 #####
46 #                               END OF YOUR CODE
47 #####
48
49 # Make sure that you have 1,000 training images and 100 testing images before moving on
```

```

50 print('Number of training images {}, number of testing images {}'.format(len(tr_dt), len(te_dt)))
51
52

```

- Task 2: Training Pix2Pix

Q1 : Please set up G_optimizer and D_optimizer in the train function.

A1 :

```

1      def train(G, D, num_epochs = 20):
2          hist_D_losses = []
3          hist_G_losses = []
4          hist_G_L1_losses = []
5          #####
6          # TODO: Add Adam optimizer to generator and discriminator      #
7          # You will use lr=0.0002, beta=0.5, beta2=0.999                #
8          #####
9
10         lr = 0.0002
11         beta1 = 0.5
12         beta2 = 0.999
13         G_optimizer = optim.Adam(G.parameters(), lr=lr, betas=(beta1, beta2))
14         D_optimizer = optim.Adam(D.parameters(), lr=lr, betas=(beta1, beta2))
15
16         #####
17         #                               END OF YOUR CODE                #
18         #####
19

```

Q2 : Implement the code for the function train as instructed by the notebook.

A2 :

```

1          # Hint: you could use following loss to complete following function
2 BCE_loss = nn.BCELoss().cuda()
3 L1_loss = nn.L1Loss().cuda()
4
5 def train(G, D, num_epochs = 20):
6     hist_D_losses = []
7     hist_G_losses = []
8     hist_G_L1_losses = []
9     #####
10    # TODO: Add Adam optimizer to generator and discriminator      #
11    # You will use lr=0.0002, beta=0.5, beta2=0.999                #
12    #####
13
14    lr = 0.0002
15    beta1 = 0.5
16    beta2 = 0.999
17    G_optimizer = optim.Adam(G.parameters(), lr=lr, betas=(beta1, beta2))
18    D_optimizer = optim.Adam(D.parameters(), lr=lr, betas=(beta1, beta2))
19
20    #####
21    #                               END OF YOUR CODE                #
22    #####
23
24    print('training start!')
25    start_time = time.time()
26    for epoch in range(num_epochs):
27        print('Start training epoch %d' % (epoch + 1))
28        D_losses = []
29        G_losses = []
30        epoch_start_time = time.time()
31        num_iter = 0
32        for x_ in train_loader:
33            y_ = x_[ :, :, :, img_size:]
34            x_ = x_[ :, :, :, 0:img_size]
35
36            x_, y_ = x_.cuda(), y_.cuda()
37            #####
38            # TODO: Implement training code for the discriminator.      #
39            # Recall that the loss is the mean of the loss for real images and fake #
40            # images, and made by some calculations with zeros and ones    #

```

```

41 # We have defined the BCE_loss, which you might would like to use. #
42 # #
43 # NOTE: While training the Discriminator, the output of the generator #
44 # must be detached from the computational graph. Refer to the method #
45 # torch.Tensor.detach() #
46 #####
47
48 #reset the optimizer
49 D_optimizer.zero_grad()
50
51 # train on real data
52 x_y_ = torch.hstack((x_, y_))
53 D_real_preds = D(x_y_)
54 D_real_preds_shape = D_real_preds.shape
55 the_real_label = torch.ones(D_real_preds_shape, device = device)
56
57 D_real_loss = BCE_loss(D_real_preds, the_real_label)
58
59 # train on fake data
60 fake_data = G(x_).detach()
61 D_fake_preds = D(torch.hstack((x_, fake_data)))
62 D_fake_preds_shape = D_fake_preds.shape
63 the_fake_label = torch.zeros(D_fake_preds_shape, device = device)
64 D_fake_loss = BCE_loss(D_fake_preds, the_fake_label)
65
66 loss_D = (D_real_loss + D_fake_loss) / 2
67
68 loss_D.backward()
69 D_optimizer.step()
70
71 #####
72 # END OF YOUR CODE #
73 #####
74
75 #####
76 # TODO: Implement training code for the Generator. #
77 #####
78
79 # 1. Train the generator
80 # 2. Append the losses to the lists 'hist_G_L1_losses' and 'hist_D_losses'
81 # (Only append the data to the list, not the complete tensor, refer
82 # torch.Tensor.item()).
83
84 G_optimizer.zero_grad()
85
86
87 generator_data = G(x_)
88 x_gen_data = torch.hstack((x_, generator_data))
89 prediction = D(x_gen_data)
90 G_L1_loss = L1_loss(generator_data, y_)
91 lambda_L1 = 100
92 loss_G = BCE_loss(prediction, the_real_label) + lambda_L1 * G_L1_loss
93
94
95 loss_G.backward()
96 G_optimizer.step()
97 #####
98 # END OF YOUR CODE #
99 #####
100
101 D_losses.append(loss_D.detach().item())
102 hist_D_losses.append(loss_D.detach().item())
103 G_losses.append(loss_G)
104
105 hist_G_losses.append(loss_G.detach().item())
106 hist_G_L1_losses.append(G_L1_loss.detach().item())
107 num_iter += 1
108
109
110
111 epoch_end_time = time.time()
112 per_epoch_ptime = epoch_end_time - epoch_start_time
113
114 print('[%d/%d] - using time: %.2f seconds' % ((epoch + 1), num_epochs, per_epoch_ptime))
115 print('loss of discriminator D: %.3f' % (torch.mean(torch.FloatTensor(D_losses))))
116 print('loss of generator G: %.3f' % (torch.mean(torch.FloatTensor(G_losses))))

```

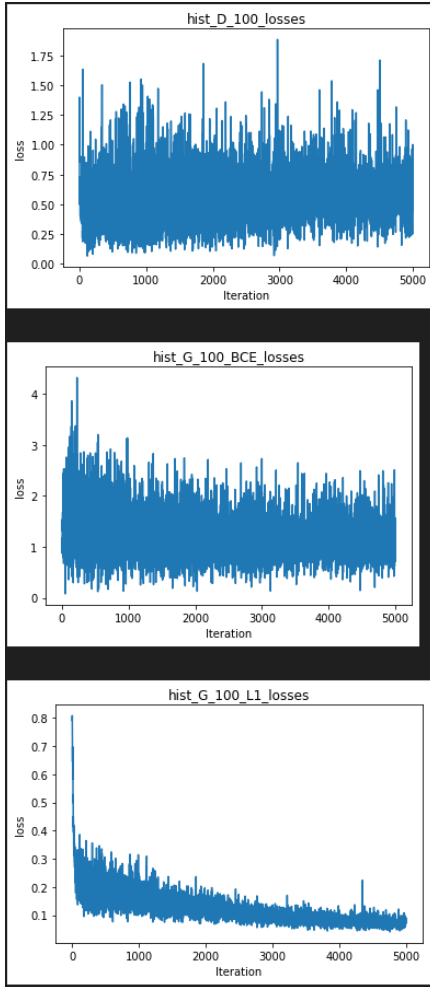
```

117     if epoch == 0 or (epoch + 1) % 5 == 0:
118         with torch.no_grad():
119             show_result(G, fixed_x_, fixed_y_, (epoch+1))
120
121     end_time = time.time()
122     total_ptime = end_time - start_time
123
124     return hist_D_losses, hist_G_losses, hist_G_L1_losses
125

```

Q3 : In your report, include these plots.

A3 :



2 Section 2: Diffusion Models

- : Task 3.1: Implement the method `get_named_beta_schedule` with linear and cosine schedules.

Q1 :

A1 :

```

1     def get_named_beta_schedule(schedule_name, num_diffusion_timesteps, beta_min=0.0001,
2         beta_max=0.02):
3         """
4         Get a pre-defined beta schedule for the given name.
5
6         Args:
7             schedule_name: str, name of the variance schedule, 'linear' or 'cosine'
8             num_diffusion_timesteps: int, number of the entire diffusion timesteps
9             beta_min: float, minimum value of beta
10            beta_max: float, maximum value of beta
11
12         Returns:
13

```

```

12     betas: np.ndarray, a 1-d array of size num_diffusion_timesteps, contains all the beta for
    each timestep
13
14     """
15     betas = [0]
16     if schedule_name == "linear":
17         ##### START TODO #####
18         # Implement the linear schedule
19         # Uniformly divide the [beta_min, beta_max) to num_diffusion_timesteps values.
20
21         betas = np.arange(beta_min, beta_max, (beta_max - beta_min) / num_diffusion_timesteps)
22         ##### END TODO #####
23     elif schedule_name == "cosine":
24         ##### START TODO #####
25         # Implement the cosine schedule
26         # Assume s = 0.008 and beta_clip=0.999
27         s = 0.008
28         beta_clip = 0.999
29
30         # betas = np.zeros(num_diffusion_timesteps)
31         old_alpha = 1
32         normalize_factor = math.cos((s/(1+s))*(math.pi/2))**2
33         T = num_diffusion_timesteps
34         for t in range(1,T):
35             alpha = math.cos((((t/T)+s)/(1+s))*(math.pi/2))**2 / normalize_factor
36             min_beta = min(beta_clip, 1-(alpha/old_alpha))
37             betas.append(min_beta)
38             #update the new alpha
39             old_alpha = alpha
40         del betas[0]
41         betas.append(0.999)
42
43         ##### End TODO #####
44
45         # Add a value to the end of betas
46
47     else:
48         raise NotImplementedError(f"unknown beta schedule: {schedule_name}")
49
50     return betas
51
52
53

```

- : Task 4.1: p_sample and p_sample_loop of class DDPMDiffusion

Q1 : Fill the TODO sections of the class DDPMDiffusion of the file guided_diffusion/simple_diffusion.py. Complete the methods p_sample and _sample_loop.

A1 :

```

1     ##### START TODO #####
2     # Calculate the values of alpha
3     # Also we will need the cumulated product of alpha.
4     # And during sampling we need the value of cumulated product of alpha from
5     # previous or next timestep.
6     self.alphas = 1 - self.betas
7     self.alphas_cumprod = np.cumprod(self.alphas) # cumulated product of alphas
8     self.alphas_cumprod_prev = np.concatenate(([1.0], self.alphas_cumprod[:-1]))
9     self.alphas_cumprod_next = np.concatenate((self.alphas_cumprod[1:], [0.0]))
10
11     ##### END TODO #####
12
13
14     def p_sample_loop(self,
15                       model,
16                       x_start,
17                       record,
18                       save_root,
19                       measurement=None,
20                       measurement_cond_fn=None,
21                       uncond=False):
22
23     """
24     The function used for sampling from noise.

```

```

11
12     Args:
13         model: nn.Module, the pretrained model that is used to predict the score and variance
14         x_start: torch.Tensor, random noise input
15         measurement: torch.Tensor, our corrupted observation
16         measurement_cond_fn: conditional function used to perform conditional sampling, is
17         None for unconditional sampling
18         record: Bool, save intermediate results if True
19         save_root: str, root of the directory to save the results
20         uncond: Bool, perform unconditional sampling if True, else perform conditional
21         sampling
22         """
23         if not uncond:
24             assert measurement is not None and measurement_cond_fn is not None, \
25                 "measurement and measurement conditional function is required for conditional
26                 sampling"
27
28         img = x_start # start from random noise
29         device = x_start.device
30
31         ##### Start TODO #####
32         # Implement the sample loop
33         # Call p_sample for every iteration
34         # It requires only one line of code implementation here
35
36         pbar = tqdm(list(range(self.num_timesteps))[:-1])
37         for idx in pbar:
38             time = torch.tensor([idx] * img.shape[0], device=device)
39
40             img = self.p_sample(model, img, time)["x_t_minus_1"]
41
42             img = img.detach_()
43
44             #####
45             if record:
46                 if idx % 10 == 0:
47                     file_path = os.path.join(save_root, f"progress/x_{str(idx).zfill(4)}.png")
48                     plt.imshow(file_path, clear_color(img))
49
50         return img

```

```

1         def p_sample(self, model, x, t):
2             """
3             Posterior sampling process, when given the model, x_t and timestep t, it returns
4             predicted
5             x_0 and x_t_minus_1
6
7             We have already provide you with the function to get the log of the variance.
8             Use self.mean_processor.get_mean_and_xstart(var_values, t), where var_values is
9             the 3:6 channels of the direct output of the model.
10            example usage: log_variance = self.var_processor.get_variance(var_values, t)
11
12            You can also use the helper function extract_and_expand() to extract the value
13            corresponding to timestep and expand it to the save size as the target for broadcast.
14            example usage: coef1 = extract_and_expand(self.posterior_mean_coef1, t, x_start)
15
16            Args:
17                model: nn.Module, the UNet model, you can call model(x, t) to get the output tensor
18                with size (B, 6, H, W)
19                x: torch.Tensor, shape (1, 3, H, W), x_t
20                t: torch.Tenosr, shape (1,), timestep
21
22            Returns:
23                output_dict: dict, contains predicted x_t_minus_1 and x_0
24                """
25                #####Start TODO#####
26                ##### Get the predicted score and variance of the pretrained model #####
27                pred_noise = model(x,t)[:,:3]
28                var_values = model(x,t)[:,:3:6]
29                ##### End TODO #####
30
31                log_variance = self.var_processor.get_variance(var_values, t) # get the log of
32                variance
33
34                ##### Start TODO #####

```

```

32     ##### get predicted x_0 and x_t_minus_1 #####
33     ##### don't forget to add noise for all the steps, except for the last one #####
34     exp_log_variance = torch.exp(log_variance)
35     sigma = torch.sqrt(exp_log_variance)
36
37     coef1 = extract_and_expand(self.alphas, t, x).cuda()
38     coef2 = extract_and_expand(self.alphas_cumprod, t, x).cuda()
39
40     if t > 0:
41         z = torch.randn(x.size()).cuda() #test
42     else:
43         z = torch.zeros_like(x)
44
45
46
47     first_part = z * sigma
48     numerator = x - ((1-coef1) / torch.sqrt(1-coef2)) * pred_noise
49     denominator = torch.sqrt(coef1)
50     second_part = numerator / denominator
51
52
53     x_t_minus_1 = first_part + second_part
54
55     ##### End TODO #####
56
57     assert x_t_minus_1.shape == log_variance.shape == x.shape
58
59     output_dict = {'x_t_minus_1': x_t_minus_1}
60     return output_dict
61

```

- : Task 4.2: Sampled Image

Q1 :

A1 :



- : Task 5.1: p_sample of class DDIMDiffusion

Q1 :

A1 :

```

1     def p_sample(self, model, x, t, eta=0.0):
2         #####
3         ##### TODO #####
4         ##### Get the predicted score and variance of the pretrained model #####
5         ##### Don't forget to use _scale_timesteps to scale the timestep for calling the model
6         prediction.
7         ##### You don't need to scale the timestep for further computations of x_t_minus_1.
8         ##### NOTE: Since this version of the model learns the variance along with the score
9         function,
10        ##### the output of the model would have double the number of channels as that of the
11        input.
12        ##### So assign the predicted score and variance values to the variables below. Refer to
13        ##### torch.split method.
14        #####

```

```

12     ##### Start TODO #####
13
14     scaled_t = self._scale_timesteps(t)
15     mode_output = model(x,scaled_t)
16     pred_noise = mode_output[:, :x.shape[1]]
17     ##### End TODO #####
18
19     model_mean, pred_xstart = self.mean_processor.get_mean_and_xstart(x, t, pred_noise)
20
21     ##### TODO #####
22     ##### Step 1: Implement the variance parameter 'sigma' for DDIM sampling. #####
23     ##### Step 2: Implement x_t_minus_1 using the pred_xstart. Don't forget #####
24     ##### to add noise for all the steps, except for the t=0. #####
25     ##### #####
26     ##### You may use the function 'extract_and_expand' to expand the timestep #####
27     ##### variable 't' to the input's shape. #####
28     ##### Assign them to the variables x_t_minus_1. #####
29
30     ##### Start TODO #####
31     a_t = extract_and_expand(self.alphas_cumprod,t,x).cuda()
32     a_t_minus_1 = extract_and_expand(self.alphas_cumprod_prev,t,x).cuda()
33     if t > 0:
34         z = torch.randn_like(x).cuda()
35     else:
36         z = torch.zeros_like(x)
37     sigma = eta*torch.sqrt(torch.sqrt((1-a_t_minus_1)/(1-a_t))*torch.sqrt(1-(a_t)/(
a_t_minus_1)))
38
39     predicted_x0 = torch.sqrt(a_t_minus_1) * ((x-torch.sqrt(1-a_t)*pred_noise)/(torch.sqrt(
a_t)))
40     direction_xt = (torch.sqrt(1-a_t_minus_1-(sigma**2))*pred_noise)
41     x_t_minus_1 = predicted_x0 + direction_xt
42
43     if t > 0:
44         x_t_minus_1 += sigma * z
45     ##### End TODO #####
46
47     return {"x_t_minus_1": x_t_minus_1, "pred_xstart": pred_xstart}
48     #####
49
50
51     def predict_eps_from_x_start(self, x_t, t, pred_xstart):
52         coef1 = extract_and_expand(self.sqrt_recip_alphas_cumprod, t, x_t)
53         coef2 = extract_and_expand(self.sqrt_recipm1_alphas_cumprod, t, x_t)
54         return (coef1 * x_t - pred_xstart) / coef2
55

```

- : Task 5.2:Sampled Image

Q1 :

A1 :



- : Task 7.1:PosteriorSampling

Q1 :

A1 :


```

1         def conditioning(self, x_i, x_t_minus_one, x_0_hat, measurement, **kwargs):
2             """
3             The conditioning function as shown in line 7
4
5             Args:
6                 x_i: torch.Tensor, x_i
7                 x_t, torch.Tensor, x_t_minus_1 prime
8                 x_0_hat: torch.Tensor, predicted x_0
9                 measurement: torch.Tensor, y, the corrupted image
10            """
11            # norm_grad, norm = self.grad_and_value(x_prev=x_prev, x_0_hat=x_0_hat, measurement=
12            measurement, **kwargs)
13            ##### Start TODO #####
14            ##### Implement the conditional sampling in line 7 #####
15            ##### A(x_0_hat) is already provided to you as A #####
16            ##### Also torch.autograd.grad() is provided to you to calculate the gradient of the
17            ##### norm term with respect to x_i, you can check https://pytorch.org/docs/stable/
18            generated/torch.autograd.grad.html#torch.autograd.grad
19            ##### for its detailed usage. You only need to specify the outputs and inputs here.
20            # A = self.operator.forward(x_0_hat, **kwargs)
21            A = self.operator.forward(x_0_hat, **kwargs)
22
23            difference = measurement - A
24            norm = torch.sqrt(torch.sum(difference ** 2))
25            diff_output = norm # outputs of the differentiated function
26            diff_input = x_i # Inputs w.r.t. which the gradient will be returned
27
28            ## TODO: Don't delete this line, you will use this
29            norm_grad = torch.autograd.grad(outputs=diff_output, inputs=diff_input)[0]
30            new_x_t_minus_one = x_t_minus_one - norm_grad
31            ##### END TODO #####
32            return new_x_t_minus_one

```

• : Task 7.2: Inpainted Image

Q1 :

A1 :

