

HW6

Yuzhen Chen

December 20, 2024

1 Section1: Pix2Pix

- Task 1.1: find_fundamental_matrix

Q1 : find_fundamental_matrix

A1 :

```
1  def normalize(points):
2      centroid = np.mean(points, axis=0)
3      scaled_points = points - centroid
4
5      squared_points = scaled_points ** 2
6      sum_of_squares = np.sum(squared_points, axis=1)
7      sqrt_sum_of_squares = np.sqrt(sum_of_squares)
8      mean_distance = np.mean(sqrt_sum_of_squares)
9
10     const_scale = np.sqrt(2) / mean_distance
11     transformation_matrix = np.array([[const_scale, 0, -const_scale * centroid[0]],
12                                       [0, const_scale, -const_scale * centroid[1]],
13                                       [0, 0, 1]])
14
15
16     ones_row = np.ones((points.shape[0],))
17     vstack_result = np.vstack((points.T, ones_row))
18
19     point_after_norm = np.dot(transformation_matrix, vstack_result)
20     return point_after_norm[:2].T, transformation_matrix
21
22
23
24
25
26 def find_fundamental_matrix(shape, pts1, pts2):
27     """
28     Computes Fundamental Matrix F that relates points in two images by the:
29
30         [u' v' 1] F [u v 1]^T = 0
31     or
32         l = F [u v 1]^T  -- the epipolar line for point [u v] in image 2
33         [u' v' 1] F = l'   -- the epipolar line for point [u' v'] in image 1
34
35     Where (u,v) and (u',v') are the 2D image coordinates of the left and
36     the right images respectively.
37
38     Inputs:
39     - shape: Tuple containing shape of img1
40     - pts1: Numpy array of shape (N,2) giving image coordinates in img1
41     - pts2: Numpy array of shape (N,2) giving image coordinates in img2
42
43     Returns:
44     - F: Numpy array of shape (3,3) giving the fundamental matrix F
45     """
46
47     #This will give you an answer you can compare with
48     #Your answer should match closely once you've divided by the last entry
49     FOpenCV, _ = cv2.findFundamentalMat(pts1, pts2, cv2.FM_8POINT)
50
51     F = np.eye(3)
```

```

52 #####
53 # TODO: Your code here #
54 #####
55
56 norm_pt1, T1 = normalize(pts1)
57 norm_pt2, T2 = normalize(pts2)
58 # Construct matrix A for the equation Ax = 0
59 A = np.zeros((len(pts1), 9))
60 for i in range(len(pts1)):
61
62     u1, v1 = norm_pt1[i]
63     u2, v2 = norm_pt2[i]
64     A[i] = [u2*u1, u2*v1, u2, v2*u1, v2*v1, v2, u1, v1, 1]
65
66
67 _, _, V = np.linalg.svd(A)
68 F = V[-1].reshape(3, 3)
69 U, S, V = np.linalg.svd(F)
70 S[-1] = 0 # Set smallest singular value to zero
71 temp_F = U @ np.diag(S) @ V
72 F = T2.T @ temp_F @ T1
73
74 #####
75 #                               END OF YOUR CODE #
76 #####
77 return F
78
79

```

- Task 1.2: compute_epipoles

Q1 : Please set up G_optimizer and D_optimizer in the train function.

A1 :

```

1     def compute_epipoles(F):
2
3         """
4         Given a Fundamental Matrix F, return the epipoles represented in
5         homogeneous coordinates.
6
7         Check: e2@F and F@e1 should be close to [0,0,0]
8
9         Inputs:
10         - F: the fundamental matrix
11
12         Return:
13         - e1: the epipole for image 1 in homogeneous coordinates
14         - e2: the epipole for image 2 in homogeneous coordinates
15         """
16         #####
17         # TODO: Your code here #
18         #####
19         # Compute the singular value decomposition of F
20         U, S, V = np.linalg.svd(F)
21
22         # Extract the last column of V to get the right singular vector
23         e1 = V[-1]
24
25         # Compute the singular value decomposition of F transpose to get the left singular vector
26         U, S, V = np.linalg.svd(F.T)
27
28         # Extract the last column of V to get the left singular vector
29         e2 = V[-1]
30
31         # Normalize the epipoles
32         e1 = e1 / e1[2]
33         e2 = e2 / e2[2]
34
35         #####
36         #                               END OF YOUR CODE #
37         #####
38         return e1, e2
39

```

Q2 : Implement the code for the function train as instructed by the notebook.


A2 :

- Task 1.3: Show epipolar lines for temple, reallyInwards, and another dataset of your choice.

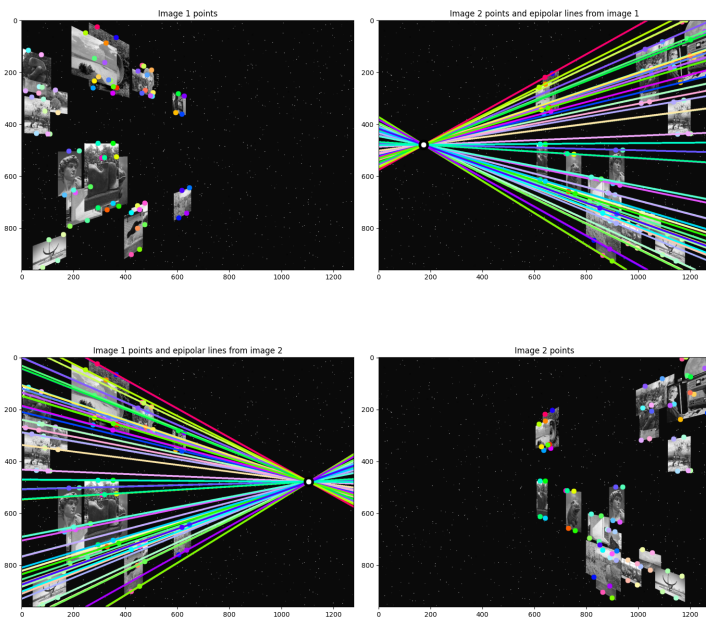
Q :Show epipolar lines for temple, reallyInwards, and another dataset of your choice.

A :

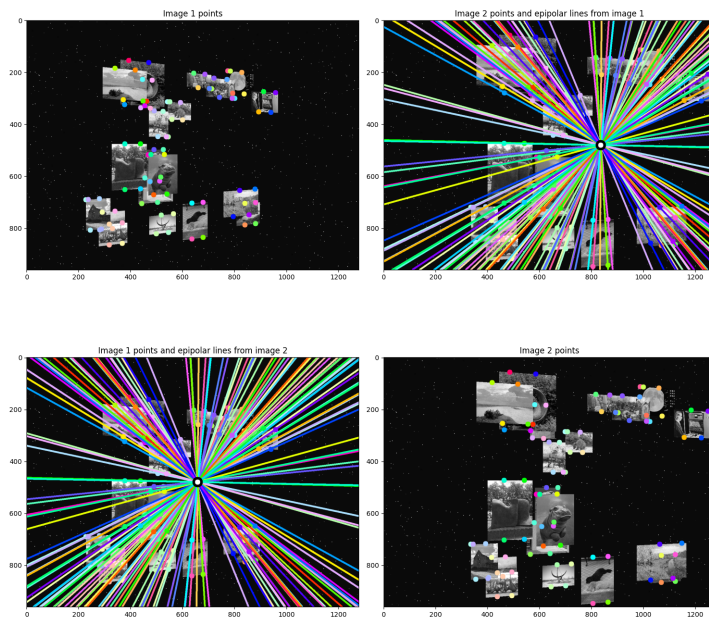
temple:



reallyInwards:



zrtransrot_us:



- Task 1.4: Report the epipoles for reallyInwards and xtrans.

Q :reallyInwards and xtrans.

A :
reallyInwards

```
[1.10594677e+03 4.79305073e+02 1.00000000e+00]
[173.14934439 479.41457107 1.          ]
[1105.94676885 479.30507302] [173.14934439 479.41457107]
```

xtrans

```
[-8.71977523e+17 2.02228728e+01 1.00000000e+00]
[-3.66017951e+17 3.93111988e+02 1.00000000e+00]
[-8.71977523e+17 2.02228728e+01] [-3.66017951e+17 3.93111988e+02]
```

- Task 2.1 Implement the function positional_encoder(x, L_embed = 6)

```
1  def positional_encoder(x, L_embed=6):
2  """
3  This function applies positional encoding to the input tensor. Positional encoding is used in NeRF
4  to allow the model to learn high-frequency details more effectively. It applies sinusoidal
5  functions
6  at different frequencies to the input.
7
8  Parameters:
9  x (torch.Tensor): The input tensor to be positionally encoded.
10 L_embed (int): The number of frequency levels to use in the encoding. Defaults to 6.
11
12 Returns:
13 torch.Tensor: The positionally encoded tensor.
14 """
15 # Initialize a list with the input tensor.
16 rets = [x]
17
18 # Loop over the number of frequency levels.
```

```

19 for i in range(L_embed):
20     #####
21     #                                TODO                                #
22     #####
23     sin_encoding = torch.sin(2.0 ** i * x)
24     cos_encoding = torch.cos(2.0 ** i * x)
25     rets.extend([sin_encoding, cos_encoding])
26     #####
27     #                                END OF YOUR CODE                                #
28     #####
29
30
31 # Concatenate the original and encoded features along the last dimension.
32 return torch.cat(rets, -1)

```

- Task 2.2 Implement the code that samples 3D points along a ray in render

```

1  def render(model, rays_o, rays_d, near, far, n_samples, rand=False):
2      """
3      Render a scene using a Neural Radiance Field (NeRF) model. This function samples points along rays,
4      evaluates the NeRF model at these points, and applies volume rendering techniques to produce an
5      image.
6
7      Parameters:
8      model (torch.nn.Module): The NeRF model to be used for rendering.
9      rays_o (torch.Tensor): Origins of the rays.
10     rays_d (torch.Tensor): Directions of the rays.
11     near (float): Near bound for depth sampling along the rays.
12     far (float): Far bound for depth sampling along the rays.
13     n_samples (int): Number of samples to take along each ray.
14     rand (bool): If True, randomize sample depths. Default is False.
15
16     Returns:
17     tuple: A tuple containing the RGB map and depth map of the rendered scene.
18     """
19
20     # Sample points along each ray, from 'near' to 'far'.
21     z = torch.linspace(near, far, n_samples).to(device)
22     if rand:
23         mids = 0.5 * (z[..., 1:] + z[..., :-1])
24         upper = torch.cat([mids, z[..., -1:]], -1)
25         lower = torch.cat([z[..., 0], mids], -1)
26         t_rand = torch.rand(z.shape).to(device)
27         z = lower + (upper - lower) * t_rand
28
29     #####
30     #                                TODO                                #
31     #####
32     # Compute 3D coordinates of the sampled points along the rays.
33     points = rays_o[..., None, :] + rays_d[..., None, :] * z[..., :, None]
34     #####
35     #                                END OF YOUR CODE                                #
36     #####
37
38     # Flatten the points and apply positional encoding.
39     flat_points = torch.reshape(points, [-1, points.shape[-1]])
40     flat_points = positional_encoder(flat_points)
41
42     # Evaluate the model on the encoded points in chunks to manage memory usage.
43     chunk = 1024 * 32
44     raw = torch.cat([model(flat_points[i:i + chunk]) for i in range(0, flat_points.shape[0], chunk)],
45                     0)
46     raw = torch.reshape(raw, list(points.shape[:-1]) + [4])
47
48     # Compute densities (sigmas) and RGB values from the model's output.
49     sigma = F.relu(raw[..., 3])
50     rgb = torch.sigmoid(raw[..., :3])
51
52     # Perform volume rendering to obtain the weights of each point.
53     one_e_10 = torch.tensor([1e10], dtype=rays_o.dtype).to(device)
54     dists = torch.cat((z[..., 1:] - z[..., :-1], one_e_10.expand(z[..., :1].shape)), dim=-1)
55     alpha = 1. - torch.exp(-sigma * dists)
56     weights = alpha * cumprod_exclusive(1. - alpha + 1e-10)
57
58     #####
59     #                                TODO                                #
60     #####

```

```

58 #####
59 # Compute the weighted sum of RGB values along each ray to get the final pixel color.
60 rgb_map = torch.sum(weights[... , None] * rgb, dim=-2)
61
62 # Compute the depth map as the weighted sum of sampled depths.
63 depth_map = torch.sum(weights * z, dim=-1)
64 #####
65 #                                     END OF YOUR CODE                                     #
66 #####
67
68 return rgb_map, depth_map

```

- Task 2.3 We will also compute, depth_map,

```

1  def render(model, rays_o, rays_d, near, far, n_samples, rand=False):
2  """
3  Render a scene using a Neural Radiance Field (NeRF) model. This function samples points along rays,
4  evaluates the NeRF model at these points, and applies volume rendering techniques to produce an
5  image.
6
7  Parameters:
8  model (torch.nn.Module): The NeRF model to be used for rendering.
9  rays_o (torch.Tensor): Origins of the rays.
10 rays_d (torch.Tensor): Directions of the rays.
11 near (float): Near bound for depth sampling along the rays.
12 far (float): Far bound for depth sampling along the rays.
13 n_samples (int): Number of samples to take along each ray.
14 rand (bool): If True, randomize sample depths. Default is False.
15
16 Returns:
17 tuple: A tuple containing the RGB map and depth map of the rendered scene.
18 """
19
20 # Sample points along each ray, from 'near' to 'far'.
21 z = torch.linspace(near, far, n_samples).to(device)
22 if rand:
23     mids = 0.5 * (z[... , 1:] + z[... , :-1])
24     upper = torch.cat([mids, z[... , -1:]], -1)
25     lower = torch.cat([z[... , :1], mids], -1)
26     t_rand = torch.rand(z.shape).to(device)
27     z = lower + (upper - lower) * t_rand
28
29 #####
30 #                                     TODO                                     #
31 #####
32 # Compute 3D coordinates of the sampled points along the rays.
33 points = rays_o[... , None, :] + rays_d[... , None, :] * z[... , :, None]
34 #####
35 #                                     END OF YOUR CODE                                     #
36 #####
37
38 # Flatten the points and apply positional encoding.
39 flat_points = torch.reshape(points, [-1, points.shape[-1]])
40 flat_points = positional_encoder(flat_points)
41
42 # Evaluate the model on the encoded points in chunks to manage memory usage.
43 chunk = 1024 * 32
44 raw = torch.cat([model(flat_points[i:i + chunk]) for i in range(0, flat_points.shape[0], chunk)],
45                 0)
46 raw = torch.reshape(raw, list(points.shape[:-1]) + [4])
47
48 # Compute densities (sigmas) and RGB values from the model's output.
49 sigma = F.relu(raw[... , 3])
50 rgb = torch.sigmoid(raw[... , :3])
51
52 # Perform volume rendering to obtain the weights of each point.
53 one_e_10 = torch.tensor([1e10], dtype=rays_o.dtype).to(device)
54 dists = torch.cat((z[... , 1:] - z[... , :-1], one_e_10.expand(z[... , :1].shape)), dim=-1)
55 alpha = 1. - torch.exp(-sigma * dists)
56 weights = alpha * cumprod_exclusive(1. - alpha + 1e-10)
57
58 #####
59 #                                     TODO                                     #
60 #####
61 # Compute the weighted sum of RGB values along each ray to get the final pixel color.
62 rgb_map = torch.sum(weights[... , None] * rgb, dim=-2)

```

```

61
62 # Compute the depth map as the weighted sum of sampled depths.
63 depth_map = torch.sum(weights * z, dim=-1)
64 #####
65 #                               END OF YOUR CODE                               #
66 #####
67
68 return rgb_map, depth_map

```

- Task2.4 Please implement part of the train(model, optimizer, n_iters) function.

```

1     mse2psnr = lambda x : -10. * torch.log(x) / torch.log(torch.Tensor([10.])).to(device)
2
3 def train(model, optimizer, n_iters=3000):
4     """
5     Train the Neural Radiance Field (NeRF) model. This function performs training over a specified
6     number of iterations,
7     updating the model parameters to minimize the difference between rendered and actual images.
8
9     Parameters:
10    model (torch.nn.Module): The NeRF model to be trained.
11    optimizer (torch.optim.Optimizer): The optimizer used for training the model.
12    n_iters (int): The number of iterations to train the model. Default is 3000.
13    """
14    psnrs = []
15    iternums = []
16
17    plot_step = 500
18    n_samples = 64 # Number of samples along each ray.
19
20    for i in tqdm(range(n_iters)):
21        # Randomly select an image from the dataset and use it as the target for training.
22        images_idx = np.random.randint(images.shape[0])
23        target = images[images_idx]
24        pose = poses[images_idx]
25
26
27        #####
28        #                               TODO                               #
29        #####
30        # Perform training. Use mse loss for loss calculation and update the model parameter using the
31        # optimizer.
32        # Hint: focal is defined as a global variable in previous section
33
34        rays_o, rays_d = get_rays(H, W, focal, pose)
35
36        # Render the scene using the NeRF model
37        rgb, _ = render(model, rays_o, rays_d, near=2., far=6., n_samples=64)
38
39        # Calculate MSE loss between rendered RGB image and target image
40        loss = torch.nn.functional.mse_loss(rgb, target)
41
42        # Clear previous gradients
43        optimizer.zero_grad()
44
45        # Backpropagation
46        loss.backward()
47
48        # Update model parameters
49        optimizer.step()
50
51        #####
52        #                               END OF YOUR CODE                               #
53        #####
54
55    if i % plot_step == 0:
56        torch.save(model.state_dict(), 'ckpt.pth')
57        # Render a test image to evaluate the current model performance.
58        with torch.no_grad():
59            rays_o, rays_d = get_rays(H, W, focal, testpose)
60            rgb, depth = render(model, rays_o, rays_d, near=2., far=6., n_samples=n_samples)
61            loss = torch.nn.functional.mse_loss(rgb, testing)
62            # Calculate PSNR for the rendered image.
63            psnr = mse2psnr(loss)

```

```

64     psnrs.append(psnr.detach().cpu().numpy())
65     iternums.append(i)
66
67
68     # Plotting the rendered image and PSNR over iterations.
69     plt.figure(figsize=(9, 3))
70
71     plt.subplot(131)
72     picture = rgb.cpu() # Copy the rendered image from GPU to CPU.
73     plt.imshow(picture)
74     plt.title(f'RGB Iter {i}')
75
76     plt.subplot(132)
77     picture = depth.cpu() * (rgb.cpu().mean(-1)>1e-2)
78     plt.imshow(picture, cmap='gray')
79     plt.title(f'Depth Iter {i}')
80
81     plt.subplot(133)
82     plt.plot(iternums, psnrs)
83     plt.title('PSNR')
84     plt.show()

```

- Task2.5 Please include the best picture (after parameter tuning) of your RGB prediction, depth prediction, and ground truth figure for different view points.

